

بسم الله الرحمن الرحيم

تکنولوژی کامپیوتر

جلسه ی سیزدهم
رپلیکیشن - همانندسازی

جلسه گذشته

اتکاپذیری، مقیاس پذیری و نگهداری پذیری

جلسه جدید

رپلیکیشن

- Replication is keeping an entire copy of the data on multiple machines
- We will assume that all of data can be stored in a single machine

چرا ریلیکیشن؟

■ مقیاس پذیری؟

- در افزایش چه لود پارامتری؟ با چه تغییری؟ روی چه پرفورمنس متریکی؟

■ اتکا پذیری؟

■ نگهداری پذیری؟

■ همچنین کمک به بهبود ریسپانس تایم؟ چطوری؟؟

اگہ دادہ مون تغییر نکنہ

روش‌های رپلیکیشن

- Single Leader Replication
- Multi-Leader Replication
- Leaderless Replication (جلسه‌ی بعد)

SINGLE LEADER REPLICATION

Leader / Follower

- Replica / Node
- Leader (Or Master / Primary)
 - *One of replicas will be Leader*
 - *All writes must go through the leader*
- Follower
 - *Catch data from leader*
 - *reads can come from any replica*

What happen on write?

- On writes, leader sends change log information to all followers
 - ***Synchronous*** - clear success/failure; clients must wait for all replicas
 - ***Asynchronous*** - lowest latency; no guarantee of durability after successful write
- Sync vs Async? Pros and Cons
 - *Single synchronous follower improves durability*
 - *Semi-Sync replication*

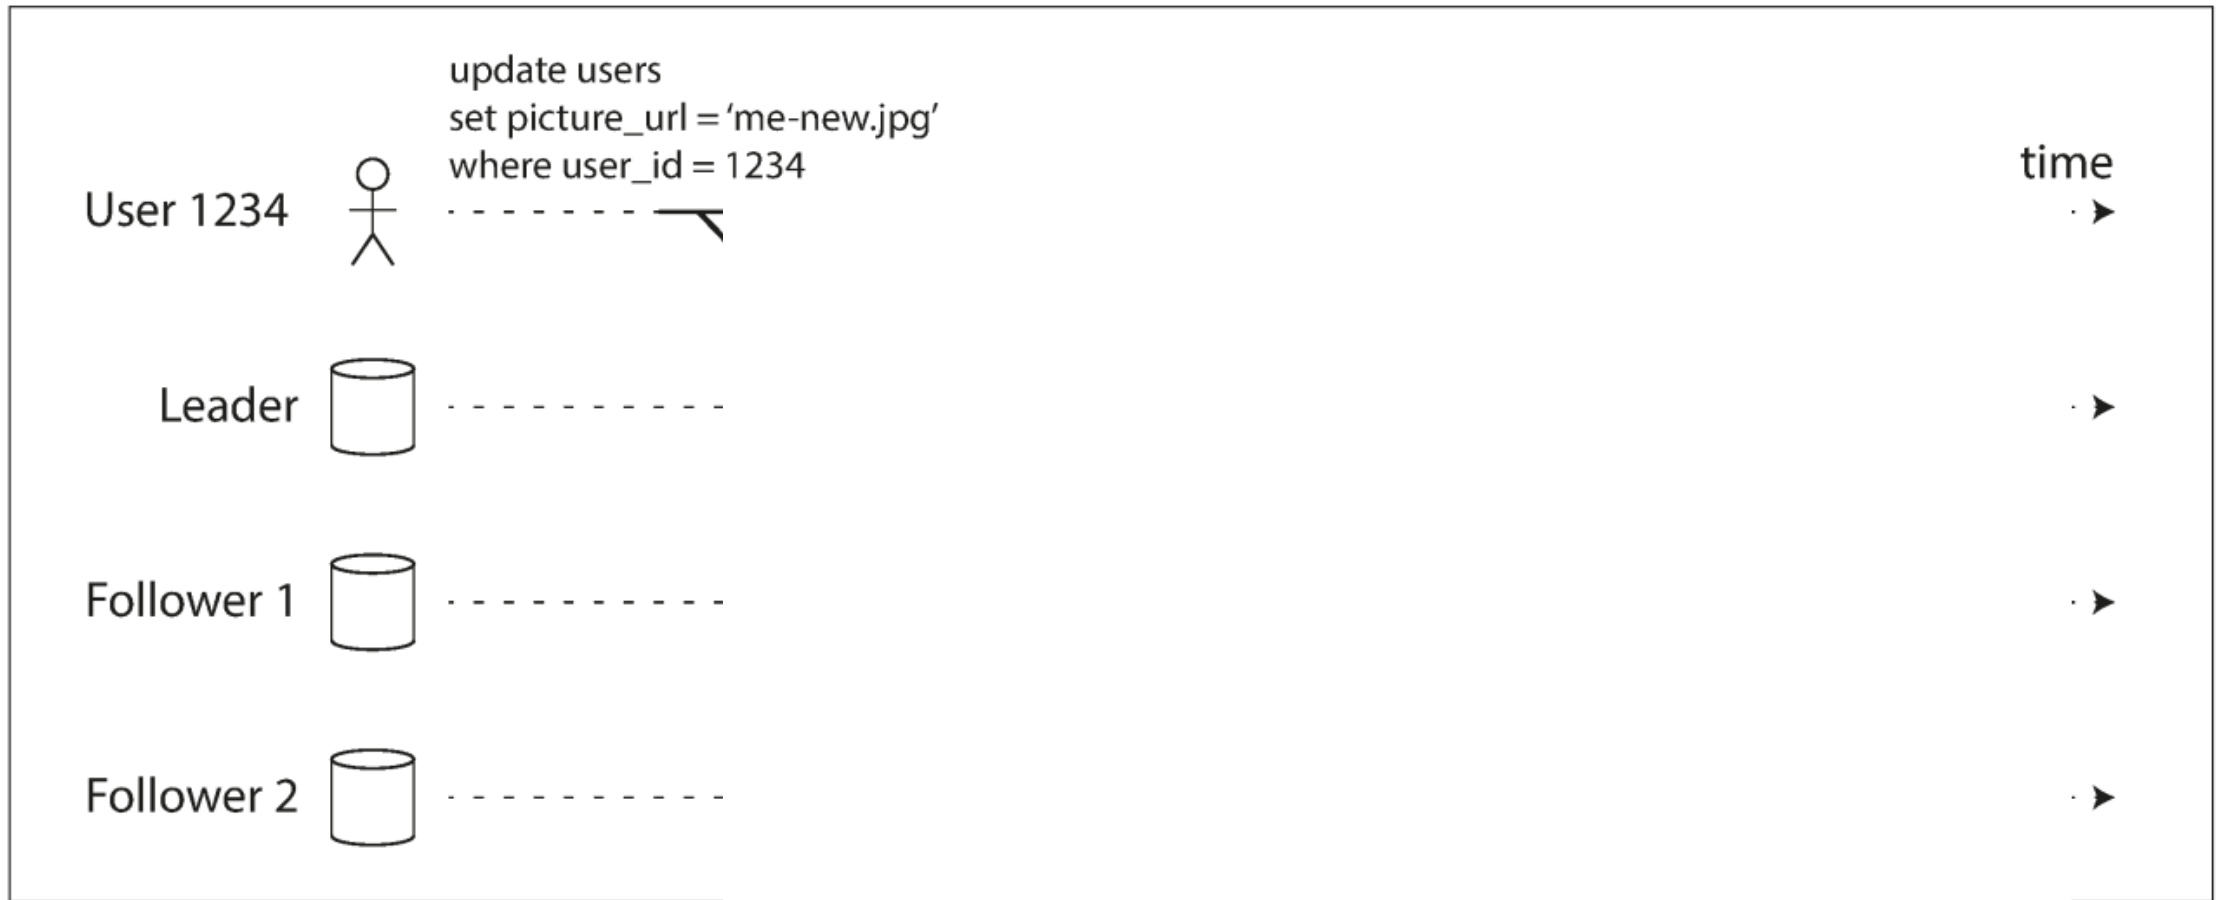


Figure 5-2. Leader-based replication with one synchronous and one asynchronous follower.




Figure 5-2. Leader-based replication with one synchronous and one asynchronous follower.




Figure 5-2. Leader-based replication with one synchronous and one asynchronous follower.

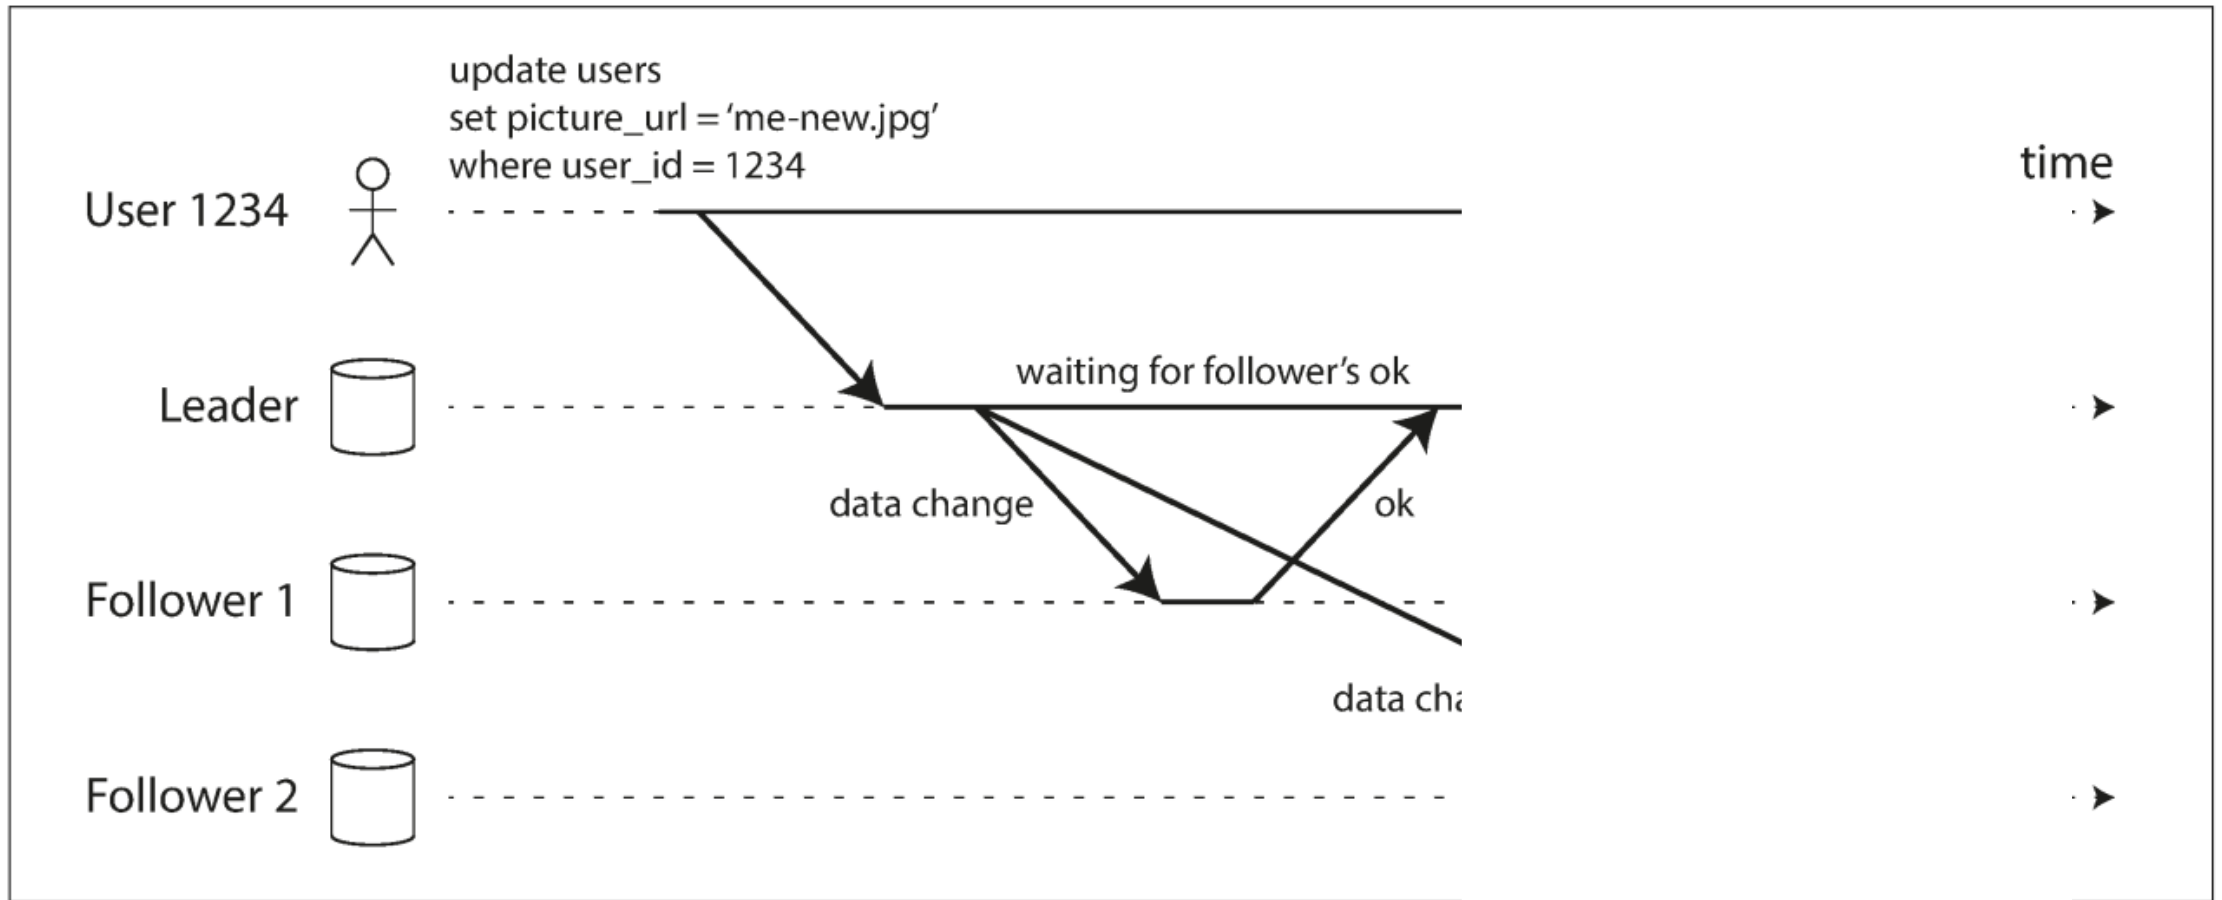


Figure 5-2. Leader-based replication with one synchronous and one asynchronous follower.

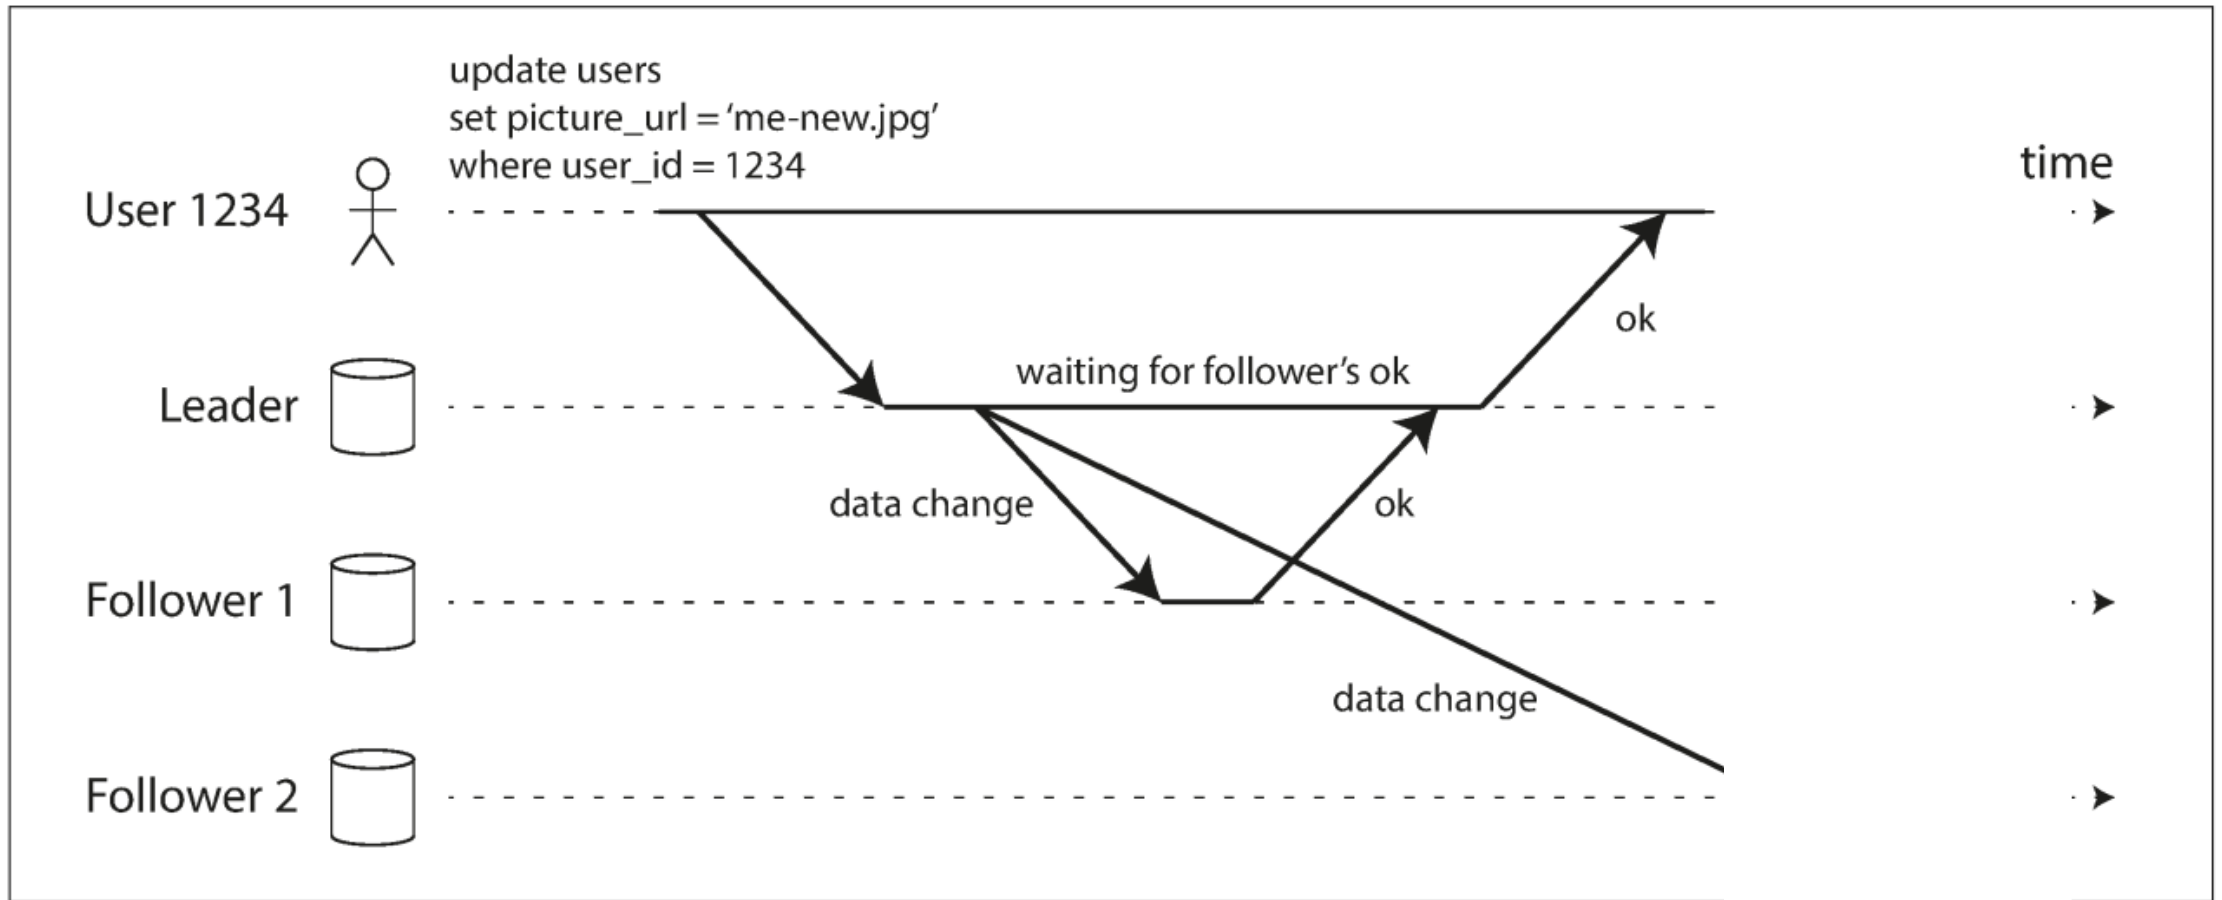


Figure 5-2. Leader-based replication with one synchronous and one asynchronous follower.

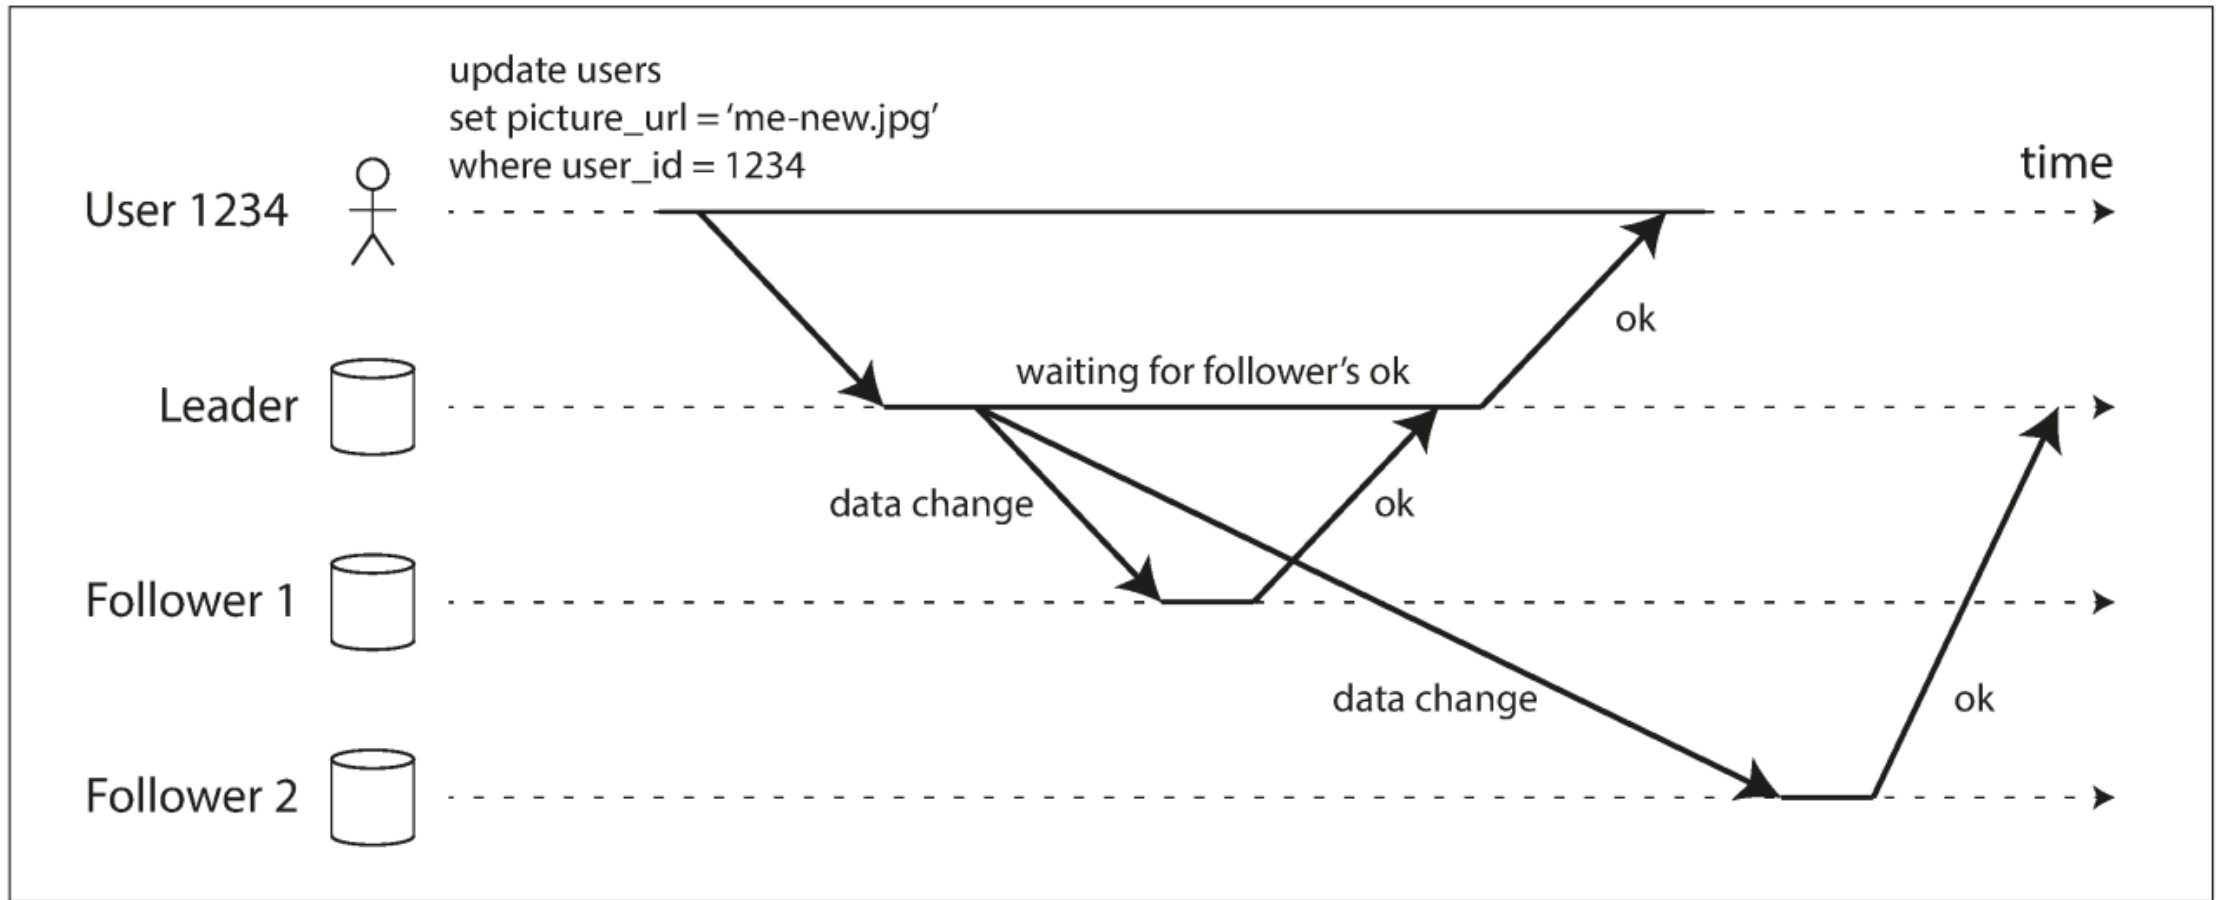


Figure 5-2. Leader-based replication with one synchronous and one asynchronous follower.

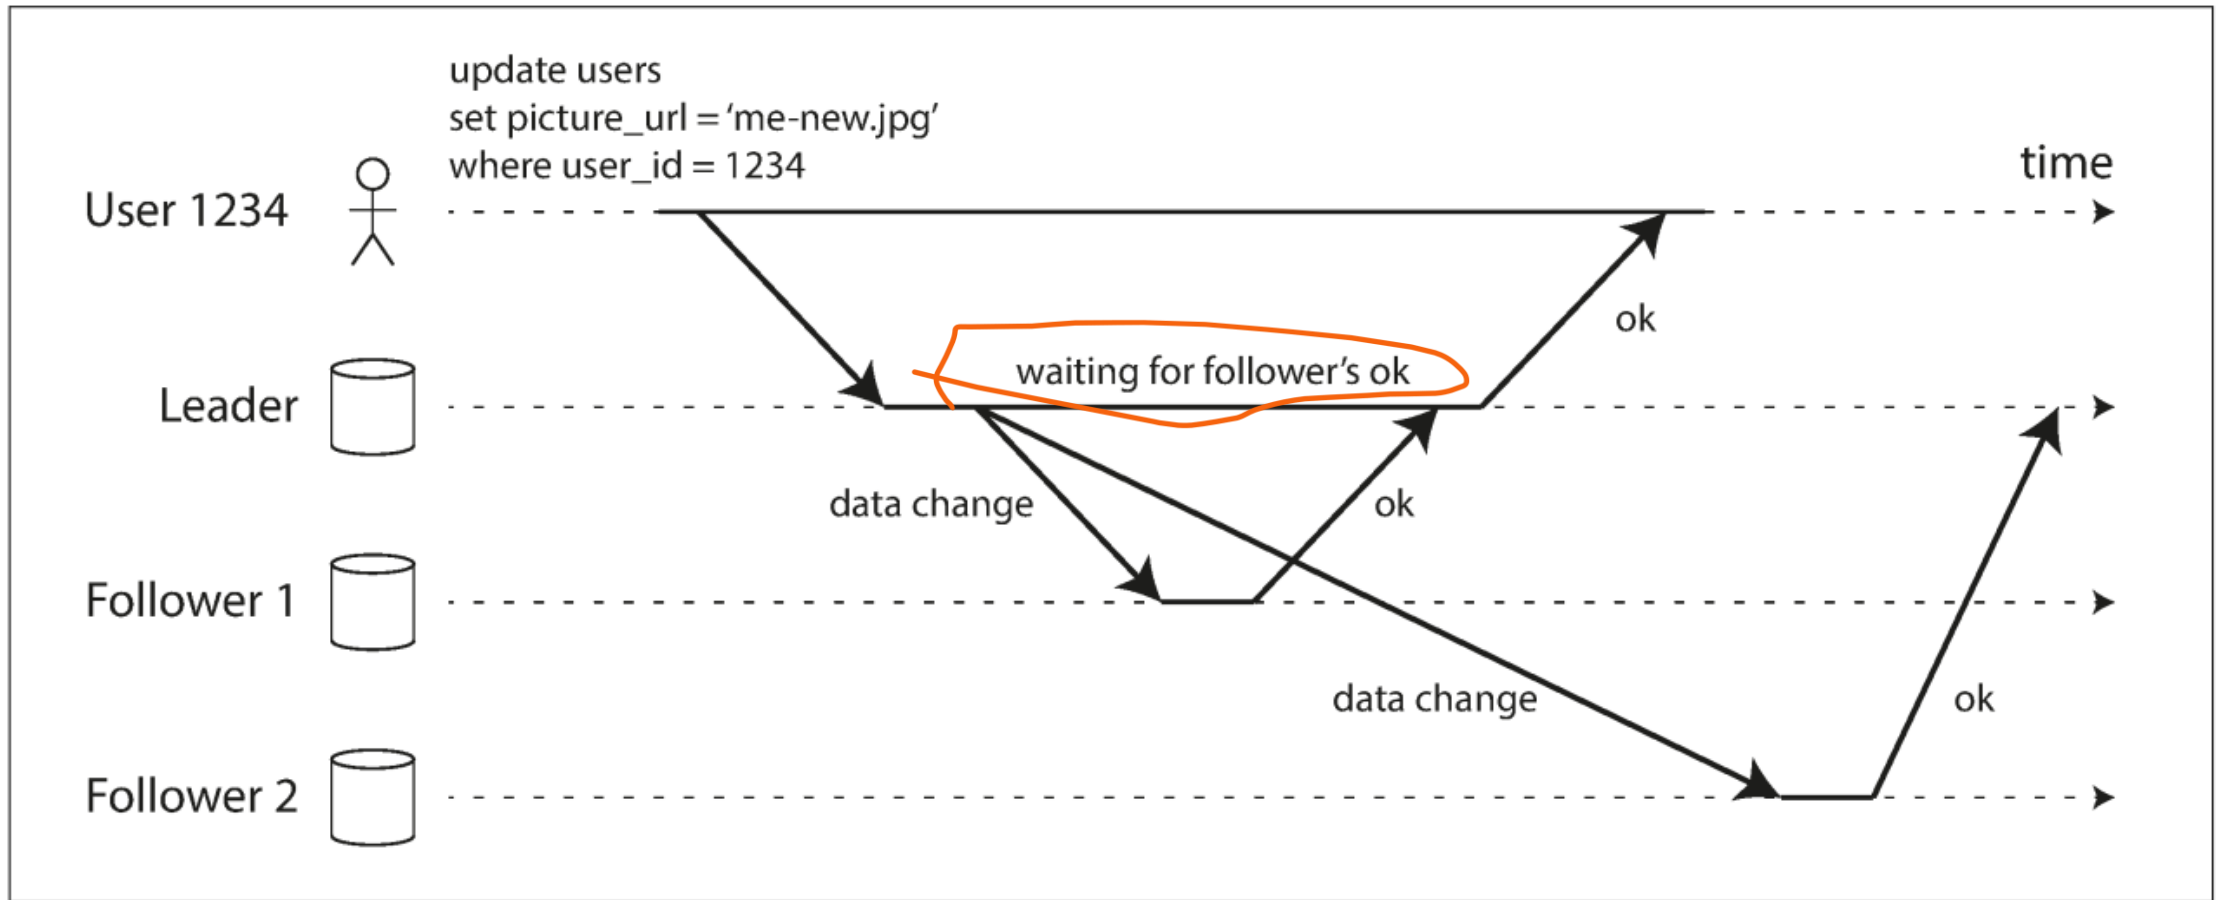


Figure 5-2. Leader-based replication with one synchronous and one asynchronous follower.

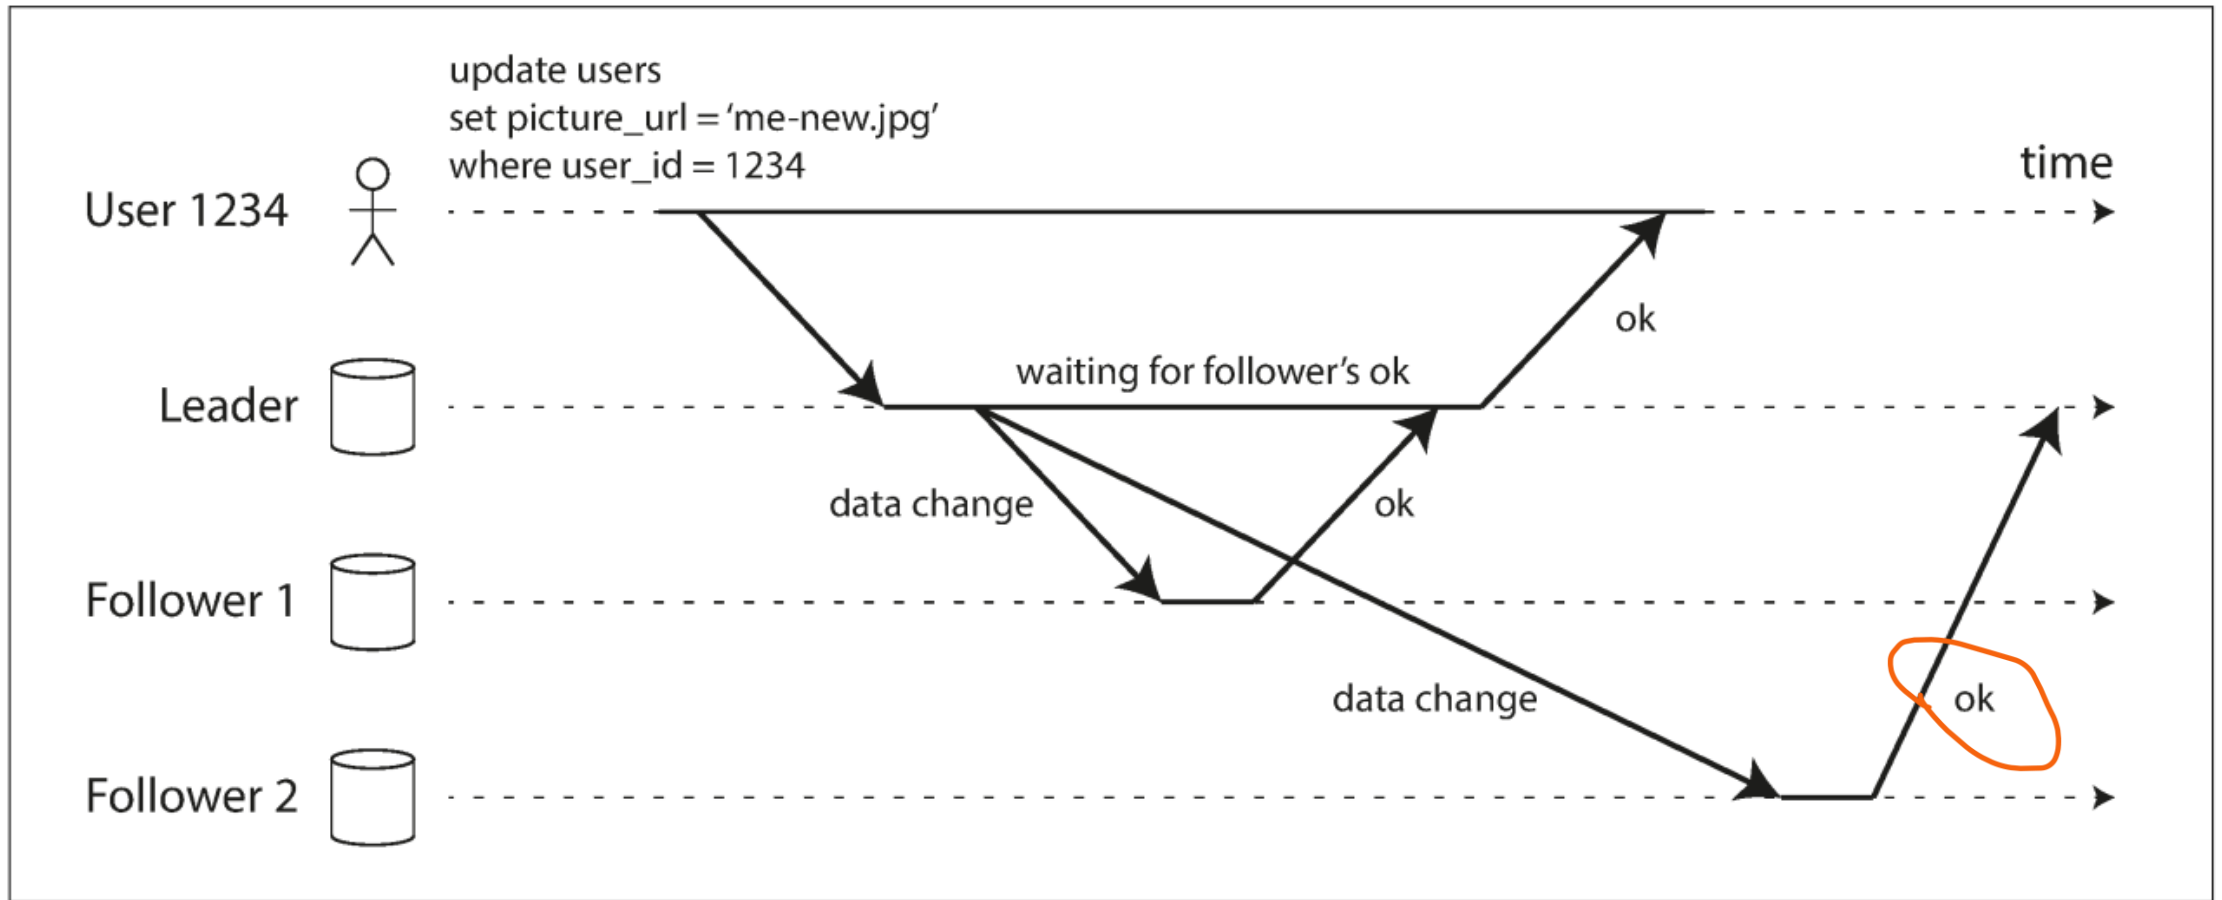


Figure 5-2. Leader-based replication with one synchronous and one asynchronous follower.

What happen if a replica fails?

- Leader
- Sync Replica
- Async Replica

How to add new replica?

- Why it is a problem?
- Snapshot / Backup of data?
 - *Some models:*
 - MVCC (Multi-Version Concurrency Control) model
 - *Most of SQL dbs*
 - Copy on write + Log-Based
 - *Redis*
 - ...

Leader Failure - Failover

- If leader fails -> promote one of the followers to be leader.
- Is it easy?
- How we determine the leader fails?
 - *Who determines?*
 - *Timeout? How long?*

Leader Failure - Failover

- How to choose the new leader? Which node is up to date?
- How to reconfigure system for new leader?
 - *Writes*
 - *Other replicas?*
- What happen to writes that may not in new leader?
 - *The Github failure example*

Leader Failure - Failover

- What happen when the failed leader come back?
 - *Split Brain problem!*
- Manual operation or automatic operation?

Replication Log

- Statement based replication
 - repeat the original command
 - nondeterministic commands!
 - RAND(), NOW()
 - Side effects, ...

Replication Log

- Write-ahead log (WAL) shipping
 - send detailed physical write change log
 - فلان بایت از فلان بلاک از بهمان فایل رو از این به اون تغییر دادم.
 - Problem with version upgrade!?

Replication Log

- Logical (row based) log shipping
 - send row level key and data
 - *On insert?*
 - *On update?*
 - *On delete?*
- Logical Replication Vs. Physical Replication?

Replication Log

- Trigger-based replication
 - *Manually coded*
 - *Why we can need it?*
 - *Flexible*
 - *Higher overhead*
 - *Risk bugs!*
- Postgresql Example

REPLICATION LAG

Replication Lag

- How Long?
 - *Seconds?*
 - *Minutes?*
 - *Hours / Days ??!*

Consistency

- Strong Consistency
- Eventual Consistency

Read your own writes

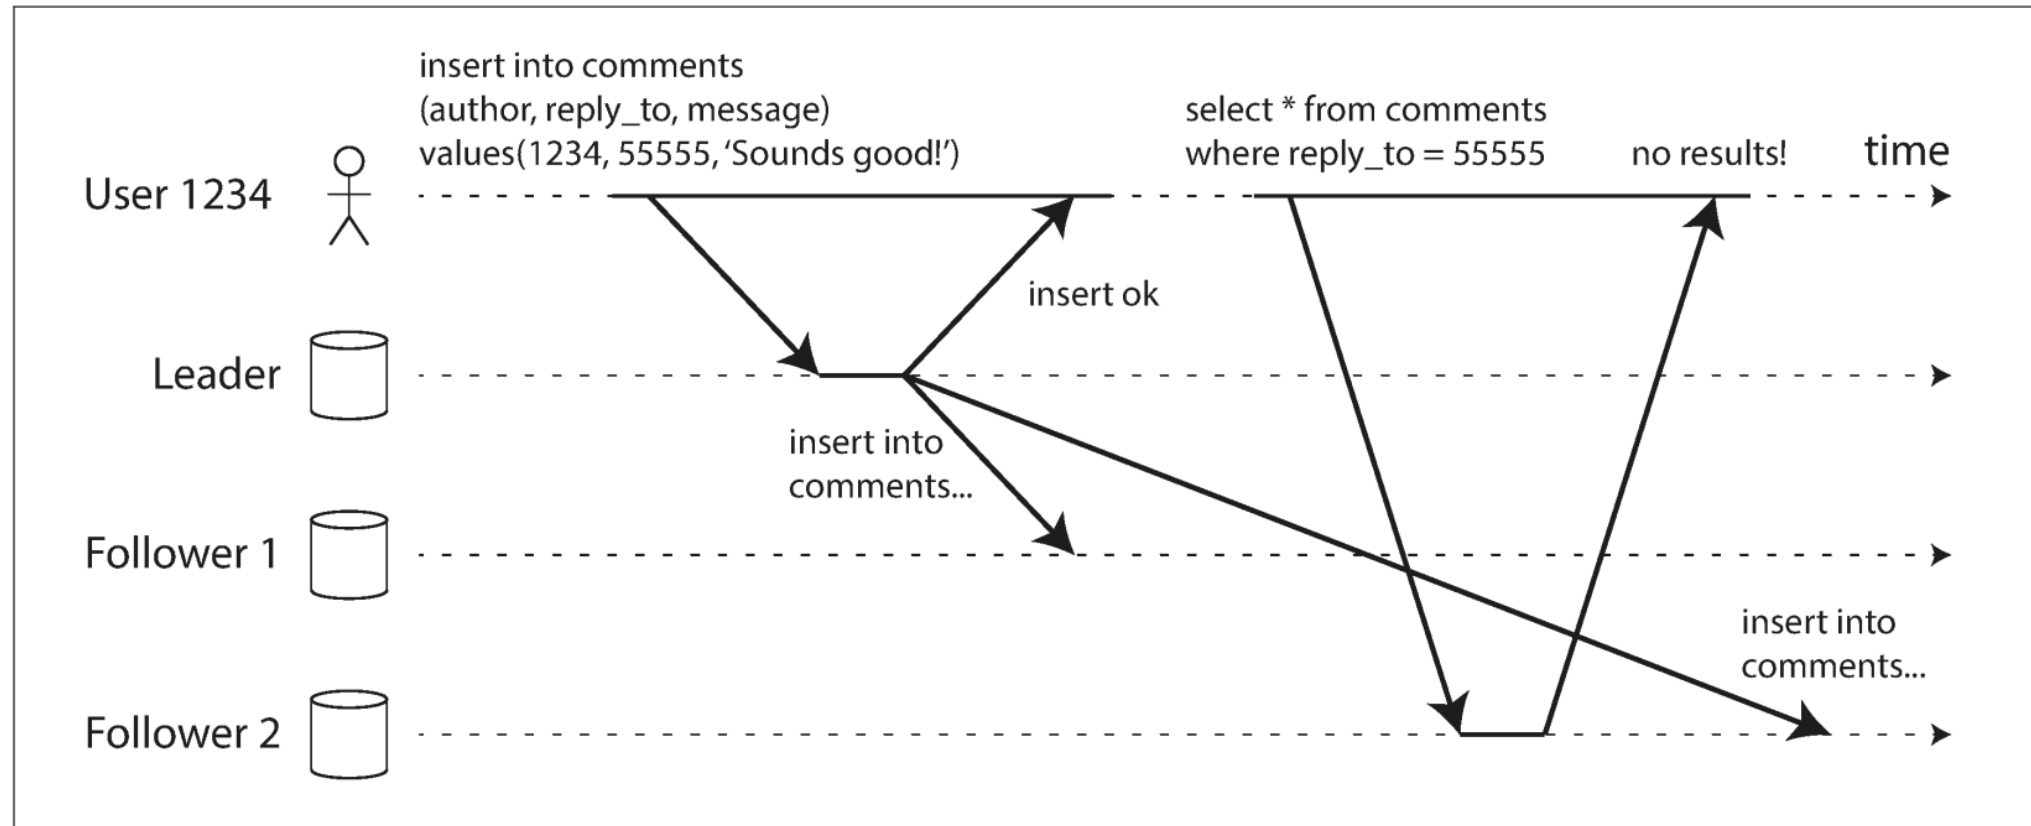


Figure 5-3. A user makes a write, followed by a read from a stale replica. To prevent this anomaly, we need read-after-write consistency.

Read-After-Write Consistency

- Or read-your-writes consistency
- Solutions?
 - *if few things modifiable, read them all from the leader*
 - *Track last update and read everything from leader for a short period*
 - *Save last timestamp of write*
- Cross-device read-after-write consistency?

Monotonic reads

- Comments of a site
- user makes two reads, the second is from a more lagged replica
 - *This will look to the user like information went backward in time*
- Solution?
 - *Each user, read from same replica...*

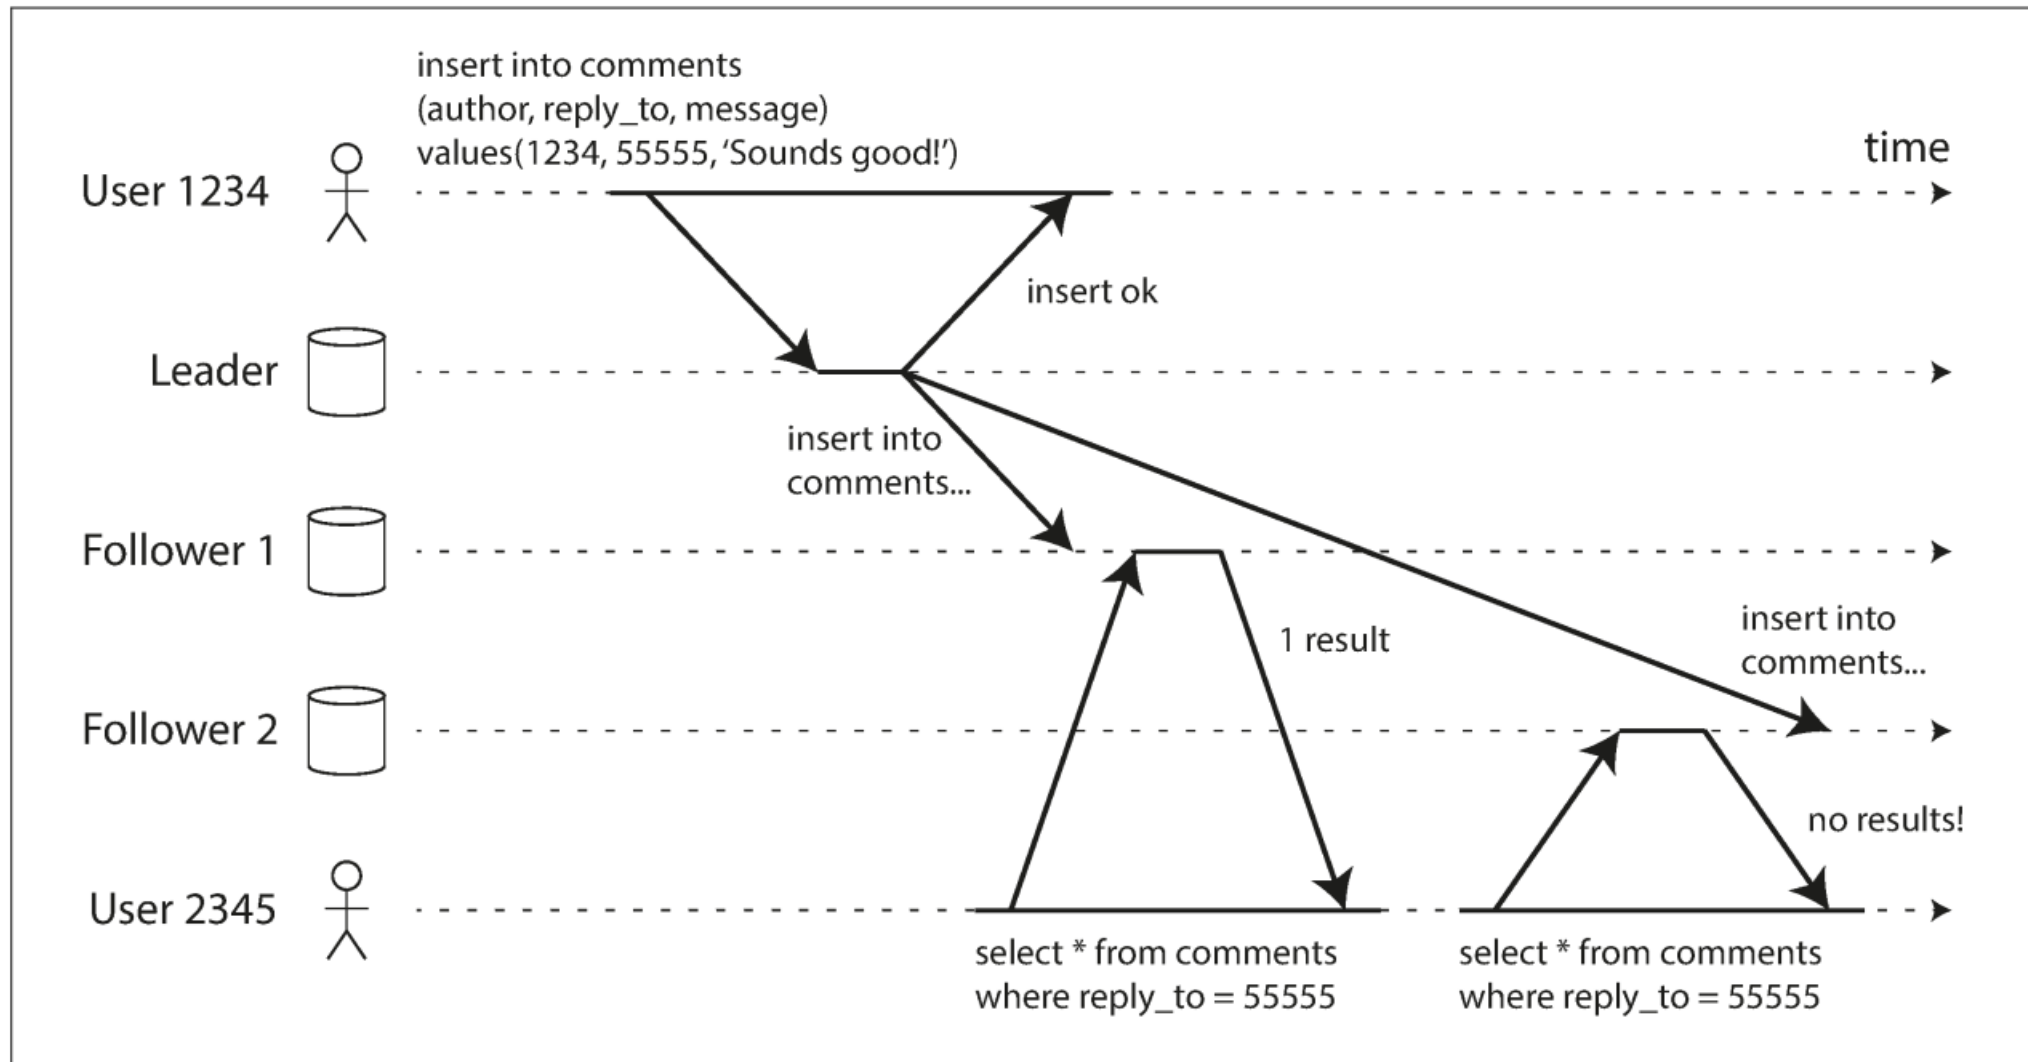


Figure 5-4. A user first reads from a fresh replica, then from a stale replica. Time appears to go backward. To prevent this anomaly, we need monotonic reads.

Monotonic reads

- Solution?

- *Each user, read from same replica...*

Consistent prefix reads

- Assume we have multiple databases / partitions

Mr. Poons

How far into the future can you see, Mrs. Cake?

Mrs. Cake

About ten seconds usually, Mr. Poons.

Consistent prefix reads


- Now a third-person, read data and observe following:

Mrs. Cake

About ten seconds usually, Mr. Poons.

Mr. Poons

How far into the future can you see, Mrs. Cake?



Solutions for Replication Lag?

- Transactions?

MULTI-LEADER REPLICATION

Multi-Leader Replication

- each leader also acts as a follower to the other leaders
- Use cases?
 - *Multi-data center operation*
 - *Clients with offline operation*
 - *Collaborative editing*

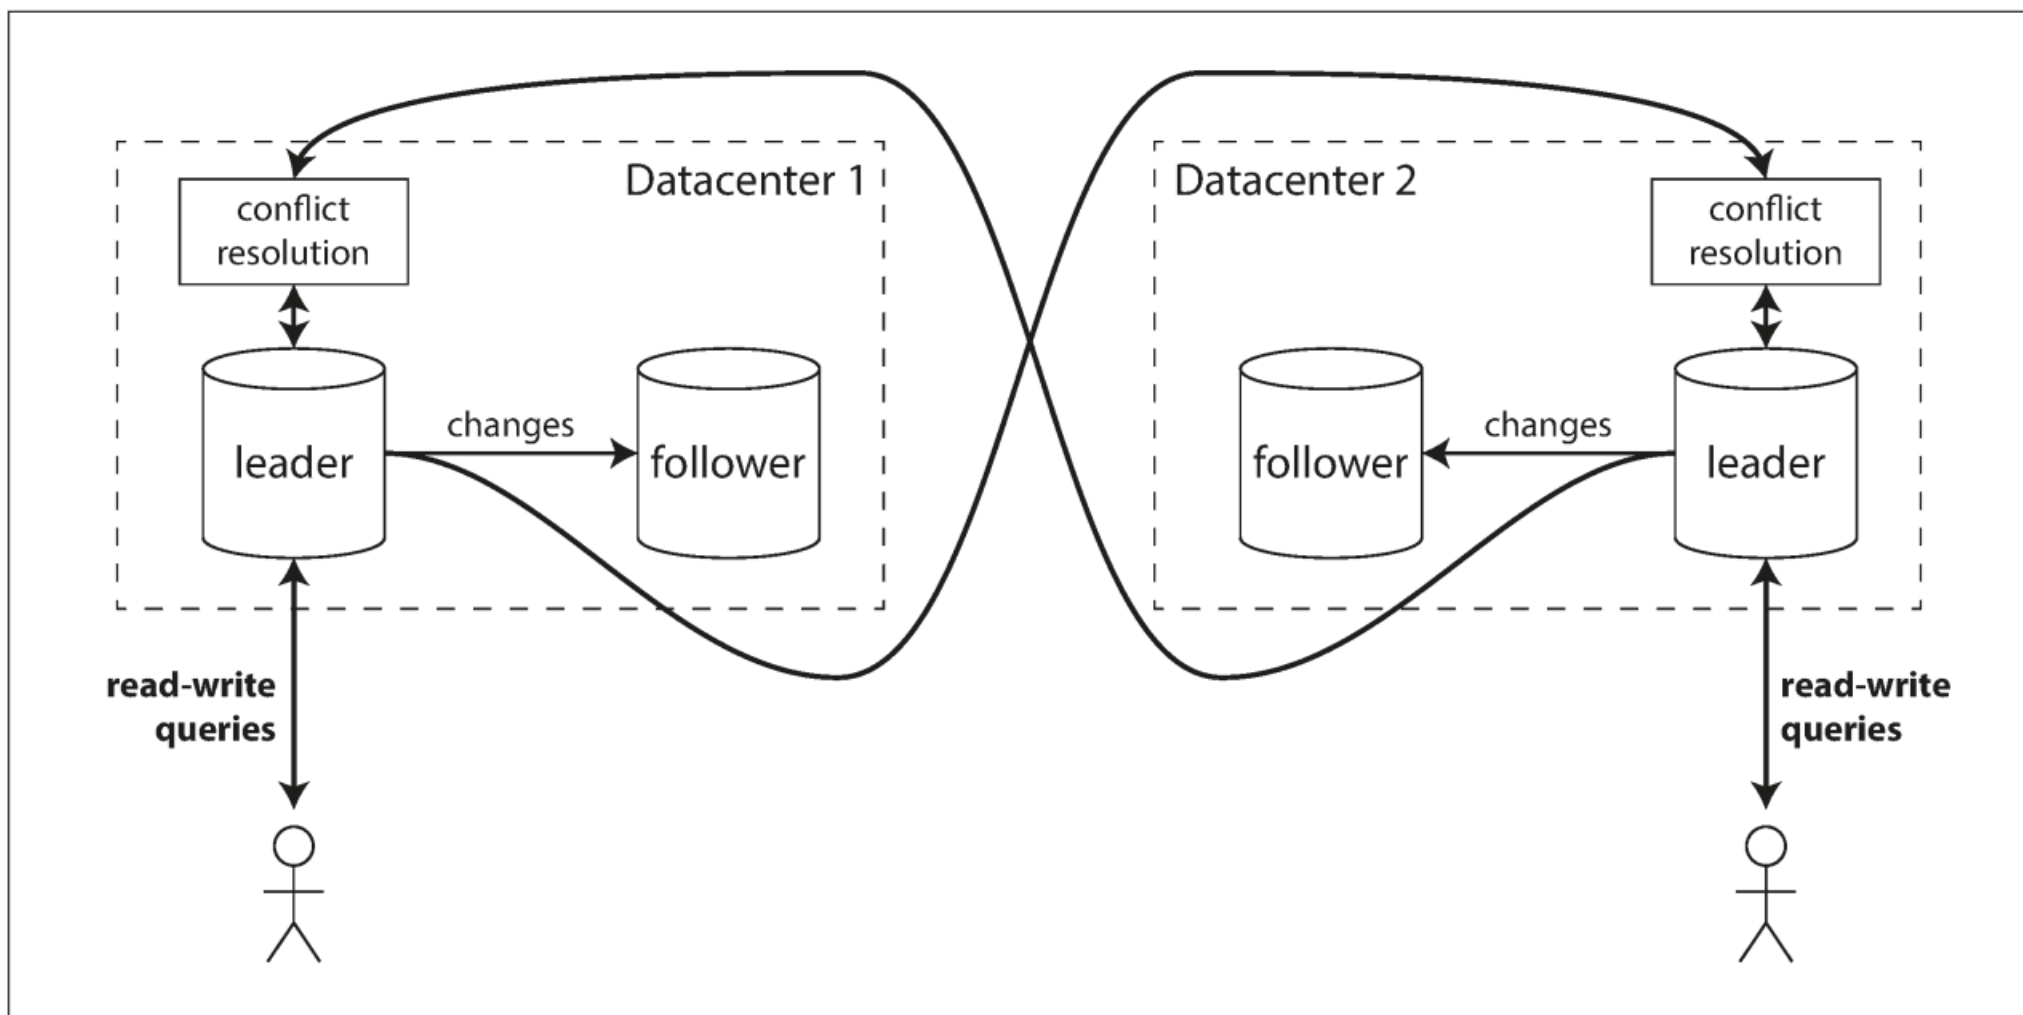


Figure 5-6. Multi-leader replication across multiple datacenters.

Handling Write Conflicts




Figure 5-7. A write conflict caused by two leaders concurrently updating the same record.

Handling Write Conflicts

- Downside of multi-leader is that write conflicts can happen, requiring conflict resolution
- Why it's hard and have bad User Experience?
 - Conflicts are generally detected asynchronously, so too late to prompt user

Handling Write Conflicts

■ Conflict avoidance

- *Assigning each record a home datacenter works until a failure or they move*
- *If you need to change the designated leader for a record -> conflict may happen 😞*

جلسه‌ی بعد