

بسم الله الرحمن الرحيم

تکنولوژی کامپیوتر

جلسه‌ی چهاردهم
رپلیکیشن بدون رهبر

جلسه گذشته

رپلیکیشن

رپلیکیشن

- Replication is keeping an entire copy of the data on multiple machines
- We will assume that all of data can be stored in a single machine

Leader / Follower

- Replica / Node
- Leader (Or Master / Primary)
 - *One of replicas will be Leader*
 - *All writes must go through the leader*
- Follower
 - *Catch data from leader*
 - *reads can come from any replica*

Multi-Leader Replication

- each leader also acts as a follower to the other leaders
- Use cases?
 - *Multi-data center operation*
 - *Clients with offline operation*
 - *Collaborative editing*

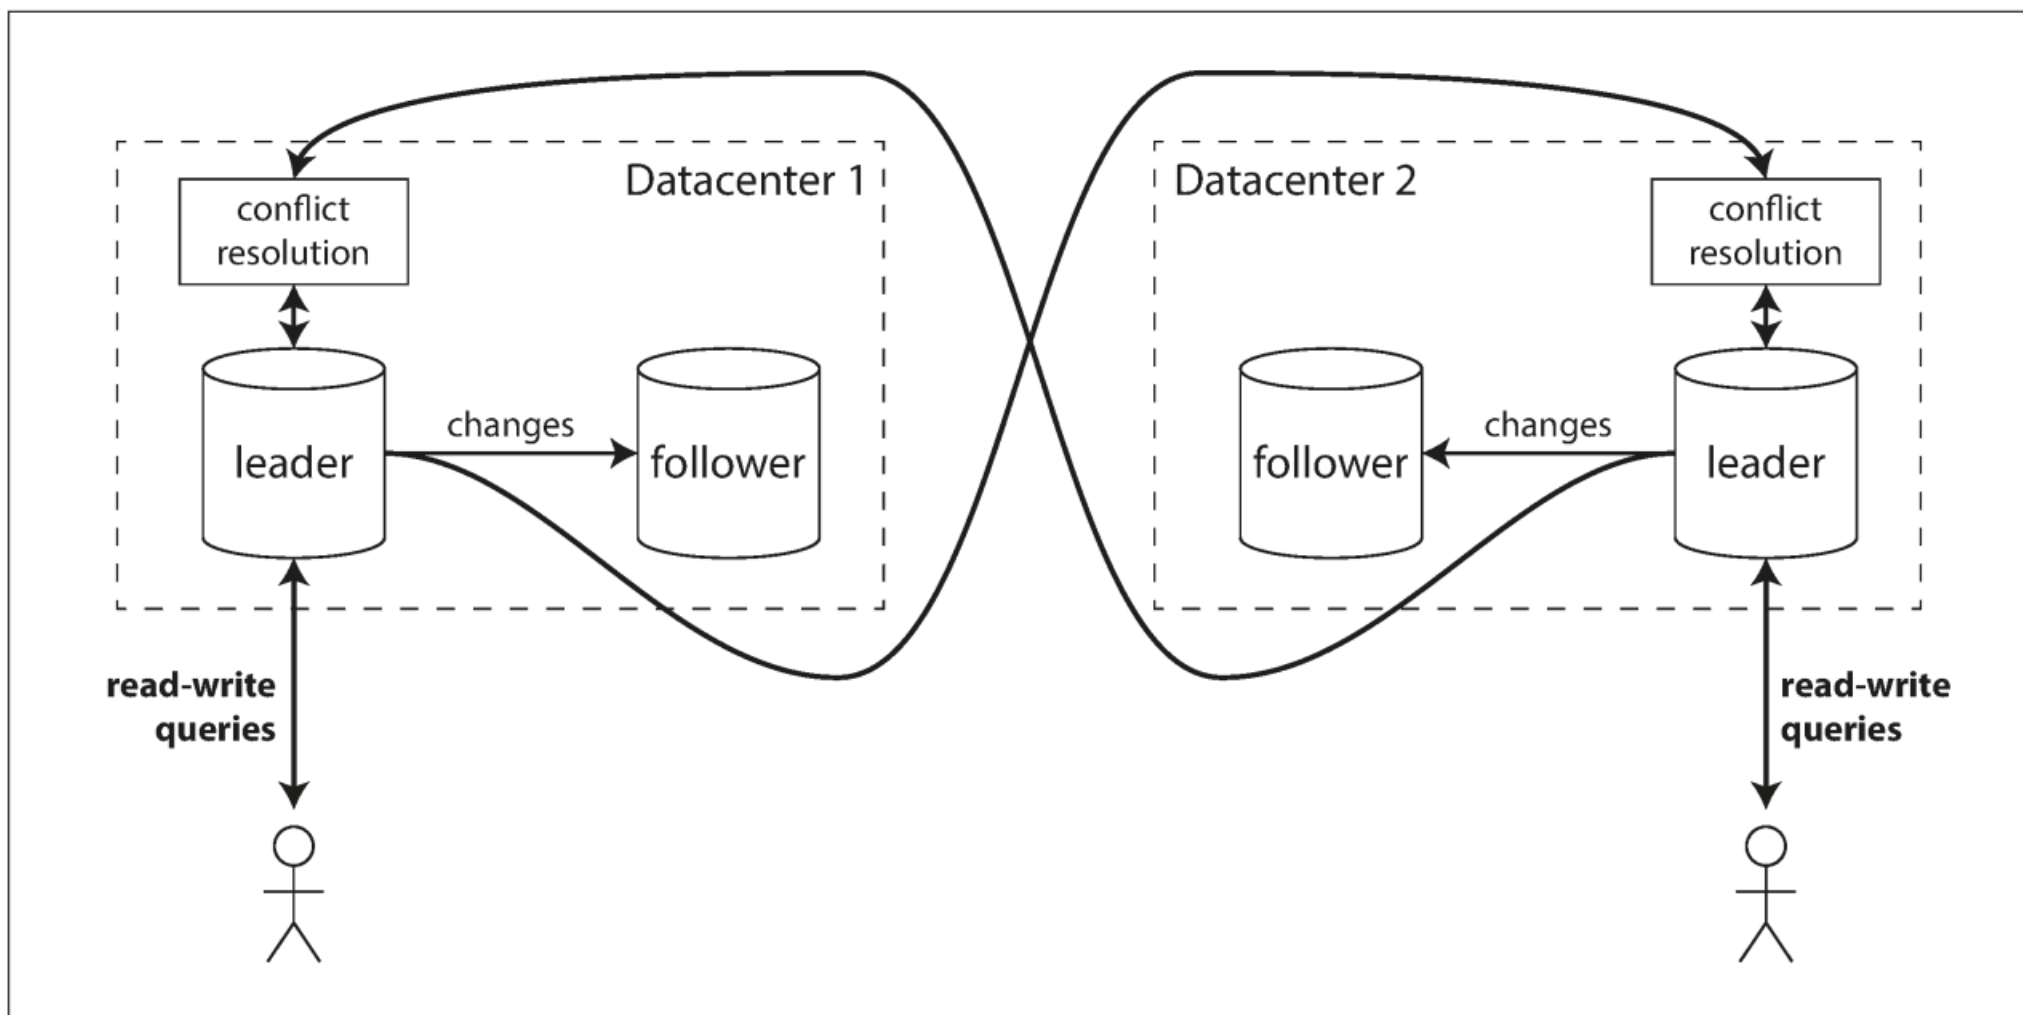


Figure 5-6. Multi-leader replication across multiple datacenters.

Handling Write Conflicts

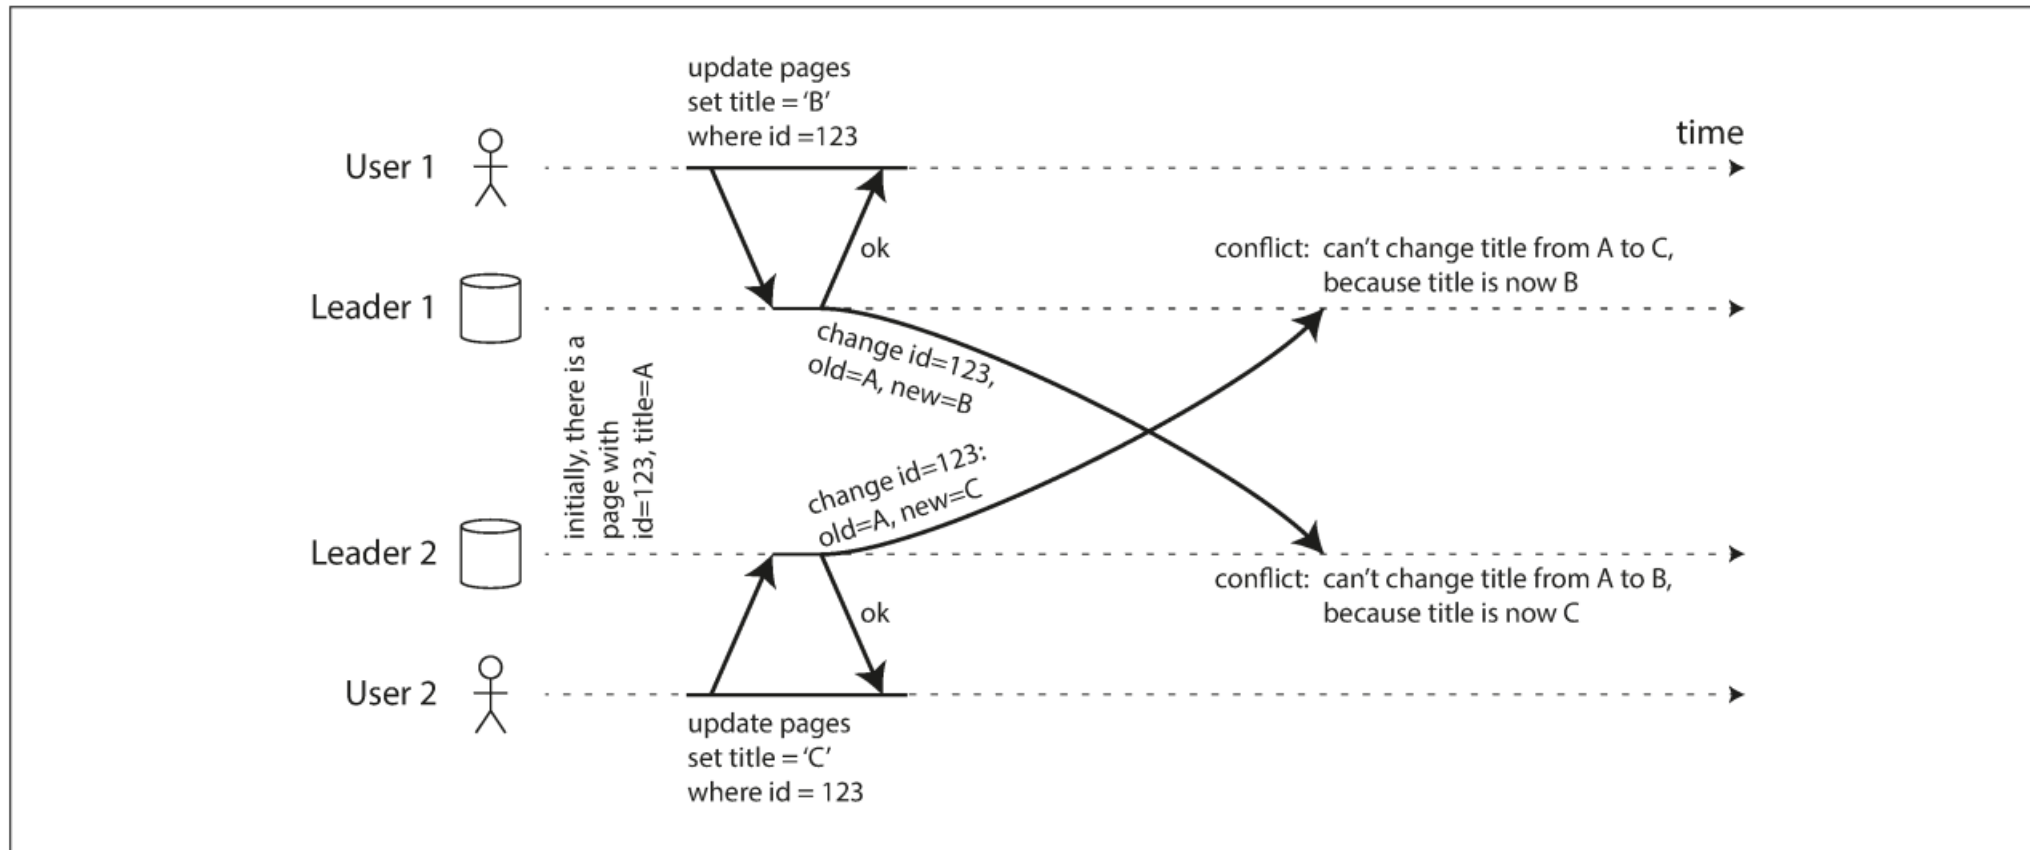


Figure 5-7. A write conflict caused by two leaders concurrently updating the same record.

Handling Write Conflicts

- Downside of multi-leader is that write conflicts can happen, requiring conflict resolution
- Why it's hard and have bad User Experience?
 - Conflicts are generally detected asynchronously, so too late to prompt user

Handling Write Conflicts

■ Conflict avoidance

- *Assigning each record a home datacenter works until a failure or they move*
- *If you need to change the designated leader for a record -> conflict may happen 😞*

جلسه‌ی جدید

WRITE CONFLICTS

Handling Write Conflicts

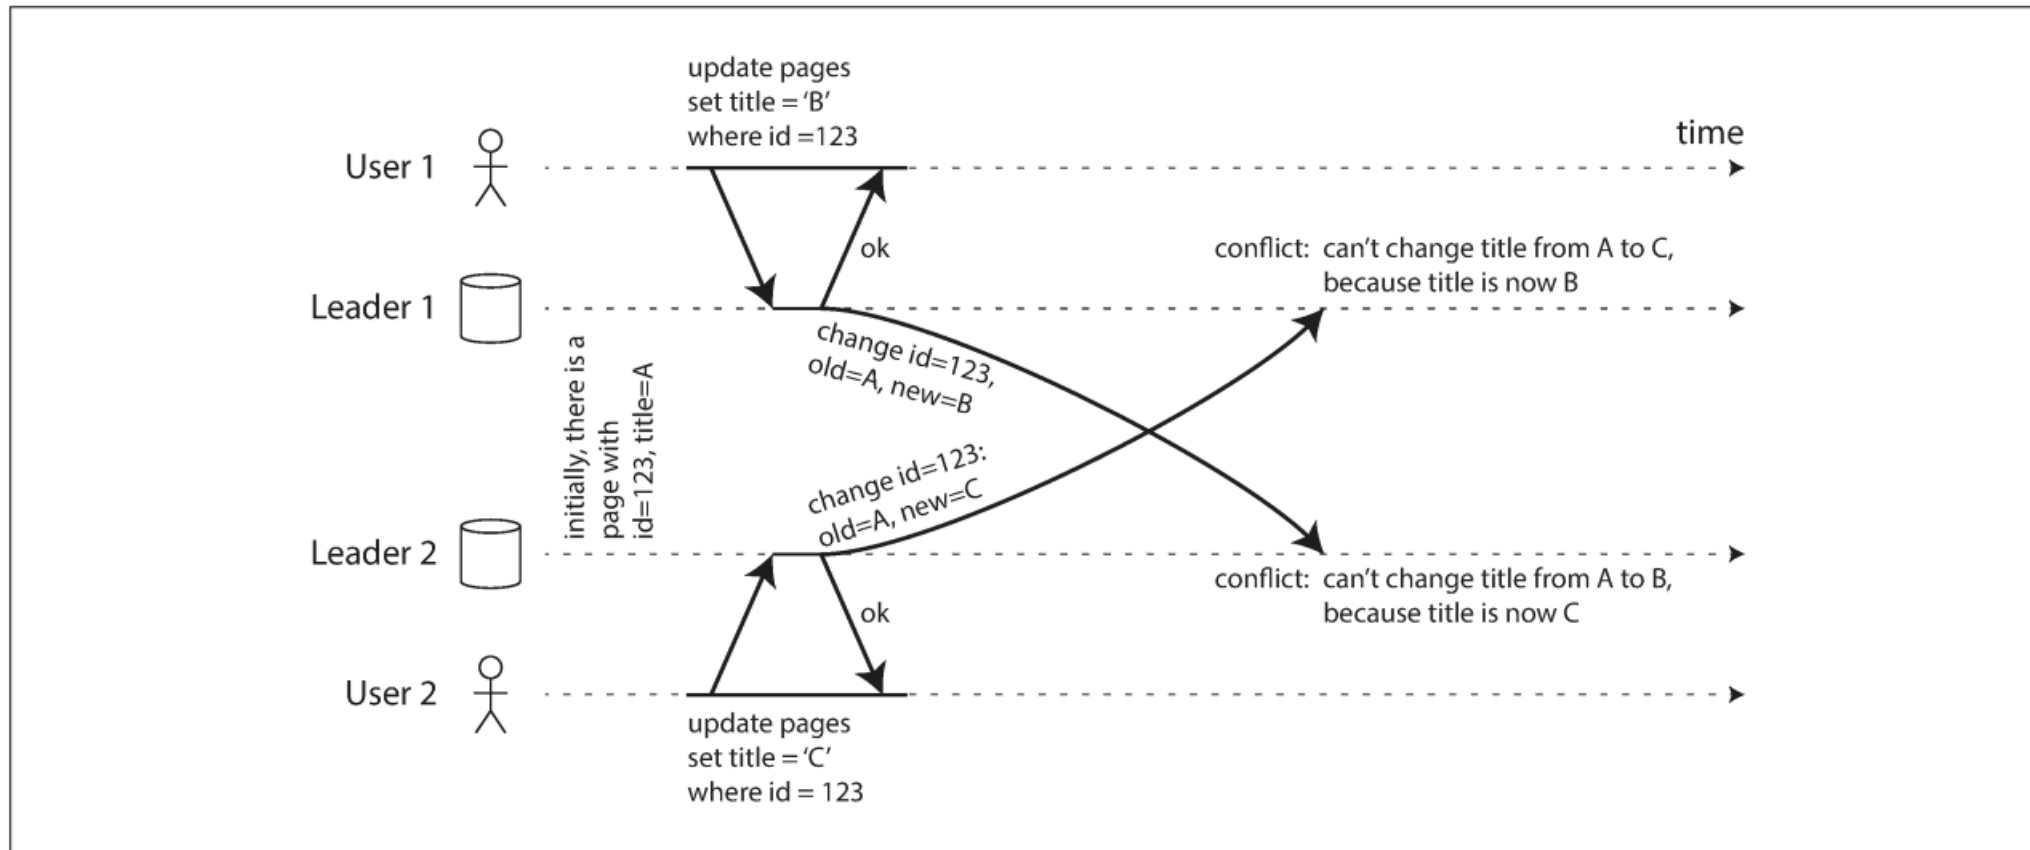


Figure 5-7. A write conflict caused by two leaders concurrently updating the same record.

Converging toward a consistent state

- What happens if each replica simply applied writes in the order that it saw the write?

Converging toward a consistent state

■ Last Write Wins (LWW)

- *Each write request contains a unique id/timestamp*
- *The request with bigger id wins*
- *It is dangerously prone to data loss*
- *if there are several concurrent writes to the same key:*
 - all reported as successful to the client
 - But only one of the writes will survive and the others will be silently discarded

LWW

- some situations, such as caching, in which lost writes are perhaps acceptable
- If losing data is not acceptable, LWW is a poor choice for conflict resolution
- The only safe way of using a database with LWW is to ensure that a key is only written once and thereafter treated as immutable

Pre-define arbitrary replica precedence rules

- Give each replica a unique ID
- Writes that originated at a higher-numbered replica always take precedence over writes that originated at a lower-numbered replica.
- This approach also implies data loss.

Merge Values together?

- Somehow merge the values together
 - *e.g., order them alphabetically*
 - *then concatenate them*
- Example:
 - *Concurrent writes: B and C*
 - *Conflict resolution: B/C*

Record log and resolve later

Custom conflict resolution login

- On write
- On read

Automatic Conflict Resolution Datatypes?

- Conflict-free replicated datatypes (CRDTs)
 - *Family of data structures for sets, maps, ordered lists, counters, ... that can concurrently edited and automatically resolve conflicts in sensible ways.*
- Mergeable persistent data structures
 - *Track history and merge like git.*
- Operational transformation
 - *Conflict resolution of collaborative editing (google docs)*
 - *Concurrent editing of an ordered list of items*

فقط نوشتن همزمان یک فیلد کانفلیکت هست؟

- Meeting room booking system
- در جلسات آینده در مورد روش‌های حلش بیشتر حرف می‌زنیم.

Multi-Leader Replication Topologies

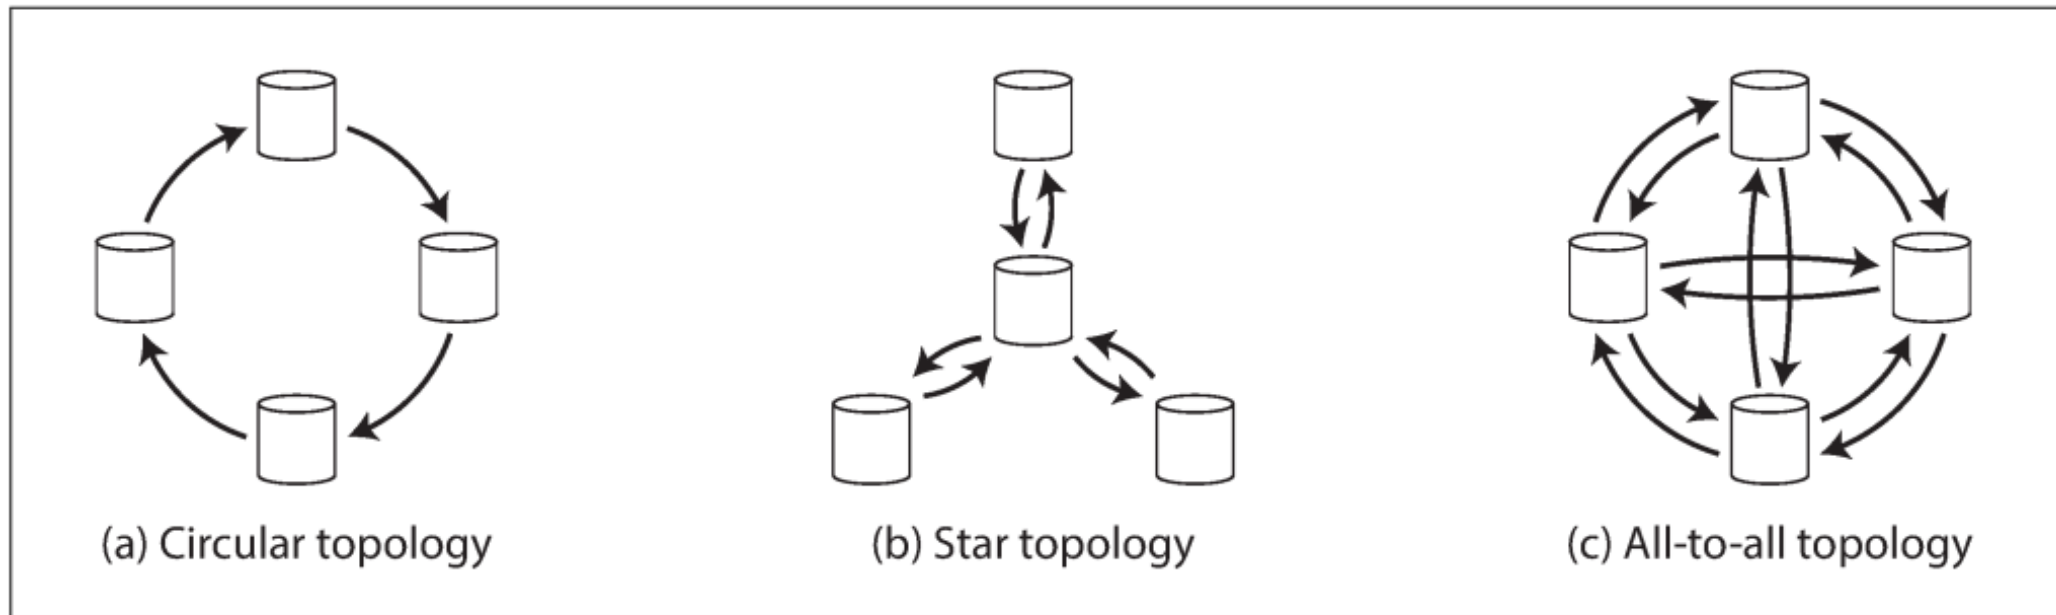


Figure 5-8. Three example topologies in which multi-leader replication can be set up.

Multi-Leader Replication Topologies

- In circular and star
 - *node receive writes from one node should forward it to other node.*
 - *How to prevent infinite replication loops?*
 - *What happed if one node fails?*

Multi-Leader Replication Topologies

■ All-to-all topologies

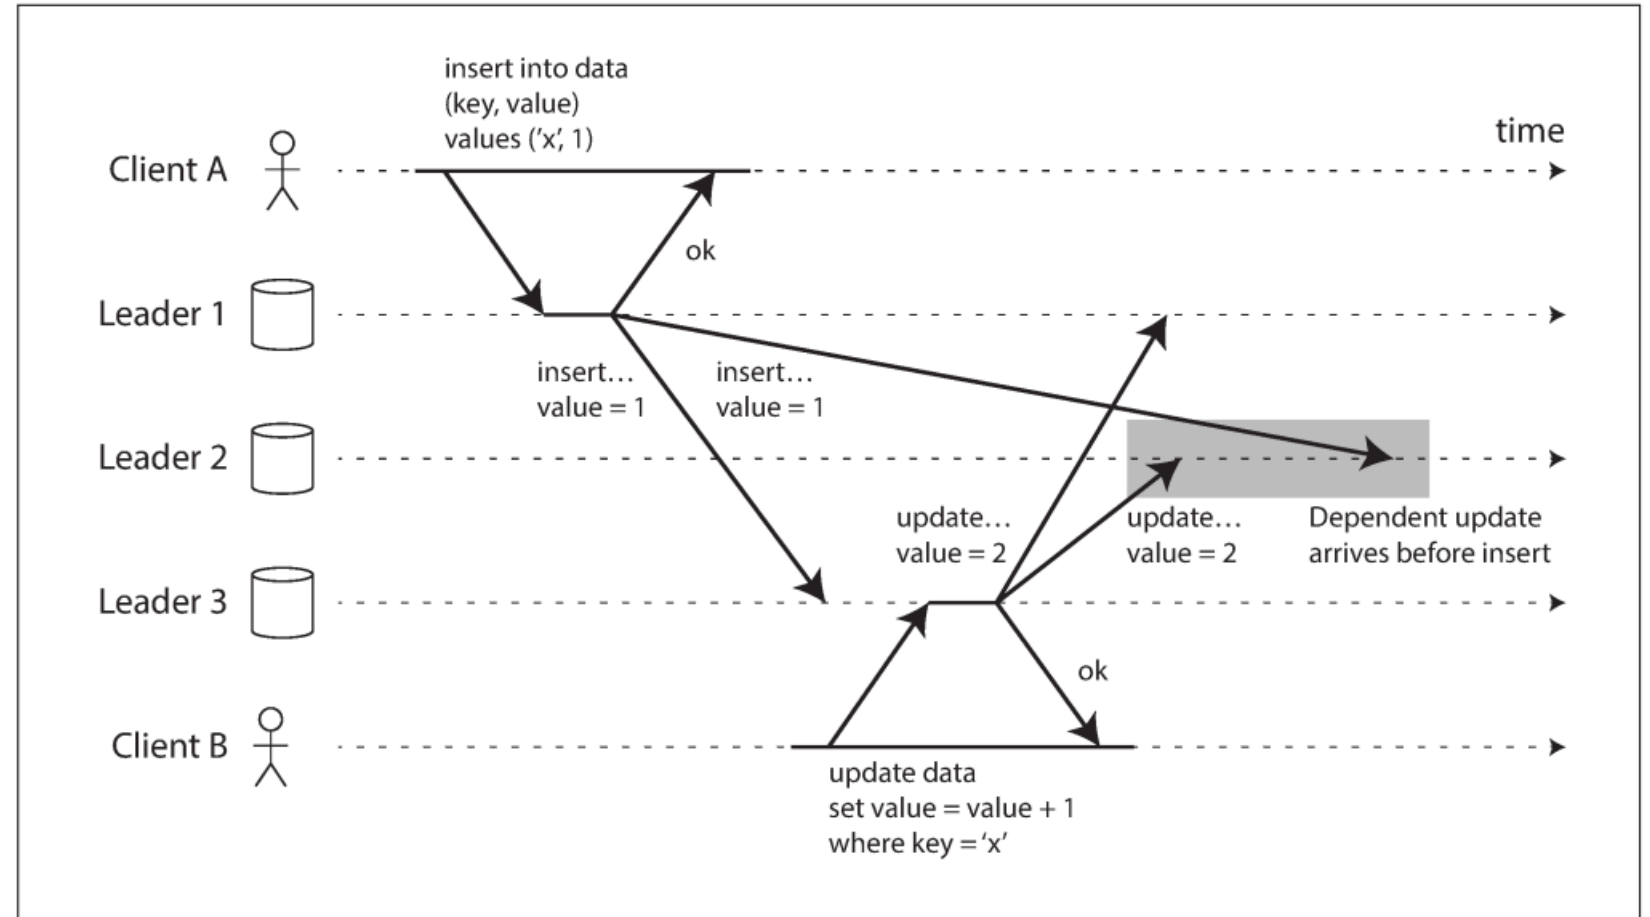


Figure 5-9. With multi-leader replication, writes may arrive in the wrong order at some replicas.

LEADERLESS REPLICATION

- Leader based replication: client send data to one node, and the database takes care of copying that write to the other replicas.
- Writes send to multiple replicas
 - *Client directly send writes to replicas*
 - *Or a coordinator handles that*

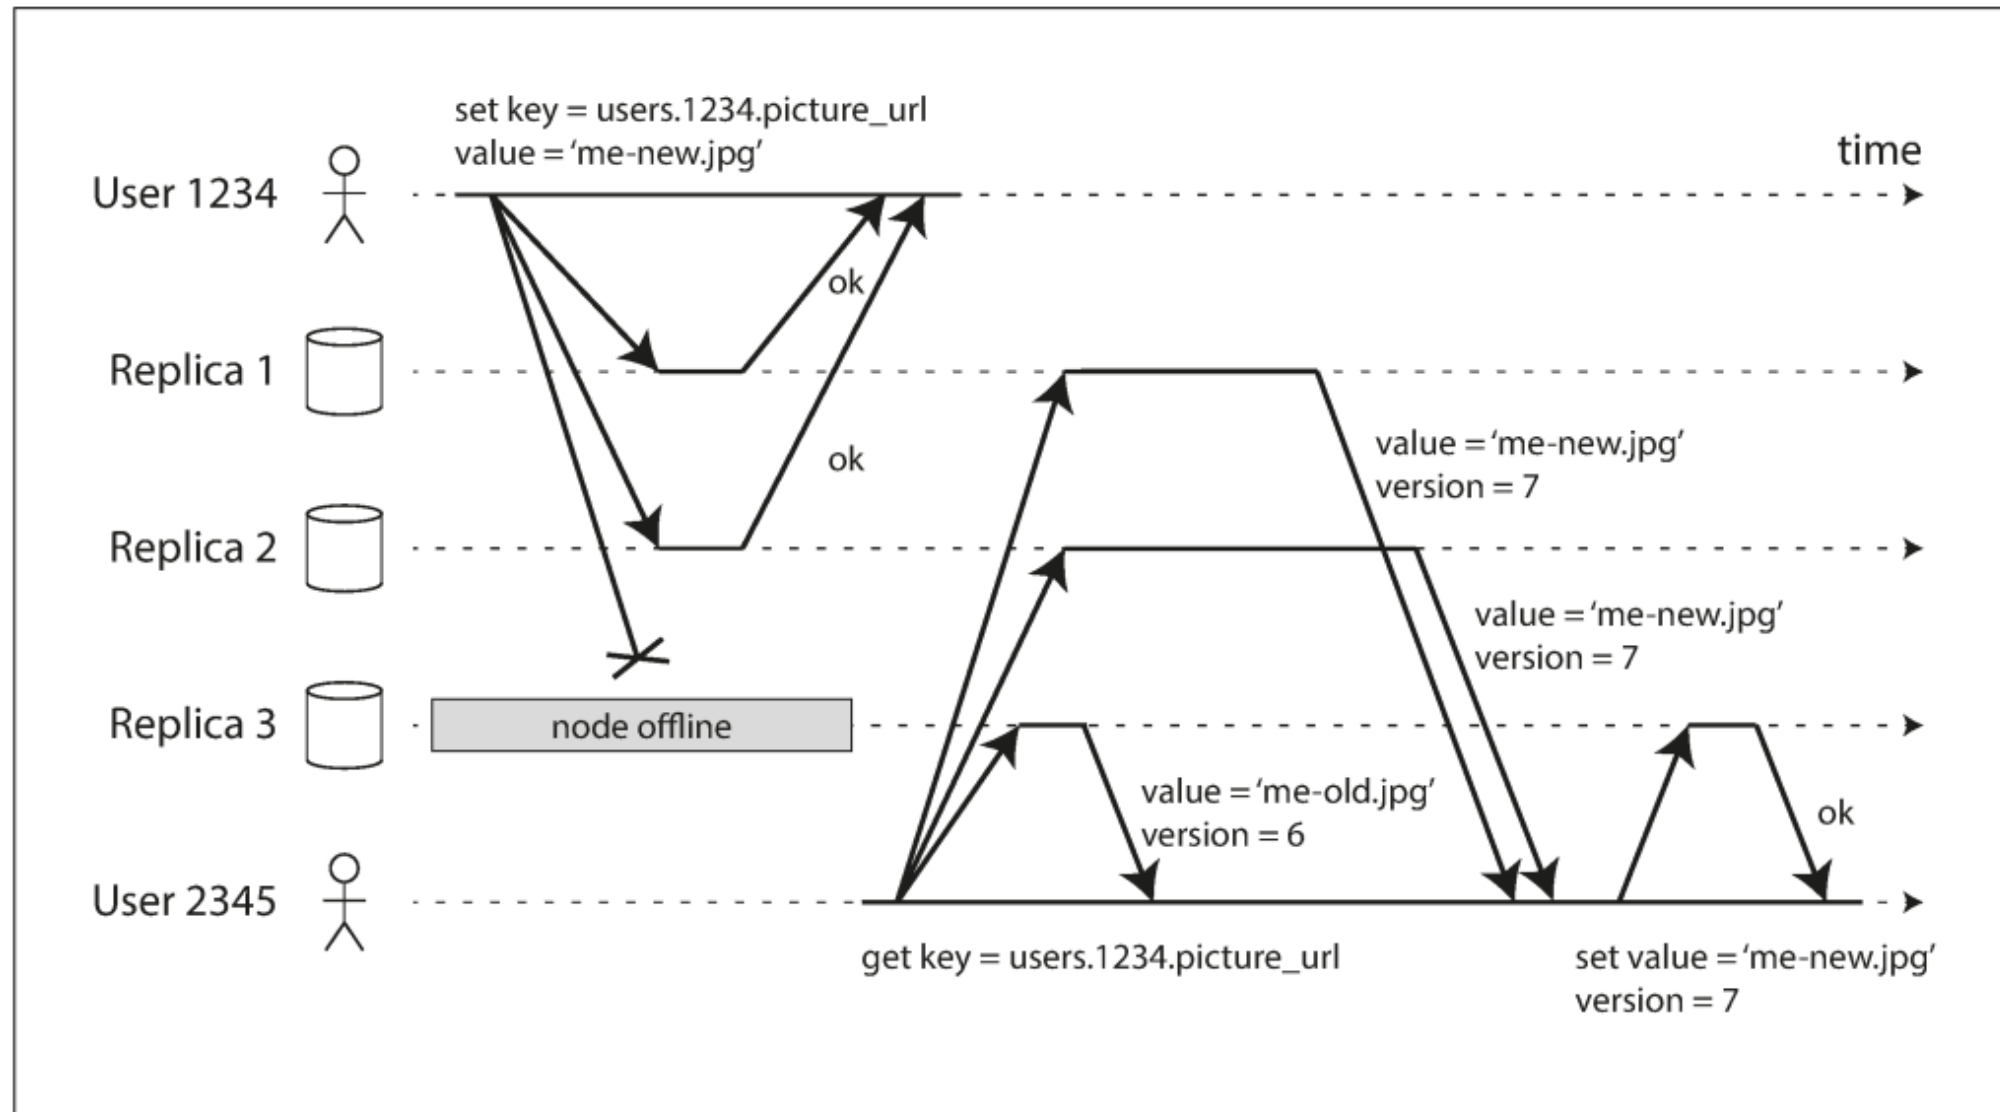


Figure 5-10. A quorum write, quorum read, and read repair after a node outage.

Writing to the Database When a Node Is Down

- you can successfully continue
 - *When nodes come back up, some of the data is stale*
- How we can detect stale on read?

Catching up on stale data from missed writes

■ Read Repair

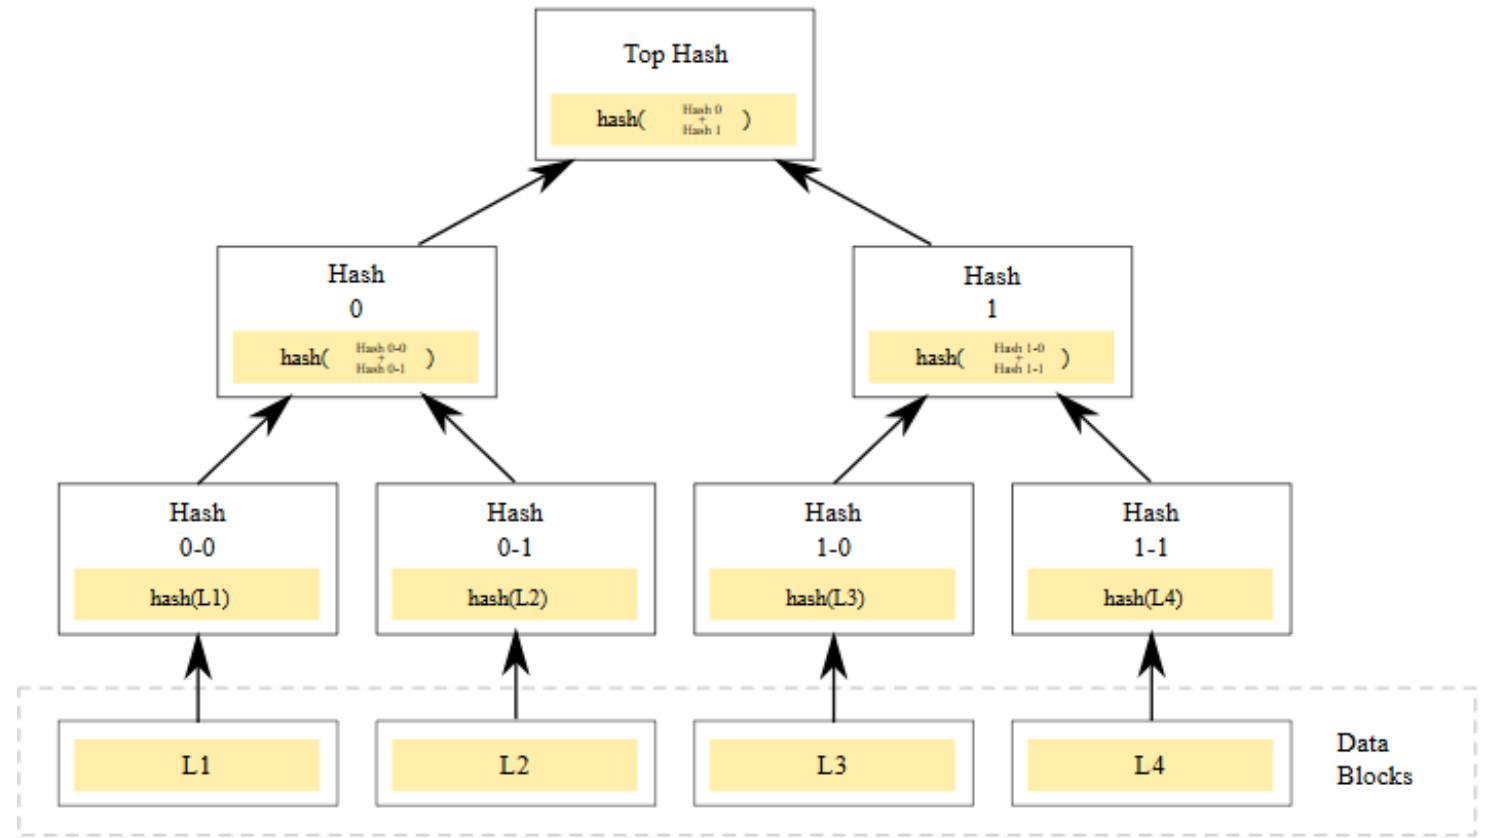
- *when a client detects a stale read, it writes the newer value back*

■ Anti-entropy process

- *background process that searches for differences between replicas and corrects them*
- *Merkle tree can be used to find differences*

Catching up on stale data from missed writes

- Anti-entropy process
 - *Merkle tree:*



Quorums for readings and writings

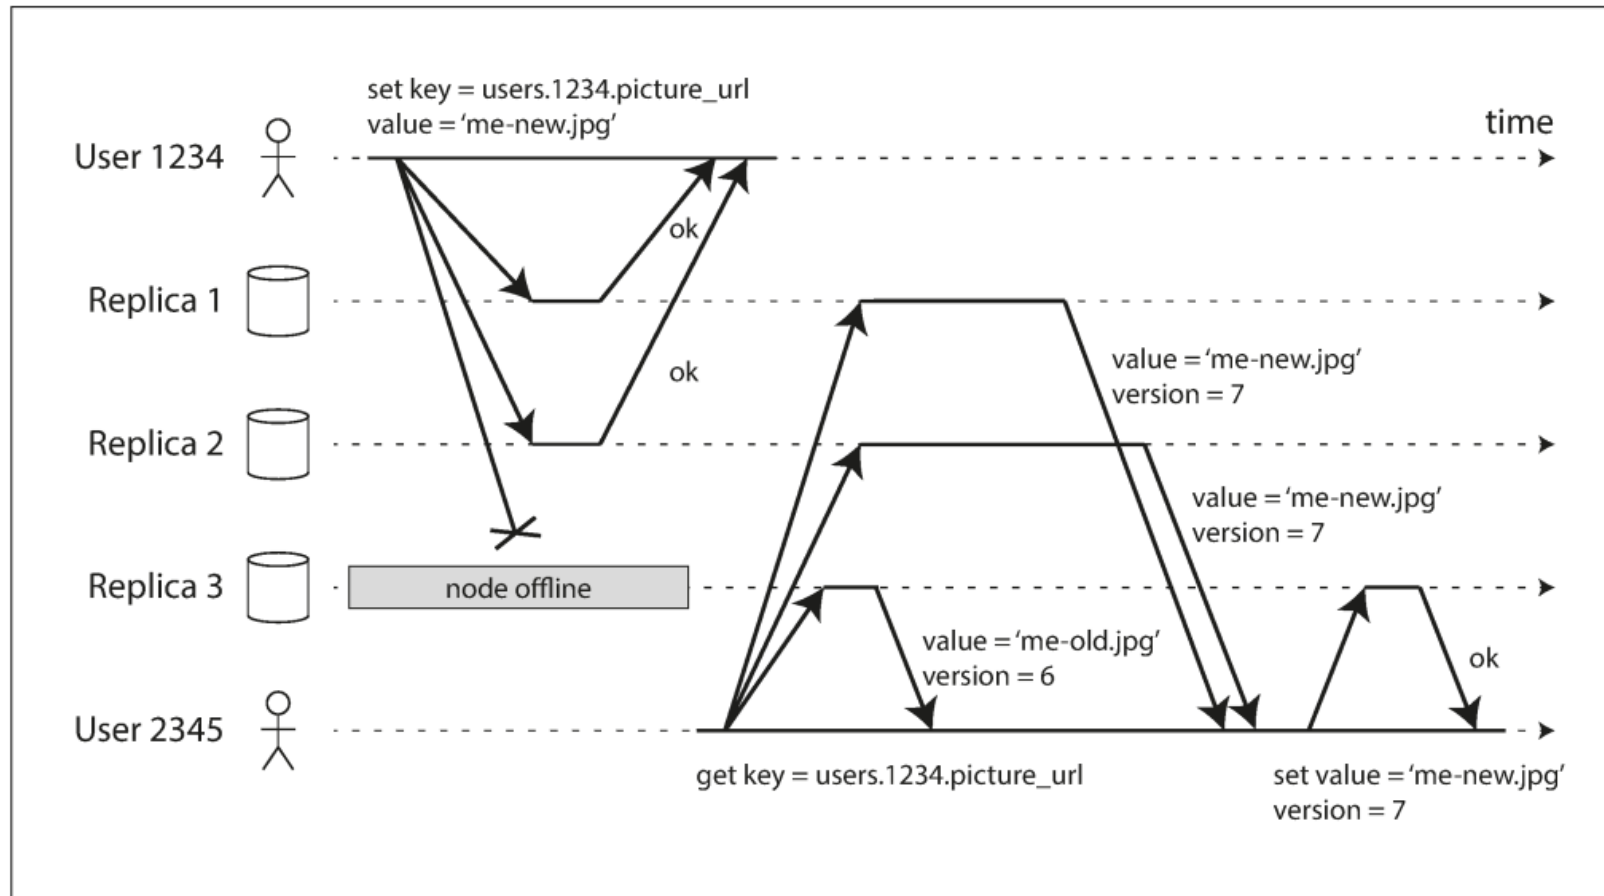


Figure 5-10. A quorum write, quorum read, and read repair after a node outage.

Quorums for readings and writings

- N : number of nodes
- n : number of replicas that particular data is stored
- w : minimum number of OKs that write considered as successful
- r : minimum number of OKs that read considered as successful

$$\mathbf{\blacksquare \quad w + r > n}$$

Quorums for readings and writings

- n, w, r ?
- Usually they are configurable.
- A common choice:
 - *n an odd number*
 - *$w, r = (n+1)/2$*

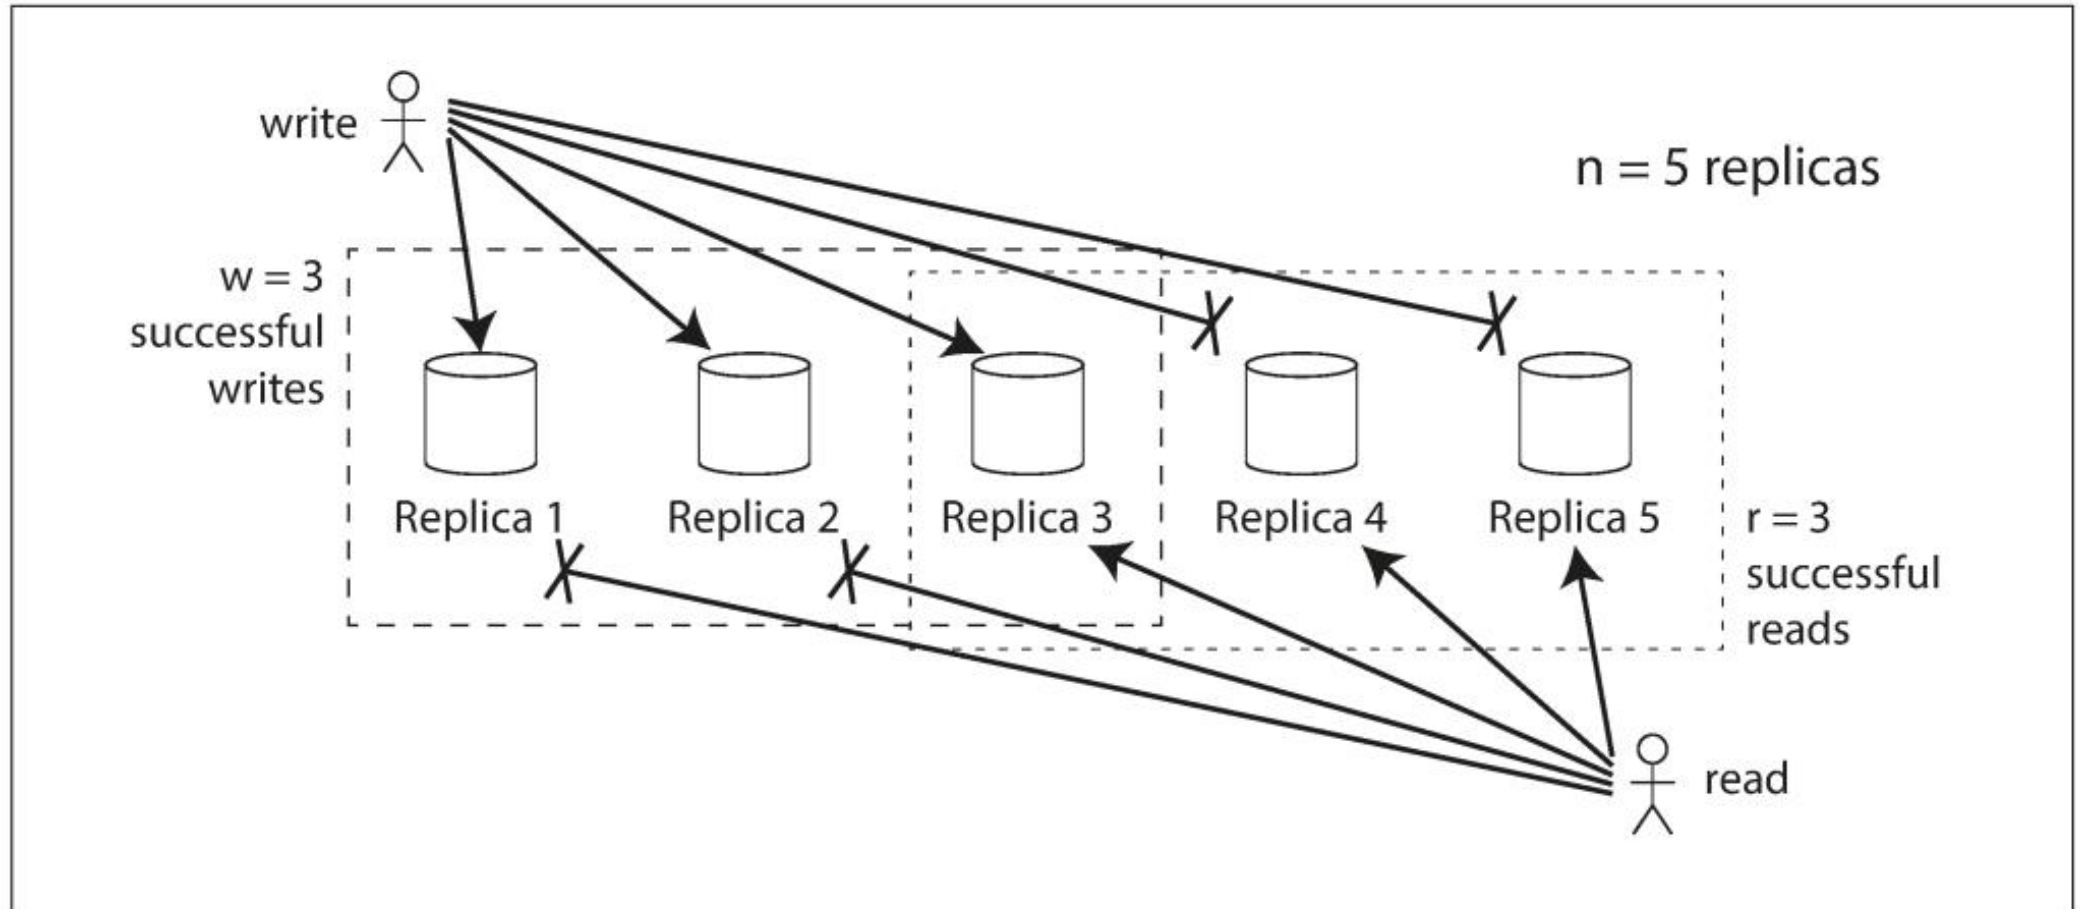


Figure 5-11. If $w + r > n$, at least one of the r replicas you read from must have seen the most recent successful write.

Limitations of Quorum Consistency

- Choosing $w + r \leq n$ risks stale reads.
- Sloppy quorums (next section)
- Concurrent writes
- If a write operation fails to get a quorum, a subset of nodes won't roll back, so will incorrectly have the newer value
- If a failed node A is restored from node B, anything stale on B is now stale on A, and both A and B can contribute toward a read quorum

Sloppy Quorums and Hinted Handoff

- In a large cluster, a network interruption could easily cause there to be not enough nodes for a quorum
- Is it better to simply return errors, or to try a workaround?
- Suppose $N > n$ and we don't have w ready nodes from n nodes.
 - ***Sloppy quorum***: use other $N-n$ nodes for write.
 - When connectivity is restored, temporary nodes use ***hinted handoff*** to write the data back to nodes where it belongs.

Multi-Datacenter Operation

- Given that communication between datacenters is expected to be slow
- Some configured to send cross-datacenter writes asynchronously
- Riak limits initial replication to be within a datacenter, and uses a multi-leader strategy between datacenters.

Concurrent Writes

- Can happen on multi-leader and leaderless replication
- Also during read repair and hinted handoff
- Last Write Wins (LWW)

Concurrent Writes

- What the meaning of concurrent write?
- If A was aware of write B, we say B happened before A.
- If B was aware of A, then A happened before B.
- In all other situations where neither A nor B were aware of the other, we define A and B as concurrent writes.

Version Vector – Single node Algorithm

- Server maintains a version with each key
- Clients must do a read, which includes version(s) and value(s), before they write
- Clients must merge multiple values read before doing a write
- Server increments max version each write
- Server receiving a write can discard data from that version or older, but keeps newer data

Version Vector – Single node Algorithm

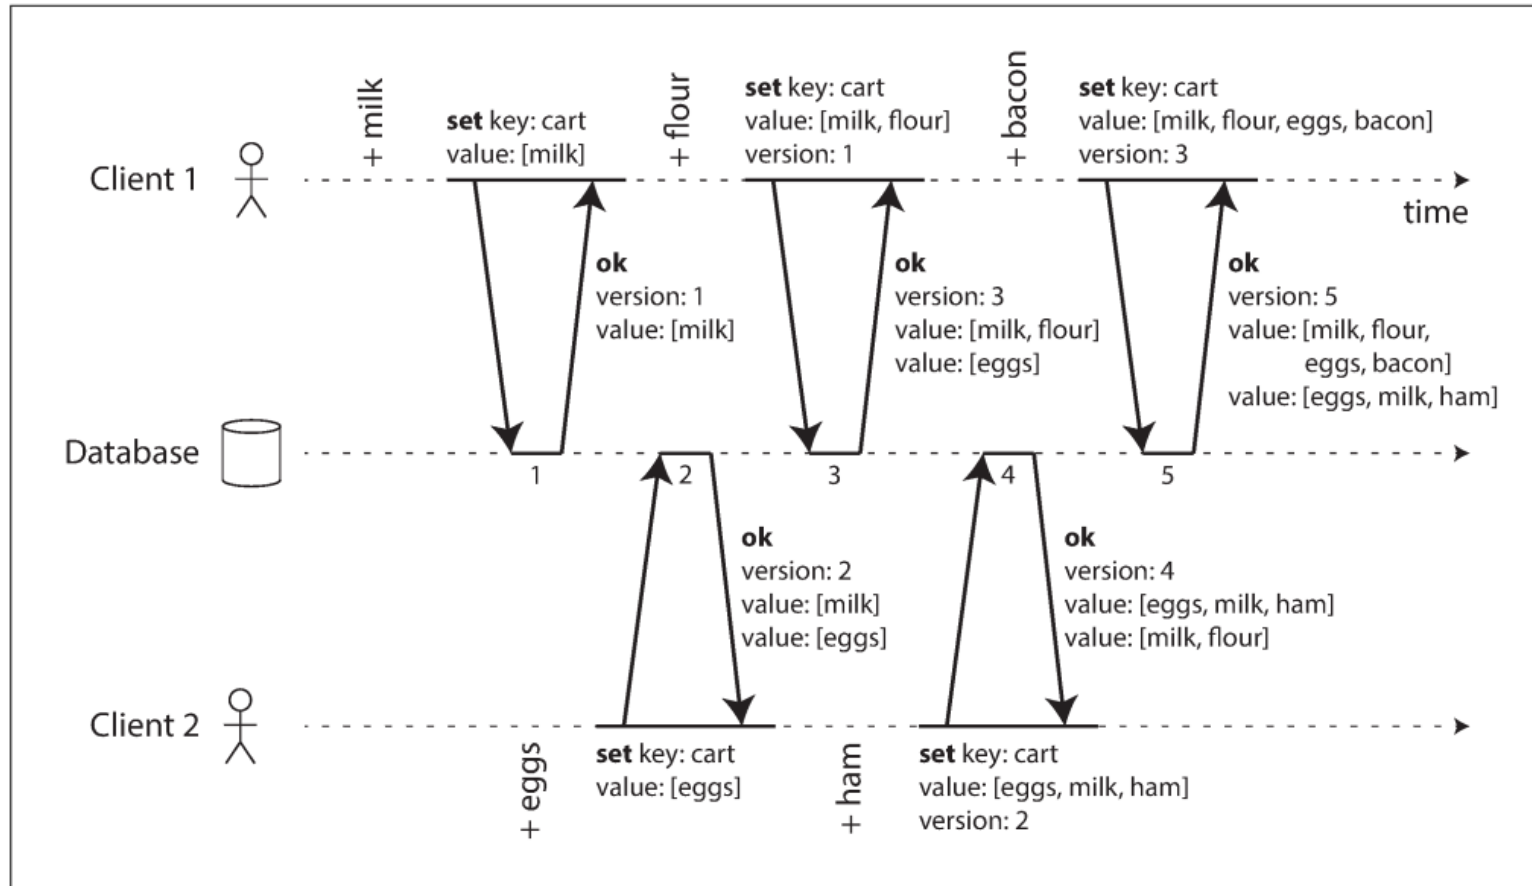


Figure 5-13. Capturing causal dependencies between two clients concurrently editing a shopping cart.

Version Vector – Single node Algorithm

- Siblings versions
- Merging concurrently written values?

Version Vector – Multiple replica

- Imagine you have N replicas: Replica A, Replica B, Replica C, ...
- Each replica keeps track of a local integer counter.
- For example:
 - *Replica A has version counter v_A .*
 - *Replica B has version counter v_B .*
 - *Replica C has version counter v_C .*

Version Vector – Multiple replica

- The version vector is a collection of each replica's version counters
- instead of a single integer, you have a small list/tuple of integer counters—one counter for each replica.
- Example: if you have three replicas, your version vector might look like [vA=5, vB=3, vC=7].

پارٹیشننگ

Partitioning

- Why?
- Partition / shard / region / vnode

Partitioning on K/V store

■ Partitioning by Key range

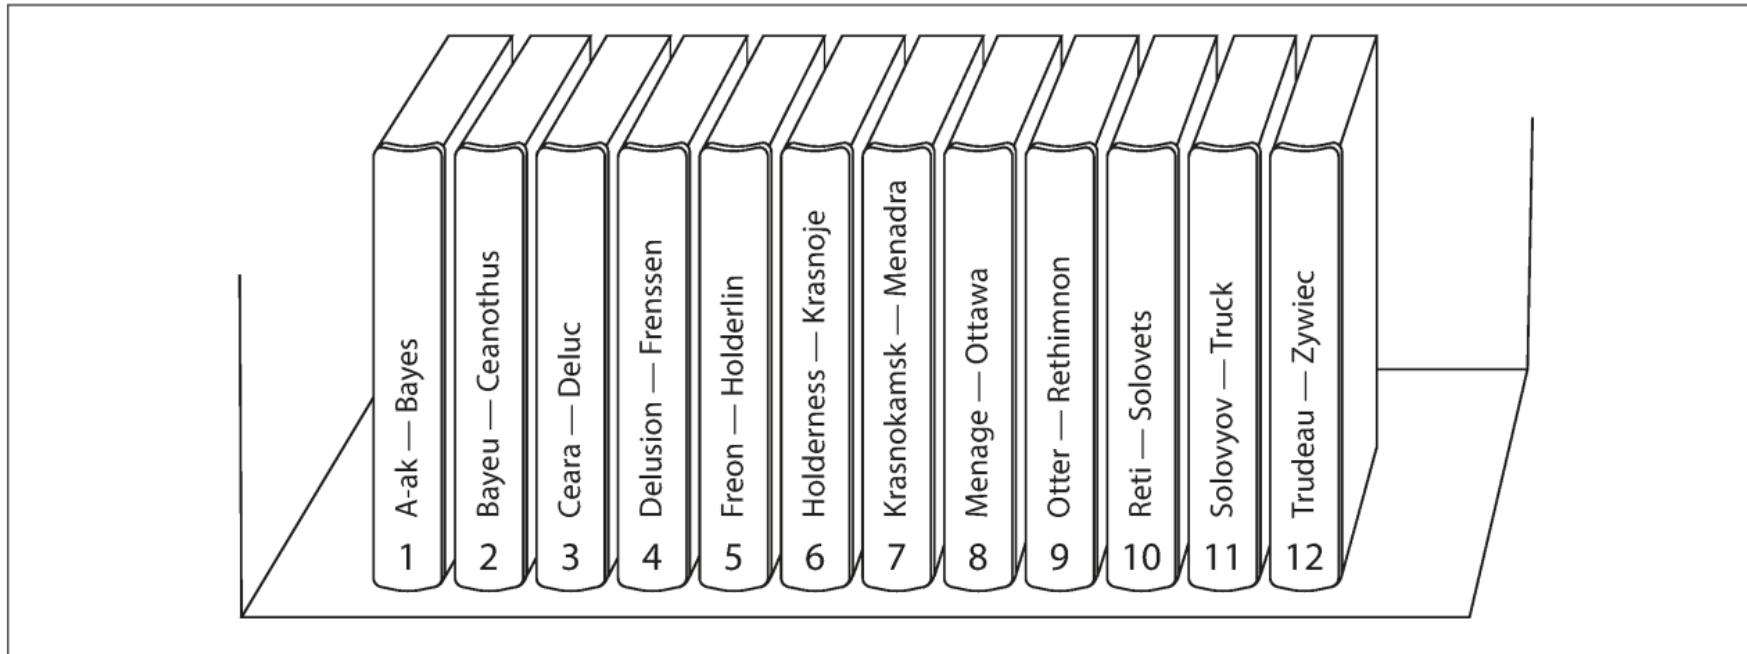


Figure 6-2. A print encyclopedia is partitioned by key range.

Partitioning on K/V store

- Partitioning by Key range
 - *Good for range queries*
 - *Example: an application that store sensors*
 - Key: timestamp
 - Value: measurement
 - *Hot Spots!*

Partitioning on K/V store

■ Partitioning by Hash key

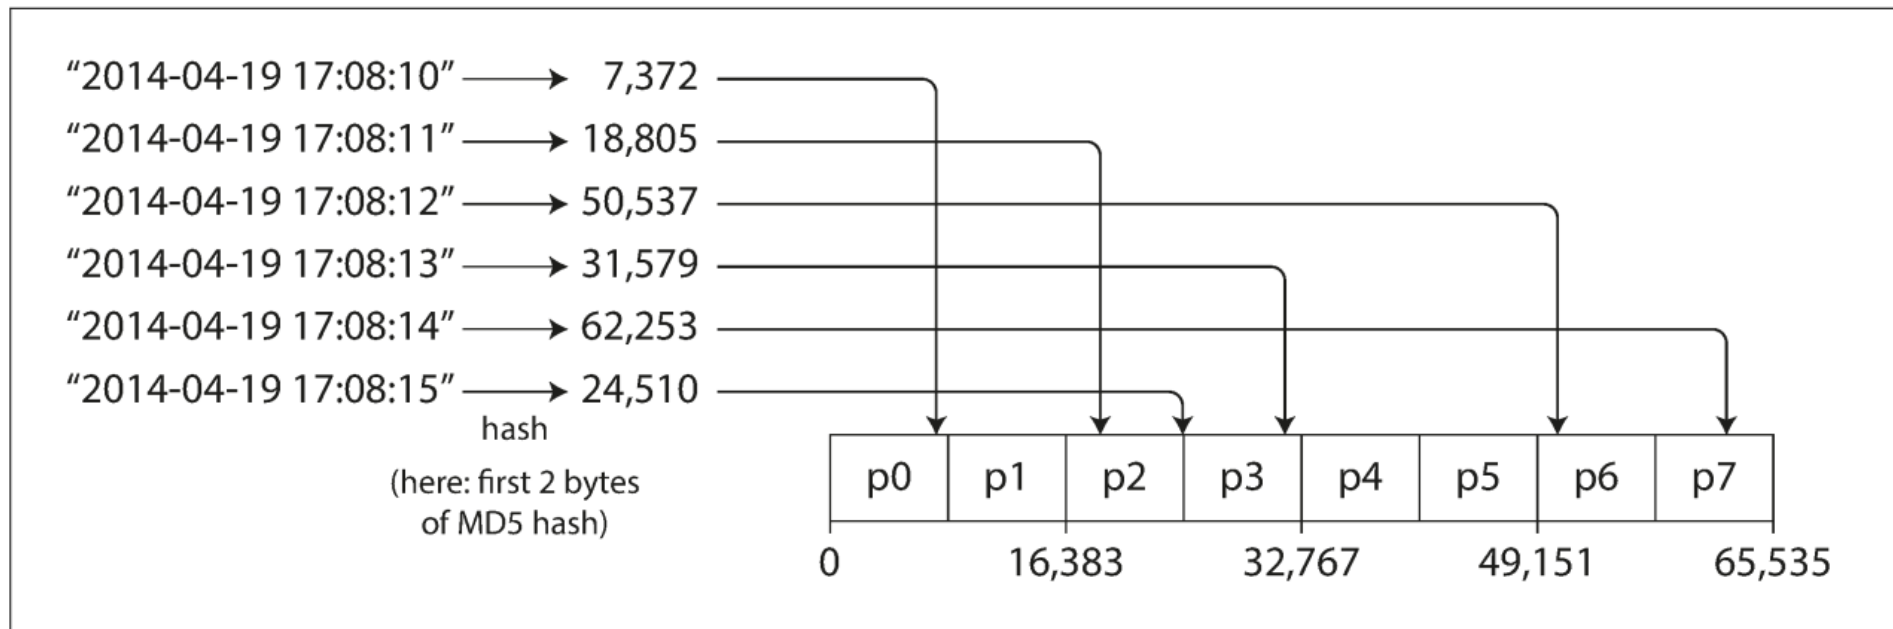


Figure 6-3. Partitioning by hash of key.

Partitioning on K/V store

- Partitioning by Hash key
 - *Hash to node?*
 - *What happen if we add new node?*