

بسم الله الرحمن الرحيم

تکنولوژی کامپیوتر

جلسه‌ی هفدهم
مصائب سیستم‌های توزیع شده

این جلسه

سیستم‌های توزیع‌شده، تفاوت و مشکلاتش نسبت به یک کامپیوتر، مدل ریاضی مسئله، خطی‌سازی

FAULT VS PARTIAL FAULT

Single Computer

- Deterministic
- Fault
 - *Blue Screen*
 - *Usually a computer fully crash*
 - *In single-system applications with proper code an application either works or it doesn't*

Distributed systems

- Partial Fault

- *We know individual systems will fail or groups of systems will fail but we need the overall system to continue working*
- *it may sometimes work and sometimes unpredictably fail.*

- You may not even know whether something succeeded or not, as the time it takes for a message to travel across a network is also nondeterministic!

Building a Reliable System from Unreliable Components??

- TCP from unreliable IP
 - *But cannot remove delays in the network*
- Error-correcting codes
 - *But only can deal with small number of single-bit errors*

UNRELIABLE NETWORKS

shared-nothing systems

Asynchronous packet networks

- The internet and most internal networks in datacenters (often Ethernet) are asynchronous packet networks
- one node can send a message (a packet) to another node
 - *no guarantees as to when it will arrive*
 - *Or no guarantees whether it will arrive at all.*

Asynchronous packet networks

- If you send a request and expect a response, many things could go wrong
 - *Your request may have been lost*
 - *Your request may be waiting in a queue and will be delivered later*
 - *The remote node may have failed*
 - *The remote node may have temporarily stopped responding*
 - *The remote node may have processed your request, but the response has been lost on the network*
 - *The remote node may have processed your request, but the response has been delayed and will be delivered later*

Asynchronous packet networks




Figure 8-1. If you send a request and don't get a response, it's not possible to distinguish whether (a) the request was lost, (b) the remote node is down, or (c) the response was lost.

Asynchronous packet networks

- the only information you have is that you haven't received a response yet!
- If you send a request to another node and don't receive a response, it is impossible to tell why
- The usual way of handling this issue is a **timeout**

Network Faults in Practice

- Network problems can be surprisingly common! 😞
- Network partitions
 - *When one part of the network is cut off from the rest due to a network fault*
 - *called a network partition or netsplit*

Network Faults in Practice

- Handling network fault?
 - *Tolerating them?*
 - *If your network is normally fairly reliable, a valid approach may be to simply show an error message to users while your network is experiencing problems*

Detecting Faults

- Automatically detect faulty nodes
- Why?
 - *A load balancer needs to stop sending requests to a node that is dead (i.e., take it out of rotation)*
 - *In a distributed database with single-leader replication, if the leader fails, one of the followers needs to be promoted to be the new leader*

Detecting Fault?

- If the system is still up but application has run into a problem then possible to get a return from the system saying it is present but non-operational
- No way of knowing what requests it may have received and left unanswered or what data was sent to it and lost
- Possible for OS to directly announce it has failed so other systems don't have to wait for timeout
- Possible to use remote management tools to check on systems if you are closely networked

Timeouts and Unbounded Delays

- How long should the timeout be?
 - *There is unfortunately no simple answer. ☹️*
 - *Long Timeout ...*
 - *Short timeout ...*

When delay happen?

- Often delays are caused by queuing of some sort
 - *Multiple systems trying to send info to the same system, all have to be handled in whatever order they are received, but some will have to wait*
 - *If all CPU cores of a system are busy it will have to wait for cycles to address the information*
 - *In virtual systems they can be temporarily frozen while other VMs are doing things*
- Possible to have a noisy neighbor that is causing delays for you by using up networking capacity

- Have to determine appropriate timeout lengths experimentally

Synchronous Versus Asynchronous Networks

- Distributed systems would be a lot simpler if we could rely on the network to deliver packets with some fixed maximum delay, and not to drop packets
- traditional fixed-line telephone network
 - *is extremely reliable*
 - *delayed audio frames and dropped calls are very rare*

Synchronous Versus Asynchronous Networks

- When you make a call over the telephone network
 - *it establishes a circuit*
 - *a fixed, guaranteed amount of bandwidth is allocated for the call, along the entire route between the two callers.*

Synchronous Versus Asynchronous Networks

- Can we not simply make network delays predictable?
- The answer is that they are optimized for bursty traffic.
- Latency and Resource Utilization

UNRELIABLE CLOCKS

Unreliable Clocks

- 1. Has this request timed out yet?
- 2. What's the 99th percentile response time of this service?
- 3. How many queries per second did this service handle on average in the last five minutes?
- 4. How long did the user spend on our site?
- 5. When was this article published?
- 6. At what date and time should the reminder email be sent?
- 7. When does this cache entry expire?
- 8. What is the timestamp on this error message in the log file?

Clocks in your computer

■ Time-of-day clocks

- *Time-of-day clocks are usually synchronized with NTP*
- *System.currentTimeMillis() in java*

■ Monotonic clocks

- *System.nanoTime() in java*
- *NTP may adjust the frequency at which the monotonic clock moves forward*

Clock Synchronization and Accuracy

- NTP synchronization can only be as good as the network delay
- ...

Relying on Synchronized Clocks

- they seem simple and easy to use But ...
- a day may not have exactly 86,400 seconds!
- time-of-day clocks may move backward in time
- time on one node may be quite different from the time on another node
- -> robust software needs to be prepared to deal with incorrect clocks.

Timestamps for ordering events

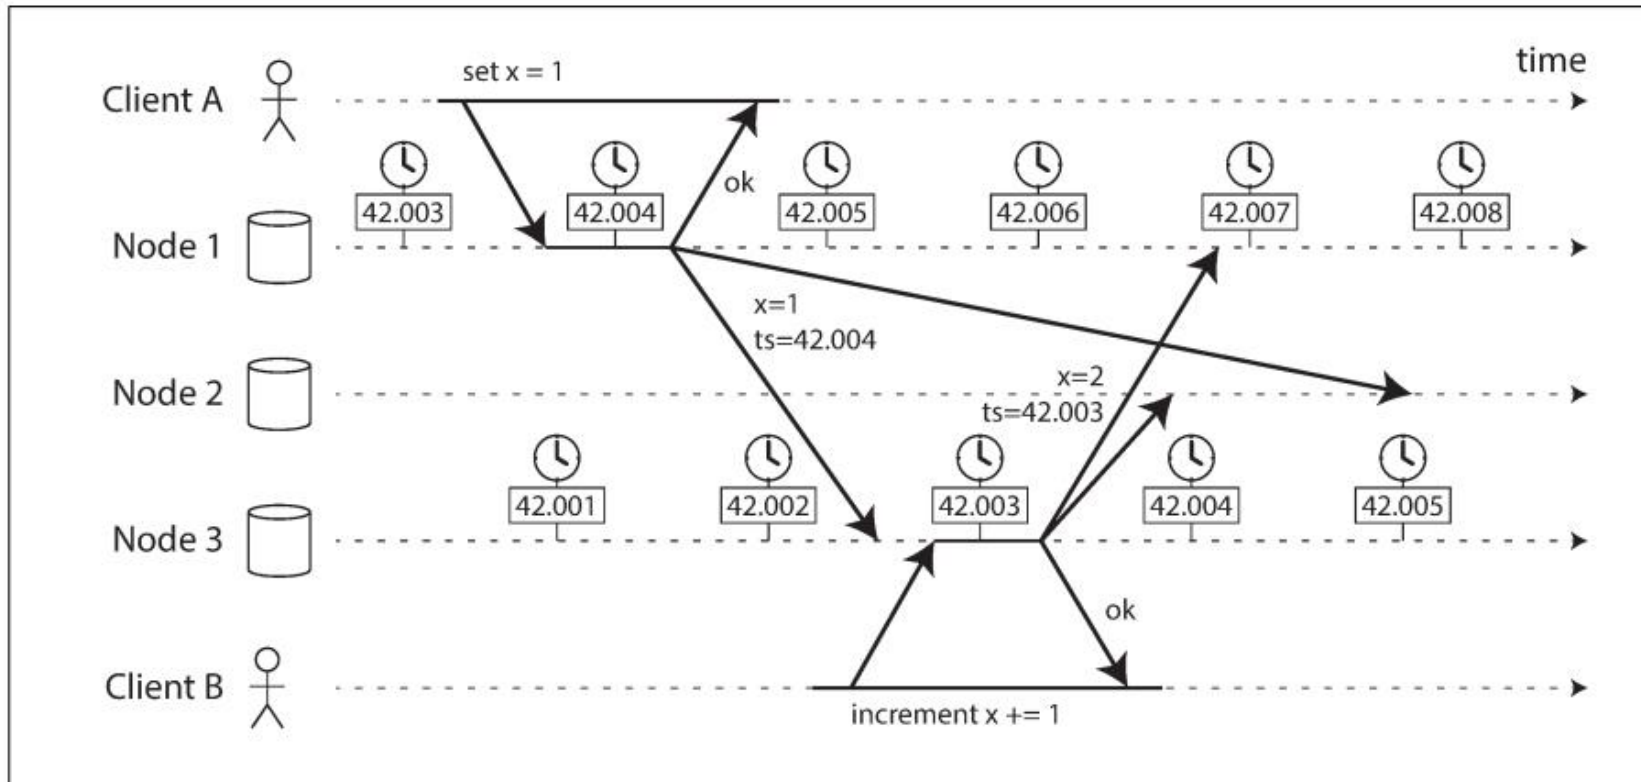


Figure 8-3. The write by client B is causally later than the write by client A, but B's write has an earlier timestamp.

Clock readings have a confidence interval

- You may be able to read a machine's time-of-day clock with microsecond or even nanosecond resolution
 - *Is it mean that the value is actually accurate to such precision?*
 - No 😞

Clock readings have a confidence interval

- Unfortunately, most systems don't expose this uncertainty
 - *for example, when you call `clock_gettime()`, the return value doesn't tell you the expected error of the timestamp,*
 - *you don't know if its confidence interval is five milliseconds or five years.*

Clock readings have a confidence interval

- The uncertainty bound can be calculated based on your time source...
- If
 - *GPS receiver*
 - *atomic (caesium) clock*
- directly attached to your computer

Clock readings have a confidence interval

- Google's Spanner TrueTime API
 - *which explicitly reports the confidence interval on the local clock*
 - *you get back two values: [earliest, latest]*
 - the earliest possible
 - And the latest possible timestamp.

Synchronized clocks for global snapshots

- Snapshot isolation...
 - *Monotonically increasing transaction ID*
 - *To implement MVVC*
- Can we use timestamp for transaction id in distributed system?
- Spanner implementation of snapshot isolation

Process Pauses

```
while (true) {  
    request = getIncomingRequest();  
  
    // Ensure that the lease always has at least 10 seconds remaining  
    if (lease.expiryTimeMillis - System.currentTimeMillis() < 10000) {  
        lease = lease.renew();  
    }  
  
    if (lease.isValid()) {  
        process(request);  
    }  
}
```

Process Pauses

- Garbage collector
- Suspend in virtual machines
- Suspend in laptops
- Context switching
- Stop process using stop signal

Preventing garbage collector pauses?

- Short lived processes
- Restart using rolling upgrade

KNOWLEDGE, TRUTH, AND LIES

Knowledge

- A node in the network cannot know anything for sure
- it can only make guesses based on the messages it receives (or doesn't receive) via the network
- If a remote node doesn't respond...
 - *Is node has a problem? Which node?*
 - *Is network has a problem?*

The truth?

■ Example

- *Node receive message from others*
- *But can't send message to others*
- *Other nodes think that it is dead*
- *the semi-disconnected node is dragged to the graveyard, kicking and screaming "I'm not dead!"*

- Example

- *Node with long stop-the-world GC*

The Truth Is Defined by the Majority

- the quorum is an absolute majority of more than half the nodes
- it is safe, because there can only be only one majority in the system
- ... consensus algorithms ...

The leader and the lock

- Lock for distributed systems
 - *Only one node is allowed to be the leader for a database partition, to avoid split brain*
 - *Only one transaction or client is allowed to hold the lock for a particular resource or object, to prevent concurrently writing to it and corrupting it.*
 - *Only one user is allowed to register a particular username, because a username must uniquely identify a user.*

The leader and the lock

- even if a node believes that it is “the chosen one”
- Is it mean that quorum of nodes agrees?
 - *No!*
 - *Example...*

The leader and the lock

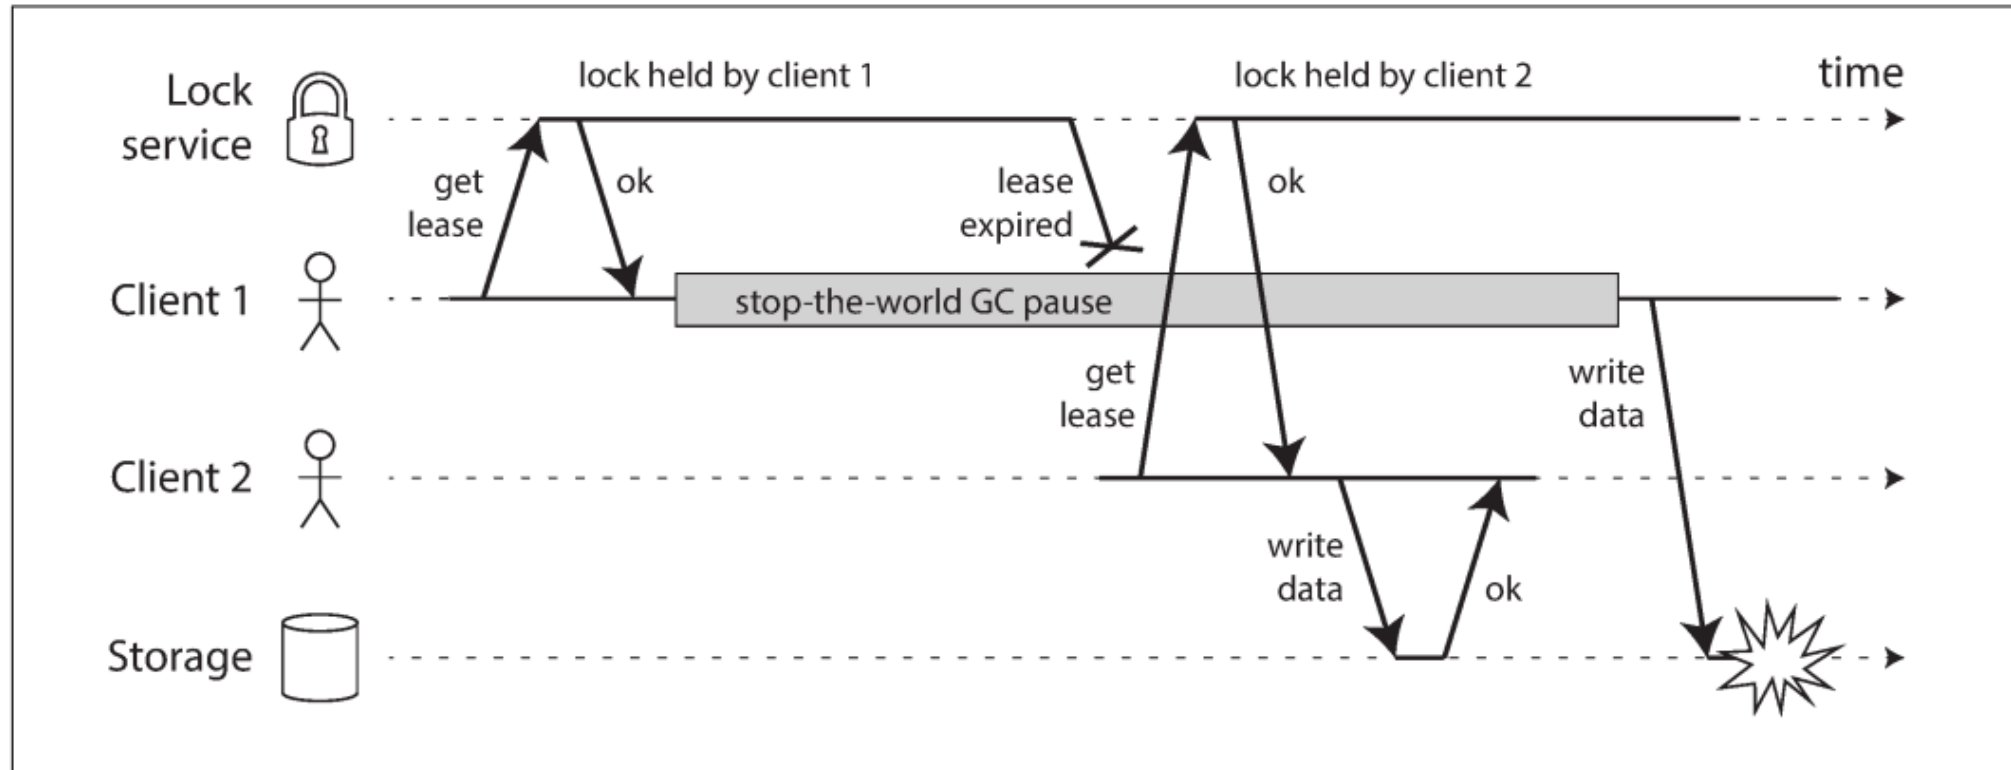


Figure 8-4. Incorrect implementation of a distributed lock: client 1 believes that it still has a valid lease, even though it has expired, and thus corrupts a file in storage.

The leader and the lock - Fencing tokens

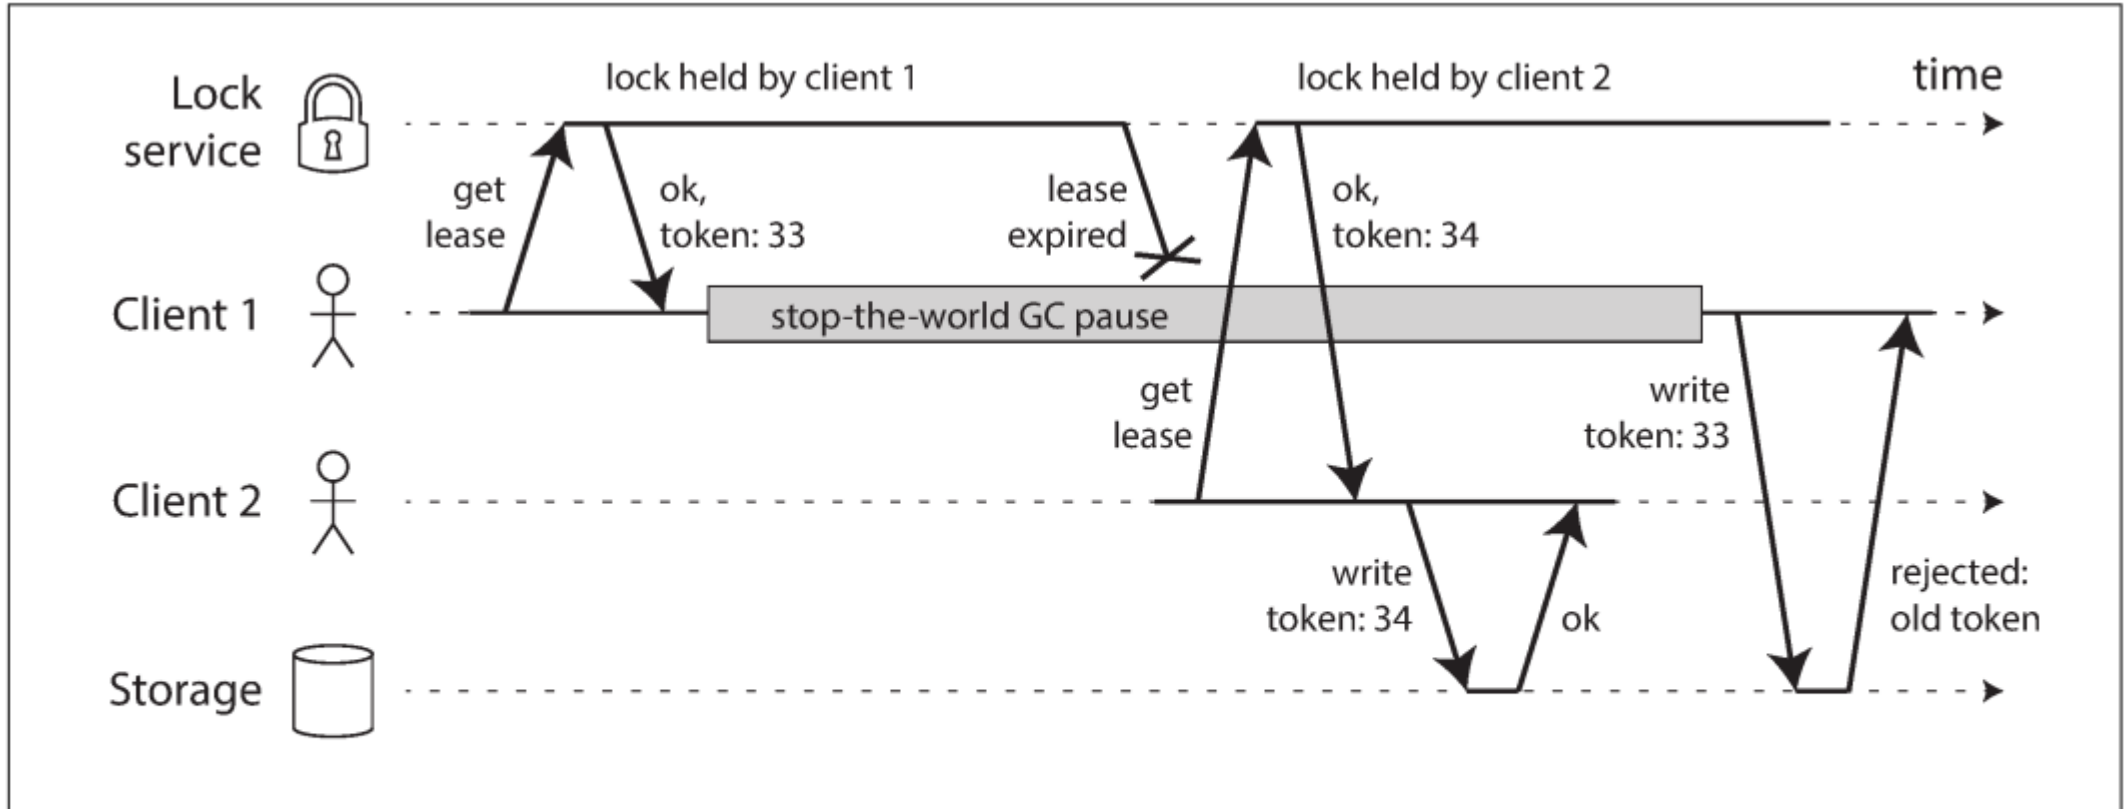


Figure 8-5. Making access to storage safe by allowing writes only in the order of increasing fencing tokens.

Byzantine Faults

SYSTEM MODEL AND REALITY

Timing assumptions

- Synchronous model

- *bounded network delay, bounded process pauses, and bounded clock error*

- Partially synchronous model

- *system behaves like a synchronous system most of the time, but it sometimes exceeds the bounds for network delay, process pauses, and clock drift*

- Asynchronous model

- *an algorithm is not allowed to make any timing assumption*

Node failures

- Crash-stop faults
- Crash-recovery faults
- Byzantine (arbitrary) faults

Correctness of an algorithm?

- Example on fencing token
- Uniqueness
 - *No two requests for a fencing token return the same value.*
- Monotonic sequence
 - *If request x returned token tx , and request y returned token ty , and x completed before y began, then $tx < ty$*
- Availability
 - *A node that requests a fencing token and does not crash eventually receives a response*

Correctness of an algorithm?

- Safety and liveness
- Safety
 - *nothing bad happens*
- Liveness
 - *something good eventually happens*

LINEARIZABILITY

Consistency Guarantees

- Eventual consistency

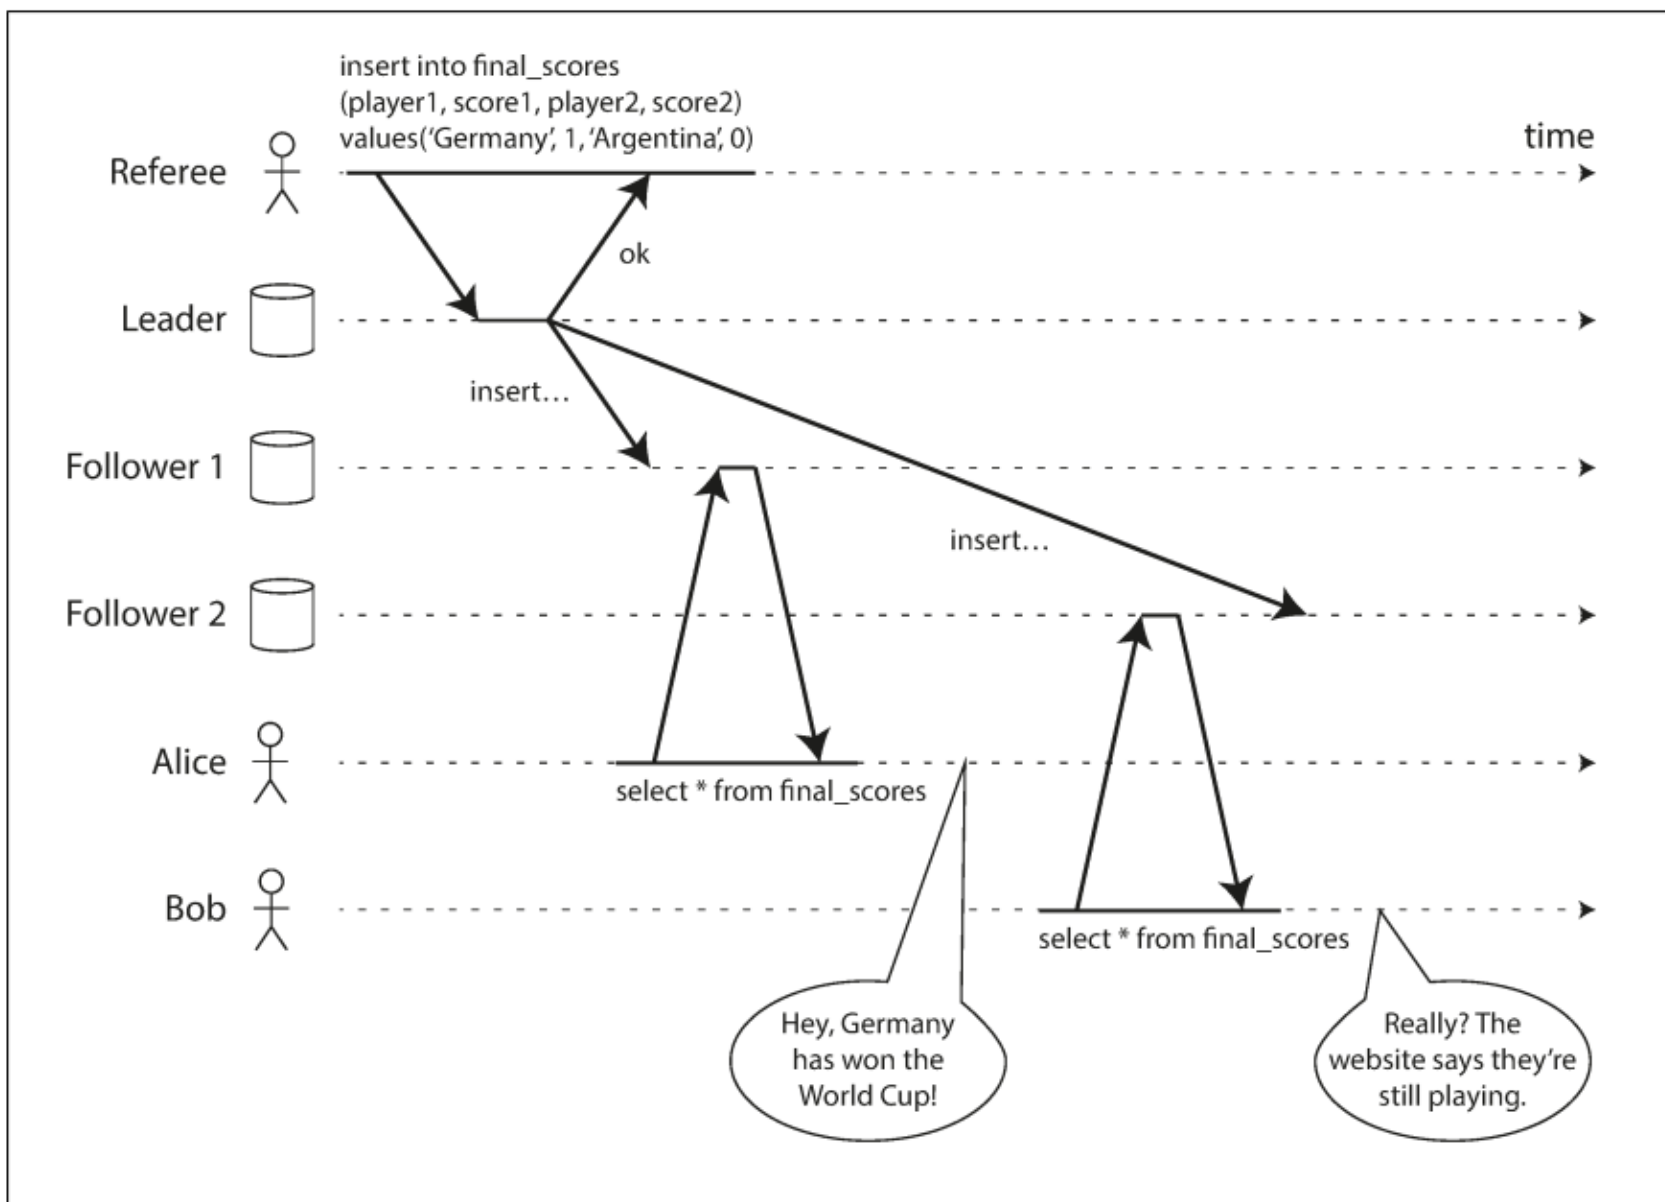


Figure 9-1. This system is not linearizable, causing football fans to be confused.

Linearizability

- atomic consistency
- strong consistency
- immediate consistency
- external consistency

Linearizability

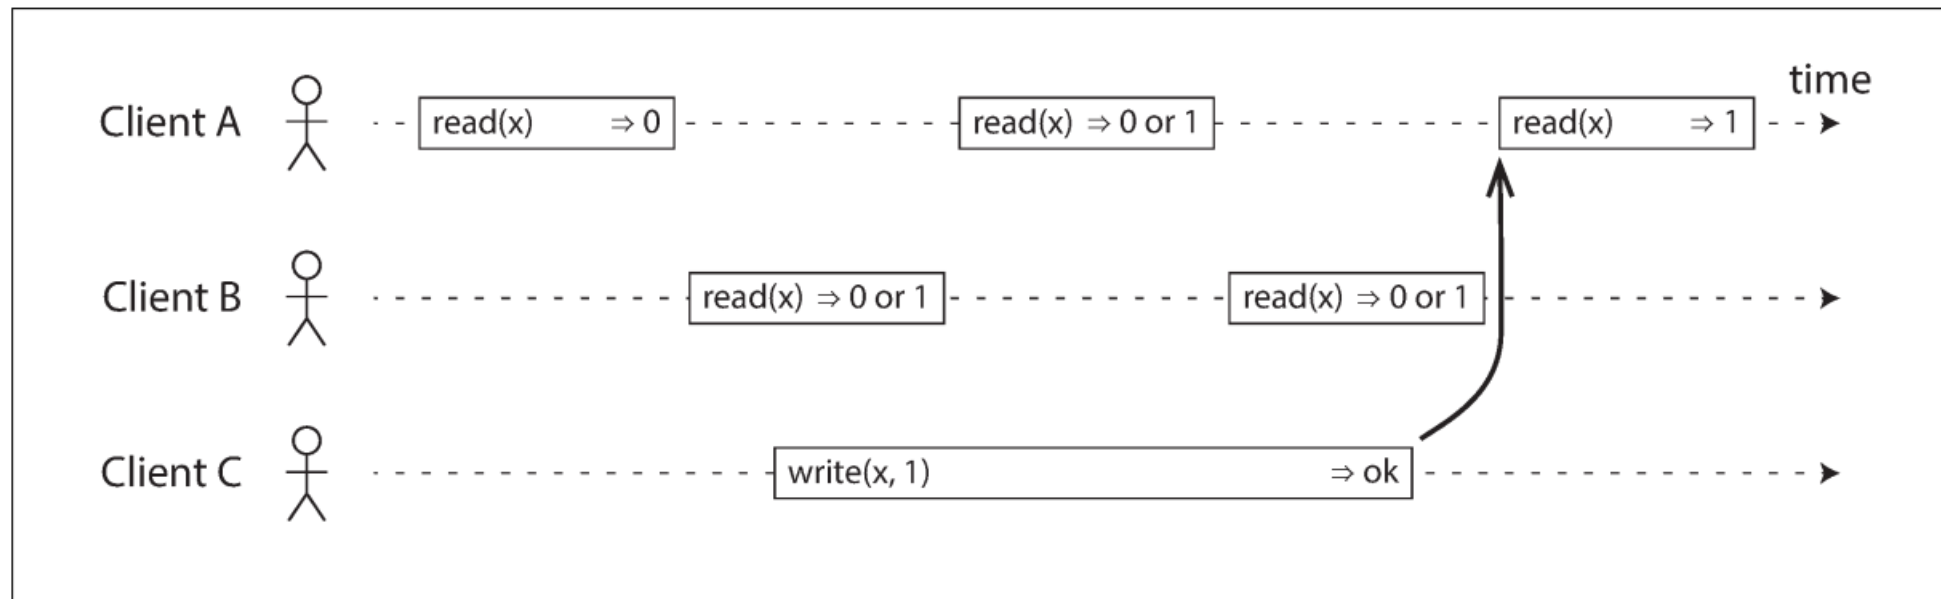


Figure 9-2. If a read request is concurrent with a write request, it may return either the old or the new value.

Linearizability

To make the system linearizable, we need to add another constraint, illustrated in **Figure 9-3**.

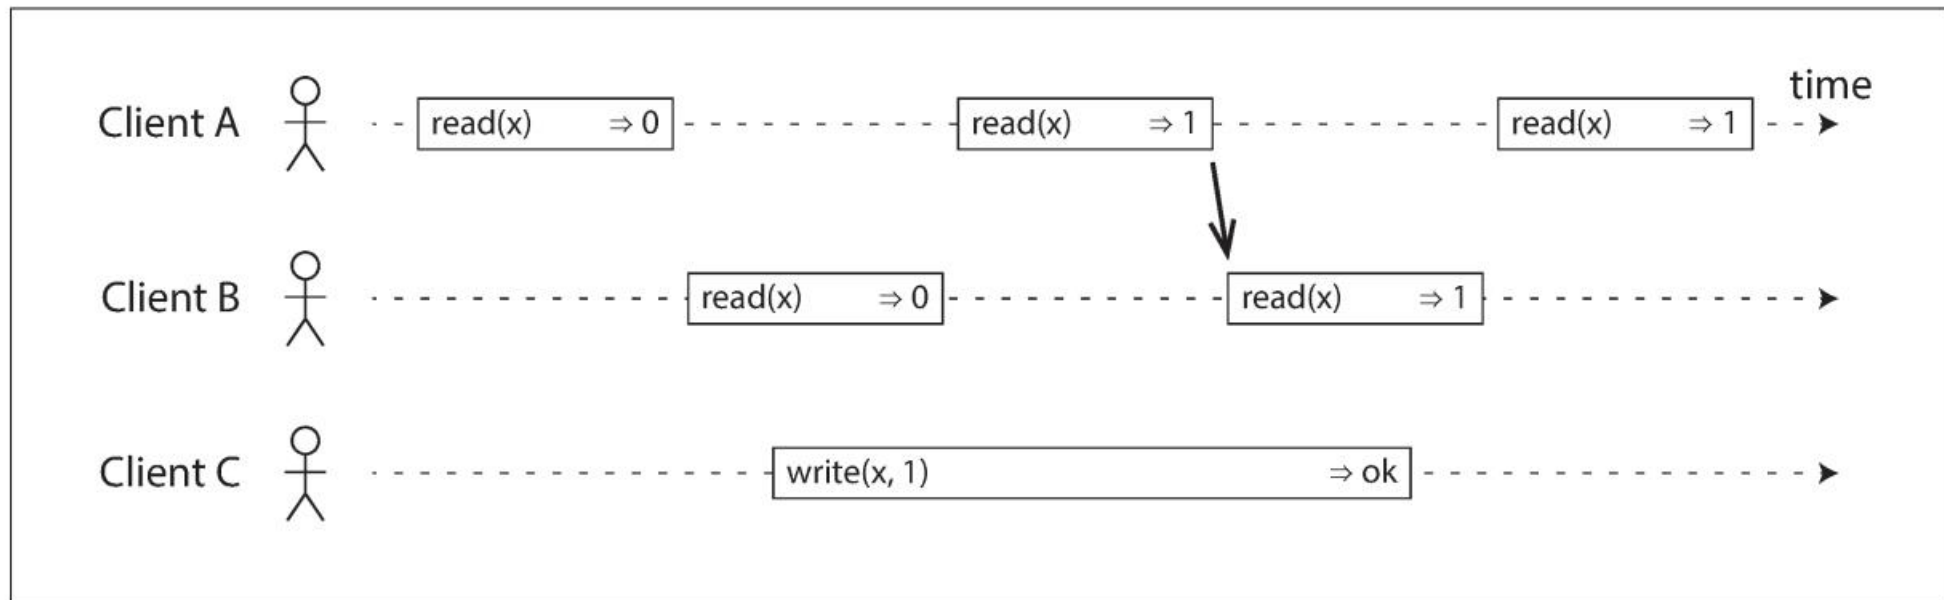


Figure 9-3. After any one read has returned the new value, all following reads (on the same or other clients) must also return the new value.

Linearizability

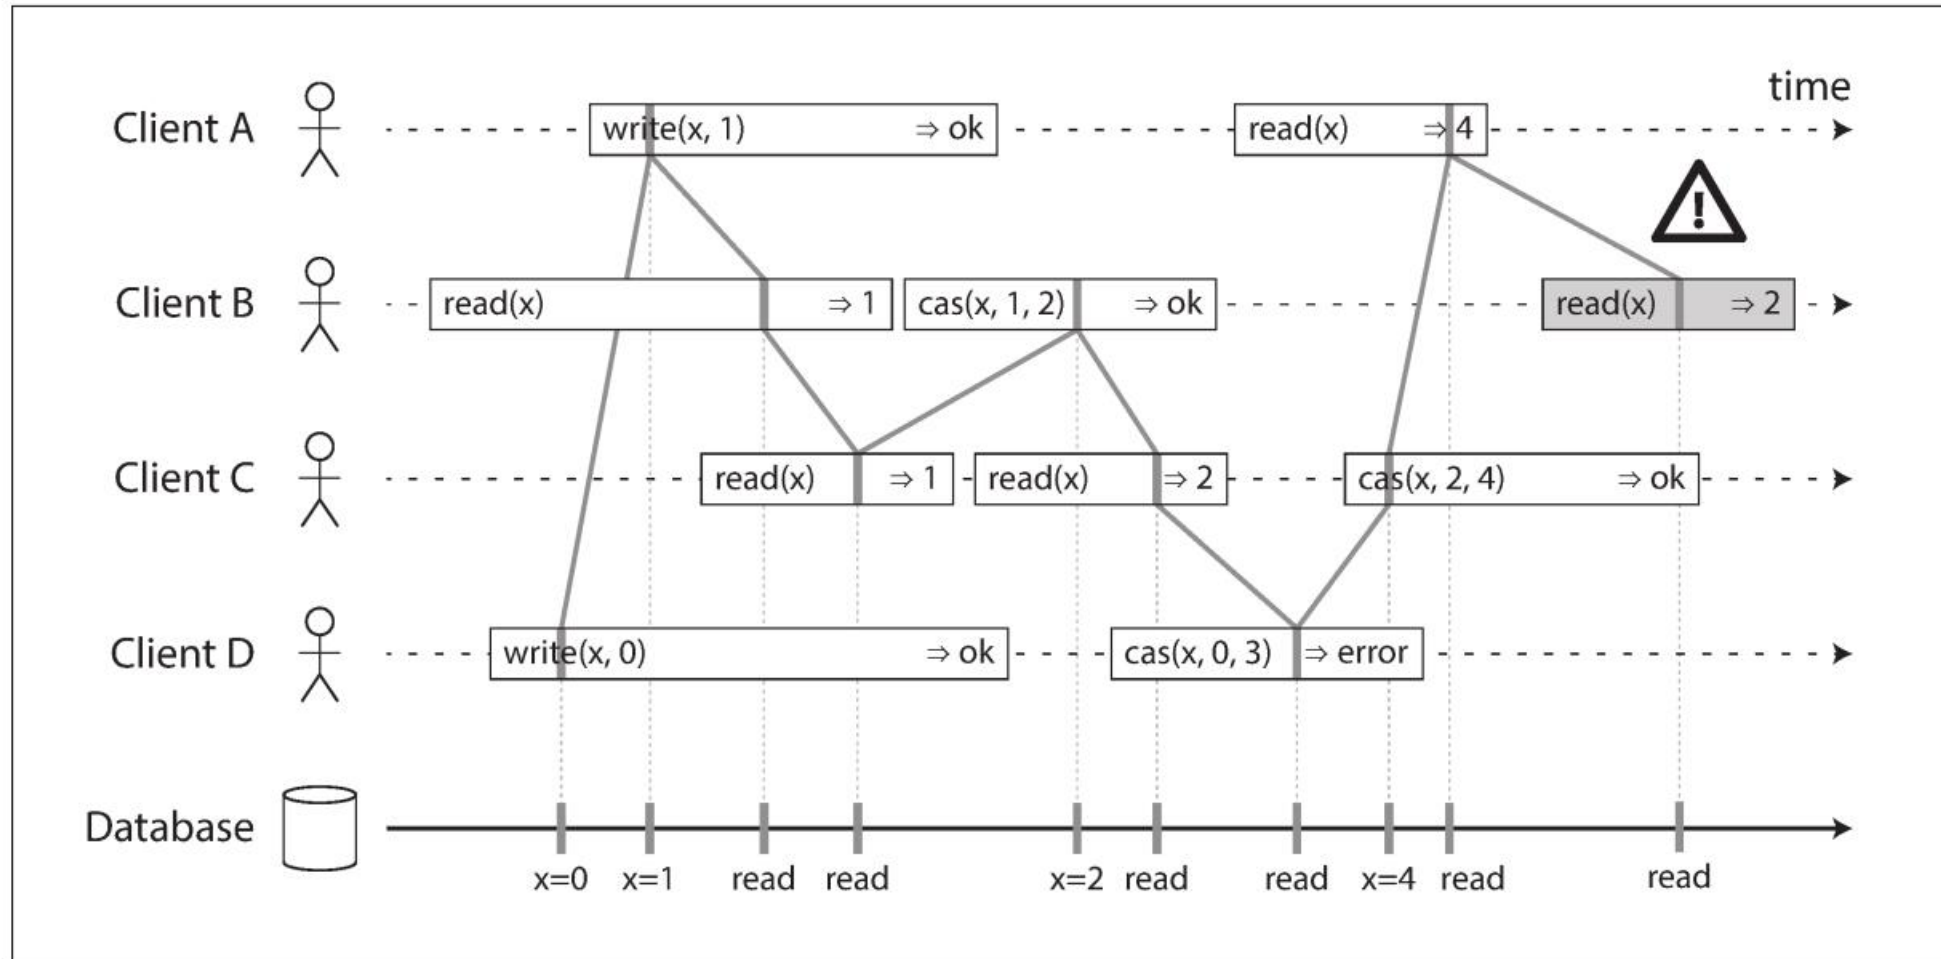


Figure 9-4. Visualizing the points in time at which the reads and writes appear to have taken effect. The final read by B is not linearizable.

Linearizability Versus Serializability

جلسه‌ی بعد