

بسم الله الرحمن الرحيم

# سیستم عامل

جلسه اول - معرفی درس

# معرفی خودم

- Alireza Tofighi
- Ph.D Student at Sharif, CTO at Balad.ir
- Email: [alirtofighim+os1403@gmail.com](mailto:alirtofighim+os1403@gmail.com)
- <https://alireza.dev/teaching/os1403/>

## ■ مشاوره و رفع اشکال؟

- به من ایمیل بزنید، ممکن است در اتاق ۱۱۴ دانشکده یا  
لابی پیدایم کنید!

# Class Goals

- Understand the basic concepts of operating systems
  - *designing & building operating systems, not using them!*
- Gain some practical experience so that it is not just words!

# قوانین درس

# تمرین‌های برنامه‌نویسی

- ۵ سری تمرین عملی
- با زبان KPL، باید مستندات زبان را هم بخوانید که یاد بگیرید و استفاده کنید. (زبان سختی نیست: دی)
- شروع تمرین اول از امروز! و مهلت تا پایان ۲۰ مهر
- قوانین تاخیر؟
  - در مجموع ۱۰ روز تاخیر مجاز، ولی هر تمرین حداکثر ۳ روز تاخیر مجاز
  - محاسبه‌ی تاخیر به شکل ساعتی
  - تاخیر بیش از زمان  $\leq$  نمره‌ای کسب نمی‌شود.

# Programming Projects

- You will read, understand and write operating system code!
- We will use the BLITZ system, written by Harry Porter
  - *CPU emulator, assembler, high-level language, operating system, and debugging environment*
  - *- Simple enough to understand in detail how everything works!*
  - *- Realistic enough to understand in detail how everything works!*
  - *- Runs on the most modern machines*

# آزمونک!


- تلاش برای دنبال کردن بخش تئوری درس
- ۵ سری
- به شکل حضوری – کتاب بسته و در پایان کلاس‌ها
- زمان آن از قبل اعلام می‌شود.

# اصل مطلب – بارم‌بندی

- ۵ سری تمرین عملی و برنامه‌نویسی - مجموعاً ۶ نمره
- ۵ سری آزمونک - مجموعاً ۳ نمره
- آزمون میان‌ترم - ۵ نمره
- آزمون پایان‌ترم - ۷ نمره



# جنگ اول به از صلح آخر – تقلب/مشابهت


- انتظار داریم تمام بخش‌های نمره‌دار درس تلاش مستقیم خودتان باشد.
- مثلاً در تمرین‌های عملی، این موارد مجاز نیست: (موارد مشابه هم مجاز نیست)
  - استفاده از کد هم‌کلاسی‌هایتان یا هر فرد دیگر
  - استفاده از ابزارهای تولید کد همچون GPT‌ها
  - استفاده از کدهای موجود در اینترنت
  - دیدن کد هم‌کلاسی یا شنیدن آن
- چرا سخت‌گیری؟
- وجود مشابهت  $\Rightarrow$  احتمال صفر گرفتن آن تمرین یا کل تمرینات درس یا افتادن درس! به تشخیص تیم درس 

# تیم درس

- علیرضا توفیقی محمدی - استاد درس  
*alirtofighim+os1403@gmail.com*
- الهه توحیدی - سردستیار درس  
*elahetohidi2017@gmail.com*
- امیرسالار صفایی قادری - دستیار درس
- سید ستار لطفی فاطمی - دستیار درس
- سوگل زمانیان - دستیار درس
- امیر طایفه - دستیار درس
- رابعه پرهیزکاری - دستیار درس
- وحید انصاری - دستیار درس
- عرفان موسویان - دستیار درس

# اطلاعات درس

<https://alireza.dev/teaching/os1403/>



# تقویم درس در سایت درس

این تقویم در طول ترم بروز خواهد شد.

جلسه	تاریخ	موضوع جلسه	آزمونک	تمرین
۱	۰۳/۰۷/۰۲	معرفی درس - شروع		تمرین ۱ آشنایی با ساختار پروژه
۲	۰۳/۰۷/۰۷	معماری کامپیوتر		
۳	۰۳/۰۷/۰۹	پردازه		
۴	۰۳/۰۷/۱۴	پردازه		
۵	۰۳/۰۷/۱۶	رلهسمان	آزمونک ۱ معماری کامپیوتر - پردازه	
۶	۰۳/۰۷/۲۱	همروندی		تمرین ۲ همروندی
۷	۰۳/۰۷/۲۳	همروندی		
۸	۰۳/۰۷/۲۸	همروندی		
۹	۰۳/۰۷/۳۰	بن‌بست	آزمونک ۲ همروندی	
۱۰	۰۳/۰۸/۰۵	بن‌بست		
۱۱	۰۳/۰۸/۰۷	بن‌بست		
۱۲	۰۳/۰۸/۱۲	زمان‌بندی		تمرین ۳ همروندی - ناخیه بحرانی
۱۳	۰۳/۰۸/۱۴	مدیریت حافظه		
۱۴	۰۳/۰۸/۱۶	مدیریت حافظه		

# کوئرای درس


رمز کلاس: os1403



# گروه تلگرامی

تیم درس رو می‌تونید در تلگرام  
هم بیابید.

لطفا ولی به من در PV تلگرام برای  
کارهای درس پیام ندید: D



# QUESTIONS?

# معرفی منبع درس





# PART ONE ■ OVERVIEW

## Chapter 1 Introduction

- 1.1 What Operating Systems Do 4
- 1.2 Computer-System Organization 7
- 1.3 Computer-System Architecture 15
- 1.4 Operating-System Operations 21
- 1.5 Resource Management 27
- 1.6 Security and Protection 33
- 1.7 Virtualization 34
- 1.8 Distributed Systems 35
- 1.9 Kernel Data Structures 36
- 1.10 Computing Environments 40
- 1.11 Free and Open-Source Operating Systems 46
- Practice Exercises 53
- Further Reading 54

## Chapter 2 Operating-System Structures

- 2.1 Operating-System Services 55
- 2.2 User and Operating-System Interface 58
- 2.3 System Calls 62
- 2.4 System Services 74
- 2.5 Linkers and Loaders 75
- 2.6 Why Applications Are Operating-System Specific 77
- 2.7 Operating-System Design and Implementation 79
- 2.8 Operating-System Structure 81
- 2.9 Building and Booting an Operating System 92
- 2.10 Operating-System Debugging 95
- 2.11 Summary 100
- Practice Exercises 101
- Further Reading 101

# منبع درس

## PART TWO ■ PROCESS MANAGEMENT

### Chapter 3 Processes

3.1 Process Concept	106	3.7 Examples of IPC Systems	132
3.2 Process Scheduling	110	3.8 Communication in Client–	
3.3 Operations on Processes	116	Server Systems	145
3.4 Interprocess Communication	123	3.9 Summary	153
3.5 IPC in Shared-Memory Systems	125	Practice Exercises	154
3.6 IPC in Message-Passing Systems	127	Further Reading	156

## منبع درس

xxi

### Contents

### Chapter 4 Threads & Concurrency

4.1 Overview	160	4.6 Threading Issues	188
4.2 Multicore Programming	162	4.7 Operating-System Examples	194
4.3 Multithreading Models	166	4.8 Summary	196
4.4 Thread Libraries	168	Practice Exercises	197
4.5 Implicit Threading	176	Further Reading	198

## Chapter 5 CPU Scheduling

5.1 Basic Concepts	200	5.7 Operating-System Examples	234
5.2 Scheduling Criteria	204	5.8 Algorithm Evaluation	244
5.3 Scheduling Algorithms	205	5.9 Summary	250
5.4 Thread Scheduling	217	Practice Exercises	251
5.5 Multi-Processor Scheduling	220	Further Reading	254
5.6 Real-Time CPU Scheduling	227		

## PART THREE ■ PROCESS SYNCHRONIZATION

# منبع درس

### Chapter 6 Synchronization Tools

- |   |     |                    |     |
|---|-----|--------------------|-----|
| 6.1 Background                              | 257 | 6.7 Monitors       | 276 |
| 6.2 The Critical-Section Problem            | 260 | 6.8 Liveness       | 283 |
| 6.3 Peterson's Solution                     | 262 | 6.9 Evaluation     | 284 |
| 6.4 Hardware Support for<br>Synchronization | 265 | 6.10 Summary       | 286 |
| 6.5 Mutex Locks                             | 270 | Practice Exercises | 287 |
| 6.6 Semaphores                              | 272 | Further Reading    | 288 |

### Chapter 7 Synchronization Examples

- |  |     |                            |     |
|--|-----|----------------------------|-----|
| 7.1 Classic Problems of<br>Synchronization | 289 | 7.5 Alternative Approaches | 311 |
| 7.2 Synchronization within the Kernel      | 295 | 7.6 Summary                | 314 |
| 7.3 POSIX Synchronization                  | 299 | Practice Exercises         | 314 |
| 7.4 Synchronization in Java                | 303 | Further Reading            | 315 |

### Chapter 8 Deadlocks

- |   |     |                            |     |
|---|-----|----------------------------|-----|
| 8.1 System Model                              | 318 | 8.6 Deadlock Avoidance     | 330 |
| 8.2 Deadlock in Multithreaded<br>Applications | 319 | 8.7 Deadlock Detection     | 337 |
| 8.3 Deadlock Characterization                 | 321 | 8.8 Recovery from Deadlock | 341 |
| 8.4 Methods for Handling Deadlocks            | 326 | 8.9 Summary                | 343 |
| 8.5 Deadlock Prevention                       | 327 | Practice Exercises         | 344 |
|   |     | Further Reading            | 346 |



## PART FOUR ■ MEMORY MANAGEMENT

## منبع درس

### Chapter 9 Main Memory

- 9.1 Background 349
- 9.2 Contiguous Memory Allocation 356
- 9.3 Paging 360
- 9.4 Structure of the Page Table 371
- 9.5 Swapping 376
- 9.6 Example: Intel 32- and 64-bit Architectures 379
- 9.7 Example: ARMv8 Architecture 383
- 9.8 Summary 384
  - Practice Exercises 385
  - Further Reading 387

### Chapter 10 Virtual Memory

- 10.1 Background 389
- 10.2 Demand Paging 392
- 10.3 Copy-on-Write 399
- 10.4 Page Replacement 401
- 10.5 Allocation of Frames 413
- 10.6 Thrashing 419
- 10.7 Memory Compression 425
- 10.8 Allocating Kernel Memory 426
- 10.9 Other Considerations 430
- 10.10 Operating-System Examples 436
- 10.11 Summary 440
  - Practice Exercises 441
  - Further Reading 444

## PART FIVE ■ STORAGE MANAGEMENT

## منبع درس

### Chapter 11 Mass-Storage Structure

- |   |                                |
|---|--------------------------------|
| 11.1 Overview of Mass-Storage Structure 449 | 11.6 Swap-Space Management 467 |
| 11.2 HDD Scheduling 457                     | 11.7 Storage Attachment 469    |
| 11.3 NVM Scheduling 461                     | 11.8 RAID Structure 473        |
| 11.4 Error Detection and Correction 462     | 11.9 Summary 485               |
| 11.5 Storage Device Management 463          | Practice Exercises 486         |
|   | Further Reading 487            |

### Chapter 12 I/O Systems

- |   |                        |
|---|------------------------|
| 12.1 Overview 489   | 12.6 STREAMS 519       |
| 12.2 I/O Hardware 490                                     | 12.7 Performance 521   |
| 12.3 Application I/O Interface 500                        | 12.8 Summary 524       |
| 12.4 Kernel I/O Subsystem 508                             | Practice Exercises 525 |
| 12.5 Transforming I/O Requests to Hardware Operations 516 | Further Reading 526    |

## PART SIX ■ FILE SYSTEM

## منبع درس

### Chapter 13 File-System Interface

- |                              |                              |
|------------------------------|------------------------------|
| 13.1 File Concept 529        | 13.5 Memory-Mapped Files 555 |
| 13.2 Access Methods 539      | 13.6 Summary 560             |
| 13.3 Directory Structure 541 | Practice Exercises 560       |
| 13.4 Protection 550          | Further Reading 561          |

### Chapter 14 File-System Implementation

- |                                     |  |
|-------------------------------------|--|
| 14.1 File-System Structure 564      | 14.7 Recovery 586                      |
| 14.2 File-System Operations 566     | 14.8 Example: The WAFL File System 589 |
| 14.3 Directory Implementation 568   | 14.9 Summary 593                       |
| 14.4 Allocation Methods 570         | Practice Exercises 594                 |
| 14.5 Free-Space Management 578      | Further Reading 594                    |
| 14.6 Efficiency and Performance 582 |  |

### Chapter 15 File-System Internals

- |                                  |                                |
|----------------------------------|--------------------------------|
| 15.1 File Systems 597            | 15.7 Consistency Semantics 608 |
| 15.2 File-System Mounting 598    | 15.8 NFS 610                   |
| 15.3 Partitions and Mounting 601 | 15.9 Summary 615               |
| 15.4 File Sharing 602            | Practice Exercises 616         |
| 15.5 Virtual File Systems 603    | Further Reading 617            |
| 15.6 Remote File Systems 605     |                                |

## PART SEVEN ■ SECURITY AND PROTECTION

## منبع درس

### Chapter 16 Security

- 16.1 The Security Problem 621
- 16.2 Program Threats 625
- 16.3 System and Network Threats 634
- 16.4 Cryptography as a Security Tool 637
- 16.5 User Authentication 648
- 16.6 Implementing Security Defenses 653
- 16.7 An Example: Windows 10 662
- 16.8 Summary 664
- Further Reading 665

### Chapter 17 Protection

- 17.1 Goals of Protection 667
- 17.2 Principles of Protection 668
- 17.3 Protection Rings 669
- 17.4 Domain of Protection 671
- 17.5 Access Matrix 675
- 17.6 Implementation of the Access Matrix 679
- 17.7 Revocation of Access Rights 682
- 17.8 Role-Based Access Control 683
- 17.9 Mandatory Access Control (MAC) 684
- 17.10 Capability-Based Systems 685
- 17.11 Other Protection Improvement Methods 687
- 17.12 Language-Based Protection 690
- 17.13 Summary 696
- Further Reading 697



## PART EIGHT ■ ADVANCED TOPICS

## منبع درس

### Chapter 18 Virtual Machines

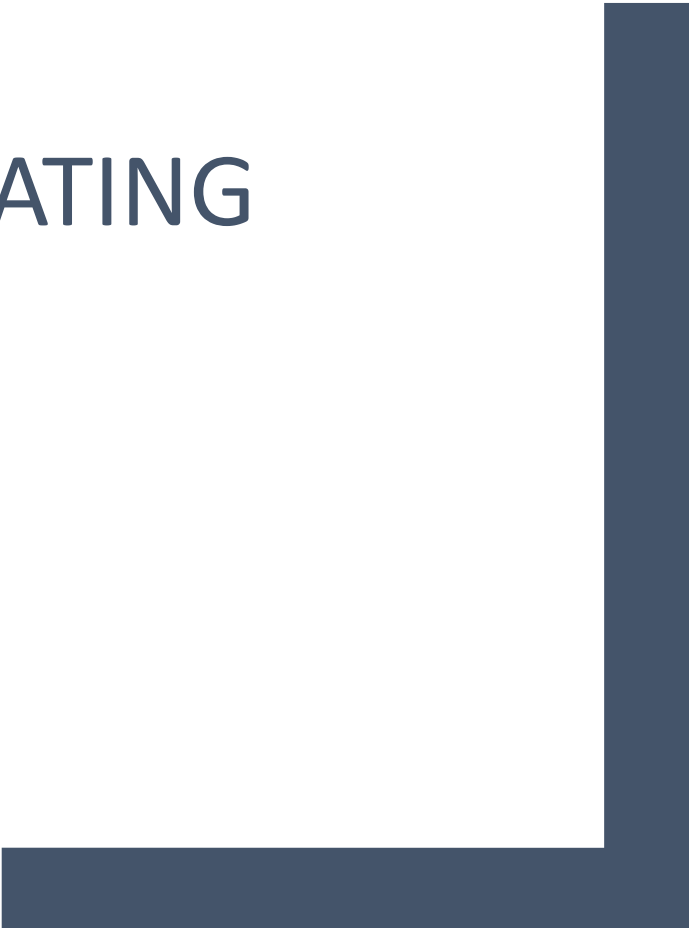
- 18.1 Overview 701
- 18.2 History 703
- 18.3 Benefits and Features 704
- 18.4 Building Blocks 707
- 18.5 Types of VMs and Their Implementations 713
- 18.6 Virtualization and Operating-System Components 719
- 18.7 Examples 726
- 18.8 Virtualization Research 728
- 18.9 Summary 729  
Further Reading 730

### Chapter 19 Networks and Distributed Systems

- 19.1 Advantages of Distributed Systems 733
- 19.2 Network Structure 735
- 19.3 Communication Structure 738
- 19.4 Network and Distributed Operating Systems 749
- 19.5 Design Issues in Distributed Systems 753
- 19.6 Distributed File Systems 757
- 19.7 DFS Naming and Transparency 761
- 19.8 Remote File Access 764
- 19.9 Final Thoughts on Distributed File Systems 767
- 19.10 Summary 768  
Practice Exercises 769  
Further Reading 770



# INTRODUCTION TO OPERATING SYSTEMS

1. What is an Operating System?
  2. Review of OS-Related Hardware
- 

# What is an Operating System?

- “A *program* ... that controls the execution of application programs and implements an interface between the user of a computer and the computer hardware”
- Runs on PCs, workstations, servers, smart phones, routers, embedded systems, etc

# Operating System Roles

## ■ Abstract Machine

- *Hides complex details of the underlying hardware*
- *Provides common API to applications and services*
- *Simplifies application writing*

## ■ Resource Manager


- *Controls accesses to shared resources*
  - CPU, memory, disks, network, ...
- *Allows for global policies to be implemented*

# ABSTRACT MACHINE

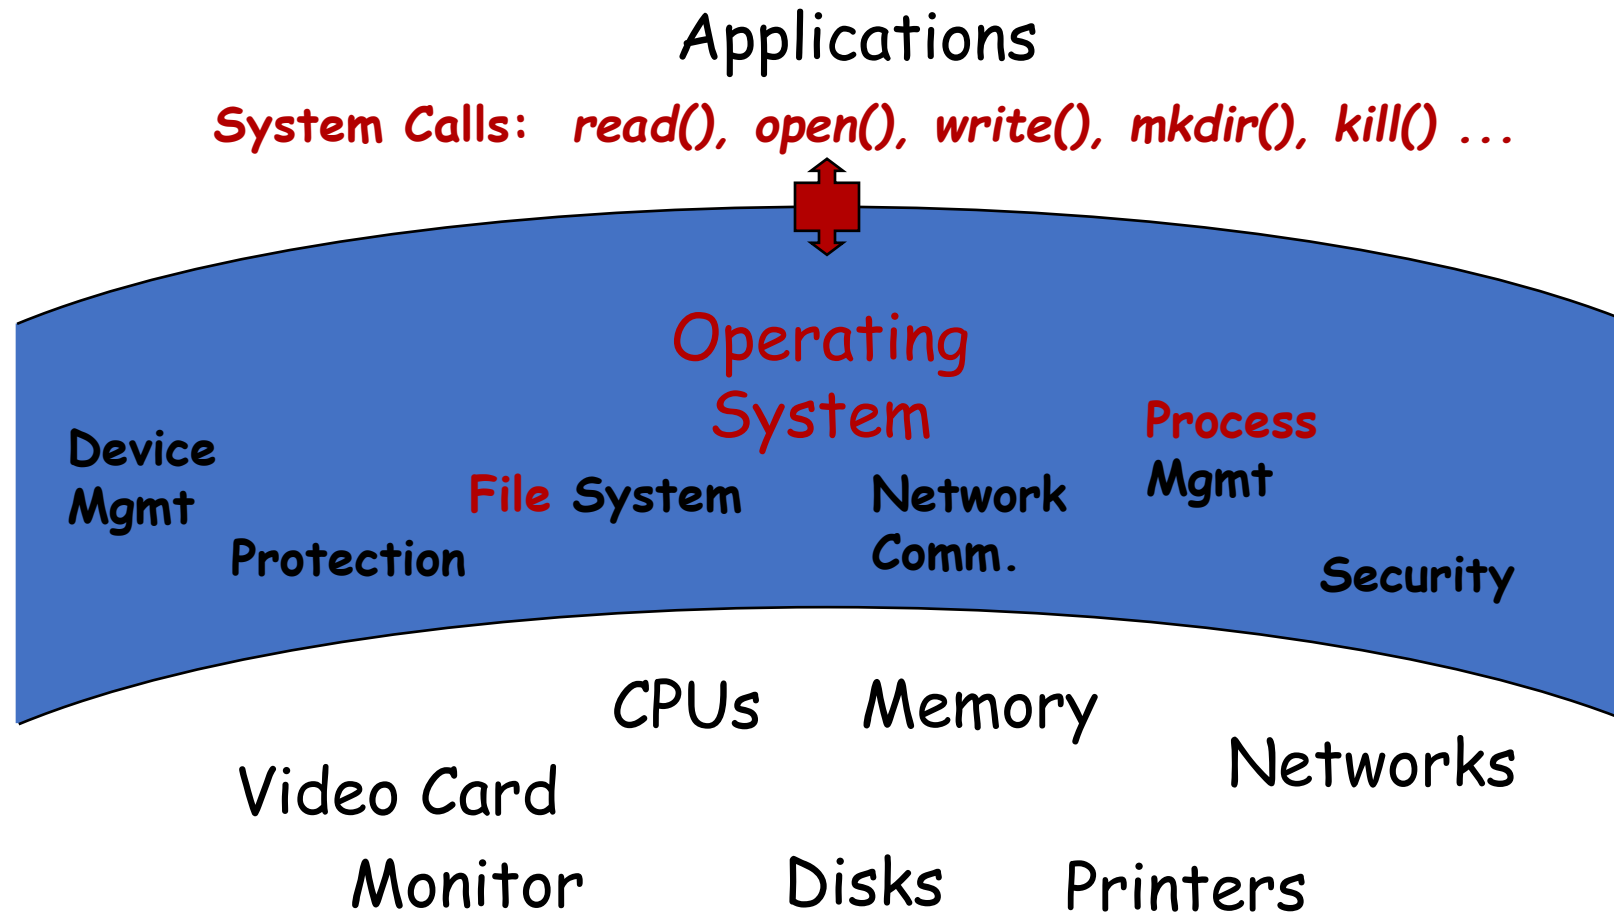
# The Abstract Machine Role

- Without operating systems, application writers would have to program all device access directly:
  - *Load device command codes into device registers*
  - *Handle initialization, recalibration, sensing, timing*
  - *Understand physical characteristics and data layout*
  - *Control motors*
  - *Interpret return codes*
  - *...*
- Application programming would be complicated
- Applications would be difficult to maintain, upgrade and port
  - *This OS code could be written just once and then shared!*

# Providing abstraction via system calls



# Providing abstraction via system calls





# RESOURCE MANAGER

# OS as a resource manager

- Allocating resources to applications
  - *time sharing resources (scheduling)*
  - *space sharing resources (allocation)*
- Making efficient use of limited resources
  - *improving utilization*
  - *minimizing overhead*
  - *improving throughput/good put*
- Protecting applications from each other
  - *enforcement of boundaries*

# Problems Solved by OS

- Time sharing the CPU among applications
- Space sharing the memory among applications
- Space sharing the disk among users
- Time sharing access to the network

# More Problems Solved by OS

- Protection of applications from each other, of user data from other users and of I/O devices
- Protection of the OS itself!
- Prevention of direct access to hardware, where this would cause problems

# QUESTIONS?

# OS NEEDS FROM HARDWARE

# OS Needs Help from Hardware

- The OS is just a program!
- When it is not running, it can't do anything!
- Its goal is to run applications, not itself!
  
- The OS needs help from the hardware in order to detect and prevent certain activities, and to maintain control

# Brief Review of Hardware

- Instruction sets define all that a CPU can do, and differ among CPU architectures
- All have **load** and **store** instructions to move data between memory and registers
- Many instructions for comparing and combining values in registers



# Basic Anatomy of a CPU

- Program Counter (PC)
  - *Holds the memory address of the next instruction*
- Instruction Register (IR)
  - *holds the instruction currently being executed*
- General Registers (Reg. 1..n)
  - *hold variables and temporary results*
- Arithmetic and Logic Unit (ALU)
  - *performs arithmetic functions and logic operations*

# Basic Anatomy of a CPU

## ■ Stack Pointer (SP)

- *holds memory address of a stack with a frame for each active procedure's parameters & local variables*


## ■ Processor Status Word (PSW)

- *contains various control bits including the **mode bit** which determines whether privileged instructions can be executed at this time*

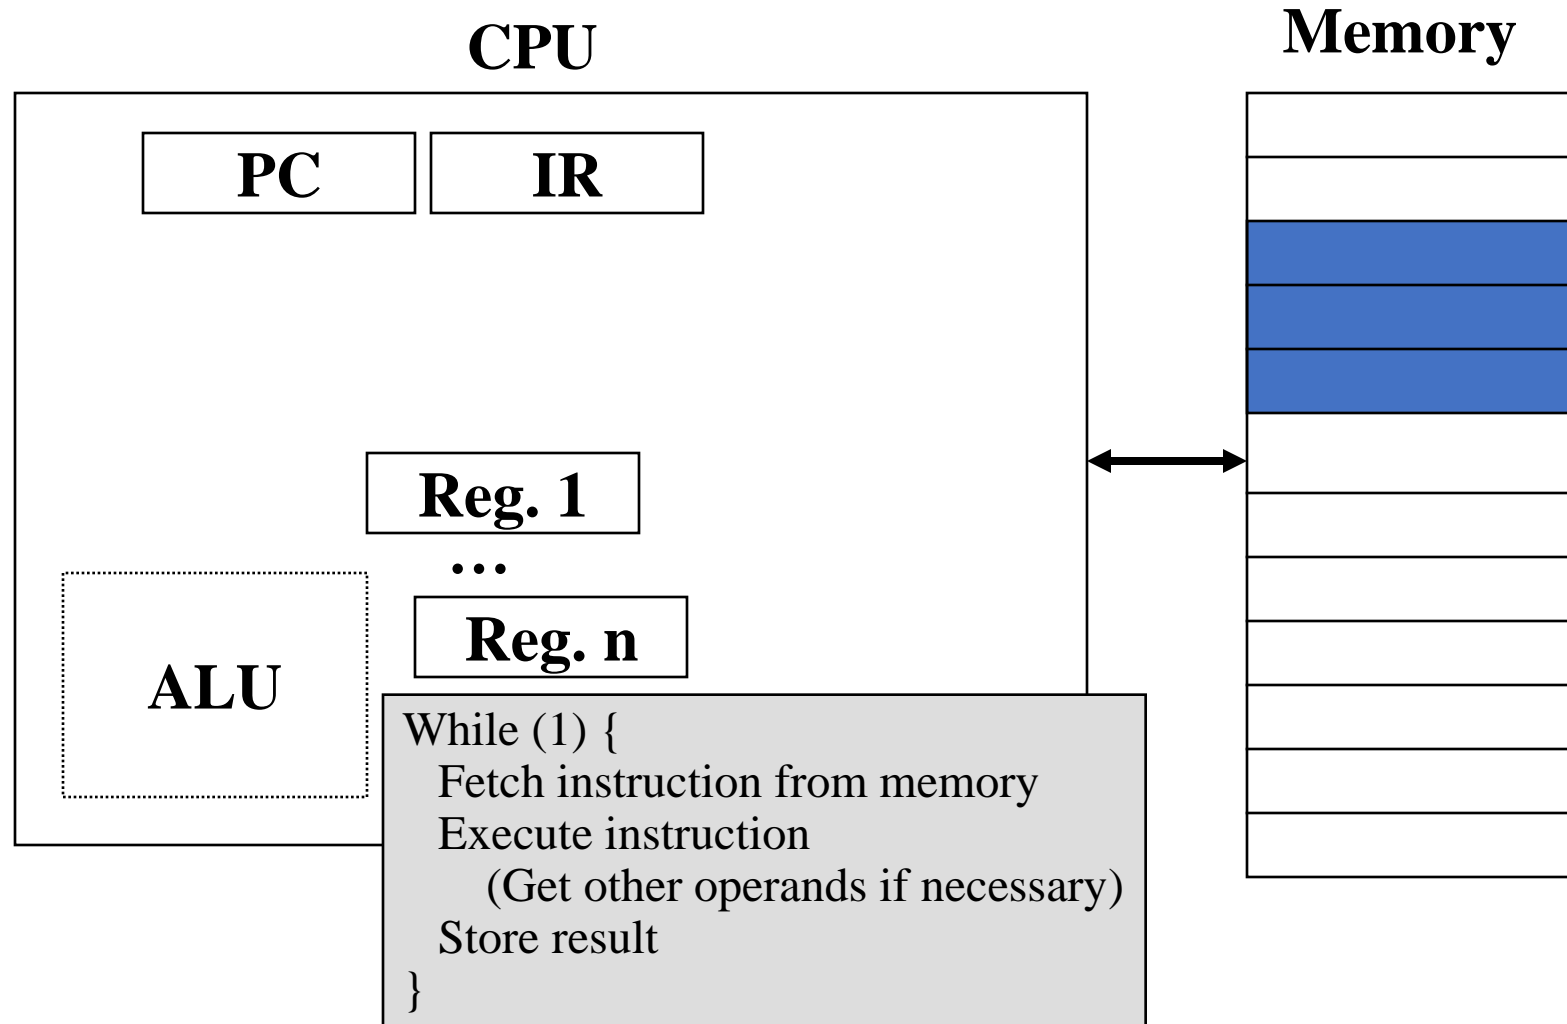
# Program Execution

- The Fetch/Decode/Execute cycle
  - *fetch next instruction pointed to by PC*
  - *decode it to find its type and operands*
  - *execute it*
  - *repeat*
- At a fundamental level, this is all a CPU does, regardless of which program it is executing

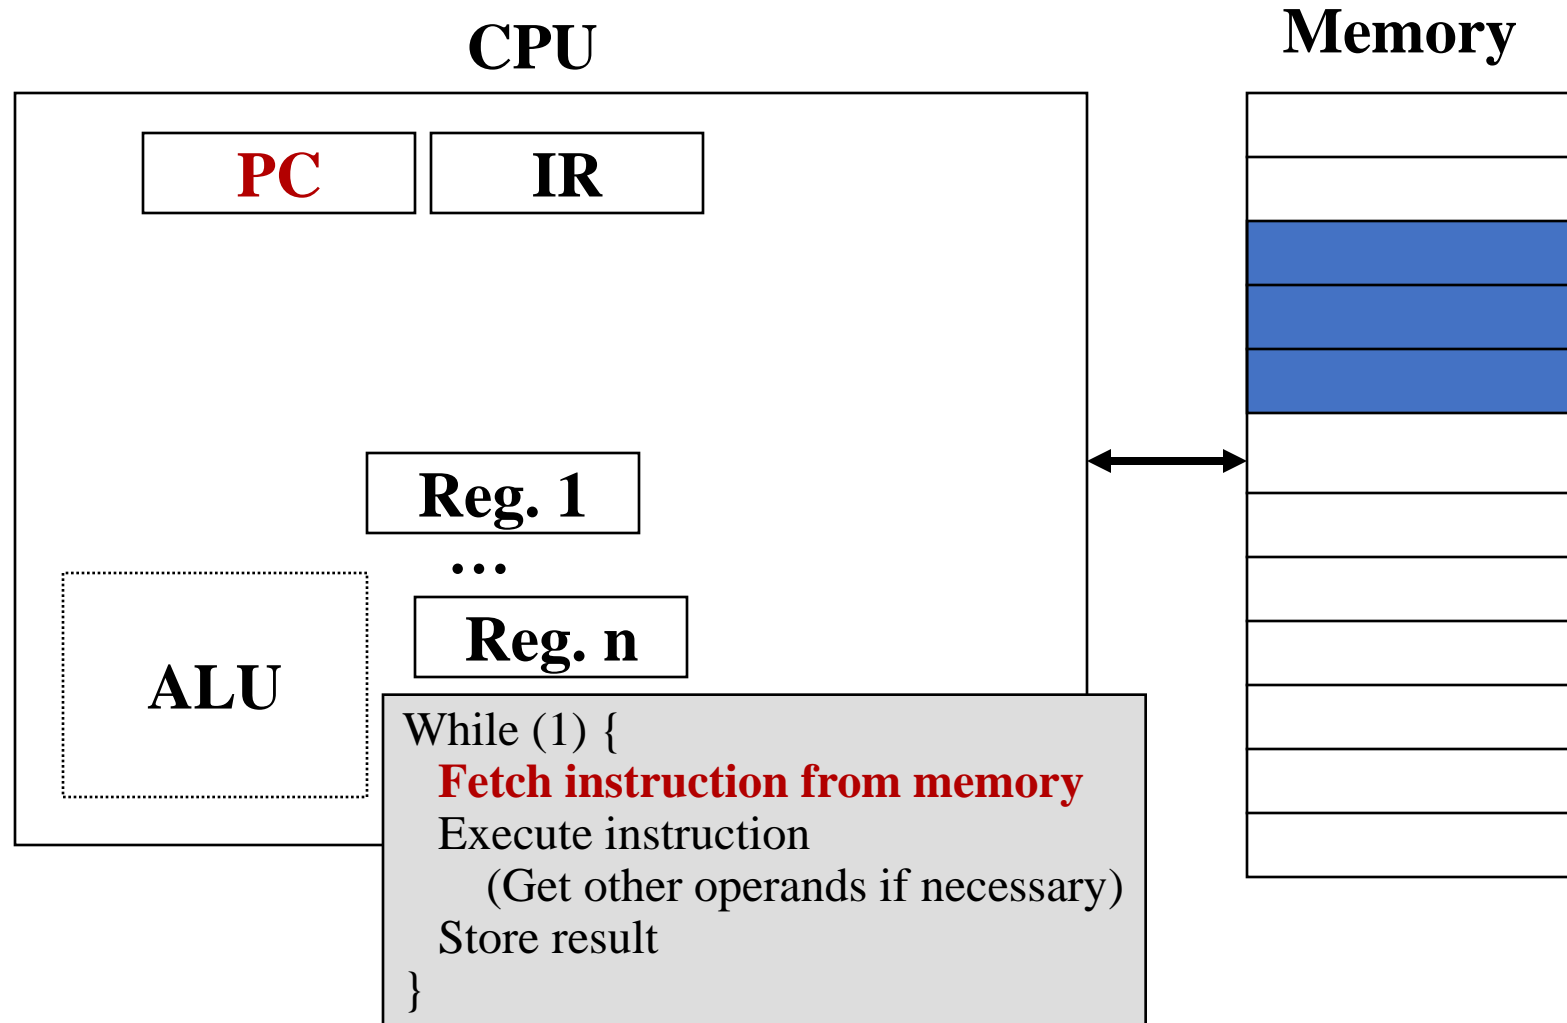
# Fetch/Decode/Execute Cycle



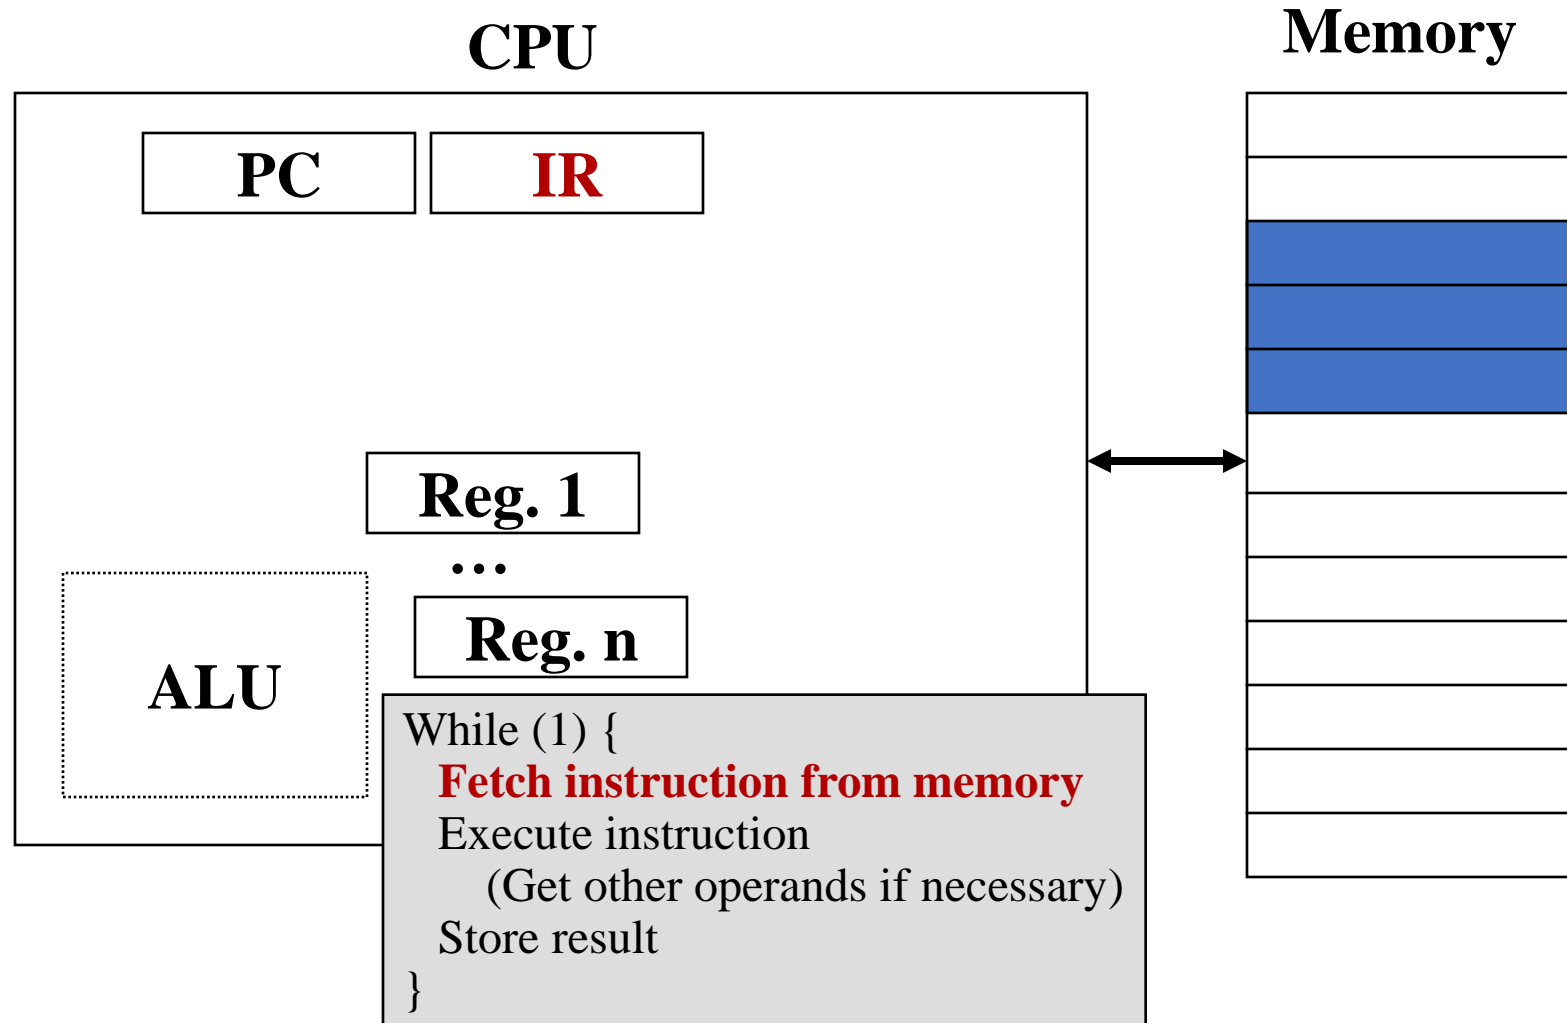
# Fetch/Decode/Execute Cycle



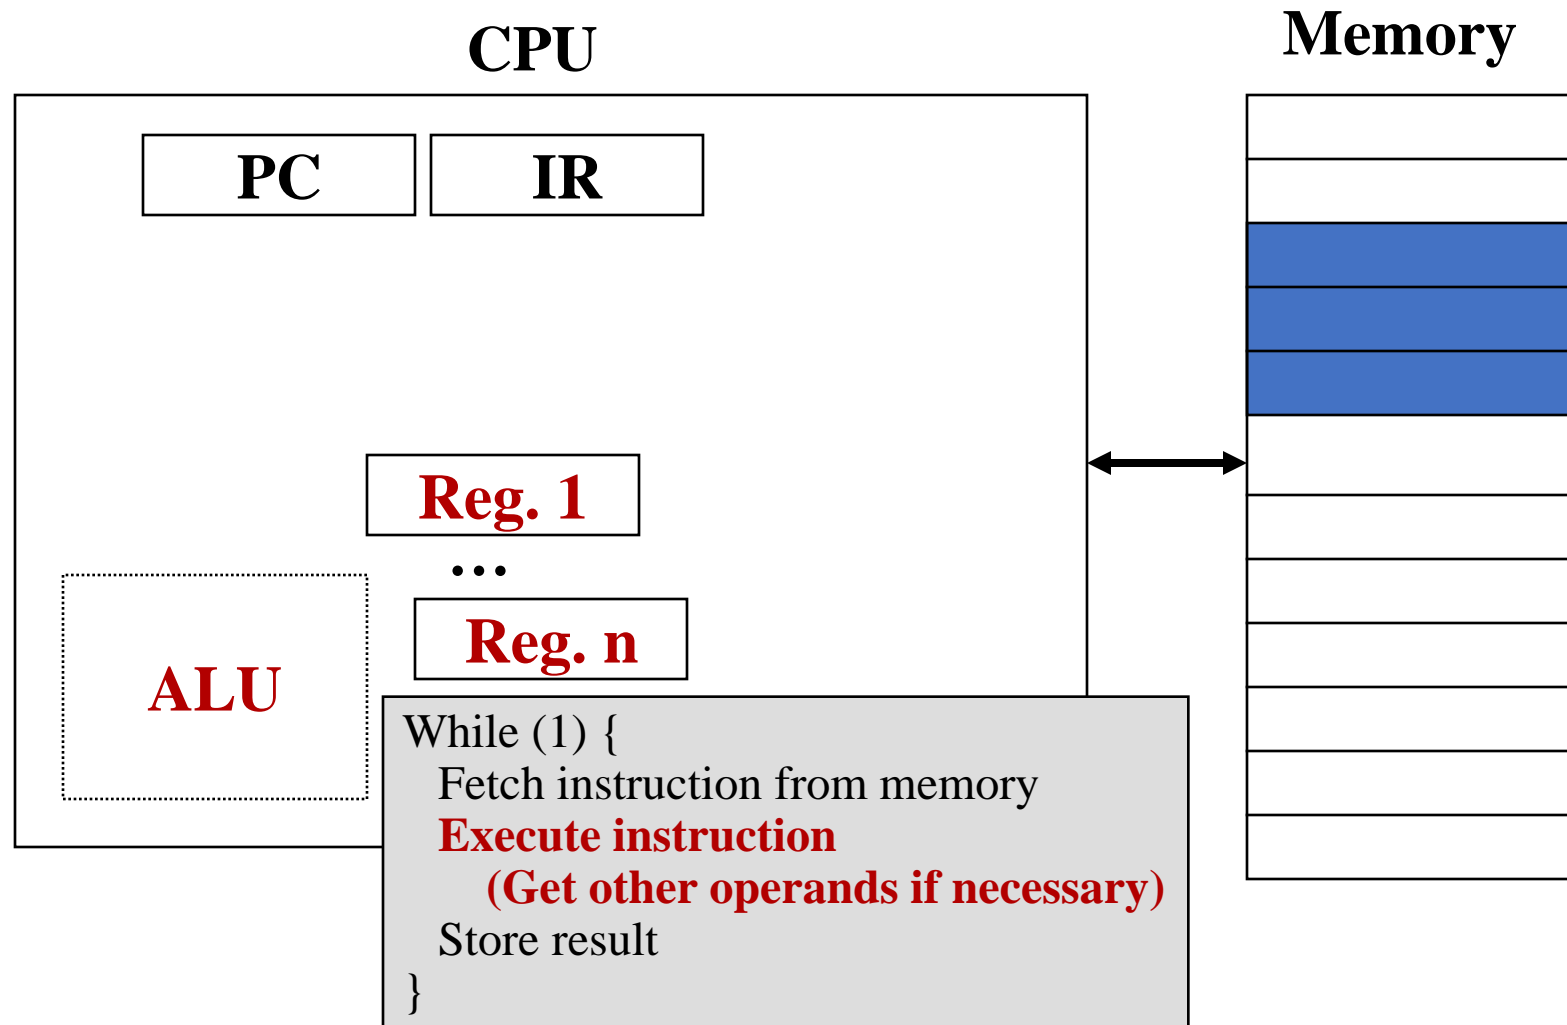
# Fetch/Decode/Execute Cycle



# Fetch/Decode/Execute Cycle

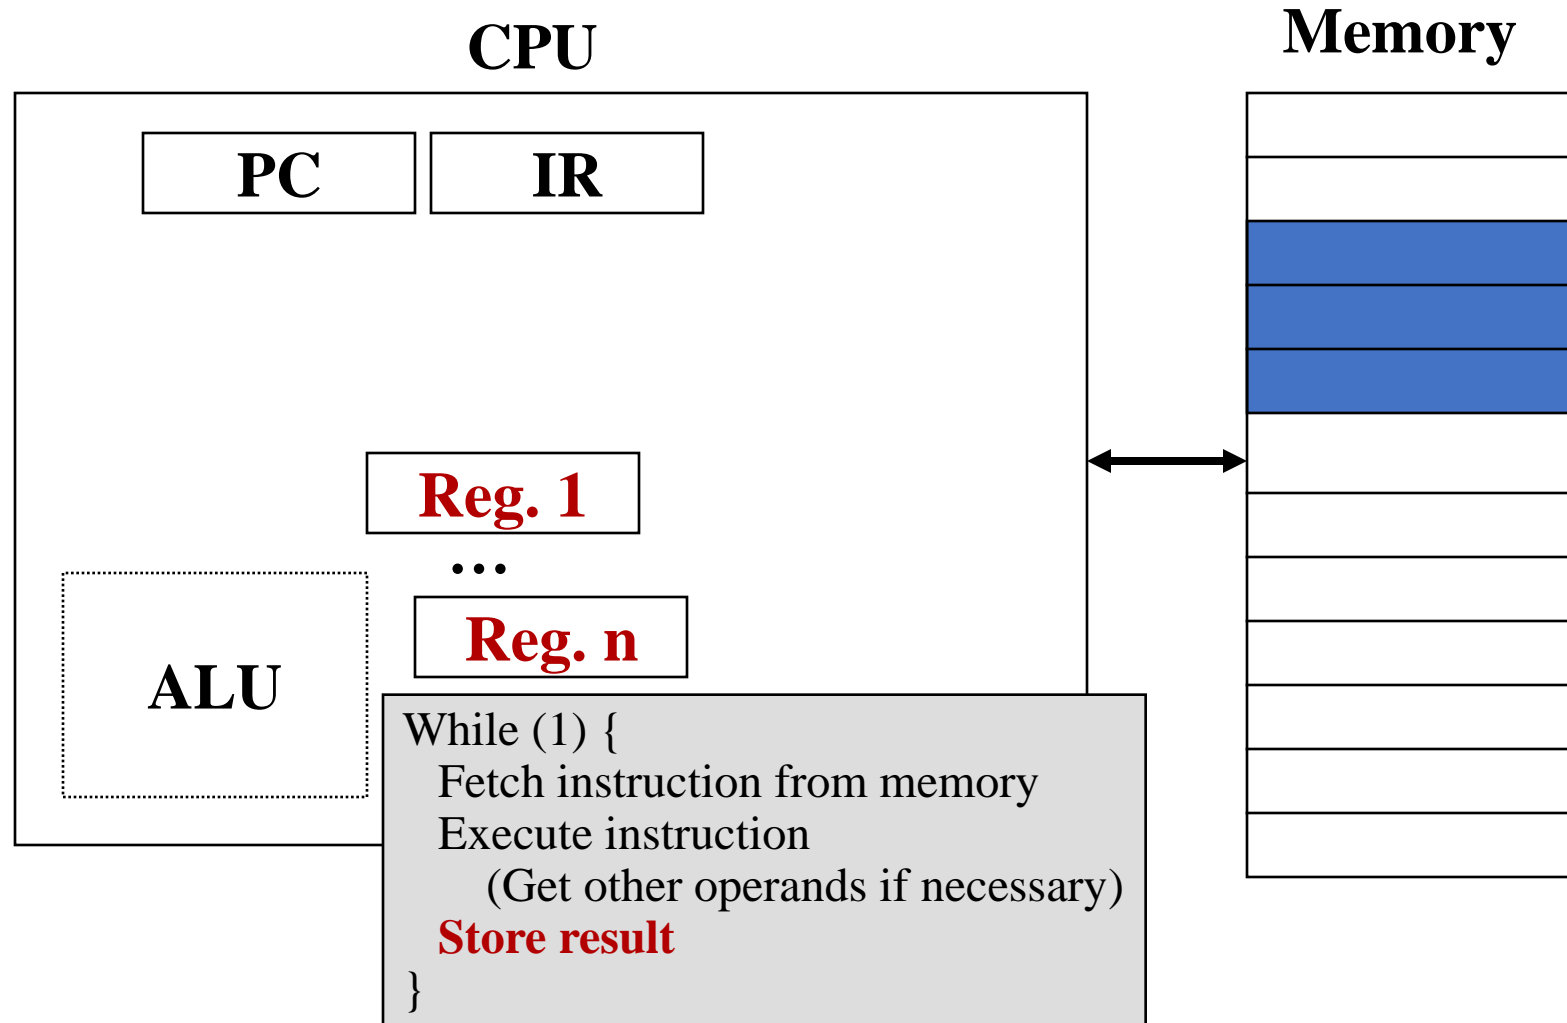


# Fetch/Decode/Execute Cycle






# Fetch/Decode/Execute Cycle



# Fetch/Decode/Execute Cycle



# The OS is Just a Program!

- The OS is just a sequence of instructions that the CPU will fetch/decode/execute
- How can the OS cause application programs to run?
- How can applications cause the OS to run?

# How Can an OS Run Applications?

- The OS must load the address of the application's starting instruction into the PC
- Example:
  - computer boots and begins running the OS
  - OS code must get into memory somehow
  - fetch/decode/execute OS instructions
  - OS requests user input to identify application program/file
  - OS loads application (executable file) into memory
  - OS loads the address of the app's first instruction into the PC
  - CPU fetches/decodes/executes the application's instructions

# The OS is Just a Program!

- How does the OS ever get to run again?
- How can the OS switch the CPU to run a new application (and later resume the first one)?
- How can the OS maintain control of what the application does when the OS is not running?
- In what ways can application try to seize control indefinitely (ie. cheat)?
- How can the OS prevent such cheating?

# How Can the OS Regain Control?

- What if an application doesn't call the OS and instead just hogs the CPU?
  - *OS needs interrupts from a timer device!*
  - *OS must register a future timer interrupt before handing control of the CPU over to an application*
  - *When the timer interrupt goes off the hardware starts running the OS at a pre-specified location called an interrupt handler*
  - *The interrupt handler is part of the OS program*
  - *The address of the interrupt handler's first instruction is placed in the PC by the h/w*

# Can the Application Cheat?

- Can the application disable the future timer interrupt so that the OS can not take control back from it?
- Disabling interrupts must be a privileged instruction that is not executable by applications
- The CPU knows whether or not to execute privileged instructions based on the value of the mode bit in the PSW!
- Privileged instructions can only be executed when the mode bit is set
  - - eg. disabling interrupts, setting the mode bit!
- - attempted execution in non-privileged mode generally causes an interrupt (trap) to occur

# Are There Other Ways to Cheat?

- What stops the running application from modifying the OS?
  - - *eg. modifying the timer interrupt handler to jump control back to the application?*



# What Stops Applications From Modifying the OS?

- Memory protection!
- Memory protection instructions must be privileged
  - - i.e., they can only be executed with the mode bit set ...
- Why must the OS clear the mode bit before it hands control to an application?

# How Can Applications Invoke the OS?

- Why not just set PC to an OS instruction address and transfer control that way?

# How Can Applications Invoke the OS?

- Special trap instruction causes a kind of interrupt
  - *changes PC to point to a predetermined OS entry point instruction*
  - *simultaneously sets the mode bit*
  - *CPU is now running in privileged mode*
- Application calls a library procedure that includes the appropriate trap instruction
- fetch/decode/execute cycle begins at a pre-specified OS entry point called a **system call handler**

# Are Traps Interrupts?

- Traps, like interrupts, are hardware events
- But traps are synchronous where as interrupts are asynchronous
- i.e. traps are caused by the executing program rather than a device external to the CPU

# Switching to a New Application?

- To suspend execution of an application the OS must run!
- After that, simply
  - - capture the application's memory state and processor state
  - - preserve all the memory values of this application
  - - copy values of all CPU registers into a data structure which is saved in memory
  - - restarting the application from the same point just requires reloading the register values

# Recap

- Why do we need a timer device?
- Why do we need an interrupt mechanism?
- Why do we need privileged instructions?
- Why are system calls different to procedure calls?
- How are system calls different to interrupts?
- Why is memory protection necessary?
- How can the OS switch among applications?