

R

بسم الله الرحمن الرحيم

سیستم عامل

جلسه هفتم – سمافور و مسائل همزمانی

آنچه گذشت

پیاده‌سازی mutex

Mutex Lock Operations

- Lock (*mutex*)
 - *Acquire the lock if it is free ... and continue*
 - *Otherwise wait until it can be acquired*
- Unlock (*mutex*)
 - *Release the lock*
 - *If there are waiting threads wake one up*

Peterson's Algorithm for Thread *i*

```
while (true) {
```

```
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j);
```

```
    /* critical section */
```

```
    flag[i] = false;
```

```
    /* remainder section */
```

```
}
```

Test-and-Set-Lock (TSL) Instruction

- ❑ Test-and-set-lock does the following *atomically*:
 - *Get the (old) value*
 - *Set the lock to TRUE*
 - *Return the old value*

If the returned value was FALSE...

- Then you got the lock!!!

If the returned value was TRUE...

- Then someone else has the lock
(so try again later)

Implementing a Mutex With TSL

```
1 repeat
2   while(TSL(mylock)) } Lock (mylock)
3   no-op;
4   critical section
5   mylock = FALSE; } Unlock (mylock)
6 remainder section
7 until FALSE
```

- Note that processes are **busy** while waiting
 - *this kind of mutex is called a **spin lock***

Busy Waiting

- Also called polling or spinning
 - *The thread consumes CPU cycles to evaluate when the lock becomes free !*
- Problem on a single CPU system...
 - *A busy-waiting thread can prevent the lock holder from running & completing its critical section & releasing the lock!*
 - *time spent spinning is wasted on a single CPU system*
 - *Why not block instead of busy wait ?*

Concurrency in the Kernel

Solution 1: Disable interrupts during critical sections

- *Ensures that interrupt handling code will not run*
- *... but what if there are multiple CPUs?*

Solution 2: Use mutex locks based on TSL for critical sections

- *Ensures mutual exclusion for all code that follows that convention*

Disabling interrupts is not enough on MPs...

- Disabling interrupts during critical sections
 - *Ensures that interrupt handling code will not run*
 - *But what if there are multiple CPUs?*
 - *A thread on a different CPU might make a system call which invokes code that manipulates the ready queue*
 - *Disabling interrupts on one CPU didn't prevent this!*
- Solution: use a mutex lock (based on TSL)
 - *Ensures mutual exclusion for all code that uses it*

Mutex is not enough

- Interrupt inside interrupt handler

جلسه‌ی جدید

سمافور و مسائل کلاسیک همروندی

مسئله‌ی نویسنده، خواننده

The Producer-Consumer Problem

- An example of the **pipelined model**
 - *One thread produces data items*
 - *Another thread consumes them*
- Use a bounded buffer between the threads
- The buffer is a shared resource
 - *Code that manipulates it is a **critical section***
- Must suspend the producer thread if the buffer is full
- Must suspend the consumer thread if the buffer is empty

Is This Solution Correct?

```
thread producer {
    while(1){
        // Produce char c
        while (count==n) {
            no_op
        }
        buf[InP] = c
        InP = InP + 1 mod n
        count++
    }
}
```

```
thread consumer {
    while(1){
        while (count==0) {
            no_op
        }
        c = buf[OutP]
        OutP = OutP + 1 mod n
        count--
        // Consume char
    }
}
```

Global variables:

```
char buf[n]
int InP = 0    // place to add
int OutP = 0   // place to get
int count
```

This Code is Incorrect!

- The “count” variable can be corrupted:
 - *Increments or decrements may be lost!*
 - *Possible Consequences:*
 - Both threads may spin forever
 - Buffer contents may be over-written
- *What is this problem called?*

This Code is Incorrect!

- *What is this problem called?*
- *Race Condition*
- Code that manipulates count must be made into a *???* and protected using *???*

This Code is Incorrect!

- *What is this problem called?*
- *Race Condition*
- Code that manipulates count must be made into a *critical section* and protected using *mutual exclusion*


More Problems With This Code

- What if buffer is full?
 - *Producer will busy-wait*
- What if buffer is empty?
 - *Consumer will busy-wait*
- We need a solution based on blocking!

A Solution Based On Blocking

```
0  thread producer {
1      while(1) {
2          // Produce char c
3          if (count==n) {
4              sleep(full)
5          }
6          buf[InP] = c;
7          InP = InP + 1 mod n
8          count++
9          if (count == 1)
10             wakeup(empty)
11      }
12 }
```

```
0  thread consumer {
1      while(1) {
2          if(count==0) {
3              sleep(empty)
4          }
5          c = buf[OutP]
6          OutP = OutP + 1 mod n
7          count--;
8          if (count == n-1)
9              wakeup(full)
10         // Consume char
11     }
12 }
```




Global variables:

```
char buf[n]
int InP = 0    // place to add
int OutP = 0   // place to get
int count
```

Use a Mutex to Fix The Problem

```
0  thread producer {
1      while(1) {
2          // Produce char c
3          if (count==n) {
4              sleep(full)
5          }
6          buf[InP] = c;
7          InP = InP + 1 mod n
8          count++
9          if (count == 1)
10             wakeup(empty)
11     }
12 }
```

```
0  thread consumer {
1      while(1) {
2          if(count==0) {
3              sleep(empty)
4          }
5          c = buf[OutP]
6          OutP = OutP + 1 mod n
7          count--;
8          if (count == n-1)
9              wakeup(full)
10         // Consume char
11     }
12 }
```



Global variables:

```
char buf[n]
int InP = 0    // place to add
int OutP = 0   // place to get
int count
```

Problems

- 1. Sleeping while holding the mutex causes deadlock !
- 2. Releasing the mutex then sleeping opens up a window during which a context switch might occur ... again risking deadlock
- 3. How can we release the mutex and sleep in a single atomic operation?
- We need a more powerful synchronization primitive

Semaphores

- An abstract data type that can be used for condition synchronization and mutual exclusion

What is the difference between mutual exclusion and condition synchronization?

Semaphores

- Condition synchronization
 - *wait* until condition holds before proceeding
 - *signal* when condition holds so others may proceed
- Mutual exclusion
 - *only one at a time in a critical section*

Semaphores

- An abstract data type
 - *containing an integer variable (S)*
 - *Two operations: $Wait(S)$ and $Signal(S)$*
- Alternative names for the two operations
 - $Wait(S) = Down(S) = P(S)$
 - $Signal(S) = Up(S) = V(S)$
- Blitz names its semaphore operations Down and Up

Classical Definition

Down (S)

{

while S <= 0 do noop; / busy wait!*

**/*

S = S - 1; / S >= 0 */*

}

Up (S)

{

S = S + 1;

}

Problems With The Definition

- Waiting threads hold the CPU
 - *Waste of time in single CPU systems*
 - *Required preemption to avoid deadlock*

Blocking Semaphores

Semaphore S has a value, S.val, and a thread list, S.list.

Down (S)

```
S.val = S.val - 1
```

```
If S.val < 0
```

```
{ add calling thread to S.list;  
  block;
```

```
}
```

```
/* negative value of S.val */
```

```
/* is # waiting threads */
```

```
/* sleep */
```

Up (S)

```
S.val = S.val + 1
```

```
If S.val <= 0
```

```
{ remove a thread T from S.list;  
  wakeup (T);
```

```
}
```

Implementing Semaphores

- Down () and Up () are assumed to be **atomic**

How can we ensure that they are atomic?

Implementing Semaphores

- Implement Down() and Up() as system calls?
 - *How can the kernel ensure Down() and Up() are completed atomically?*
 - *Same solutions as before (disable interrupts, or use TSL-based mutex)*

Semaphores With Disabling

```
struct semaphore {  
    int val;  
    list L;  
}
```

```
Down(semaphore sem)  
    DISABLE_INTS  
    sem.val--  
    if (sem.val < 0){  
        add thread to sem.L  
        sleep(thread)  
    }  
    ENABLE_INTS
```

```
Up(semaphore sem)  
    DISABLE_INTS  
    sem.val++  
    if (sem.val <= 0) {  
        th = remove next  
        thread from sem.L  
        wakeup(th)  
    }  
    ENABLE_INTS
```

Semaphores With Disabling

```
struct semaphore {  
    int val;  
    list L;  
}
```

```
Down(semaphore sem)  
    DISABLE_INTS  
    sem.val--  
    if (sem.val < 0){  
        add thread to sem.L  
        sleep(thread)  
    }  
    ENABLE_INTS
```

```
Up(semaphore sem)  
    DISABLE_INTS  
    sem.val++  
    if (sem.val <= 0) {  
        th = remove next  
        thread from sem.L  
        wakeup(th)  
    }  
    ENABLE_INTS
```


Semaphore.down in Blitz

```
method Down()  
    var oldIntStat: int  
    oldIntStat = SetInterruptsTo (DISABLED)  
    if count == 0x80000000  
        FatalError ("Semaphore count underflowed during 'Wait' operation")  
    EndIf  
    count = count - 1  
    if count < 0 waitingThreads.AddToEnd (currentThread)  
        currentThread.Sleep ()  
    endIf  
    oldIntStat = SetInterruptsTo (oldIntStat)  
endMethod
```

Semaphore.down in Blitz

```
method Down()  
    var oldIntStat: int  
    oldIntStat = SetInterruptsTo (DISABLED)  
    if count == 0x80000000  
        FatalError ("Semaphore count underflowed during 'Wait' operation")  
    EndIf  
    count = count - 1  
    if count < 0 waitingThreads.AddToEnd (currentThread)  
        currentThread.Sleep ()  
    EndIf  
    oldIntStat = SetInterruptsTo (oldIntStat)  
endMethod
```

Semaphore.down in Blitz

```
method Down()  
    var oldIntStat: int  
    oldIntStat = SetInterruptsTo (DISABLED)  
    if count == 0x80000000  
        FatalError ("Semaphore count underflowed during 'Wait' operation")  
    EndIf  
    count = count - 1  
    if count < 0 waitingThreads.AddToEnd (currentThread)  
        currentThread.Sleep ()  
    endIf  
    oldIntStat = SetInterruptsTo (oldIntStat)  
endMethod
```

Semaphore.down in Blitz

```
method Down()  
    var oldIntStat: int  
    oldIntStat = SetInterruptsTo (DISABLED)  
    if count == 0x80000000  
        FatalError ("Semaphore count underflowed during 'Wait' operation")  
    EndIf  
    count = count - 1  
    if count < 0 waitingThreads.AddToEnd (currentThread)  
        currentThread.Sleep ()  
    endIf  
    oldIntStat = SetInterruptsTo (oldIntStat)  
endMethod
```

What is `currentThread.Sleep ()` ?

- If `sleep` stops a thread from executing, how, where, and when does it return?
 - *which thread enables interrupts following `sleep`?*
 - *the thread that called `sleep` shouldn't return until another thread has called `signal` !*
 - *... but how does that other thread get to run?*
 - *... where exactly does the `thread switch` occur?*
- Trace down through the Blitz code until you find a call to `switch()`
 - *Switch is called in one thread but returns in another!*
 - *See where registers are saved and restored*

Study The Blitz Code

- Thread.c
 - *Thread.Sleep ()*
 - *Run (nextThread)*
- Switch.s
 - *Switch (prevThread, nextThread)*

Blitz Code For Semaphore.up

```
method Up ()
    var oldIntStat: int
    t: ptr to Thread
    oldIntStat = SetInterruptsTo (DISABLED)
    if count == 0x7fffffff
        FatalError ("Semaphore count overflowed during 'Signal' operation")
    endIf
    count = count + 1
    if count <= 0
        t = waitingThreads.Remove ()
        t.status = READY
        readyList.AddToEnd (t)
    endIf
    oldIntStat = SetInterruptsTo (oldIntStat)
endMethod
```

Using Atomic Instructions

- Implementing semaphores with interrupt disabling only works on uni-processors
 - *What should we do on a multiprocessor?*
- Special (hardware) atomic instructions for synchronization
 - *test and set lock (TSL)*
 - *compare and swap (CAS)*
- Semaphore can be built using atomic instructions
 1. *build mutex locks from atomic instructions*
 2. *build semaphores from mutex locks*

پیاده‌سازی سمافور با
میوتکس؟!؟

How about this solution?

```
var cnt: int = 0          -- Signal count
var m1: Mutex = unlocked  -- Protects access to "cnt"
    m2: Mutex = locked    -- Locked when waiting
```

Down () :

```
    Lock (m1)
    cnt = cnt - 1
    if cnt < 0
        Lock (m2)
        Unlock (m1)
    else
        Unlock (m1)
    endIf
```

Up () :

```
    Lock (m1)
    cnt = cnt + 1
    if cnt <= 0
        Unlock (m2)
    endIf
    Unlock (m1)
```

How about this solution?

```
var cnt: int = 0           -- Signal count
var m1: Mutex = unlocked  -- Protects access to "cnt"
    m2: Mutex = locked    -- Locked when waiting
```

Down () :

```
Lock (m1)
cnt = cnt - 1
if cnt < 0
    Lock (m2)
    Unlock (m1)
else
    Unlock (m1)
endIf
```

Up () :

```
Lock (m1)
cnt = cnt + 1
if cnt <= 0
    Unlock (m2)
endIf
Unlock (m1)
```

Contains a
Deadlock!

How about this solution then?

```
var cnt: int = 0          -- Signal count
var m1: Mutex = unlocked  -- Protects access to "cnt"
    m2: Mutex = locked    -- Locked when waiting
```

Down () :

```
    Lock (m1)
    cnt = cnt - 1
    if cnt < 0
        → Unlock (m1)
        Lock (m2)
    else
        Unlock (m1)
    endIf
```

Up () :

```
    Lock (m1)
    cnt = cnt + 1
    if cnt <= 0
        Unlock (m2)
    endIf
    Unlock (m1)
```

How about this solution then?

```
var cnt: int = 0          -- Signal count
var m1: Mutex = unlocked  -- Protects access to "cnt"
    m2: Mutex = locked    -- Locked when waiting
```

Down () :

```
Lock (m1)
cnt = cnt - 1
if cnt < 0
    Unlock (m1)
    Lock (m2)
else
    Unlock (m1)
endif
```

Up () :

```
Lock (m1)
cnt = cnt + 1
if cnt <= 0
    Unlock (m2)
endif
Unlock (m1)
```

Another solution?

```
var cnt: int = 0          -- Signal count
var m1: Mutex = unlocked -- Protects access to "cnt"
    m2: Mutex = locked    -- Locked when waiting
```

Down () :

```
    Lock (m1)
    cnt = cnt - 1
    if cnt < 0
        → Unlock (m1)
        → Lock (m2)
    → endIf
    → Unlock (m1)
```

Up () :

```
    Lock (m1)
    cnt = cnt + 1
    if cnt <= 0
        Unlock (m2)
    else
        Unlock (m1)
    endIf
```

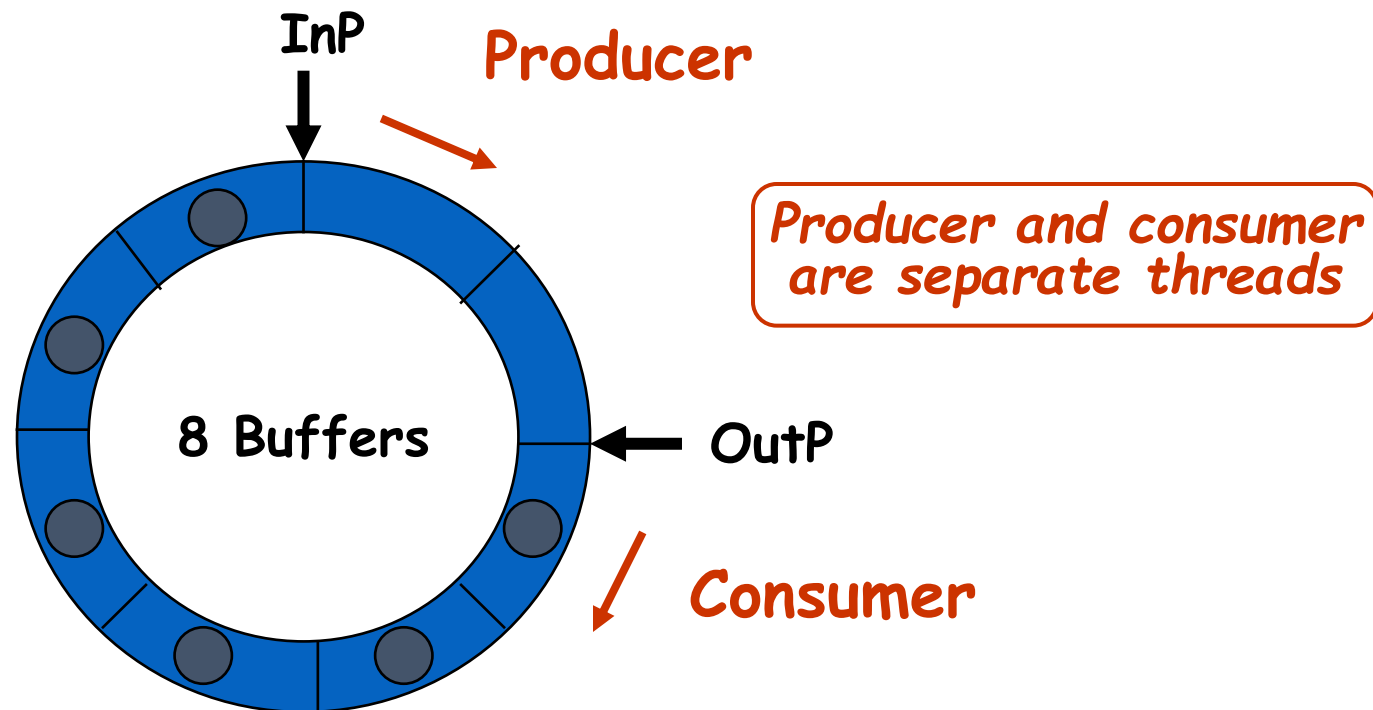
مسئله‌ی کلاسیک همزمانی

Classical Synchronization Problems

- Producer Consumer (bounded buffer)
- Dining philosophers
- Sleeping barber
- Readers and writers

Producer Consumer Problem


- Also known as the bounded buffer problem



Is This a Valid Solution?

```
thread producer {  
    while(1){  
        // Produce char c  
        while (count==n) {  
            no_op  
        }  
        buf[InP] = c  
        InP = InP + 1 mod n  
        count++  
    }  
}
```

```
thread consumer {  
    while(1){  
        while (count==0) {  
            no_op  
        }  
        c = buf[OutP]  
        OutP = OutP + 1 mod n  
        count--  
        // Consume char  
    }  
}
```



Global variables:

```
char buf[n]  
int InP = 0    // place to add  
int OutP = 0   // place to get  
int count
```

Does This Solution Work?

Global variables

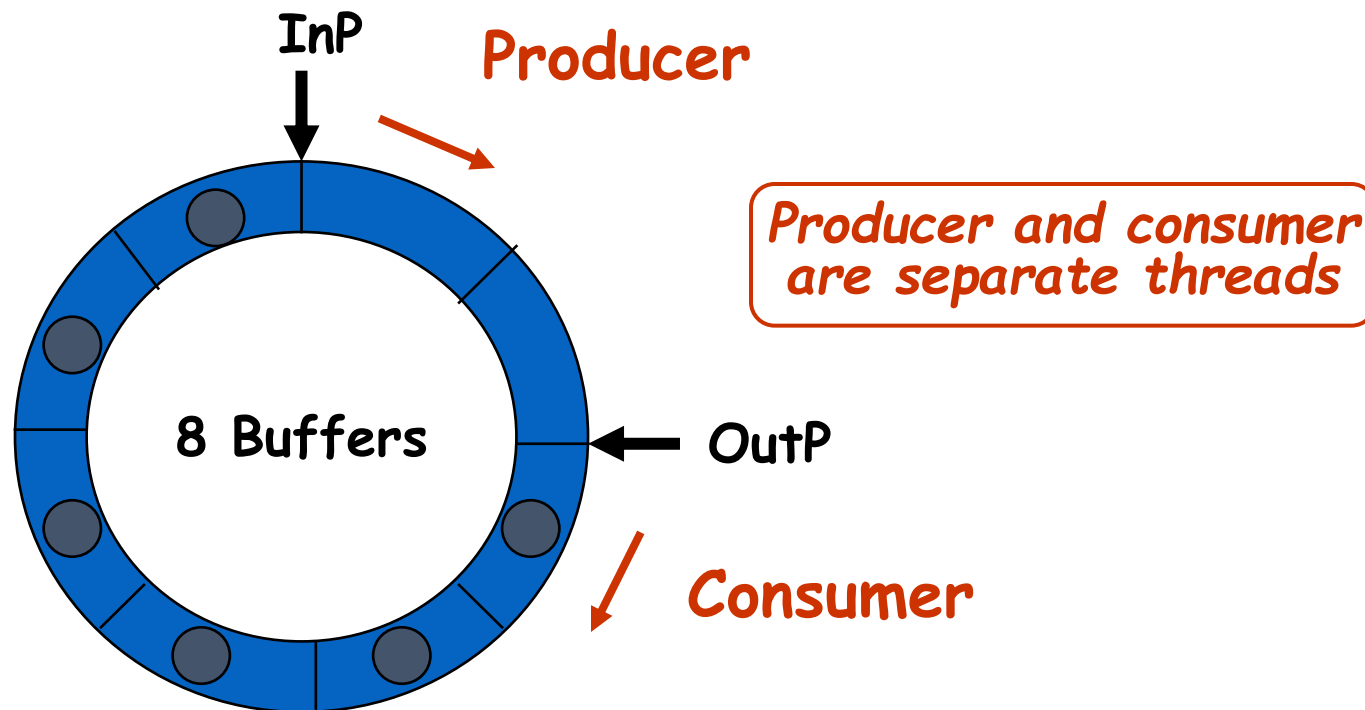
```
semaphore full_buffs = 0;  
semaphore empty_buffs = n;  
char buff[n];  
int InP, OutP;
```

```
0 thread producer {  
1     while(1){  
2         // Produce char c...  
3         down(empty_buffs)  
4         buf[InP] = c  
5         InP = InP + 1 mod n  
6         up(full_buffs)  
7     }  
8 }
```

```
0 thread consumer {  
1     while(1){  
2         down(full_buffs)  
3         c = buf[OutP]  
4         OutP = OutP + 1 mod n  
5         up(empty_buffs)  
6         // Consume char...  
7     }  
8 }
```

Producer Consumer Problem

- What is the shared state in the last solution?
- Does it apply mutual exclusion? If so, how?




Problems With Solution

- What if we have multiple producers and multiple consumers?
 - *Producer-specific and consumer-specific data becomes shared*
 - *We need to define and protect critical sections*
- *You 'll do this in the next parts of the current Blitz project, using the mutex locks you built!*

Dining Philosophers Problem

- Five philosophers sit at a table
- One chopstick between each philosopher
(need two to eat)



Each philosopher is modeled with a thread

```
while (TRUE) {  
    Think();  
    Grab first chopstick;  
    Grab second chopstick;  
    Eat();  
    Put down first chopstick;  
    Put down second chopstick;  
}
```

- *Why do they need to synchronize? How should they do it?*

Is This a Valid Solution?

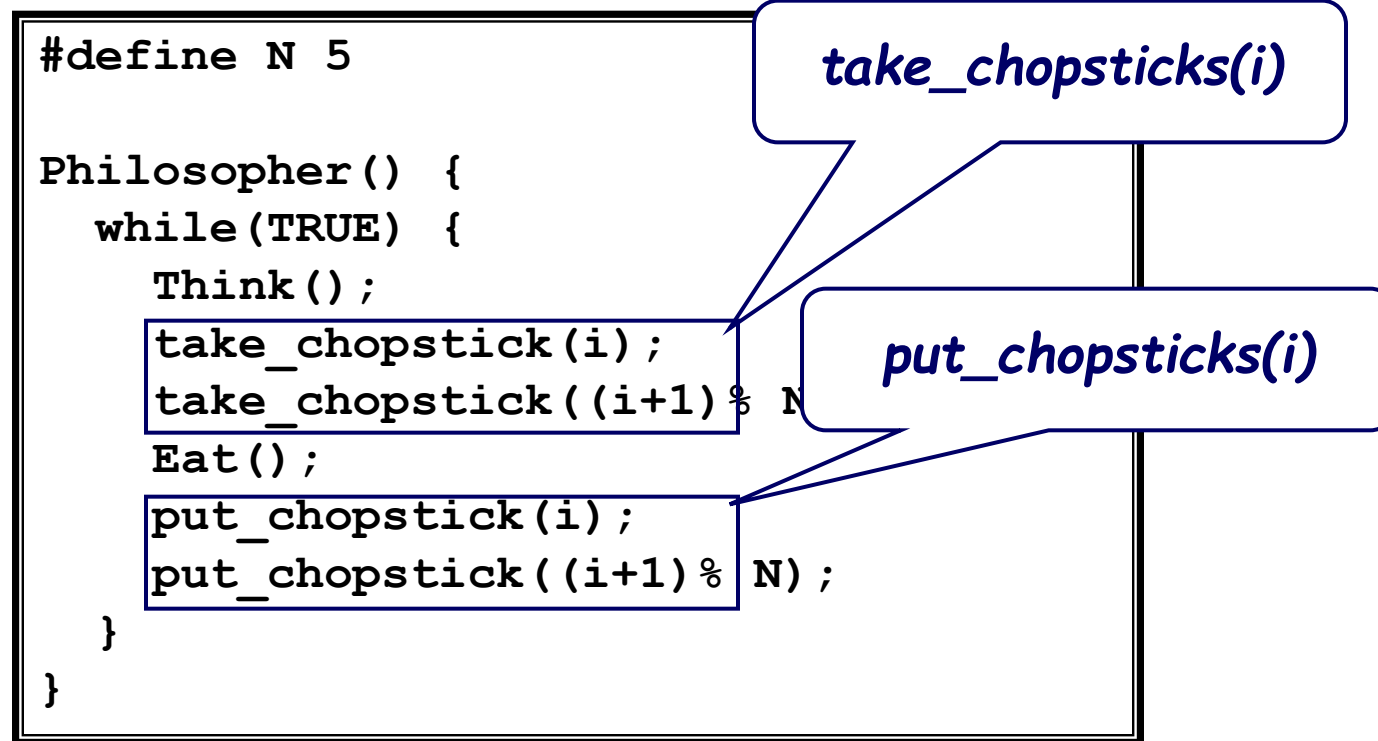
```
#define N 5

Philosopher() {
    while(TRUE) {
        Think();
        take_chopstick(i);
        take_chopstick((i+1) % N);
        Eat();
        put_chopstick(i);
        put_chopstick((i+1) % N);
    }
}
```

Problems

- Potential for deadlock !

Working Towards a Solution



Working Towards a Solution

```
#define N 5

Philosopher() {
    while(TRUE) {
        Think();
        take_chopsticks(i);
        Eat();
        put_chopsticks(i);
    }
}
```

Taking Chopsticks

```
int state[N]
semaphore mutex = 1
semaphore sem[i]
```

```
take_chopsticks(int i) {
    down(mutex);
    state[i] = HUNGRY;
    test(i);
    up(mutex);
    down(sem[i]);
}
```

```
// only called with mutex set!

test(int i) {
    if (state[i] == HUNGRY &&
        state[LEFT] != EATING &&
        state[RIGHT] != EATING) {
        state[i] = EATING;
        up(sem[i]);
    }
}
```

Putting Down Chopsticks

```
int state[N]
semaphore mutex = 1
semaphore sem[i]
```

```
put_chopsticks(int i) {
    down(mutex);
    state[i] = THINKING;
    test(LEFT);
    test(RIGHT);
    up(mutex);
}
```

```
// only called with mutex set!

test(int i) {
    if (state[i] == HUNGRY &&
        state[LEFT] != EATING &&
        state[RIGHT] != EATING) {
        state[i] = EATING;
        up(sem[i]);
    }
}
```

Dining Philosophers

- Is the previous solution correct?
- What does it mean for it to be correct?
- Is there an easier way?

Recap

- What is a race condition?
- How can we protect against race conditions?
- Can locks be implemented simply by reading and writing to a binary variable in memory?
- How can a kernel make synchronization-related system calls atomic on a uniprocessor?
 - *Why wouldn't this work on a multiprocessor?*
- Why is it better to block rather than spin on a uniprocessor?
- Why is it sometimes better to spin rather than block on a multiprocessor?

Recap

- When faced with a concurrent programming problem, what strategy would you follow in designing a solution?
- What does all of this have to do with Operating Systems?