

R

بسم الله الرحمن الرحيم

سیستم عامل

جلسه دوازدهم – مدیریت حافظه (مقدمه)

جلسه‌ی گذشته

الگوریتم‌های زمان‌بندی پردازنده‌ها

Scheduling Policies

- First-Come, First Served (FIFO)
- Shortest Job First (non-preemptive)
- Shortest Job First (with preemption)
- Round-Robin Scheduling
- Priority Scheduling
- Real-Time Scheduling

جلسه‌ی جدید

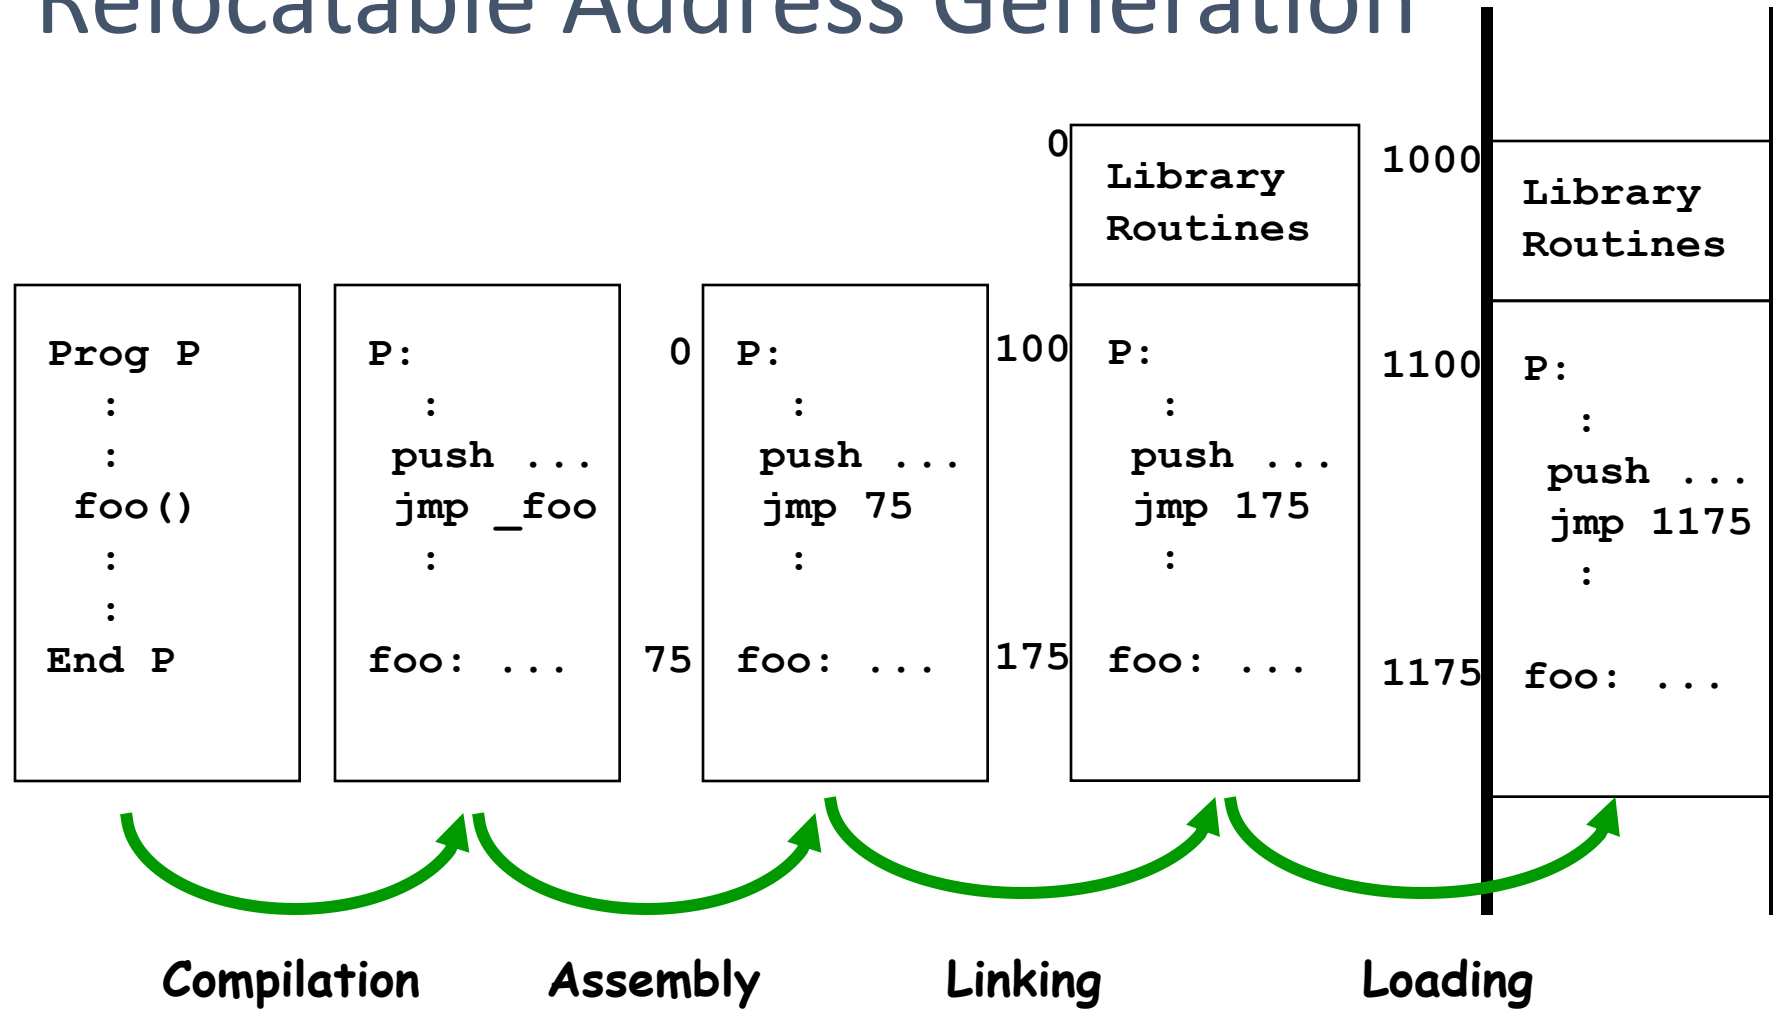
Memory Management

- Memory – a linear array of bytes
 - *Holds O.S. and programs (processes)*
 - *Each cell (byte) is named by a unique memory address*
- Recall, processes are defined by an *address space*, consisting of text, data, and stack regions
- Process execution
 - *CPU fetches instructions from the text region according to the value of the program counter (PC)*
 - *Each instruction may request additional operands from the data or stack region*

Addressing Memory

- Cannot know ahead of time where in memory a program will be loaded!
- Compiler produces code containing embedded addresses
 - *these addresses can't be absolute (physical addresses)*
- Linker combines pieces of the program
 - *Assumes the program will be loaded at address 0*
- We need to **bind** the compiler/linker generated addresses to the actual memory locations

Relocatable Address Generation



Address Binding

- Address binding
 - *fixing a physical address to the logical address of a process' address space*

Address Binding

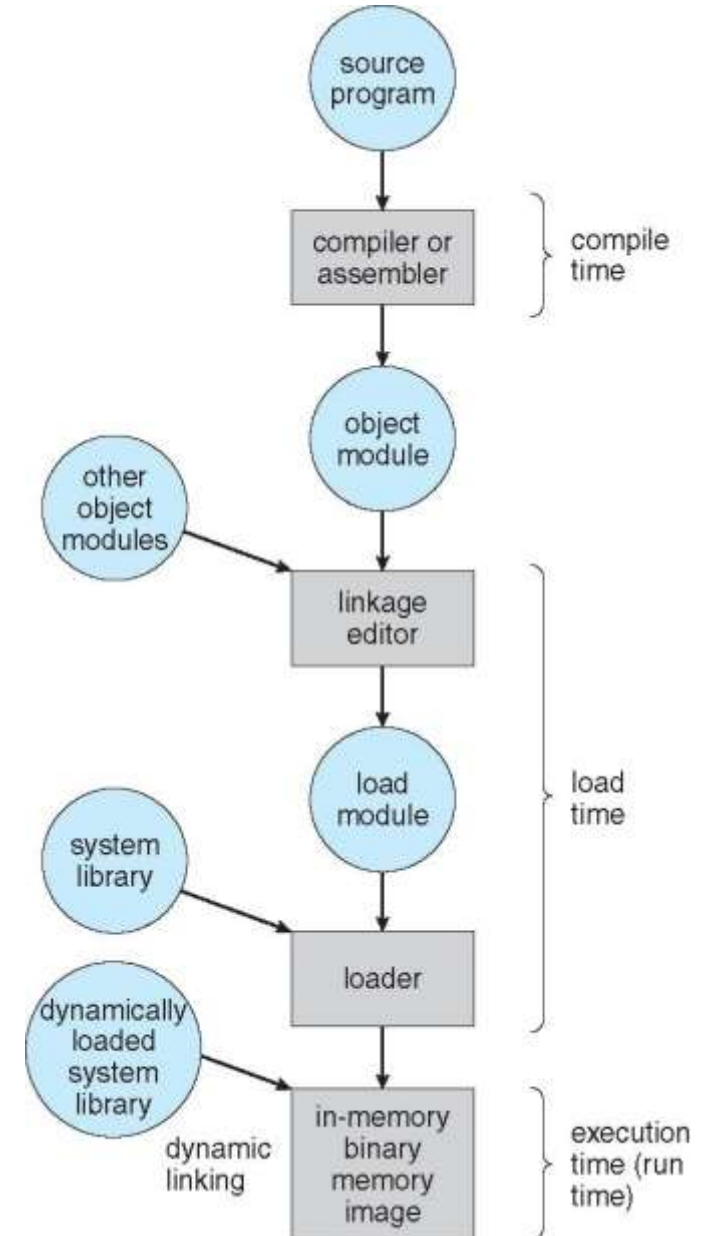
Address binding of instructions and data to memory addresses can happen at three different stages:

- **Compile time:** If memory location known a priori, **absolute code** can be generated; must recompile code if starting location changes
- **Load time:** Must generate **relocatable code** if memory location is not known at compile time
- **Execution time:** Binding delayed until run time if the process can be moved during its execution from one memory segment to another
 - Need hardware support for address maps (e.g., base and limit registers)

Base and Limit Registers

- Simple runtime relocation scheme
 - *Use 2 registers to describe a partition*
- For every address generated, at runtime...
 - *Compare to the **limit** register (& abort if larger)*
 - *Add to the **base** register to give **physical** memory address*

Multistep Processing of a User Program

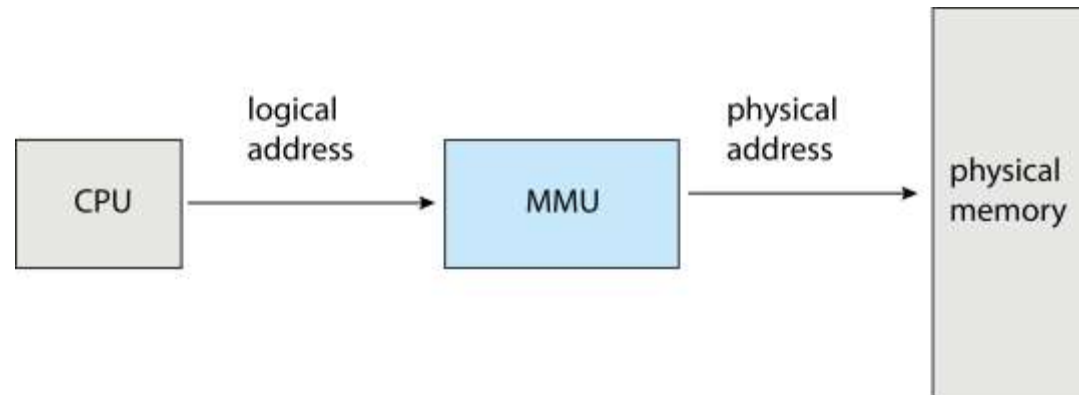


Logical vs. Physical Address Space

- The concept of a logical address space that is bound to a separate **physical address space** is central to proper memory management
 - **Logical address** – *generated by the CPU; also referred to as **virtual address***
 - **Physical address** – *address seen by the memory unit*
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes; logical (virtual) and physical addresses differ in execution-time address-binding scheme
- **Logical address space** is the set of all logical addresses generated by a program
- **Physical address space** is the set of all physical addresses generated by a program

Memory-Management Unit (MMU)

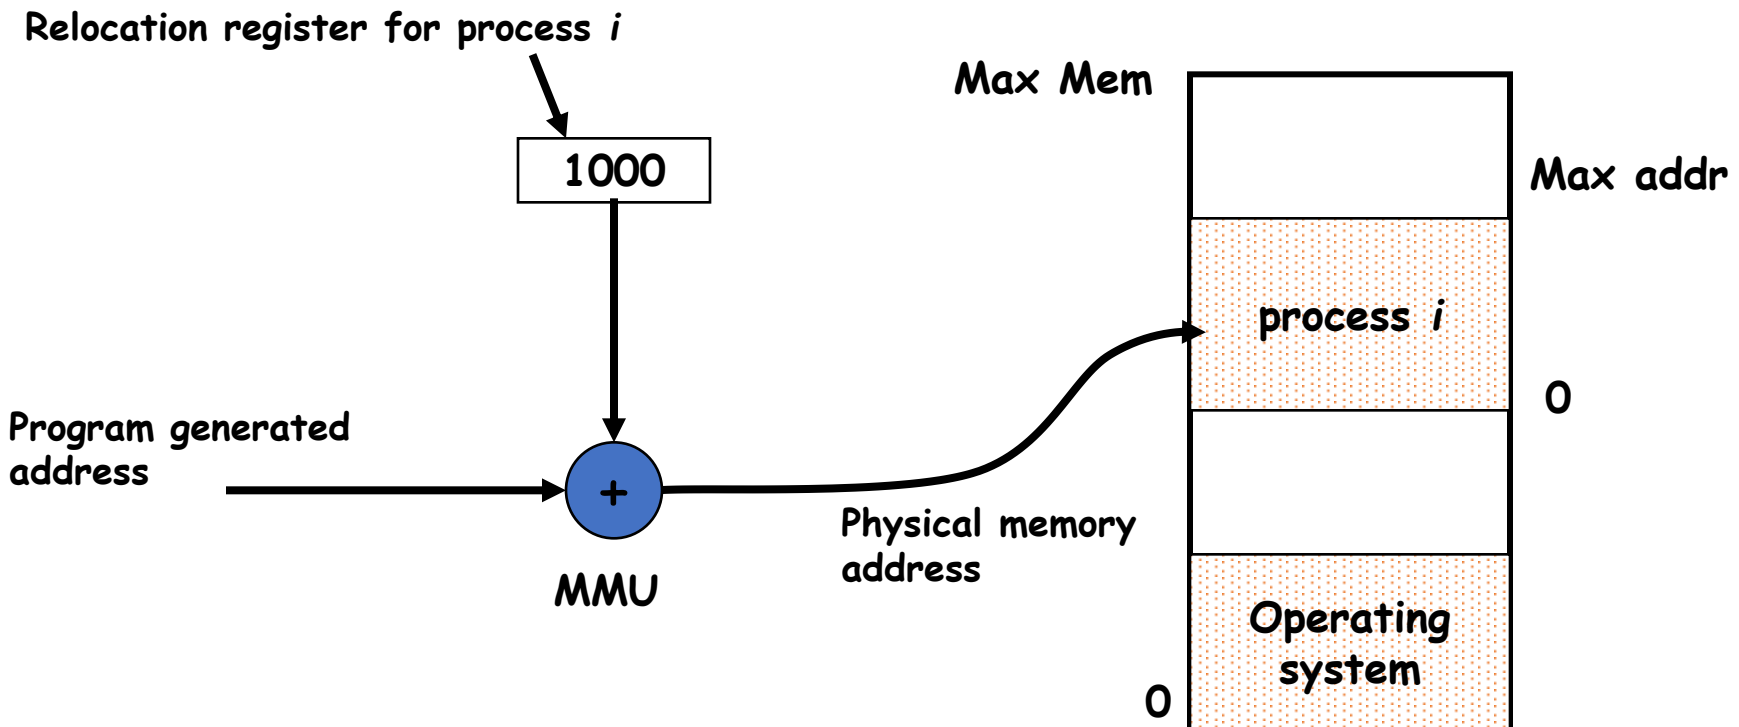
- Hardware device that at run time maps virtual to physical address



- Many methods possible, covered in the rest of this lecture

Dynamic Relocation

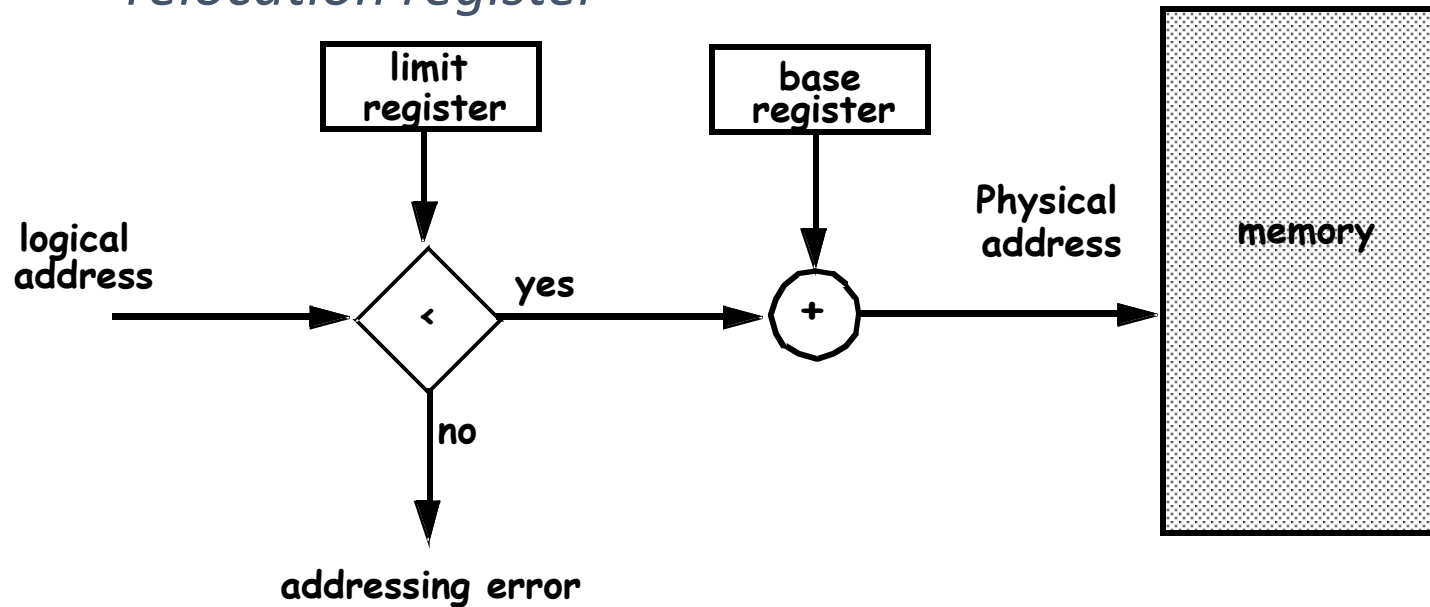
- **Memory Management Unit (MMU)**
- - Dynamically converts logical to physical address
- - Contains base address register for running process



Protection

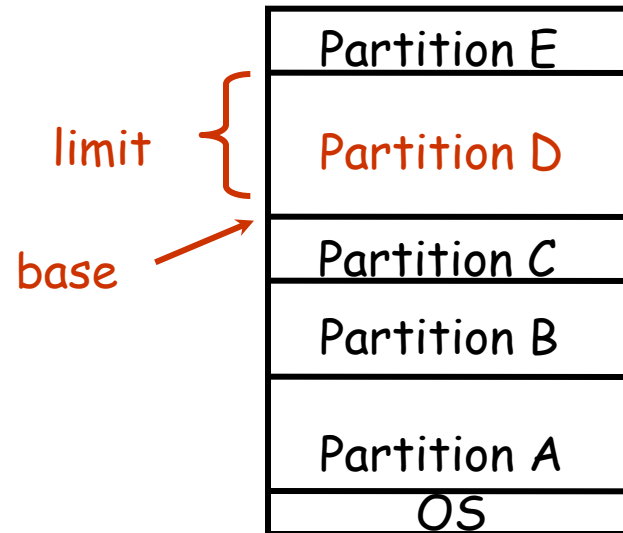
■ Memory protection

- *Base* register gives starting address for process
- *Limit* register limits the offset accessible from the relocation register



Multiprogramming

- Multiprogramming: a separate partition per process
- What happens on a context switch?
 - Store process *base* and *limit* register values
 - Load new values into *base* and *limit* registers




Swapping

- When a program is running...
 - *The entire program must be in memory*
 - *Each program is put into a single **partition***
- When the program is not running...
 - *May remain resident in memory*
 - *May get “swapped” out to disk*
- Over time...
 - *Programs come into memory when they get swapped in*
 - *Programs leave memory when they get swapped out*

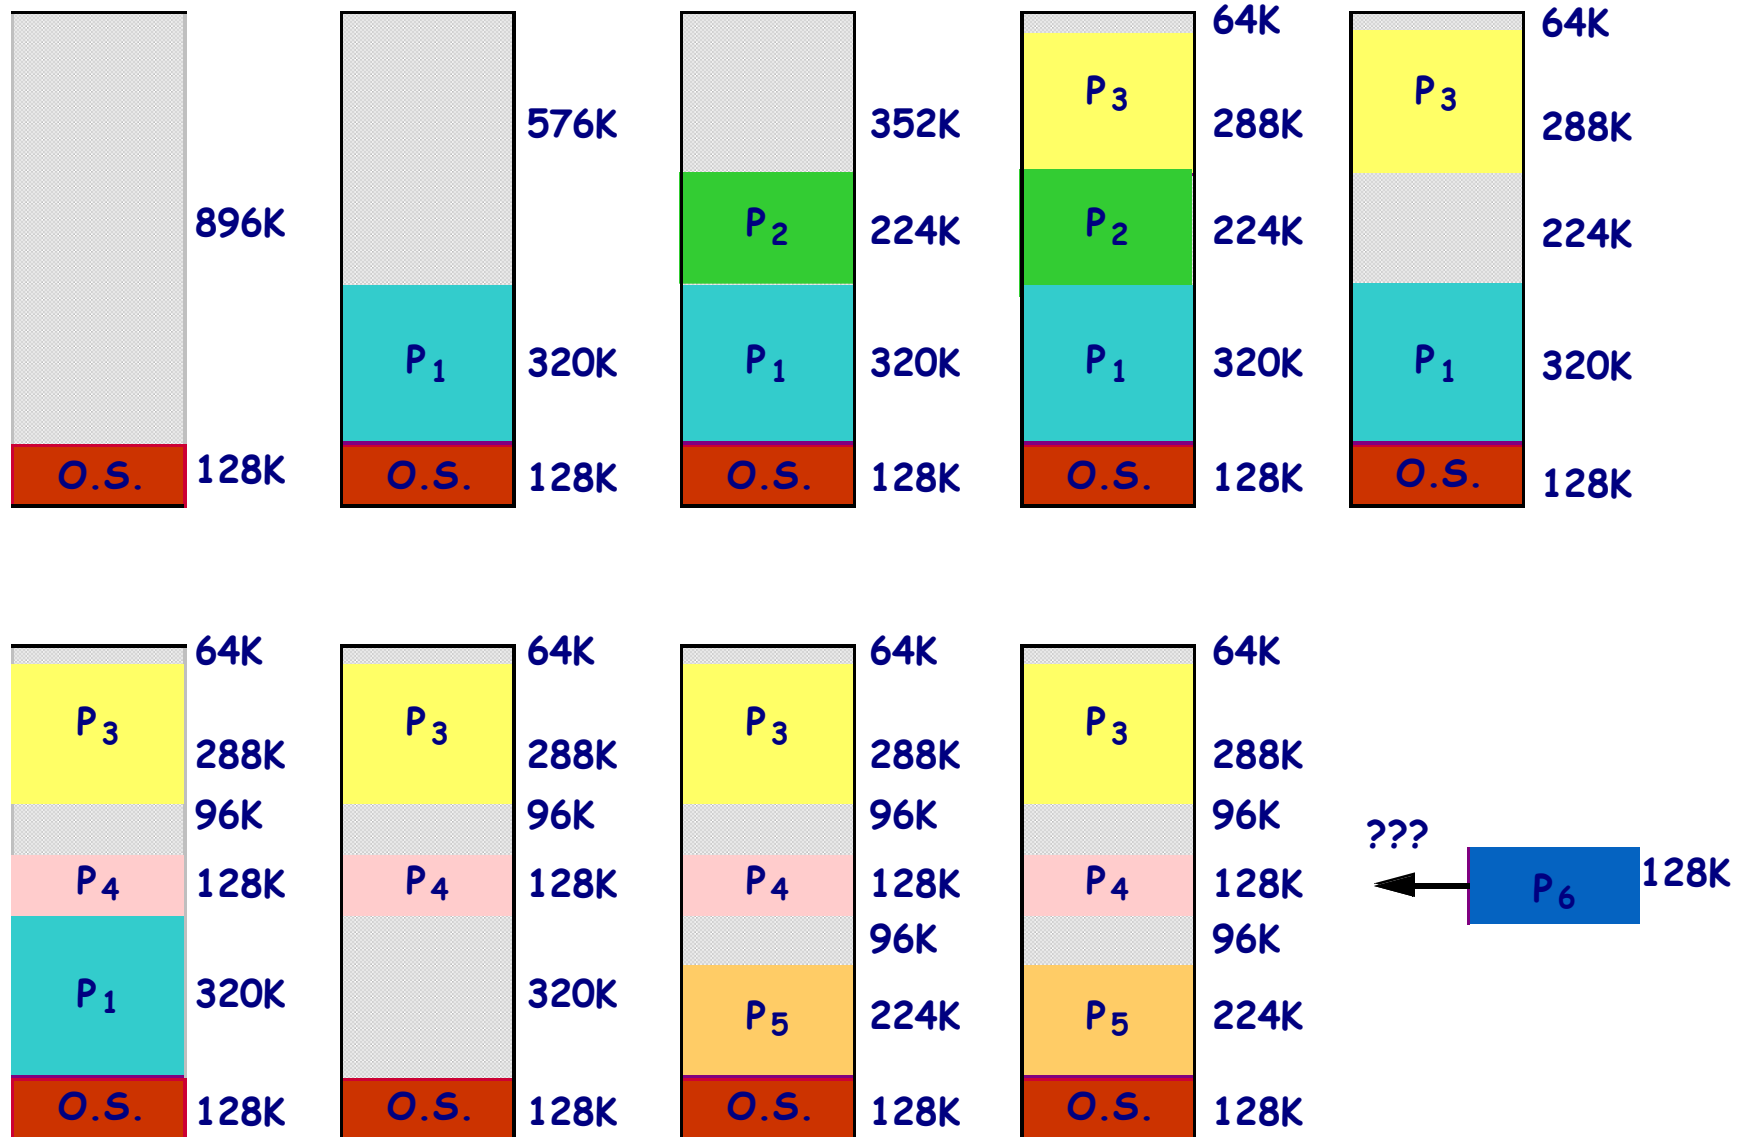
Swapping

- Benefits of swapping:
 - *Allows multiple programs to be run concurrently*
 - *... more than will fit in memory at once*




FRAGMENTATION





Dealing With Fragmentation

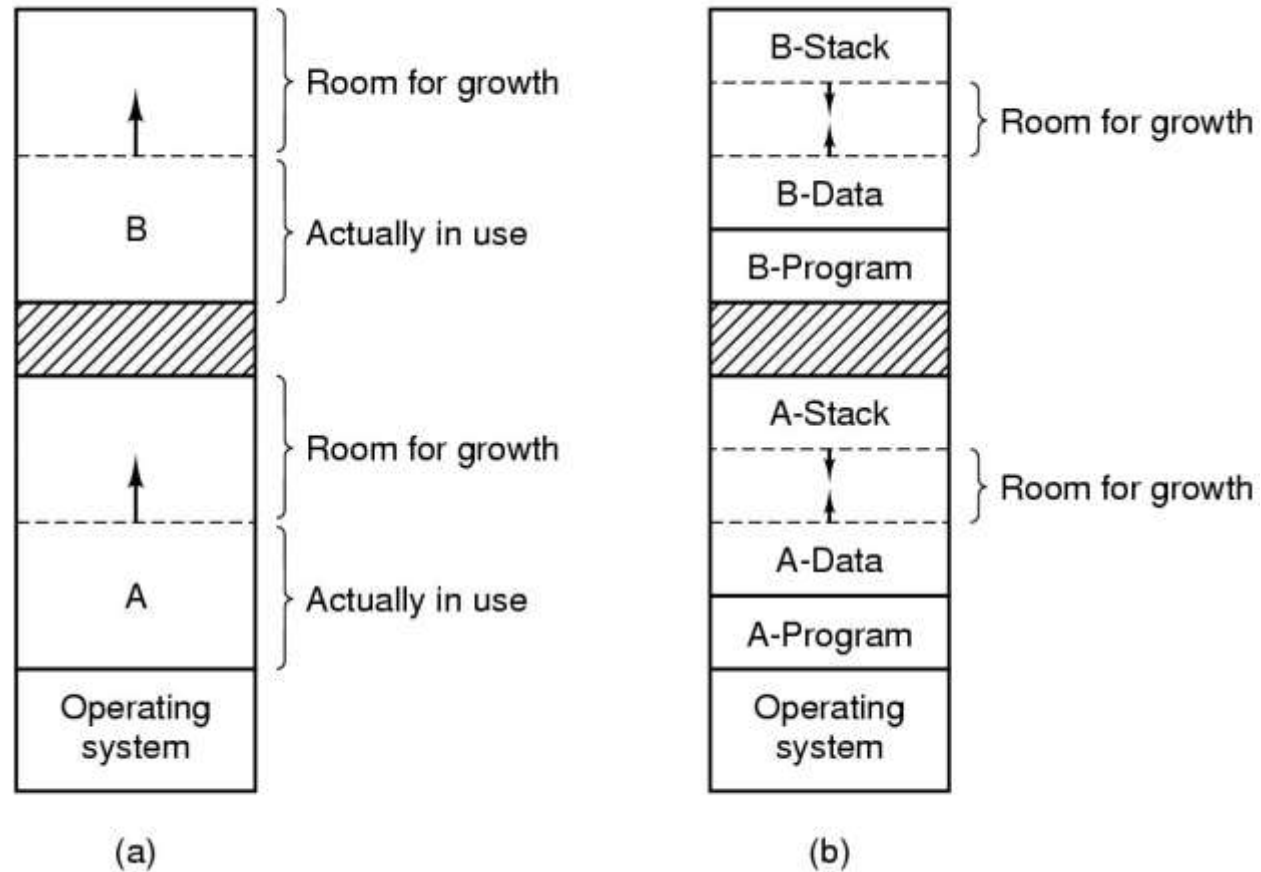
- *Compaction* – from time to time shift processes around to collect all free space into one contiguous block
 - *Memory to memory copying overhead*
 - *Memory to disk to memory for compaction via swapping!*



How Big Should Partitions Be?

- Programs may want to grow during execution
 - *More room for stack, heap allocation, etc*
- Problem:
 - *If the partition is too small, programs must be moved*
 - *Requires copying overhead*
 - *Why not make the partitions a little larger than necessary to accommodate “some” cheap growth?*

Allocating Extra Space Within



Fragmentation

- **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous
- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used

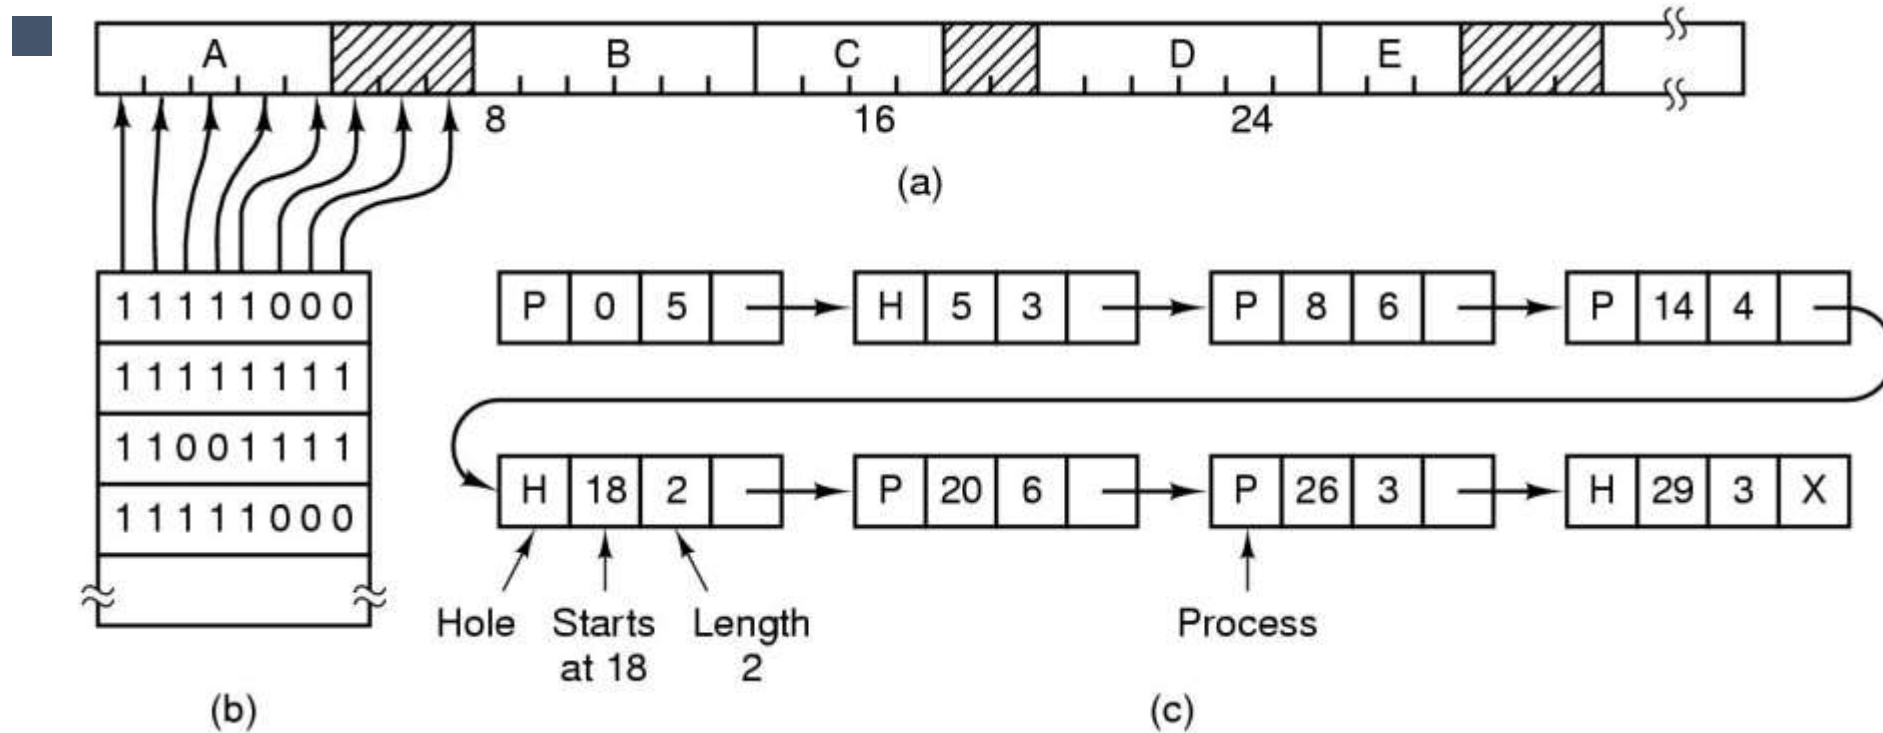
Management Data Structures

- Each chunk of memory is either
 - *Used by some process or unused (free)*
- Operations
 - *Allocate* a chunk of unused memory big enough to hold a new process
 - *Free* a chunk of memory by returning it to the *free pool* after a process terminates or is swapped out

Management With Bit Maps

- Problem - how to keep track of used and unused memory?
- **Technique 1** - Bit Maps
 - *A long bit string*
 - *One bit for every chunk of memory*
 - 1 = in use
 - 0 = free
 - *Size of allocation unit influences space required*
 - Example: unit size = 32 bits
 - *overhead for bit map: $1/33 = 3\%$*
 - Example: unit size = 4Kbytes
 - *overhead for bit map: $1/32,769$*

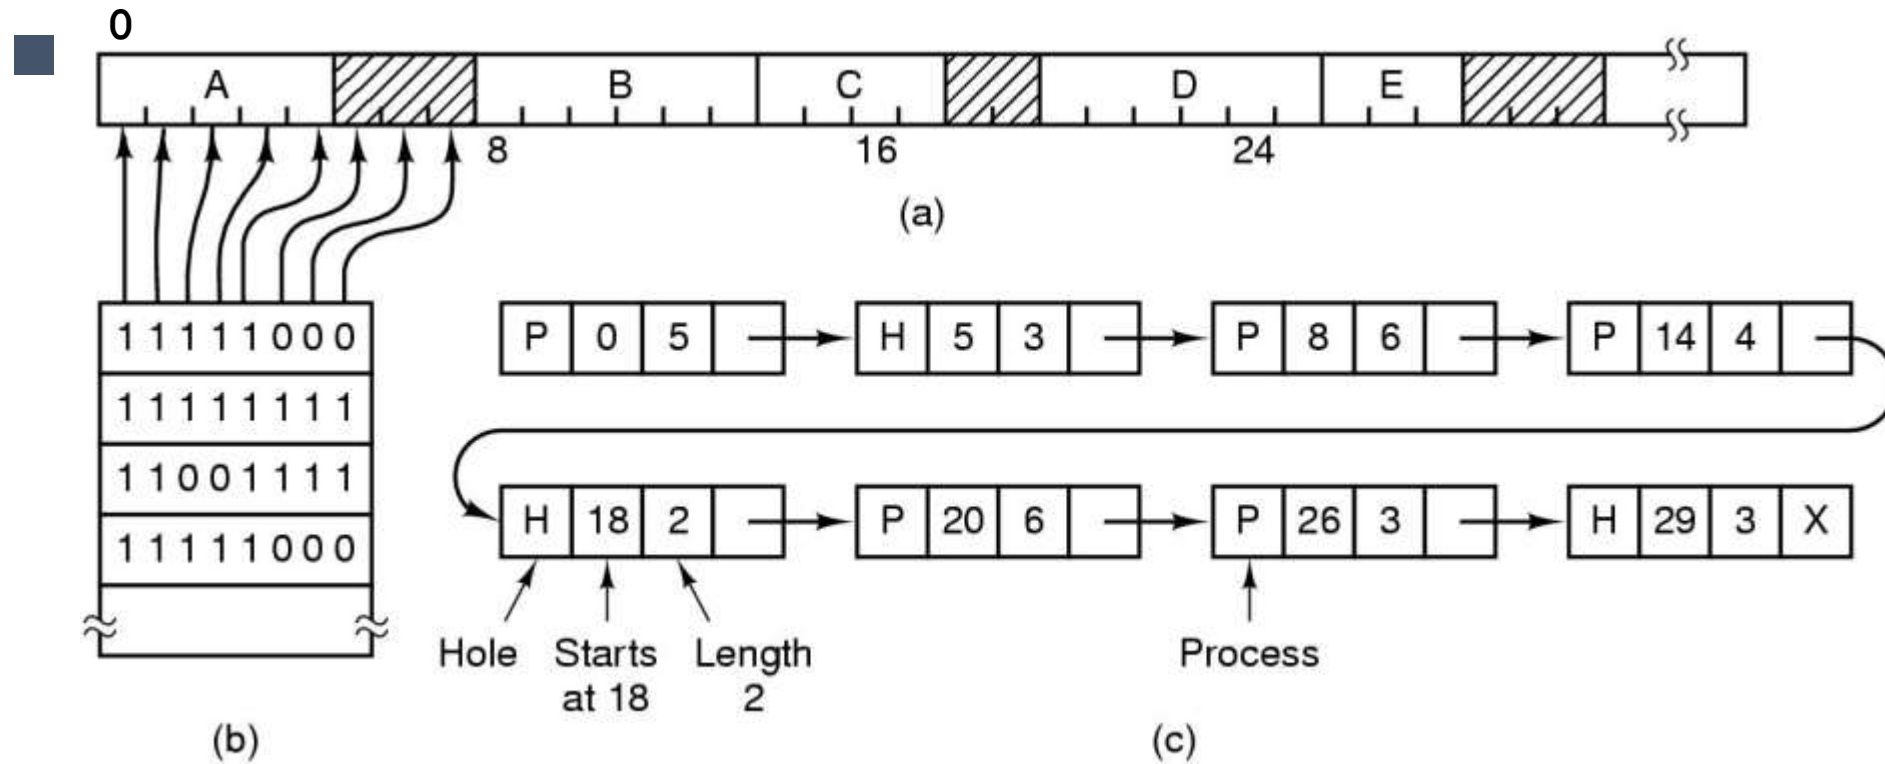
Management With Bit Maps



Management With Linked Lists

- **Technique 2 - Linked List**
- Keep a list of elements
- Each element describes one unit of memory
 - *Free / in-use Bit (“P=process, H=hole”)*
 - *Starting address*
 - *Length*
 - *Pointer to next element*

Management With Linked Lists



Management With Linked Lists

Searching the list for space for a new process

- First Fit
- Next Fit
 - *Start from current location in the list*
- Best Fit
 - *Find the smallest hole that will work*
 - *Tends to create lots of really small holes*
- Worst Fit
 - *Find the largest hole*
 - *Remainder will be big*
- Quick Fit
 - *Keep separate lists for common sizes*

Fragmentation Revisited

- Memory is divided into partitions
- Each partition has a different size
- Processes are allocated space and later freed
- After a while memory will be full of small holes!
 - *No free space large enough for a new process even though there is enough free memory in total*
- If we allow free space within a partition we have fragmentation
 - *External fragmentation* = unused space between partitions
 - *Internal fragmentation* = unused space within partitions

Solutions to Fragmentation

- Compaction requires high copying overhead
- Why not allocate memory in non-contiguous equal fixed size units?
 - *No external fragmentation!*
 - *Internal fragmentation < 1 unit per process*
- How big should the units be?
 - *The smaller the better for internal fragmentation*
 - *The larger the better for management overhead*
- The key challenge for this approach

How can we do secure dynamic address translation?