

R

بسم الله الرحمن الرحيم

سیستم عامل

جلسه شانزدهم – الگوریتم‌های جایگزینی صفحه (۲)

جلسه‌ی گذشته

Page Replacement

- Assume a normal page table (e.g., BLITZ)
- User-program is executing
- A *PageInvalidFault* occurs!
 - *The page needed is not in memory*
- Select some frame and remove the page in it
 - *If it has been modified, it must be written back to disk*
 - the “dirty” bit in its page table entry tells us if this is necessary
- Figure out which page was needed from the faulting addr
- Read the needed page into this frame
- Restart the interrupted process by retrying the same instruction

Page Replacement Algorithms

Which frame to replace? ■

Algorithms? ■

- ورودی: لیست صفحه‌هایی که در حافظه هستند و درخواست‌های مربوط به گرفتن صفحه
- خروجی: کدام صفحه را از حافظه خارج کنیم
- هدف: کاهش تعداد *page fault* ها

The optimal page replacement algorithm

- **Idea:** Given all the data, how to find the optimal page replacement?

Time	0	1	2	3	4	5	6	7	8	9	10
Requests		c	a	d	b	e	b	a	b	c	d

Page	0	a
Frames	1	b
	2	c
	3	d

Page faults

LFD

- Replace the page that will not be needed for the longest
- Example:

Time		0	1	2	3	4	5	6	7	8	9	10
Requests			c	a	d	b	e	b	a	b	c	d
Page	0	a										
Frames	1	b	a	a	a	a						
	2	c	b	b	b	b						
	3	d	c	c	c	c						
			d	d	d	d						
Page faults							X					

LFD

- Replace the page that will not be needed for the longest
- Example:


Time		0	1	2	3	4	5	6	7	8	9	10
Requests			c	a	d	b	e	b	a	b	c	d
<hr/>												
Page	0	a	a	a	a	a	a	a	a	a	a	
Frames	1	b	b	b	b	b	b	b	b	b	b	
	2	c	c	c	c	c	c	c	c	c	c	
	3	d	d	d	d	d	e	e	e	e	e	
Page faults							X				X	

Is LFD Optimal?

- **OPT:** Optimum with longest prefix equal to LFD
 - *First non-equal position: $i+1$*
- Case 1) $i+1$ is not a page fault.
 - *Is it possible?*
- Case 2) $i+1$ is page fault
 - *LFD evicts p*
 - *OPT evicts p'*

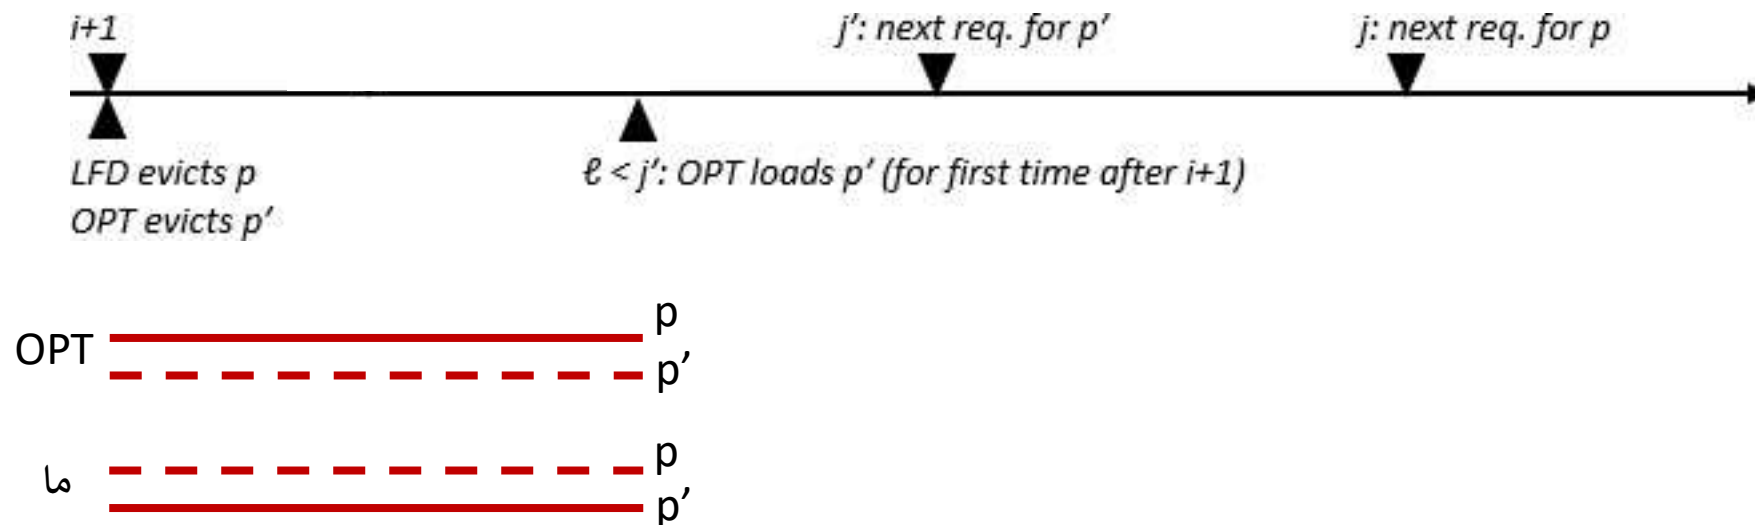
Is LFD Optimal?

- **OPT:** Optimum with longest prefix equal to LFD
 - *First non-equal position: $i+1$*
- Case 2) $i+1$ is page fault
 - *LFD evicts p*
 - *OPT evicts p'*




Is LFD Optimal?

- Proof (by contradiction):
 - **OPT**: Optimum with longest prefix equal to LFD Case
 - 2) $i+1$ is page fault
 - Case 2-A) OPT keeps p until l
 - Case 2-B) OPT evicts p at $l' < l$



Is LFD Optimal?

- Proof (by contradiction):
 - **OPT**: Optimum with longest prefix equal to LFD Case
 - 2) $i+1$ is page fault
 - Case 2-A) OPT keeps p until l
 - Case 2-B) OPT evicts p at $l' < l$



$$\text{LFD} = \text{OPT}$$

Optimal Page Replacement

■ Idea:

- *Select the page that will not be needed for the longest time*

■ Problem:

- *Can't know the future of a program*
- *Can't know when a given page will be needed next*
- *The optimal algorithm is unrealizable*

Optimal Page Replacement

■ However:

- *We can use it as a control case for simulation studies*
 - Run the program once
 - Generate a log of all memory references
 - Do we need all of them?
 - Use the log to simulate various page replacement algorithms
 - Can compare others to “optimal” algorithm

جلسه‌ی جدید

تعداد PAGE FAULT ها

چه پارامترهایی روی تعداد پیج فالت‌ها موثر هستند؟

■ الگوریتم جایگزینی صفحه

■ اندازه‌ی تعداد frame‌های قابل قرارگیری در حافظه

- این ممکنه *global* برای کل پردازنده‌ها مشترک باشه

- ممکنه *local* به ازای هر پردازنده باشه.

■ ??




Figure 10.11 Graph of page faults versus number of frames.

FIFO ALGORITHM

FIFO Algorithm

- Always replace the oldest page ...
 - *Replace the page that has been in memory for the longest time*

- Replace the page that was first brought into memory
- Example: Memory system with 4 frames:

Time	0	1	2	3	4	5	6	7	8	9	10
Requests		c	a	d	b	e	b	a	b	c	a
Page	0	a									
Frames	1	b		a	a	a					
	2	c				b					
	3	d	c	c	c	c					
				d	d						
Page faults						x					

FIFO Algorithm

- Replace the page that was first brought into memory
- Example: Memory system with 4 frames:

Time	0	1	2	3	4	5	6	7	8	9	10
Requests		c	a	d	b	e	b	a	b	c	a
<hr/>											
Page 0	a		a	a	a	a	a	a	a		
Frames 1	b				b	b	b	b	b		
2	c										
3	d	c	c	c	c	e	e	e	e		
				d	d	d	d	d	d		
Page faults						x				x	

FIFO Algorithm

- Replace the page that was first brought into memory
- Example: Memory system with 4 frames:

Time	0	1	2	3	4	5	6	7	8	9	10
Requests		c	a	d	b	e	b	a	b	c	a
<hr/>											
Page 0	a		a	a	a	a	a	a	a	c	
Frames 1	b				b	b	b	b	b	b	
2	c										
3	d	c	c	c	c	e	e	e	e	e	
				d	d	d	d	d	d	d	
Page faults						x				x	x

FIFO Algorithm

- Replace the page that was first brought into memory
- Example: Memory system with 4 frames:

Time	0	1	2	3	4	5	6	7	8	9	10
Requests		c	a	d	b	e	b	a	b	c	a
<hr/>											
Page 0	a		a	a	a	a	a	a	a	c	c
Frames 1	b				b	b	b	b	b	b	b
2	c										
3	d	c	c	c	c	e	e	e	e	e	e
				d	d	d	d	d	d	d	a
Page faults						x				x	x

FIFO Algorithm

- Always replace the oldest page.
 - *Replace the page that has been in memory for the longest time*
- Implementation
 - *Maintain a linked list of all pages in memory*
 - *Keep it in order of when they came into memory*
 - *The page at the tail of the list is oldest*
 - *Add new page to head of list*

FIFO Algorithm

■ Disadvantage?


- *The oldest page may be needed again soon*
- *Some page may be important throughout execution*
- *It will get old, but replacing it will cause an immediate page fault*

سوال؟

- در الگوریتم FIFO، اگر تعداد frame ها رو زیاد کنیم، اوضاع تعداد page fault ها چطور می شود؟
 - لزوما بهتر می شود؟!
 - متأسفانه لزوما نه!!

Belady's Anomaly

- Adding more frames can cause more page faults!



NRU ALGORITHM

How Can We Do Better?

- Need an approximation of how likely each frame is to be accessed in the future
 - *If we base this on past behavior we need a way to track past behavior*
 - *Tracking memory accesses requires hardware support to be efficient*

مسئله: یکی از فریم‌هایی که اخیرا استفاده نشده رو بریزیم دور

- چطوری انتخاب کنیم کدوم فریم؟
- علی‌الحساب فرض کنید نمی‌خوایم داده‌ساختار جدیدی نگه داریم...
- مثلا از اطلاعات توی page table استفاده کنیم.
- چه اطلاعاتی داریم؟
- Reference bit
- Dirty bit

Referenced and Dirty Bits

- Each page table entry (and TLB entry!) has a
 - *Referenced bit - set by TLB when page read / written*
 - *Dirty / modified bit - set when page is written*
 - *If TLB entry for this page is valid, it has the most up to date version of these bits for the page*
 - *OS must copy them into the page table entry during fault handling*
- Idea: use the information contained in these bits to drive the page replacement algorithm

Referenced and Dirty Bits

- Some hardware does not have support for the dirty bit
- Instead, memory protection can be used to emulate it
- Idea:
 - *Software sets the protection bits for all pages to “read only”*
 - *When program tries to update the page...*
 - A trap occurs
 - Software sets the Dirty Bit in the page table and clears the ReadOnly bit
 - Resumes execution of the program

Not Recently Used Algorithm

- Uses the Referenced Bit and the Dirty Bit
- Initially, all pages have
 - *Referenced Bit = 0*
 - *Dirty Bit = 0*
- Periodically... (e.g. whenever a timer interrupt occurs)
 - *Clear the Referenced Bit*
 - *Referenced bit now indicates “recent” access*

Not Recently Used Algorithm

- When a page fault occurs...
- Categorize each page...
 - Class 1: Referenced = 0 Dirty = 0
 - Class 2: Referenced = 0 Dirty = 1
 - Class 3: Referenced = 1 Dirty = 0
 - Class 4: Referenced = 1 Dirty = 1
- Choose a victim page from ...
 - class 1 ... why?
- If none, choose a page from ...
 - class 2 ... why?
- If none, choose a page from ...
 - class 3 ... why?
- If none, choose a page from ...
 - class 4 ... why?

مسئله: یکی از فریم‌هایی که اخیرا استفاده نشده رو بریزیم دور

- چطوری انتخاب کنیم کدوم فریم؟
- علی‌الحساب فرض کنید نمی‌خواهیم داده‌ساختار جدیدی نگه داریم...
- مثلا از اطلاعات توی page table استفاده کنیم.
- راه دیگه‌ای هم داریم؟
- یه طوری می‌خواهیم اونی که اخیرا reference نشده رو بندازیم بیرون...


SECOND CHANCE ALGORITHM

Second Chance Algorithm

- An implementation of NRU based on FIFO
- Pages kept in a linked list (oldest at the front)
- Look at the oldest page
 - *If its “referenced bit” is 0...*
 - Select it for replacement
 - *Else*
 - It was used recently; don't want to replace it
 - Clear its “referenced bit”
 - Move it to the end of the list
 - *Repeat*
- What if every page was used in last clock tick?

Implementation of Second Chance

- Maintain a circular list of pages in memory
- Set a bit for the page when a page is referenced
- Search list looking for a victim page that does not have the referenced bit set
 - *If the bit is set, clear it and move on to the next page*
 - *Replaces pages that haven't been referenced for one complete clock revolution*



Enhanced Second Chance algorithm

- همزمان از هر دو ایده‌ی قبلی استفاده کنیم...
- در ۴ گروه براساس reference bit و dirty bit در نظر بگیریم.
- حالا second chance رو براساس گروه‌بندی قبلی انجام بدیم...



LRU

Least Recently Used Algorithm

- A refinement of NRU that orders how recently a page was used
 - *Keep track of when a page is used*
 - *Replace the page that has been used least recently*

Least Recently Used Algorithm

- Replace the page that hasn't been referenced in the longest time

Time		0	1	2	3	4	5	6	7	8	9	10
Requests			c	a	d	b	e	b	a	b	c	d
<hr/>												
Page	0	a	a	a	a	a	a	a	a	a	a	a
Frames	1	b	b	b	b	b	b	b	b	b	b	b
	2	c	c	c	c	e	e	e	e	e	e	d
	3	d	d	d	d	d	d	d	d	d	c	c
<hr/>												
Page faults							x				x	x

Least Recently Used Algorithm

- But how can we implement LRU?

Least Recently Used Algorithm

- But how can we implement LRU?
- Idea #1:
 - *Keep a linked list of all pages*
 - *On every memory reference, Move that page to the front of the list*
 - *The page at the tail of the list is replaced*

Least Recently Used Algorithm

- But how can we implement LRU?
 - ... *without requiring every access to be recorded?*
- Idea #2:
 - *MMU (hardware) maintains a counter*
 - *Incremented on every clock cycle*
 - *Every time a page table entry is used*
 - MMU writes the value to the page table entry
 - This *timestamp* value is the *time-of-last-use*
 - *When a page fault occurs*
 - OS looks through the page table
 - Identifies the entry with the oldest timestamp

Least Recently Used Algorithm

- What if we don't have hardware support for a counter?
- Idea #3:
 - *Maintain a counter in software*
 - *One every timer interrupt...*
 - Increment counter
 - Run through the page table
 - For every entry that has "ReferencedBit" = 1
 - * *Update its timestamp*
 - * *Clear the ReferencedBit*
 - *Approximates LRU*
 - *If several have oldest time, choose one arbitrarily*

NFU

Not Frequently Used Algorithm

- Bases decision of frequency of use rather than recency
- Associate a counter with each page
- On every clock interrupt, the OS looks at each page.
 - *If the reference bit is set increment that page's counter & clear the bit*
- The counter approximates how often the page is used
- For replacement, choose the page with lowest counter

Not Frequently Used Algorithm

■ *Problem:*

- *Some page may be heavily used*
 - Its counter is large
- *The program's behavior changes*
 - Now, this page is not used ever again (or only rarely)
- *This algorithm never forgets!*
 - - *This page will never be chosen for replacement!*
- *We may want to combine frequency and recency somehow*

NFU With Aging

- Associate a counter with each page
- On every clock tick, the OS looks at each page.
 - *Shift the counter right 1 bit (divide its value by 2)*
 - *If the reference bit is set...*
 - Set the most-significant bit
 - Clear the Referenced Bit

– T	$100000 = 32$
– T^1	$010000 = 16$
– T^2	$001000 = 8$
– T^3	$000100 = 4$
– T^4	$100010 = 34$
– T_5	