

R

بسم الله الرحمن الرحيم

سیستم عامل

جلسه هجدهم – الگوریتم‌های جایگزینی صفحه (۳)

جلسه‌ی گذشته

FIFO Algorithm

- Replace the page that was first brought into memory
- Example: Memory system with 4 frames:


Time	0	1	2	3	4	5	6	7	8	9	10
Requests		c	a	d	b	e	b	a	b	c	a
<hr/>											
Page 0	a		a	a	a	a	a	a	a	c	c
Frames 1	b				b	b	b	b	b	b	b
2	c										
3	d	c	c	c	c	e	e	e	e	e	e
				d	d	d	d	d	d	d	a
Page faults						x				x	x

FIFO Algorithm

- Always replace the oldest page.
 - *Replace the page that has been in memory for the longest time*
- Implementation
 - *Maintain a linked list of all pages in memory*
 - *Keep it in order of when they came into memory*
 - *The page at the tail of the list is oldest*
 - *Add new page to head of list*

Belady's Anomaly

- Adding more frames can cause more page faults!



آزمونک قبل

پرسش ۱: فرض کنید صفحه‌ای به شماره‌ی 7F35 (آدرس در مبنای ۱۶)، در حافظه‌ی مجازی داریم. فرض کنید از یک صفحه‌بندی یک مرحله‌ای استفاده کردیم و اندازه‌ی صفحات ۲ کیلوبایت هستند. (۸ نمره)

P#	PF#	P/A present / Absent
0	4	0
1	3	1

الف. Page number و offset را برای این آدرس مجازی مشخص کنید.
(۱ نمره)

ب. جدول روبرو، Page table ما است که ستون اول آن، page number، ستون دوم Page Frame Number و ستون سوم آن valid bit است. براساس آن، شماره‌ی آدرس فیزیکی متناظر با آدرس مجازی 7F35 را محاسبه کنید. (۲ نمره)

P#	PF#	P/A present / Absent
0	4	0
1	3	1
2	2	0
12	14	1
13	20	1
14	11	1
15	18	1

پ. براساس این جدول صفحه، حافظه‌ی مجازی چند بیتی است؟
(۱ نمره)

P#	PF#	P/A present / Absent
0	4	0
1	3	1
2	2	0

30	2	0
31	4	0

ت. آدرس یا آدرس‌های مجازی متناظر با آدرس فیزیکی 1A0E را بیابید.
(۲ نمره)

P#	PF#	P/A present / Absent
0	4	0
1	3	1
2	2	0
3	3	0
4	7	1

31

4

0

ث. اگر حافظه‌ی مجازی ۵۰ بیتی بود، برای هر سطر از جدول صفحه ۱۰ بایت ذخیره می‌کردیم و ۲۰ پردازشی در حال اجرا داشتیم، آیا این روش صفحه بندی مناسب است؟ برای کم کردن سربار، چه راه حل‌هایی مطرح است؟ (۲ نمره)

پرسش ۲: می‌خواهیم عمل‌کرد دو سیستم صفحه‌بندی با داشتن TLB را مقایسه کنیم. فرض کنید زمان دسترسی به یک خانه از حافظه‌ی اصلی 2 micro second باشد. (۴ نمره)

۱. سیستم S1، دارای TLB با زمان پاسخ 120 nano seconds، که در ۶۰٪ مواقع صفحه در TLB است و از صفحه‌بندی ۱ مرحله‌ای استفاده می‌کند.

۲. سیستم S2، دارای TLB بزرگتر با زمان پاسخ 150 nano seconds، که در ۸۰٪ مواقع صفحه در TLB است و از صفحه‌بندی ۳ مرحله‌ای استفاده می‌کند.

در سیستم S1، یک درخواست همزمان به TLB و page table می‌خورد ولی در سیستم S2، ابتدا درخواست به TLB می‌خورد و در صورت TLB miss، درخواست به page table می‌رود.

برای هر کدام از این دو سیستم، کمینه و بیشینه زمان انتظار برای خواندن مقدار یک خانه از حافظه را به دست آورید.

جلسه‌ی جدید

NRU ALGORITHM

How Can We Do Better?

- Need an approximation of how likely each frame is to be accessed in the future
 - *If we base this on past behavior we need a way to track past behavior*
 - *Tracking memory accesses requires hardware support to be efficient*

مسئله: یکی از فریم‌هایی که اخیرا استفاده نشده رو بریزیم دور

- چطوری انتخاب کنیم کدوم فریم؟
- علی‌الحساب فرض کنید نمی‌خوایم داده‌ساختار جدیدی نگه داریم...
- مثلا از اطلاعات توی page table استفاده کنیم.
- چه اطلاعاتی داریم؟
- Reference bit
- Dirty bit

Referenced and Dirty Bits

- Each page table entry (and TLB entry!) has a
 - *Referenced bit - set by TLB when page read / written*
 - *Dirty / modified bit - set when page is written*
 - *If TLB entry for this page is valid, it has the most up to date version of these bits for the page*
 - *OS must copy them into the page table entry during fault handling*
- Idea: use the information contained in these bits to drive the page replacement algorithm

Referenced and Dirty Bits

- Some hardware does not have support for the dirty bit
- Instead, memory protection can be used to emulate it
- Idea:
 - *Software sets the protection bits for all pages to “read only”*
 - *When program tries to update the page...*
 - A trap occurs
 - Software sets the Dirty Bit in the page table and clears the ReadOnly bit
 - Resumes execution of the program

Not Recently Used Algorithm

- Uses the Referenced Bit and the Dirty Bit
- Initially, all pages have
 - *Referenced Bit = 0*
 - *Dirty Bit = 0*
- Periodically... (e.g. whenever a timer interrupt occurs)
 - *Clear the Referenced Bit*
 - *Referenced bit now indicates “recent” access*

Not Recently Used Algorithm

- When a page fault occurs...
- Categorize each page...
 - Class 1: Referenced = 0 Dirty = 0
 - Class 2: Referenced = 0 Dirty = 1
 - Class 3: Referenced = 1 Dirty = 0
 - Class 4: Referenced = 1 Dirty = 1
- Choose a victim page from ...
 - class 1 ... why?
- If none, choose a page from ...
 - class 2 ... why?
- If none, choose a page from ...
 - class 3 ... why?
- If none, choose a page from ...
 - class 4 ... why?


SECOND CHANCE ALGORITHM

Second Chance Algorithm

- An implementation of NRU based on FIFO
- Pages kept in a linked list (oldest at the front)
- Look at the oldest page
 - *If its “referenced bit” is 0...*
 - Select it for replacement
 - *Else*
 - It was used recently; don't want to replace it
 - Clear its “referenced bit”
 - Move it to the end of the list
 - *Repeat*
- What if every page was used in last clock tick?

Implementation of Second Chance

- Maintain a circular list of pages in memory
- Set a bit for the page when a page is referenced
- Search list looking for a victim page that does not have the referenced bit set
 - *If the bit is set, clear it and move on to the next page*
 - *Replaces pages that haven't been referenced for one complete clock revolution*



Enhanced Second Chance algorithm

- همزمان از هر دو ایده‌ی قبلی استفاده کنیم...
- در ۴ گروه براساس reference bit و dirty bit در نظر بگیریم.
- حالا second chance رو براساس گروه‌بندی قبلی انجام بدیم...



LRU

Least Recently Used Algorithm

- A refinement of NRU that orders how recently a page was used
 - *Keep track of when a page is used*
 - *Replace the page that has been used least recently*

Least Recently Used Algorithm

- Replace the page that hasn't been referenced in the longest time

Time		0	1	2	3	4	5	6	7	8	9	10
Requests			c	a	d	b	e	b	a	b	c	d
<hr/>												
Page	0	a	a	a	a	a	a	a	a	a	a	a
Frames	1	b	b	b	b	b	b	b	b	b	b	b
	2	c	c	c	c	e	e	e	e	e	e	d
	3	d	d	d	d	d	d	d	d	d	c	c
<hr/>												
Page faults							x				x	x

Least Recently Used Algorithm

- But how can we implement LRU?
- Idea #1:
 - *Keep a linked list of all pages*
 - *On every memory reference, Move that page to the front of the list*
 - *The page at the tail of the list is replaced*

Least Recently Used Algorithm

- But how can we implement LRU?
 - ... *without requiring every access to be recorded?*
- Idea #2:
 - *MMU (hardware) maintains a counter*
 - *Incremented on every clock cycle*
 - *Every time a page table entry is used*
 - MMU writes the value to the page table entry
 - This *timestamp* value is the *time-of-last-use*
 - *When a page fault occurs*
 - OS looks through the page table
 - Identifies the entry with the oldest timestamp

Least Recently Used Algorithm

- What if we don't have hardware support for a counter?
- Idea #3:
 - *Maintain a counter in software*
 - *One every timer interrupt...*
 - Increment counter
 - Run through the page table
 - For every entry that has "ReferencedBit" = 1
 - * *Update its timestamp*
 - * *Clear the ReferencedBit*
 - *Approximates LRU*
 - *If several have oldest time, choose one arbitrarily*

NFU

Not Frequently Used Algorithm

- Bases decision of frequency of use rather than recency
- Associate a counter with each page
- On every clock interrupt, the OS looks at each page.
 - *If the reference bit is set increment that page's counter & clear the bit*
- The counter approximates how often the page is used
- For replacement, choose the page with lowest counter

Not Frequently Used Algorithm

■ *Problem:*

- *Some page may be heavily used*
 - Its counter is large
- *The program's behavior changes*
 - Now, this page is not used ever again (or only rarely)
- *This algorithm never forgets!*
 - - *This page will never be chosen for replacement!*
- *We may want to combine frequency and recency somehow*

NFU With Aging

- Associate a counter with each page
- On every clock tick, the OS looks at each page.
 - *Shift the counter right 1 bit (divide its value by 2)*
 - *If the reference bit is set...*
 - Set the most-significant bit
 - Clear the Referenced Bit

– T	$100000 = 32$
– T^1	$010000 = 16$
– T^2	$001000 = 8$
– T^3	$000100 = 4$
– T^4	$100010 = 34$
– T_5	


THRASHING

Thrashing

- If a process does not have “enough” pages, the page-fault rate is very high
 - *Page fault to get page*
 - *Replace existing frame*
 - *But quickly need replaced frame back*
 - *This leads to:*
 - Low CPU utilization
 - Operating system thinking that it needs to increase the degree of multiprogramming
 - Another process added to the system

Thrashing (Cont.)

- **Thrashing.** A process is busy swapping pages in and out



WORKING SET

The Working Set

- *Demand paging*
 - *Pages are only loaded when accessed*
 - *When process begins, all pages marked INVALID*

The Working Set

■ *Demand paging*

- *Pages are only loaded when accessed*
- *When process begins, all pages marked INVALID*

■ *Locality of reference*

- *Processes tend to use only a small fraction of their pages*

■ *Working Set*

- *The set of pages a process needs*
- *If working set is in memory, no page faults*
- *What if you can't get working set into memory?*

The Working Set

■ *Thrashing*

- *If you can't get working set into memory page faults occur every few instructions*
- *Little work gets done*
- *Most of the CPU's time is going on overhead*

Working Set Algorithm

- *Based on prepaging (prefetching)*
 - *Load pages before they are needed*
- *Main idea:*
 - *Try to identify the process's working set based on time*
 - *Keep track of each page's time since last access*
 - *Assume working set valid for T time units*
 - *Replace pages older than T*

Working Set Algorithm


- *Current Virtual Time*

- *Only consider how much CPU time this process has seen*

- *Implementation*


- *On each clock tick, look at each page*
 - *Was it referenced since the last check?*
 - Yes: make a note of Current Virtual Time
 - *If a page has not been used in the last T msec,*
 - Assume it is not in the working set!
 - Evict it
 - Write it out if it is dirty

Working Set Algorithm




Page Fault Frequency

- *If T is too small, page fault frequency will be high*
- *If you make it bigger page fault frequency will decline*



Page Fault Frequency



Page Fault Frequency

- Measure the page fault frequency of each process
- Count the number of faults every second
- May want to consider the past few seconds as well
- Aging:
 - *Keep a running value*
 - *Every second*
 - Count number of page faults
 - Divide running value by 2
 - Add in the count for this second

WHICH ALGORITHM IS
BEST?



Modeling Algorithm Performance

- Run a program


- *Look at all memory references*
- *Don't need all this data*
- *Look at which pages are accessed*
- *0000001222333300114444001123444*
- *Eliminate duplicates*
- *012301401234*

- This defines the *Reference String*

- *Use this to evaluate different algorithms*
- *Count page faults given the same reference string*

Proactive Replacement

- Replacing victim frame on each page fault typically requires two disk accesses per page fault
- Alternative → the O.S. can keep several pages free in anticipation of upcoming page faults.
- In Unix: low and high water marks



$\text{low} < \# \text{ free pages} < \text{high}$

UNIX Page Replacement

- Clock algorithm for page replacement
 - *If page has not been accessed move it to the free list for use as allocatable page*
 - *If modified/dirty → write to disk (still keep stuff in memory though)*
 - *If unmodified → just move to free list*
- High and low water marks for free pages
 - *Pages on the free-list can be re-allocated if they are accessed again before being overwritten*

Local vs. Global Replacement

- Assume several processes: A, B, C, ...
- Some process gets a page fault (say, process A)
- Choose a page to replace.
- *Local page replacement*
 - *Only choose one of A's pages*
- *Global page replacement*
 - *Choose any page*

Local vs. Global Replacement

	Age
A0	10
A1	7
A2	5
A3	4
A4	6
A5	3
B0	9
B1	4
B2	6
B3	2
B4	5
B5	6
B6	12
C1	3
C2	5
C3	6

Original

A0
A1
A2
A3
A4
A6
B0
B1
B2
B3
B4
B5
B6
C1
C2
C3

Local

A0
A1
A2
A3
A4
A5
B0
B1
B2
A6
B4
B5
B6
C1
C2
C3

Global

Local vs. Global Replacement

- Assume we have
 - *5,000 frames in memory and 10 processes*
- Idea: Give each process 500 frames
- Is this fair?
 - *Small processes do not need all those pages*
 - *Large processes may benefit from even more frames*
- Idea:
 - *Look at the size of each process (... but how?)*
 - *Give them a pro-rated number of frames with some minimum*

Load Control

- Assume:

- *The best page replacement algorithm*
- *Optimal global allocation of page frames*

- *Thrashing is still possible!*

- *Too many page faults!*
- *No useful work is getting done!*
- *Demand for frames is too great!*

- Solution:

- *Get rid of some processes (temporarily swap them out)*
- *Two-level scheduling (swapping with paging)*