

R

بسم الله الرحمن الرحيم

سیستم عامل

جلسه نوزدهم – الگوریتم‌های جایگزینی صفحه (۴)

جلسه‌ی گذشته

FIFO Algorithm

- Always replace the oldest page.
 - *Replace the page that has been in memory for the longest time*
- Implementation
 - *Maintain a linked list of all pages in memory*
 - *Keep it in order of when they came into memory*
 - *The page at the tail of the list is oldest*
 - *Add new page to head of list*

How Can We Do Better?

- Need an approximation of how likely each frame is to be accessed in the future
 - *If we base this on past behavior we need a way to track past behavior*
 - *Tracking memory accesses requires hardware support to be efficient*

Not Recently Used Algorithm

- When a page fault occurs...
- Categorize each page...
 - Class 1: Referenced = 0 Dirty = 0
 - Class 2: Referenced = 0 Dirty = 1
 - Class 3: Referenced = 1 Dirty = 0
 - Class 4: Referenced = 1 Dirty = 1
- Choose a victim page from ...
 - class 1 ... why?
- If none, choose a page from ...
 - class 2 ... why?
- If none, choose a page from ...
 - class 3 ... why?
- If none, choose a page from ...
 - class 4 ... why?

NRU

- Class 1: *Referenced = 0 Dirty = 0*
- Class 2: *Referenced = 0 Dirty = 1*
- Class 3: *Referenced = 1 Dirty = 0*
- Class 4: *Referenced = 1 Dirty = 1*


Time	0	1	2	3	4	5	6	7	8	9	10
Requests		c	a	d	b	e	b	a	b	c	d

Page	0	a
Frames	1	b
	2	c
	3	d


Page faults

Second Chance

- Maintain a circular list of pages in memory
- Set a bit for the page when a page is referenced
- Search list looking for a victim page that does not have the referenced bit set
 - *If the bit is set, clear it and move on to the next page*
 - *Replaces pages that haven't been referenced for one complete clock revolution*
- Second Chance with circular queue also referred as Clock Algorithm



Second Chance



Time	0	1	2	3	4	5	6	7	8	9	10
Requests		c	a	d	b	e	b	a	b	c	d

Page	0	a
Frames	1	b
	2	c
	3	d

Page faults

Least Recently Used Algorithm

- Replace the page that hasn't been referenced in the longest time

Time	0	1	2	3	4	5	6	7	8	9	10
Requests		c	a	d	b	e	b	a	b	c	d
<hr/>											
Page 0	a	a	a	a	a	a	a	a	a	a	a
Frames 1	b	b	b	b	b	b	b	b	b	b	b
2	c	c	c	c	c	e	e	e	e	e	d
3	d	d	d	d	d	d	d	d	d	c	c
Page faults						x				x	x

Least Recently Used Algorithm



- But how can we implement LRU?
- Idea #1:
 - *Keep a linked list of all pages*
 - *On every memory reference, Move that page to the front of the list*
 - *The page at the tail of the list is replaced*

Least Recently Used Algorithm

- But how can we implement LRU?
 - ... *without requiring every access to be recorded?*
- Idea #2:
 - *MMU (hardware) maintains a counter*
 - *Incremented on every clock cycle*
 - *Every time a page table entry is used*
 - MMU writes the value to the page table entry
 - This *timestamp* value is the *time-of-last-use*
 - *When a page fault occurs*
 - OS looks through the page table
 - Identifies the entry with the oldest timestamp

Least Recently Used Algorithm

- What if we don't have hardware support for a counter?

- Idea #3:

- *Maintain a counter in software*
- *One every timer interrupt...*
 - Increment counter
 - Run through the page table
 - For every entry that has "ReferencedBit" = 1
 - * *Update its timestamp*
 - * *Clear the ReferencedBit*
- *Approximates LRU*
- *If several have oldest time, choose one arbitrarily*

$C = 1$

10
11

	RB	C
A	0	1
B	0	1
C	0	1
D	0	0

Not Frequently Used Algorithm

- Bases decision of frequency of use rather than recency
- Associate a counter with each page
- On every clock interrupt, the OS looks at each page.
 - *If the reference bit is set increment that page's counter & clear the bit*
- The counter approximates how often the page is used
- For replacement, choose the page with lowest counter

NFU With Aging

- Associate a counter with each page
- On every clock tick, the OS looks at each page.
 - *Shift the counter right 1 bit (divide its value by 2)*
 - *If the reference bit is set...*
 - Set the most-significant bit
 - Clear the Referenced Bit

- T_1 100000 = 32
- T_1^1 010000 = 16
- T_2^2 001000 = 8
- T_3^3 000100 = 4
- T_5^4 100010 = 34

010001
110001

جلسه‌ی جدید

FRAME ALLOCATION



Local vs. Global Allocation

- Assume we have
 - *5,000 frames in memory and 10 processes*
- Idea: Give each process 500 frames
- Is this fair?
 - *Small processes do not need all those pages*
 - *Large processes may benefit from even more frames*
- Idea:
 - *Look at the size of each process (... but how?)*
 - *Give them a pro-rated number of frames with some minimum*

Local vs. Global Replacement

- Assume several processes: A, B, C, ...
- Some process gets a page fault (say, process A)
- Choose a page to replace.
- *Local page replacement*
 - *Only choose one of A's pages*
- *Global page replacement*
 - *Choose any page*

Local vs. Global Replacement

	Age
A0	10
A1	7
A2	5
A3	4
A4	6
A5	3
B0	9
B1	4
B2	6
B3	2
B4	5
B5	6
B6	12
C1	3
C2	5
C3	6

Original

A0
A1
A2
A3
A4
A6
B0
B1
B2
B3
B4
B5
B6
C1
C2
C3

Local

A0
A1
A2
A3
A4
A5
B0
B1
B2
A6
B4
B5
B6
C1
C2
C3

Global

PROACTIVE REPLACEMENT

Proactive Replacement

- Replacing victim frame on each page fault typically requires two disk accesses per page fault
- Alternative → the O.S. can keep several pages free in anticipation of upcoming page faults
- Free List (Inactive list): List of frames that ready for replacement

Proactive Replacement

- Page Faults:

- *Major Page Fault:*

- a page is referenced and the page is not in memory.

- *Minor Page Fault:*

- process does not have a logical mapping to a page, yet that page is in memory

- Minor Page faults:

- *process may reference a shared library that is in memory, but the process does not have a mapping to it in its page table*
 - *page is reclaimed from a process and placed on the free-frame list, but the page has not yet been zeroed out and allocated to another process*

Proactive Page Replacement

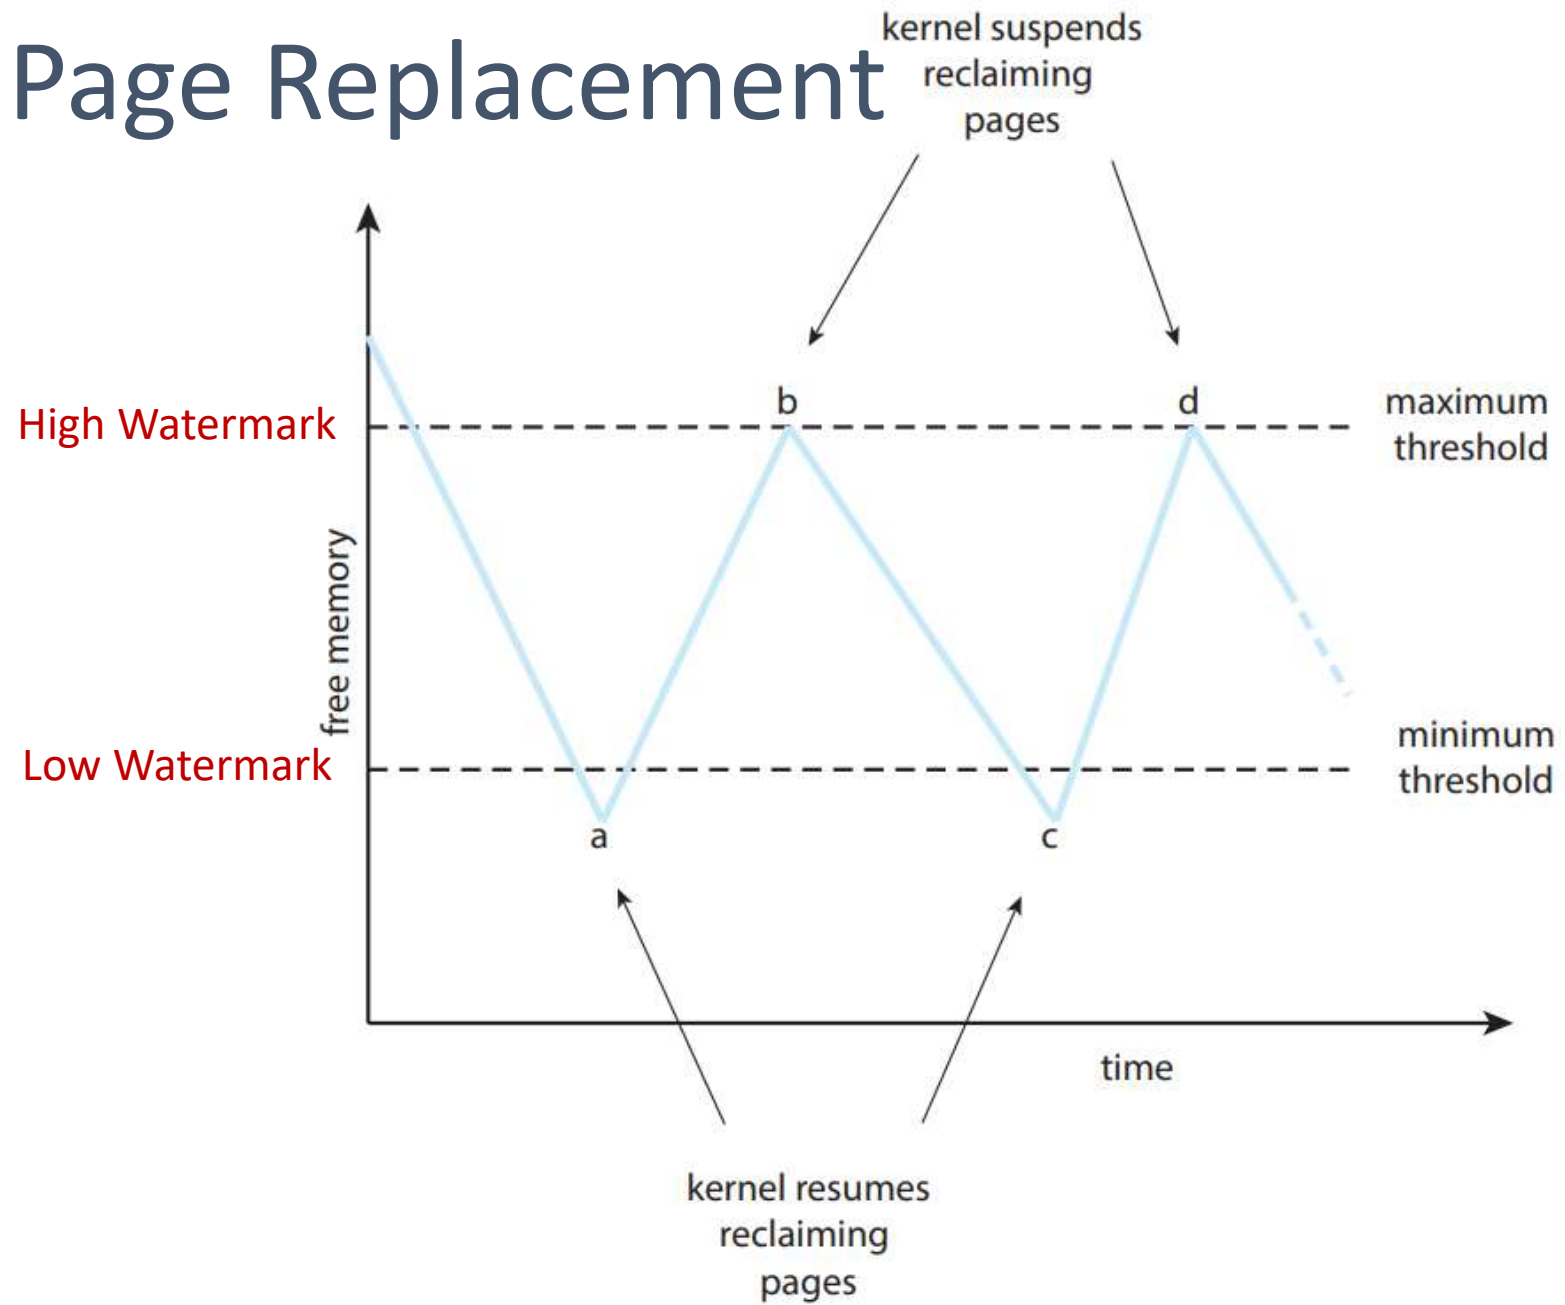


Figure 10.18 Reclaiming pages.

PAGE REPLACEMENT IN LINUX

UNIX Page Replacement

- Some Proactive Page Replacement
- Enable Swapping daemon (kwapd) on low watermark

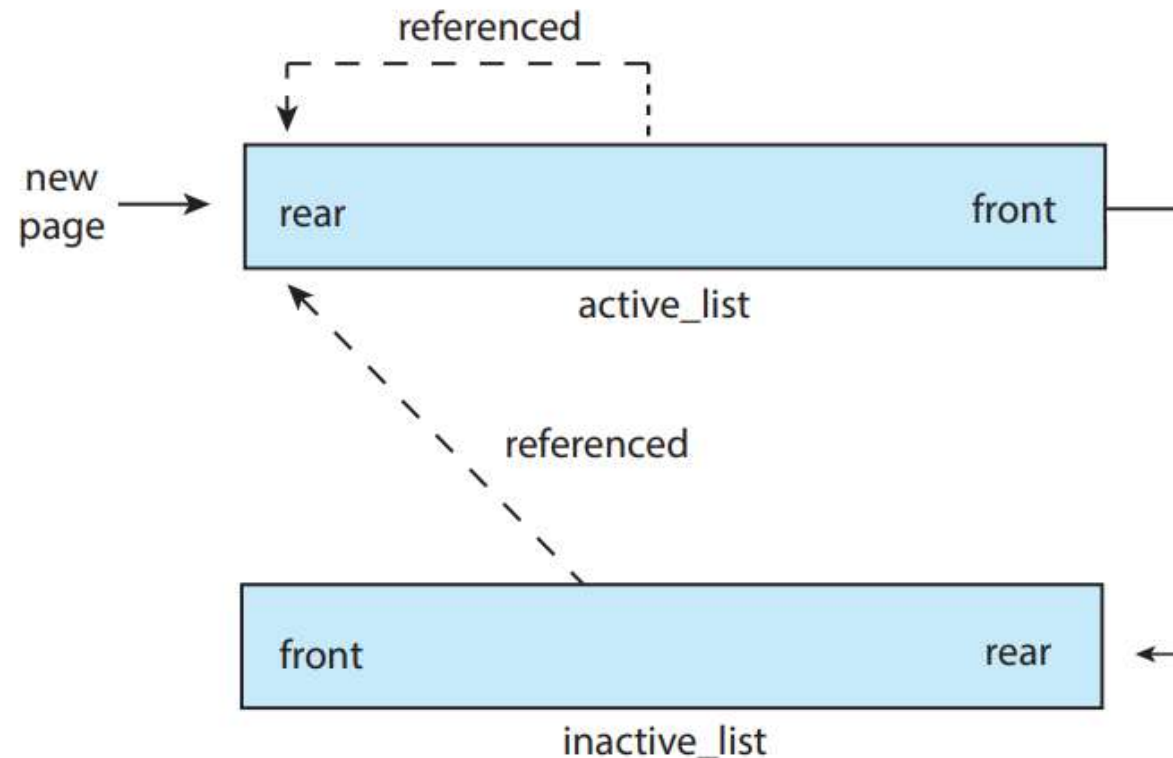


Figure 10.29 The Linux active_list and inactive_list structures.

UNIX Page Replacement

- Second Chance (Clock) Page replacement used as Approximation of LRU

WORKING SET

The Working Set

- *Locality of reference*
 - *Processes tend to use only a small fraction of their pages*
- *Working Set*
 - *The set of pages a process needs*
 - *If working set is in memory, no page faults*
 - *What if you can't get working set into memory?*





Figure 10.21 Locality in a memory-reference pattern.

Thrashing

- If a process does not have “enough” pages, the page-fault rate is very high
 - *Page fault to get page*
 - *Replace existing frame*
 - *But quickly need replaced frame back*
 - *This leads to:*
 - Low CPU utilization
 - Operating system thinking that it needs to increase the degree of multiprogramming
 - Another process added to the system

Thrashing (Cont.)

- **Thrashing.** A process is busy swapping pages in and out




Working Set Algorithm


- *Based on prepaging (prefetching)*
 - *Load pages before they are needed*
- *Main idea:*
 - *Try to identify the process's working set based on time*
 - *Keep track of each page's time since last access*
 - *Assume working set valid for T time units*
 - *Replace pages older than T*

page reference table

. . . 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 4 1 3 2 3 4 4 4 3 4 4 4 . . .



$WS(t_1) = \{1, 2, 5, 6, 7\}$



$WS(t_2) = \{3, 4\}$

Figure 10.22 Working-set model.

Working Set Algorithm

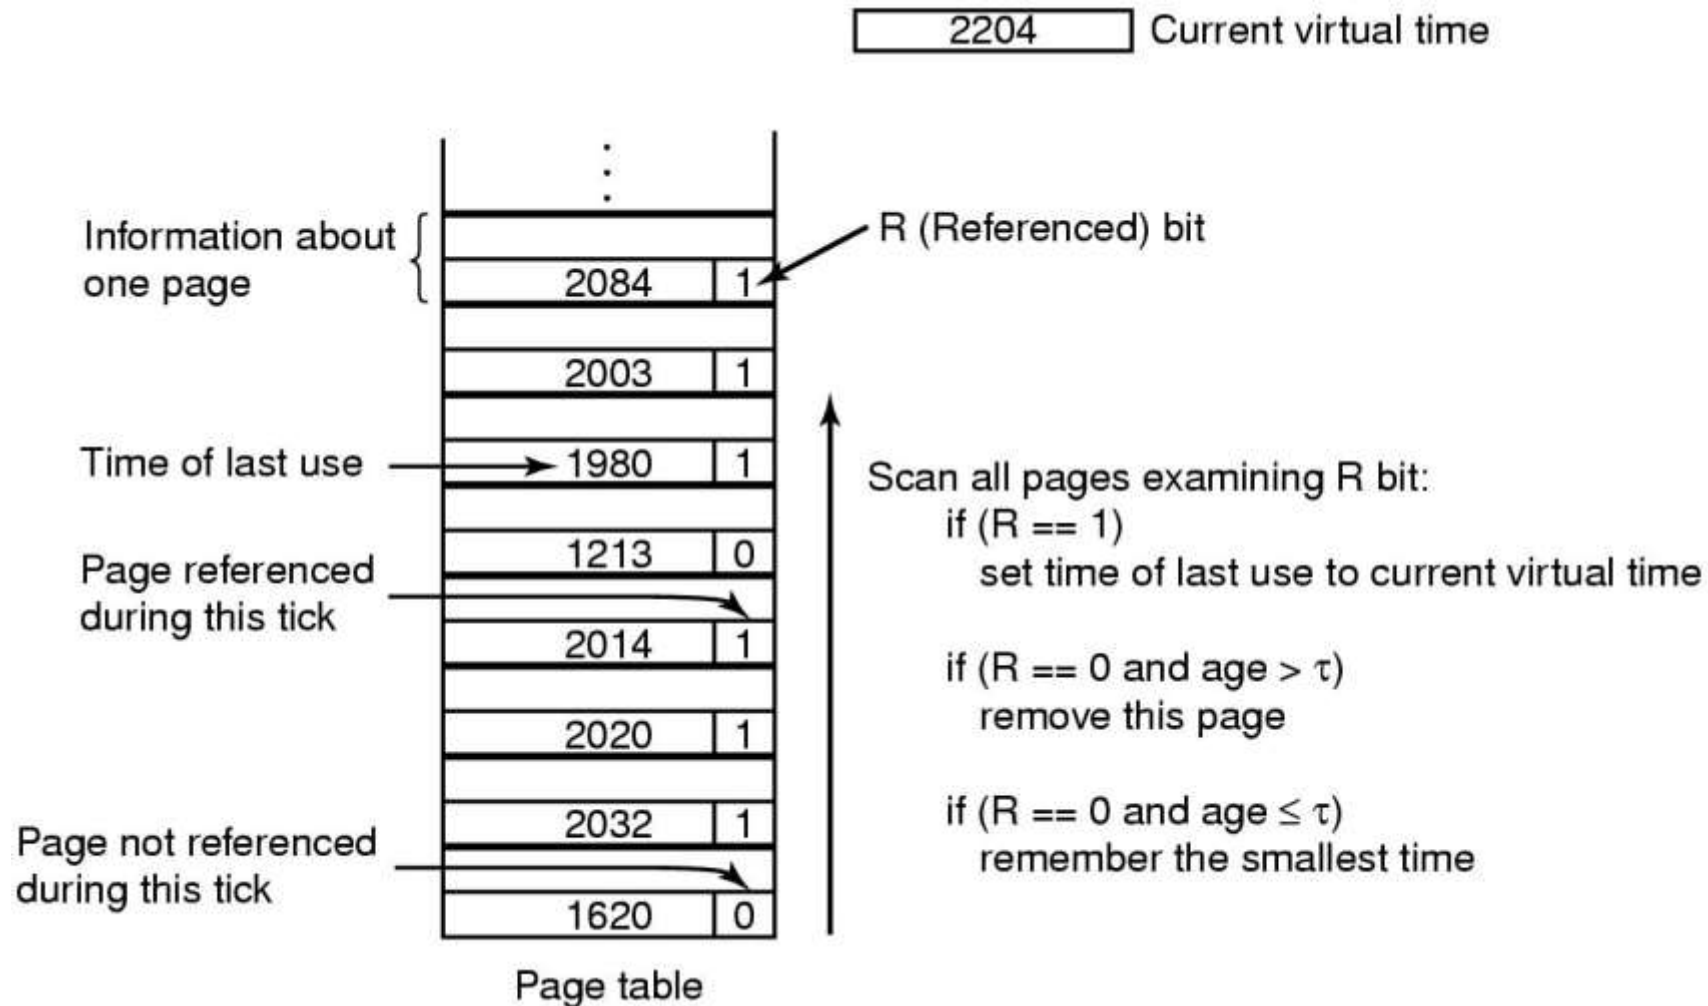
- *Current Virtual Time*

- *Only consider how much CPU time this process has seen*

- *Implementation*


- *On each clock tick, look at each page*
 - *Was it referenced since the last check?*
 - Yes: make a note of Current Virtual Time
 - *If a page has not been used in the last T msec,*
 - Assume it is not in the working set!
 - Evict it
 - Write it out if it is dirty

Working Set Algorithm

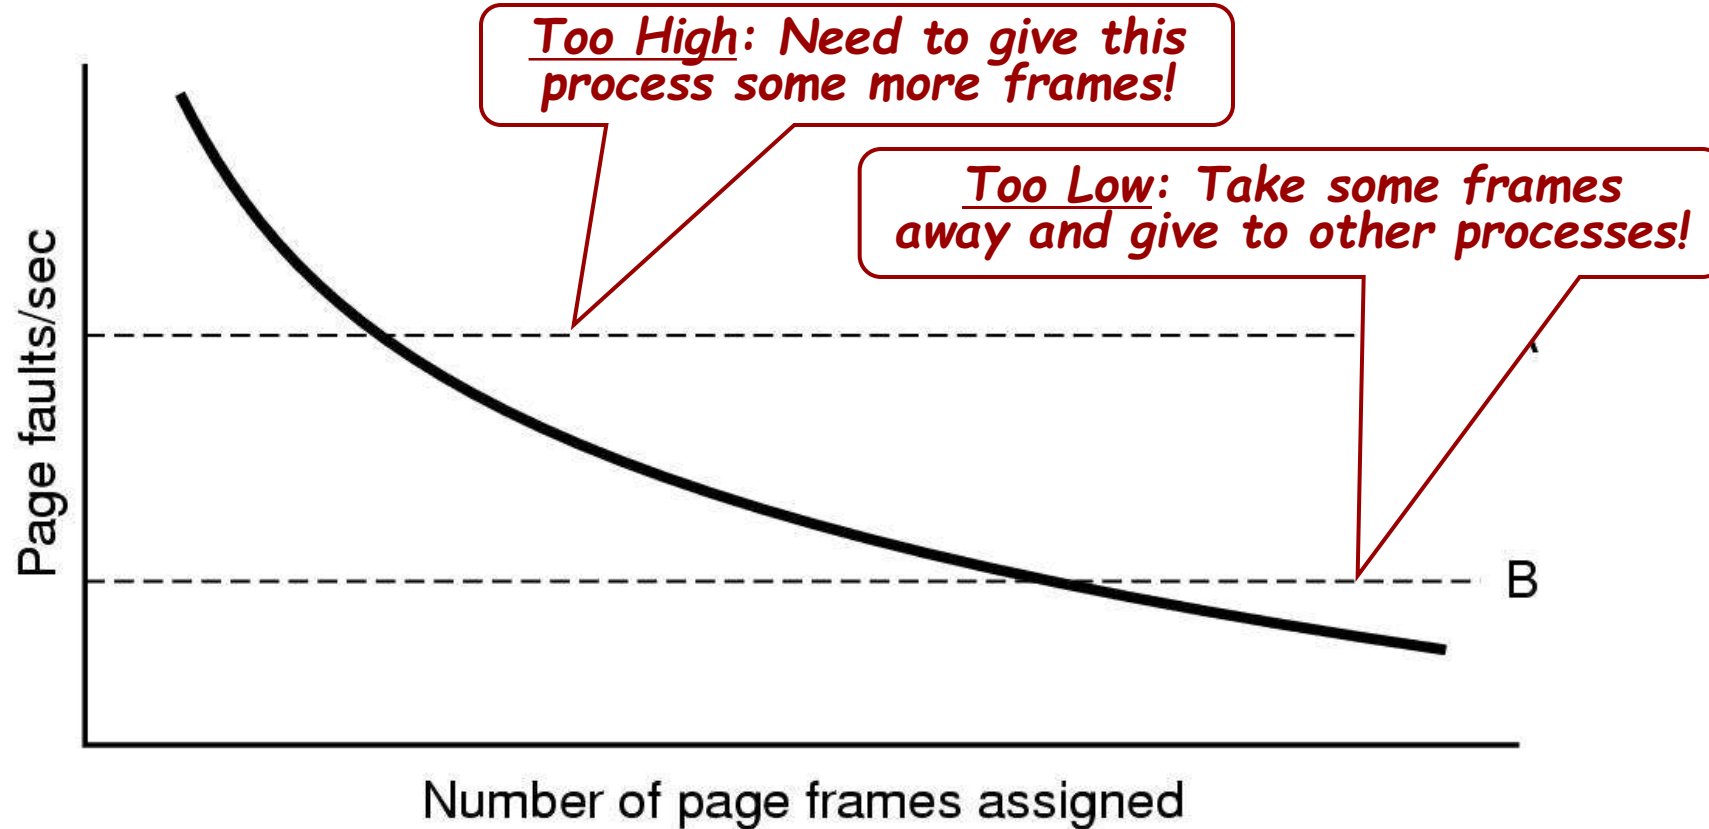


Page Fault Frequency

- *If T is too small, page fault frequency will be high*
- *If you make it bigger page fault frequency will decline*



Page Fault Frequency



Page Fault Frequency

- Measure the page fault frequency of each process
- Count the number of faults every second
- May want to consider the past few seconds as well
- Aging:
 - *Keep a running value*
 - *Every second*
 - Count number of page faults
 - Divide running value by 2
 - Add in the count for this second

WHICH ALGORITHM IS
BEST?



Modeling Algorithm Performance

- Run a program

- *Look at all memory references*
- *Don't need all this data*
- *Look at which pages are accessed*
- *0000001222333300114444001123444*
- *Eliminate duplicates*
- *012301401234*

- This defines the *Reference String*

- *Use this to evaluate different algorithms*
- *Count page faults given the same reference string*

Load Control

- Assume:

- *The best page replacement algorithm*
- *Optimal global allocation of page frames*

- *Thrashing is still possible!*

- *Too many page faults!*
- *No useful work is getting done!*
- *Demand for frames is too great!*

- Solution:

- *Get rid of some processes (temporarily swap them out)*
- *Two-level scheduling (swapping with paging)*

PAGE FAULT AND PROGRAMMING



Program Structure

■ Program structure

- `int[128,128] data;`
- *Each row is stored in one page*
- *Program 1*

```
for (j = 0; j < 128; j++)  
    for (i = 0; i < 128; i++)  
        data[i,j] = 0;
```

128 x 128 = 16,384 page faults

- *Program 2*

```
for (i = 0; i < 128; i++)  
    for (j = 0; j < 128; j++)  
        data[i,j] = 0;
```

128 page faults