

R

بسم الله الرحمن الرحيم

سیستم عامل

جلسه بیستم – دستگاه‌های ورودی و خروجی

جلسه‌ی گذشته

Local vs. Global Replacement

- Assume several processes: A, B, C, ...
- Some process gets a page fault (say, process A)
- Choose a page to replace.
- *Local page replacement*
 - *Only choose one of A's pages*
- *Global page replacement*
 - *Choose any page*

Proactive Replacement

- Replacing victim frame on each page fault typically requires two disk accesses per page fault
- Alternative → the O.S. can keep several pages free in anticipation of upcoming page faults
- Free List (Inactive list): List of frames that ready for replacement

UNIX Page Replacement

- Some Proactive Page Replacement
- Enable Swapping daemon (kwapd) on low watermark





Figure 10.29 The Linux active_list and inactive_list structures.

UNIX Page Replacement

- Second Chance (Clock) Page replacement used as Approximation of LRU

Thrashing (Cont.)

- **Thrashing.** A process is busy swapping pages in and out




Working Set Algorithm


- *Based on prepaging (prefetching)*
 - *Load pages before they are needed*
- *Main idea:*
 - *Try to identify the process's working set based on time*
 - *Keep track of each page's time since last access*
 - *Assume working set valid for T time units*
 - *Replace pages older than T*

page reference table

. . . 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 4 1 3 2 3 4 4 4 3 4 4 4 . . .



$WS(t_1) = \{1, 2, 5, 6, 7\}$

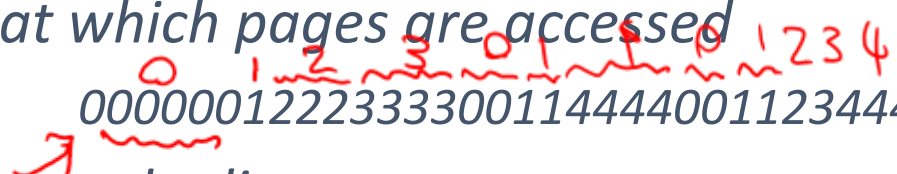


$WS(t_2) = \{3, 4\}$

Figure 10.22 Working-set model.

Modeling Algorithm Performance

■ Run a program

- *Look at all memory references*
- *Don't need all this data*
- *Look at which pages are accessed*
- 
0000001222333300114444001123444
- *Eliminate duplicates*
- 012301401234

■ This defines the *Reference String*

- *Use this to evaluate different algorithms*
- *Count page faults given the same reference string*

Program Structure

■ Program structure

- `int[128,128] data;`
- *Each row is stored in one page*
- *Program 1*

```
for (j = 0; j < 128; j++)  
  for (i = 0; i < 128; i++)  
    data[i,j] = 0;
```

128 x 128 = 16,384 page faults

- *Program 2*

```
for (i = 0; i < 128; i++)  
  for (j = 0; j < 128; j++)  
    data[i,j] = 0;
```

128 page faults



جلسه‌ی جدید

دستگاه‌های ورودی و خروجی

فصل ۱۲ کتاب

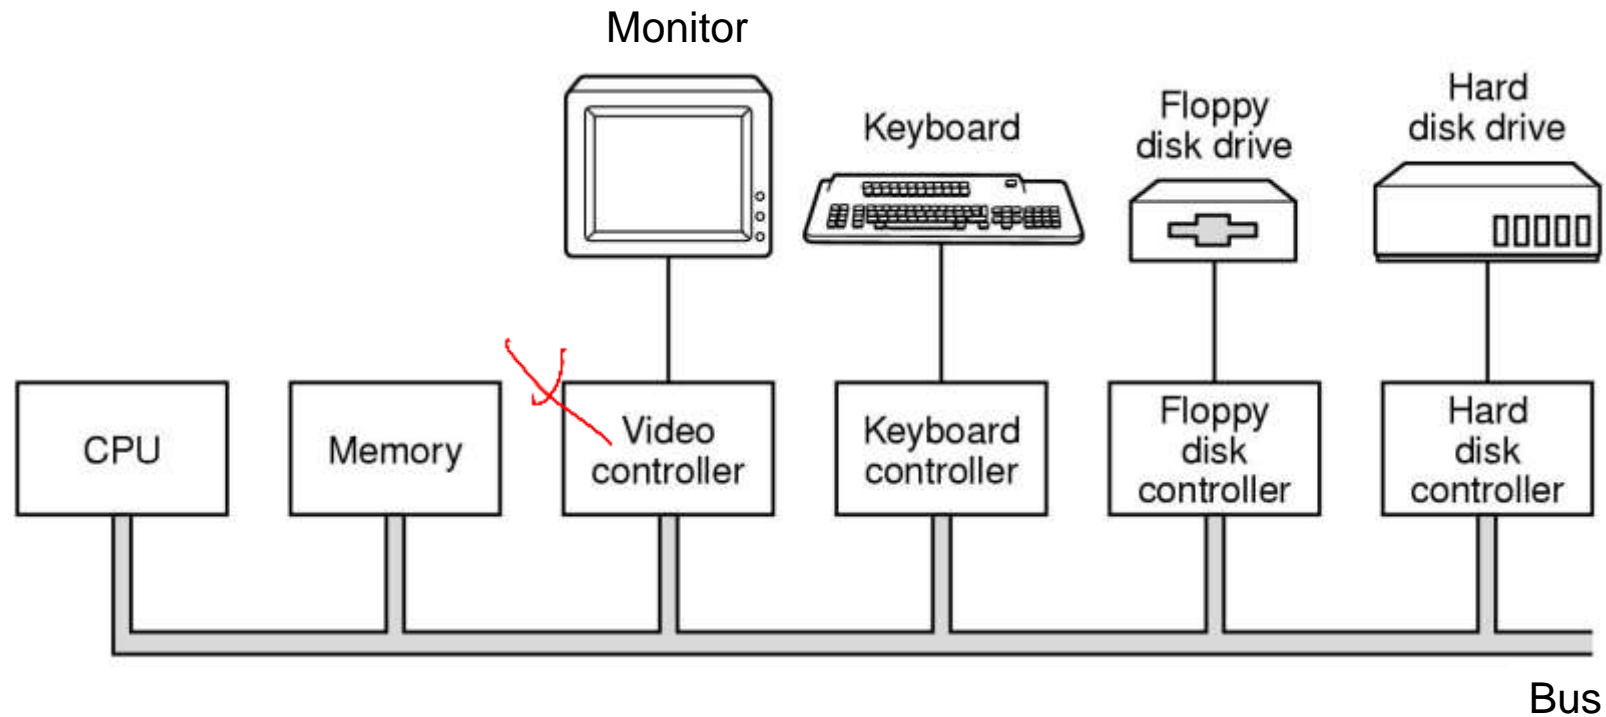
نمای کلی

Device Terminology

- Device (mechanical hardware)
- Device controller (electrical hardware)
- Device driver (software)

Devices & Controllers

■ Components of a simple personal computer



Device Controllers

■ The Device vs. its Controller


■ Some duties of a device controller:

- *Interface between CPU and the Device*
- *Start/Stop device activity*
- *Convert serial bit stream to a block of bytes*
- *Deal with error detection/correction*
- *Move data to/from main memory*

■ Some controllers may handle several (similar) devices

Communication With Devices

- Hardware supports I/O ports or memory mapped I/O for accessing device controller registers and buffers



I/O Ports

- Each port has a separate number.
- CPU has special I/O instructions

- *in* *r4, 3*


- *out* *3, r4*

The I/O Port Number

- Port numbers form an “address space” ... separate from main memory
- Contrast with
 - *load* *r4, 3*
 - *store* *3, r4*

Memory-Mapped I/O

- One address space for
 - *main memory*
 - *I/O devices*
- CPU has no special instructions
 - *load r4, addr*
 - *store addr, r4*
- I/O devices are “mapped” into
 - *very high addresses*



I/O Device Speed


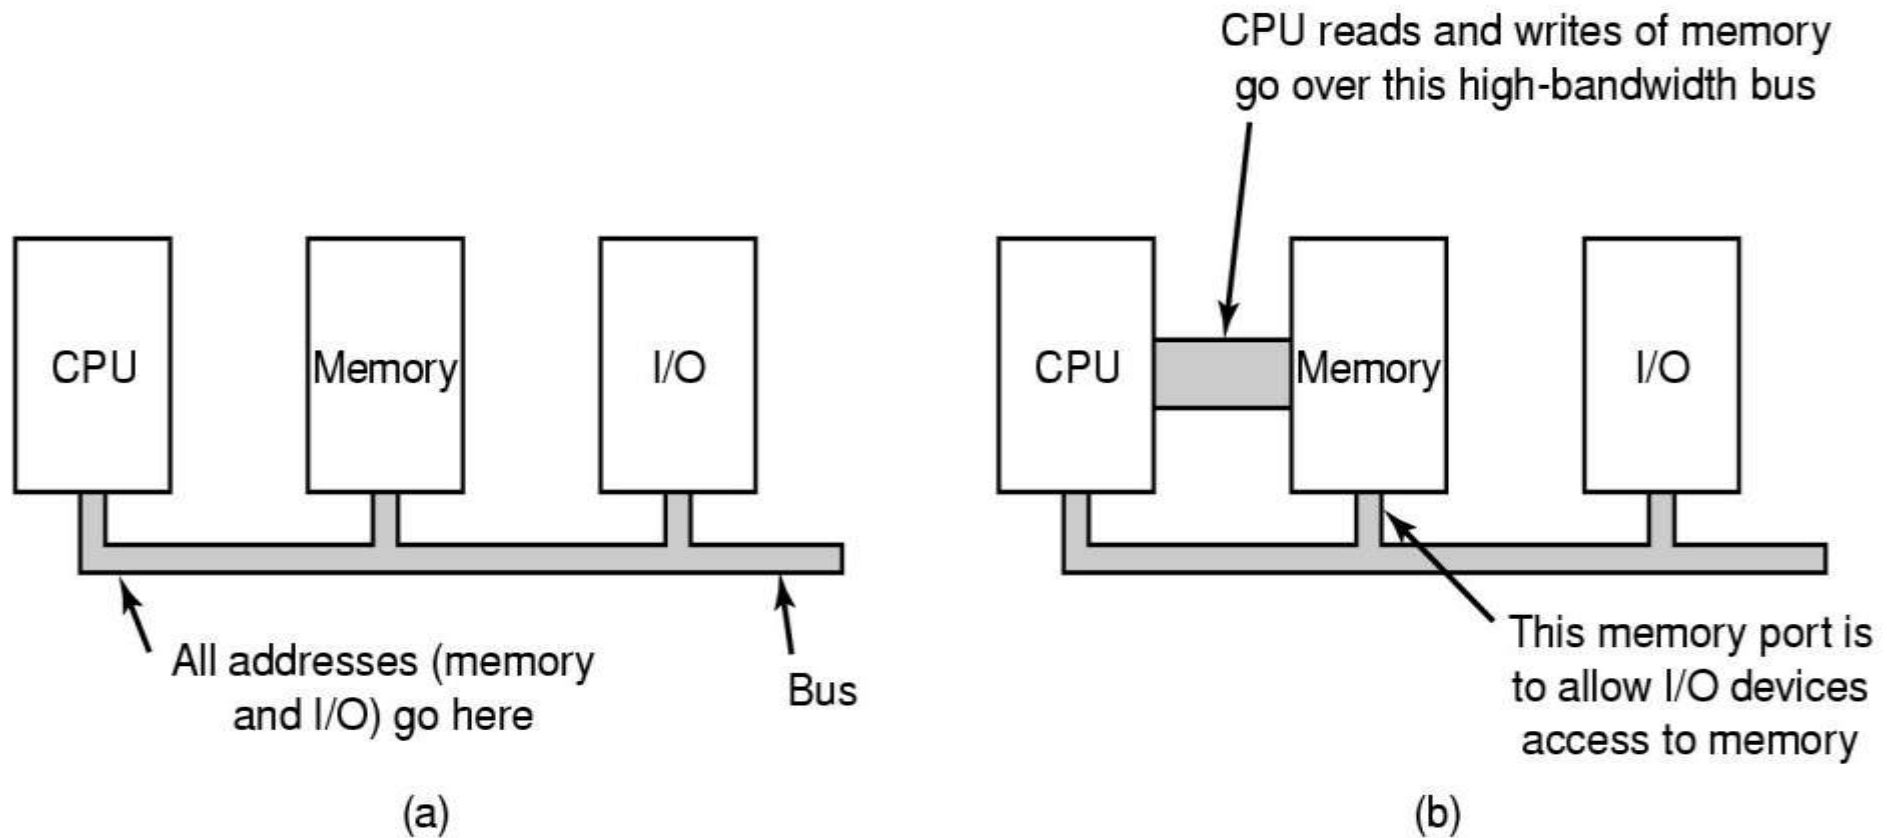


Figure 12.11 Common PC and data-center I/O device and interface speeds.

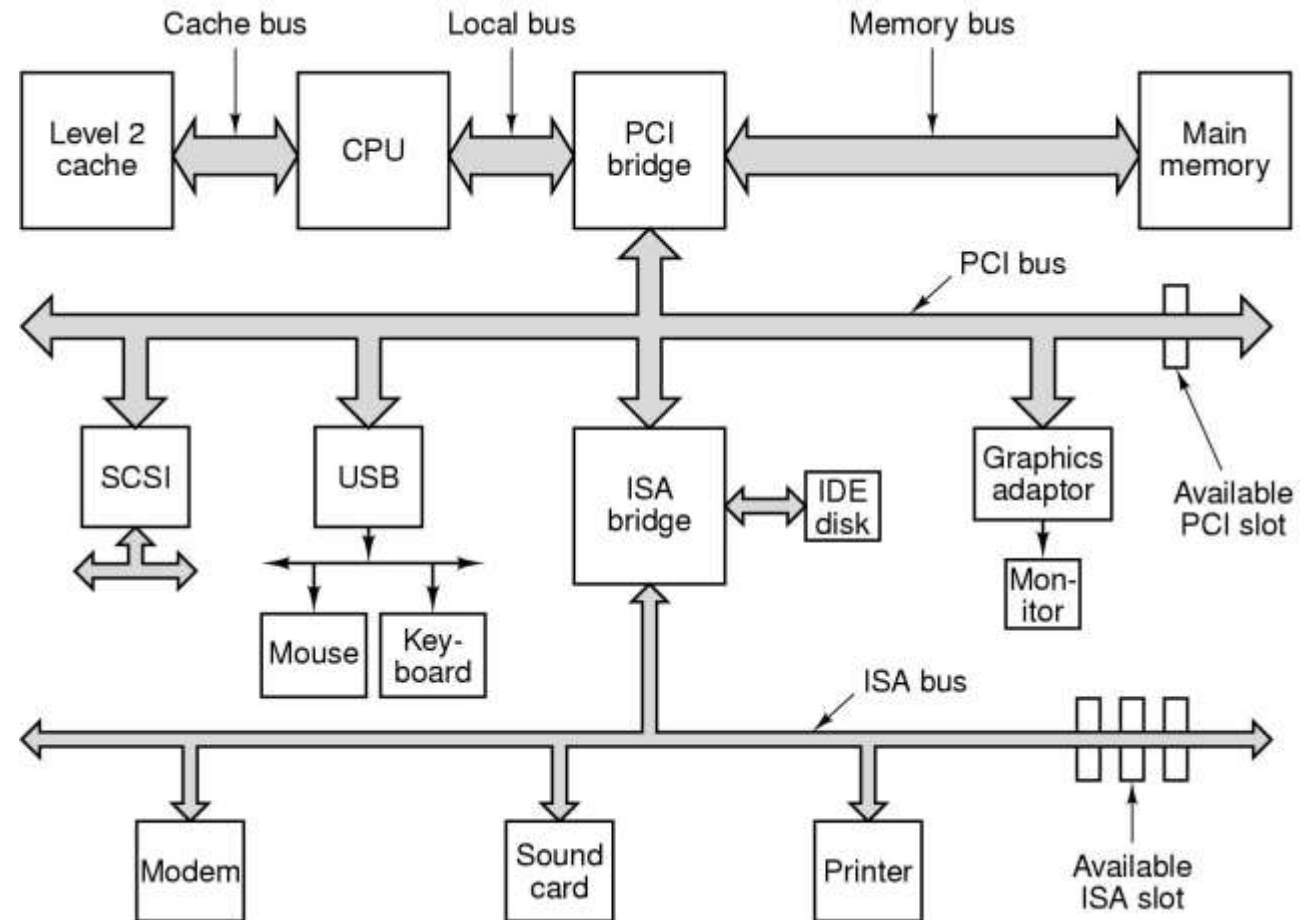
Hardware Performance Challenges

- How to prevent slow devices from slowing down memory due to bus contention
 - *What is bus contention?*
- How to access I/O addresses without interfering with memory performance

Dual Bus Architecture



Pentium Bus Architecture




Software Performance Challenges

- How to prevent CPU throughput from being limited by I/O device speed (for *slow devices*)
 - *Why would slow devices affect the CPU?*
- How to prevent I/O throughput from being limited by CPU speed (for *fast devices*)
 - *Why would device throughput be limited by the CPU?*
- How to achieve good utilization of CPU and I/O devices
- How to meet the real-time requirements of devices

ارتباط با دستگاه‌های ورودی و خروجی

Programmed I/O



Steps in printing a string

Programmed I/O

- Example: Writing a string to a serial output or printing a string

```
CopyFromUser(virtAddr, kernelBuffer, byteCount)
```

```
for i = 0 to byteCount-1
```

```
    while *serialStatusReg != READY
```

```
    endwhile
```

```
    *serialDataReg = kernelBuffer[i]
```

```
endfor
```

```
return
```

- Called “Busy Waiting” or “Polling”
- Problem: CPU is continually busy working on I/O!

Interrupt-Driven I/O

■ Getting the I/O started:

```
CopyFromUser(virtAddr, kernelBuffer,  
byteCount)
```

```
EnableInterrupts()
```

```
while *serialStatusReg != READY
```

```
endWhile
```

```
*serialDataReg = kernelBuffer[0]
```

```
Sleep ()
```

Interrupt-Driven I/O

■The Interrupt Handler:

```
if i == byteCount
```

```
    Wake up the user process
```

```
else
```

```
    *serialDataReg = kernelBuffer[i]
```

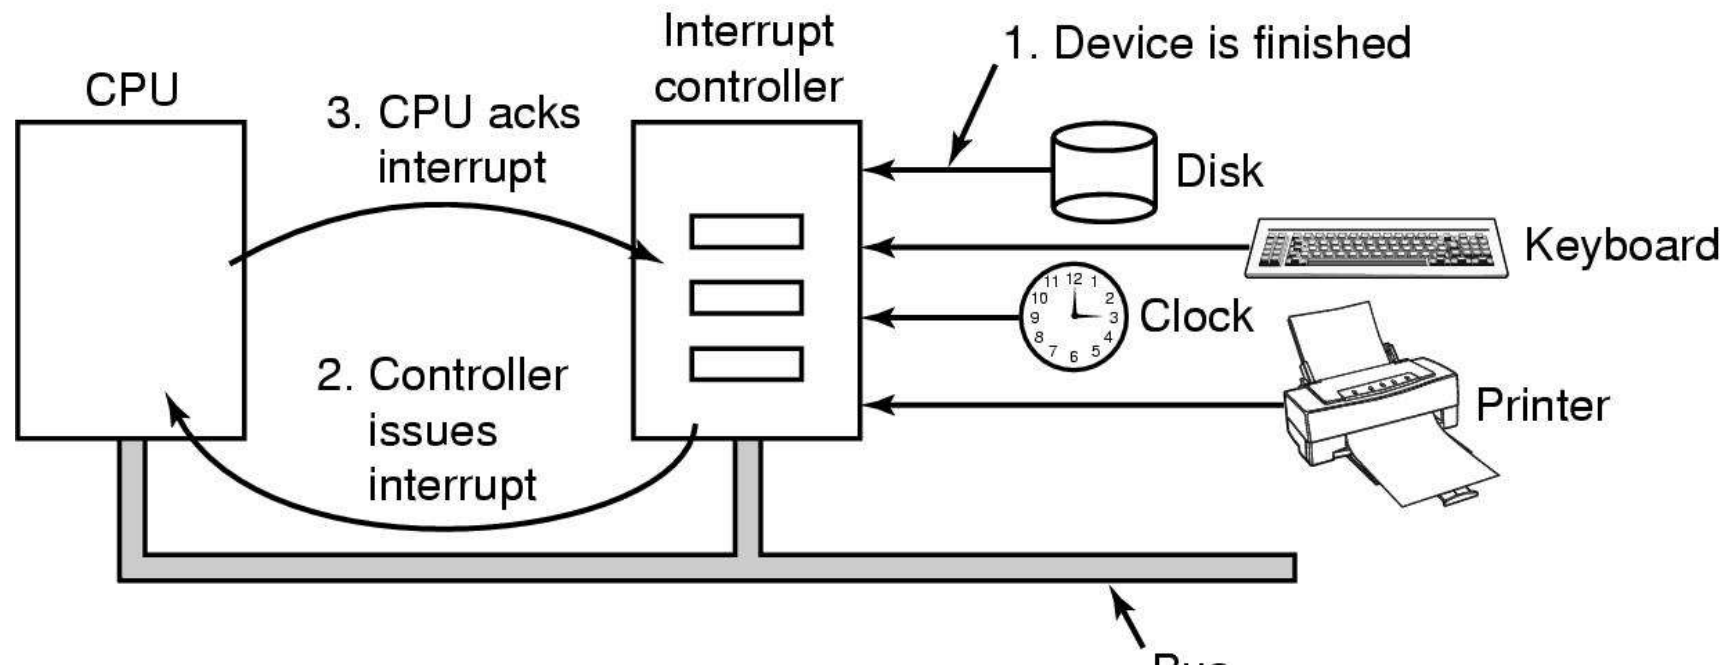
```
    i = i + 1
```

```
endIf
```

```
Return from interrupt
```

Hardware Support For Interrupts

How interrupts happen. Connections between devices and interrupt controller actually use interrupt lines on the bus rather than dedicated wires



Interrupt Driven I/O Problem

■ Problem:

- *CPU is still involved in every data transfer*
- *Interrupt handling overhead is high*
- *Overhead cost is not amortized over much data*
- *Overhead is too high for fast devices*
 - Gbps networks
 - Disk drives

Direct Memory Access (DMA)

- Data transferred from device straight to/from memory
- CPU not involved
- *The DMA controller:*
 - *Does the work of moving the data*
 - *CPU sets up the DMA controller (“programs it ”)*
 - *CPU continues*
 - *The DMA controller moves the bytes*

Sending Data Using DMA

■Getting the I/O started:

```
CopyFromUser(virtAddr, kernelBuffer,  
byteCount)
```

```
Set up DMA controller
```

```
Sleep ()
```


■The Interrupt Handler:

```
Acknowledge interrupt
```

```
Wake up the user process
```

```
Return from interrupt
```

Direct Memory Access (DMA)



Direct Memory Access (DMA)

■ Cycle Stealing

- *DMA Controller acquires control of bus*
- *Transfers a single byte (or word)*
- *Releases the bus*
- *The CPU is slowed down due to bus contention*

■ Burst Mode

- *DMA Controller acquires control of bus*
- *Transfers all the data*
- *Releases the bus*
- *The CPU operation is temporarily suspended*

Principles of I/O Software

■ Device Independence

- *Programs can access any I/O device*
 - Hard Drive, CD-ROM, Floppy,...
 - ... without specifying the device in advance

■ Uniform Naming

- *Devices / Files are named with simple strings*
- *Names should not depend on the device*

■ Error Handling

- *Should be as close to the hardware as possible because its often device-specific*

Principles of I/O Software

■ Synchronous vs. Asynchronous Transfers

- *Process is blocked vs. interrupt-driven or polling approaches*

■ Buffering

- *Data comes off a device*
- *May not know the final destination of the data*
 - e.g., a network packet... Where to put it???

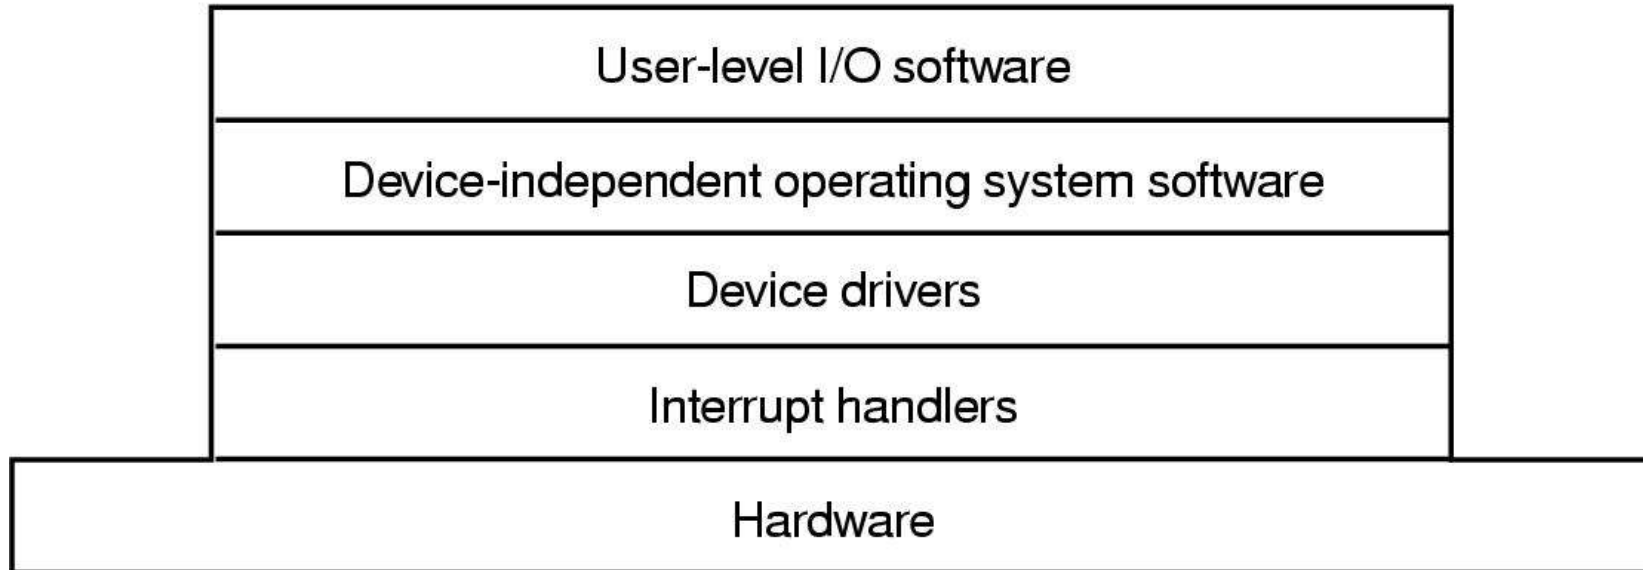
■ Sharable vs. Dedicated Devices

- *Disk should be sharable*
- *Keyboard, Screen dedicated to one process*

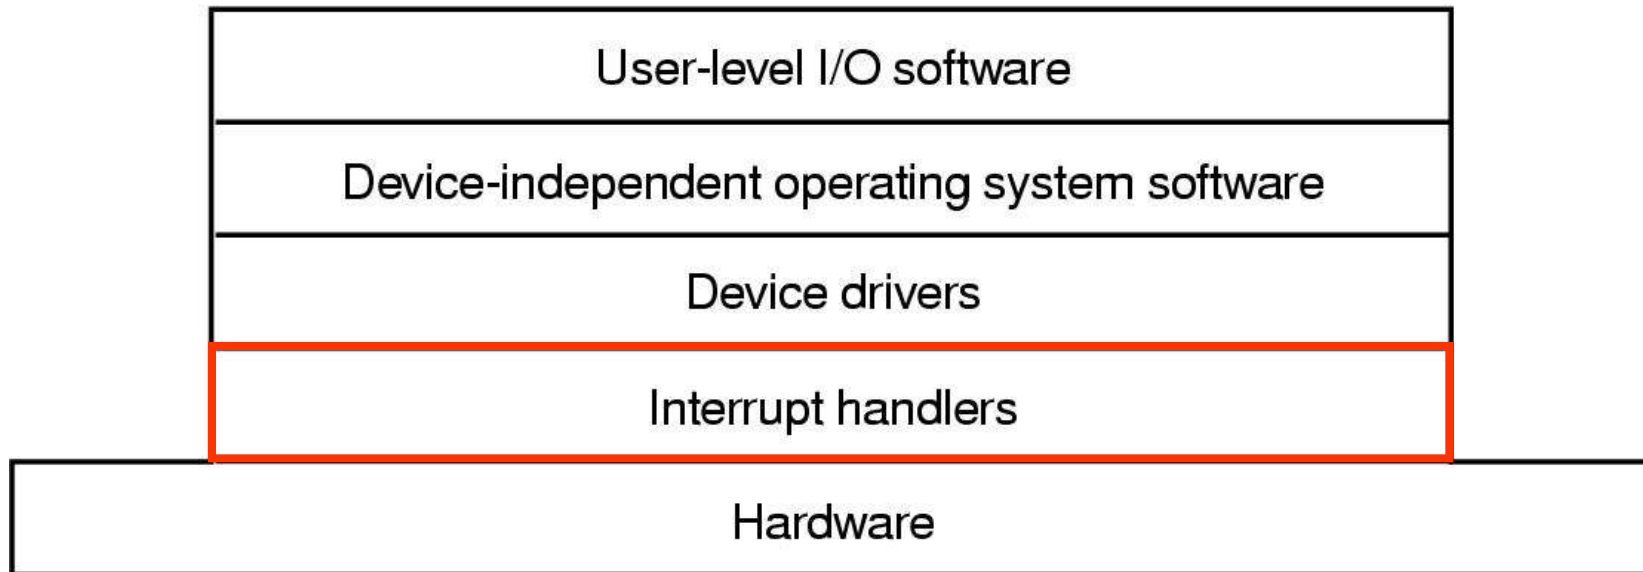
Software Engineering Challenges

- How to remove the complexities of I/O handling from application programs
 - Standard I/O APIs (libraries and system calls)
- How to support a wide range of device types on a wide range of operating systems
 - Standard interfaces for device drivers (DDI)
 - Standard/published interfaces for access to kernel facilities (DKI)

I/O Software Layers



I/O Software Layers



Interrupt Handling

- I/O Device Driver starts the operation
 - *Then blocks until an interrupt occurs*
 - *Then it wakes up, finishes, & returns*
- The Interrupt Handler
 - *Does whatever is immediately necessary*
 - *Then unblocks the driver*
- Example: The BLITZ “DiskDriver”
 - *Start I/O and block (waits on semaphore)*
 - *Interrupt routine signals the semaphore & returns*

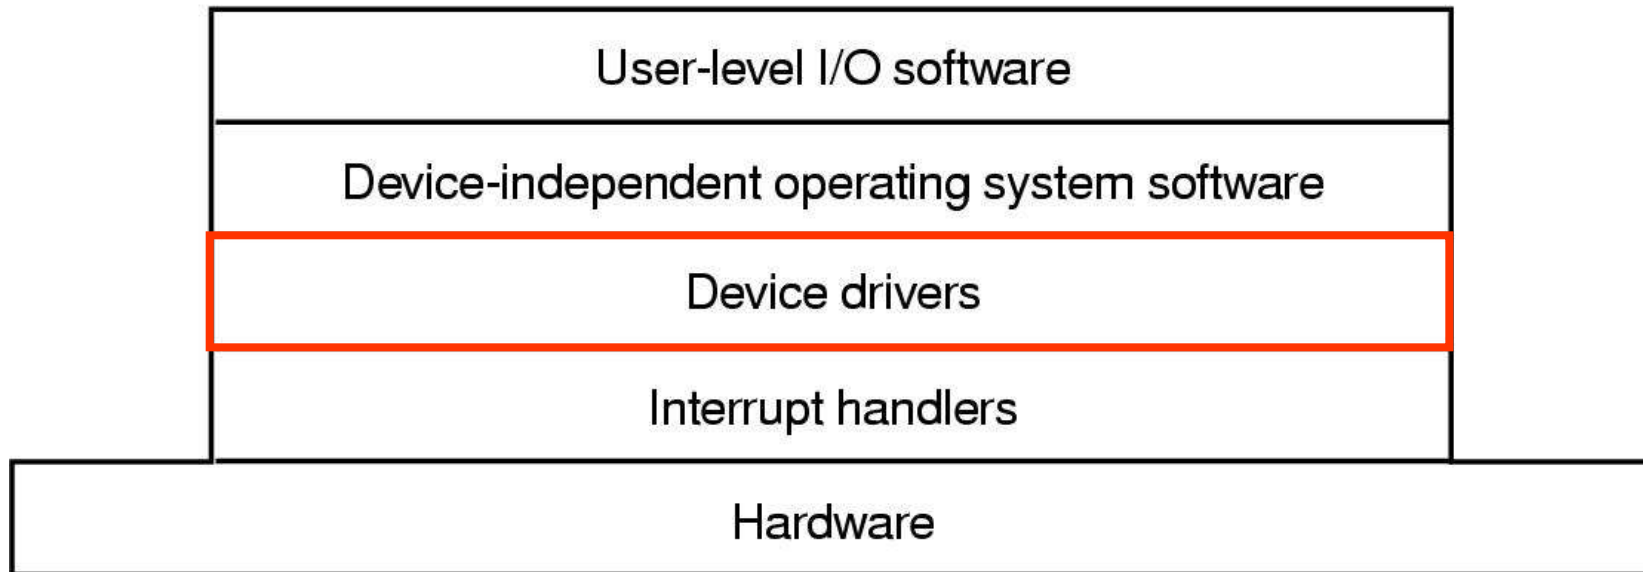
Top and Bottom Halves

- Interrupt handlers are divided into *scheduled* and *non scheduled* tasks
- Non-scheduled tasks execute immediately on interrupt and run in the context of the interrupted thread
 - *ie. There is no VM context switch*
 - *They should do a minimum amount of work so as not to disrupt progress of interrupted thread*
 - *They should minimize time during which interrupts are disabled*
- Scheduled tasks are queued for processing by a thread
 - *This thread will be scheduled to run later*
 - *May be scheduled preemptively or non-preemptively*

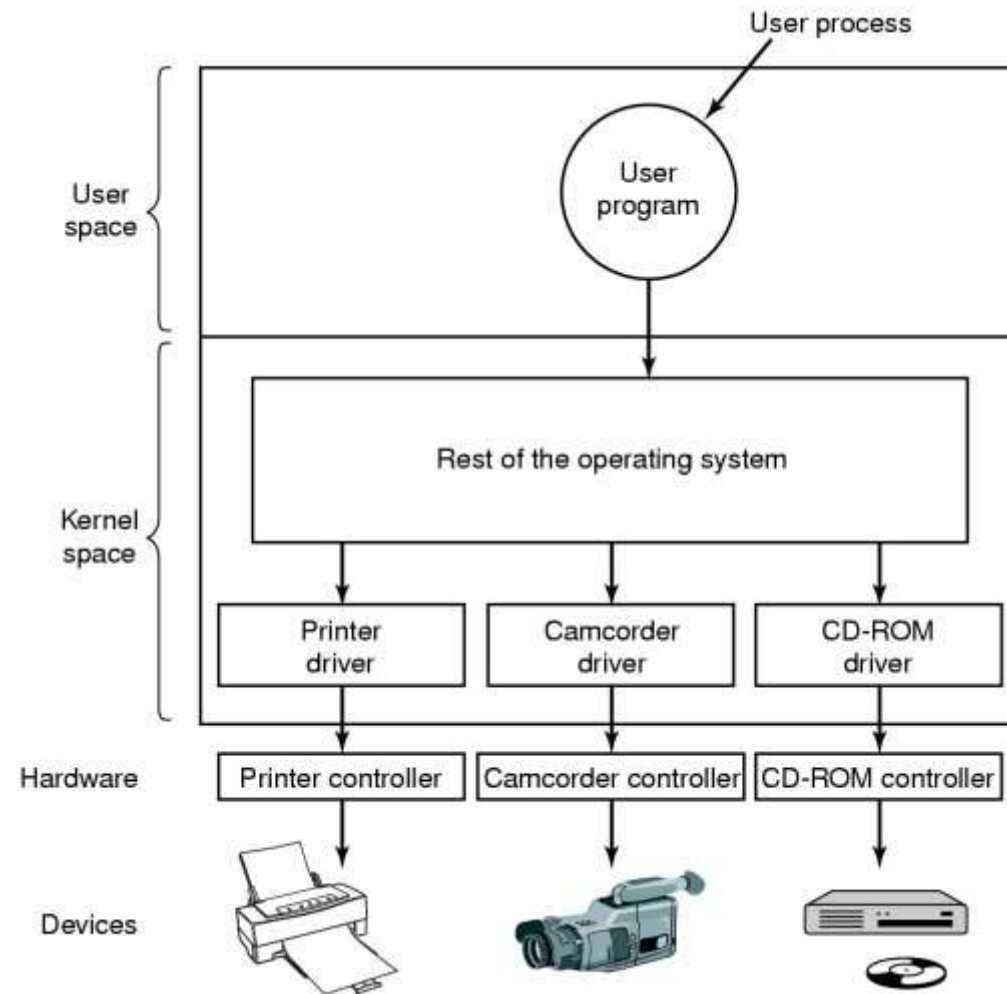
Interrupt Handler's Jobs

- Set up stack for interrupt service procedure
- Ack interrupt controller, re-enable interrupts
- Copy registers from where saved
- Run service procedure

I/O Software Layers



Device drivers in kernel space



Device Drivers

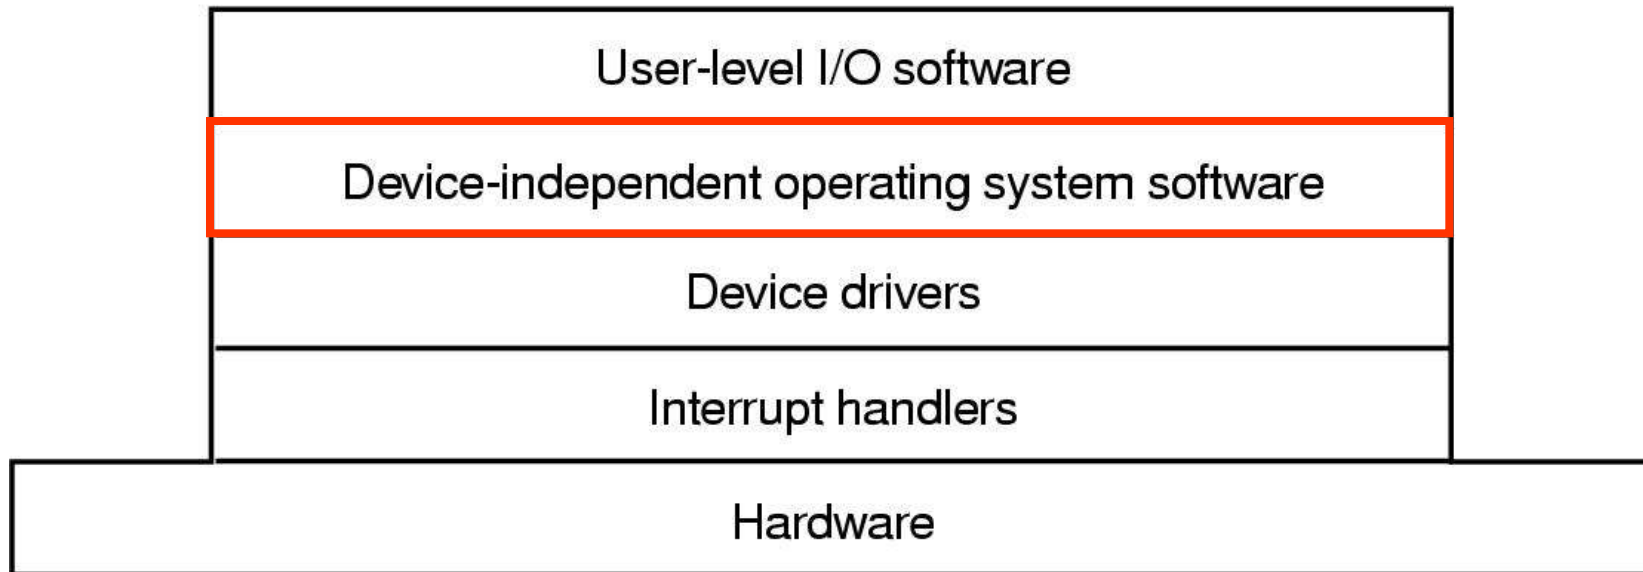
■ Device drivers are device-specific software that connects devices with the operating system

- *Typically an assembly-level job*
 - Must deal with hardware-specific details
 - Must deal with O.S. specific details
- *Goal: hide as many device-specific details as possible from higher level software*

■ Device drivers are typically given kernel privileges for efficiency

- *Bugs can bring down the O.S.!*
- *Open challenge: how to provide efficiency and safety?*

I/O Software Layers



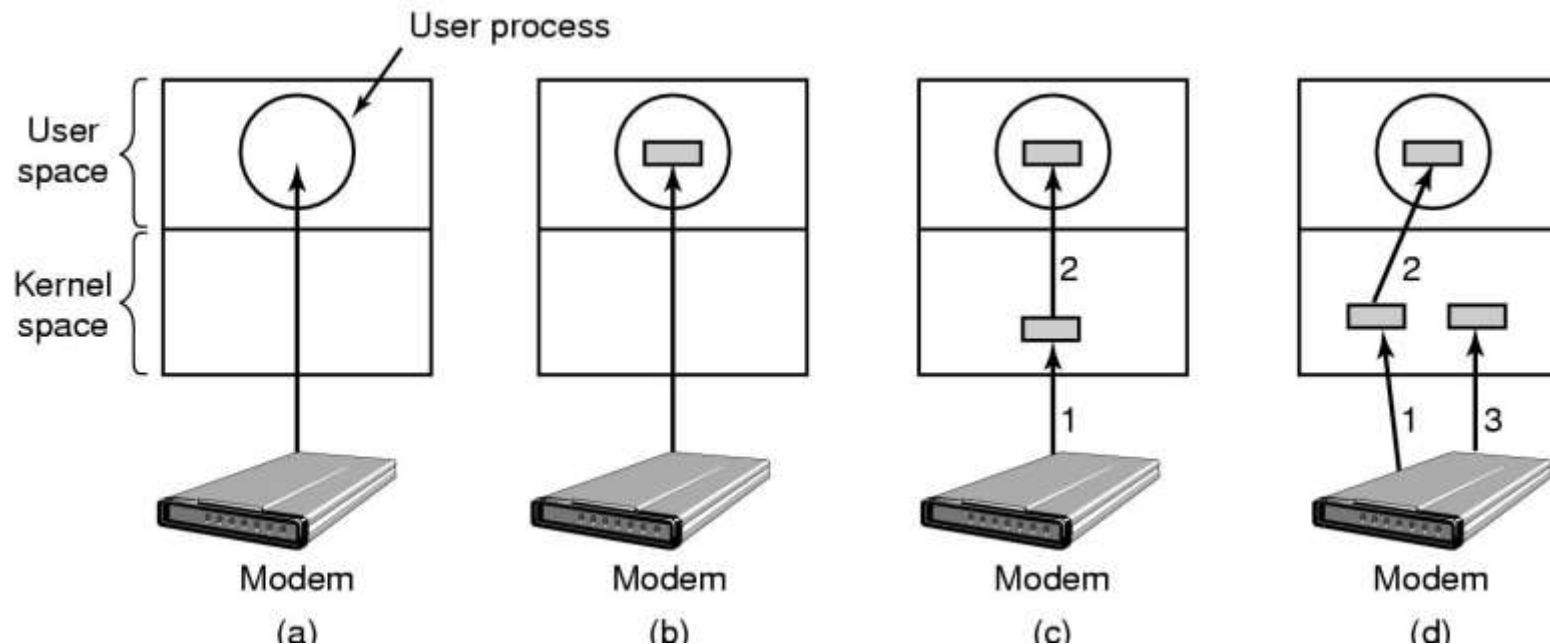
Device-Independent I/O Software

■ Functions and responsibilities

- *Uniform interfacing for device drivers*
- *Buffering*
- *Error reporting*
- *Allocating and releasing dedicated devices*
- *Providing a device-independent block size*

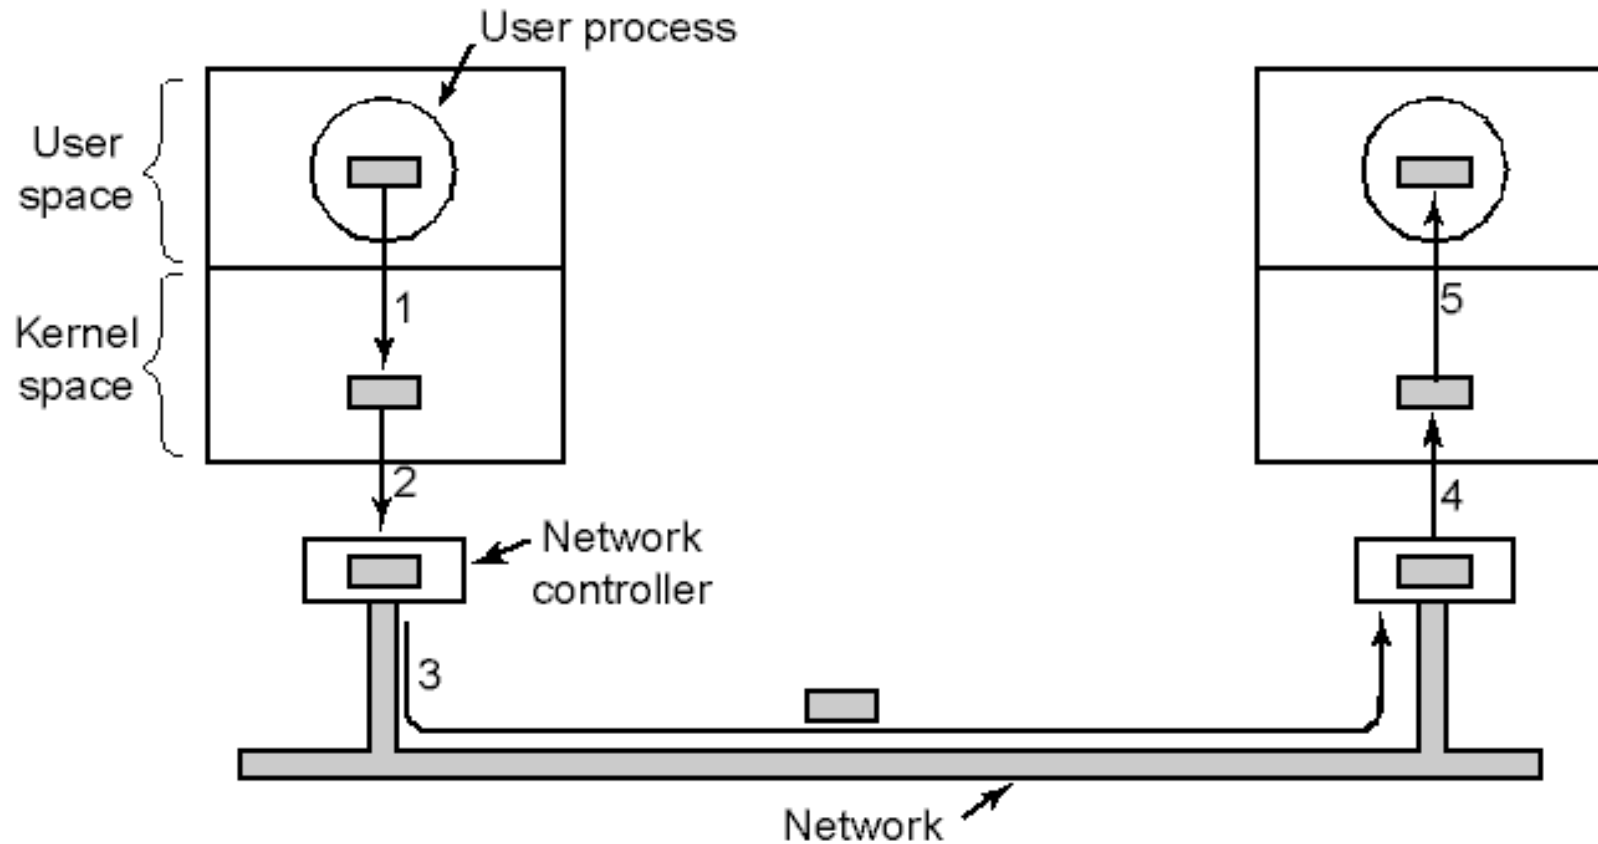
Device-Independent I/O Buffering

- (a) Unbuffered input
- (b) Buffering in user space
- (c) Buffering in the kernel followed by copying to user space
- (d) Double buffering in the kernel




Copying Overhead in Network I/O

Networking may involve many copies





Devices As Files

- Before mounting,
 - *files on floppy are inaccessible*
- After mounting floppy on b,
 - *files on floppy are part of file hierarchy*

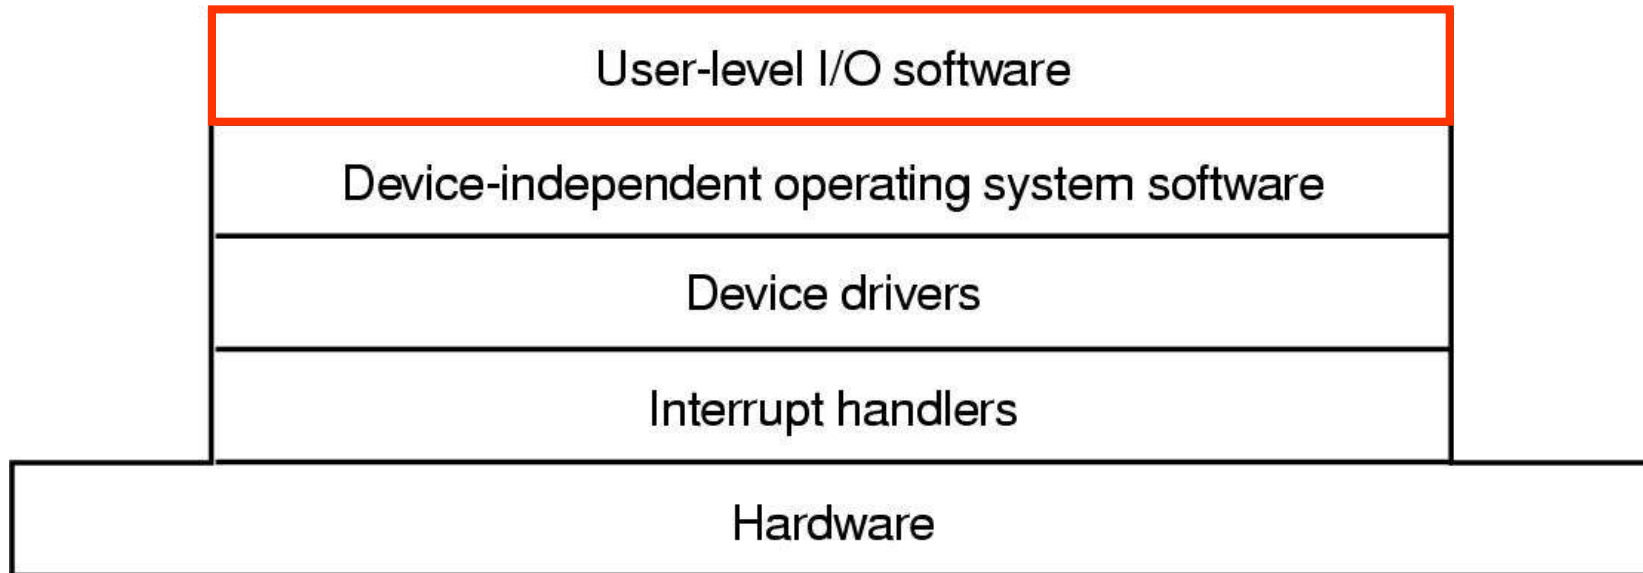


(a)



(b)

I/O Software Layers



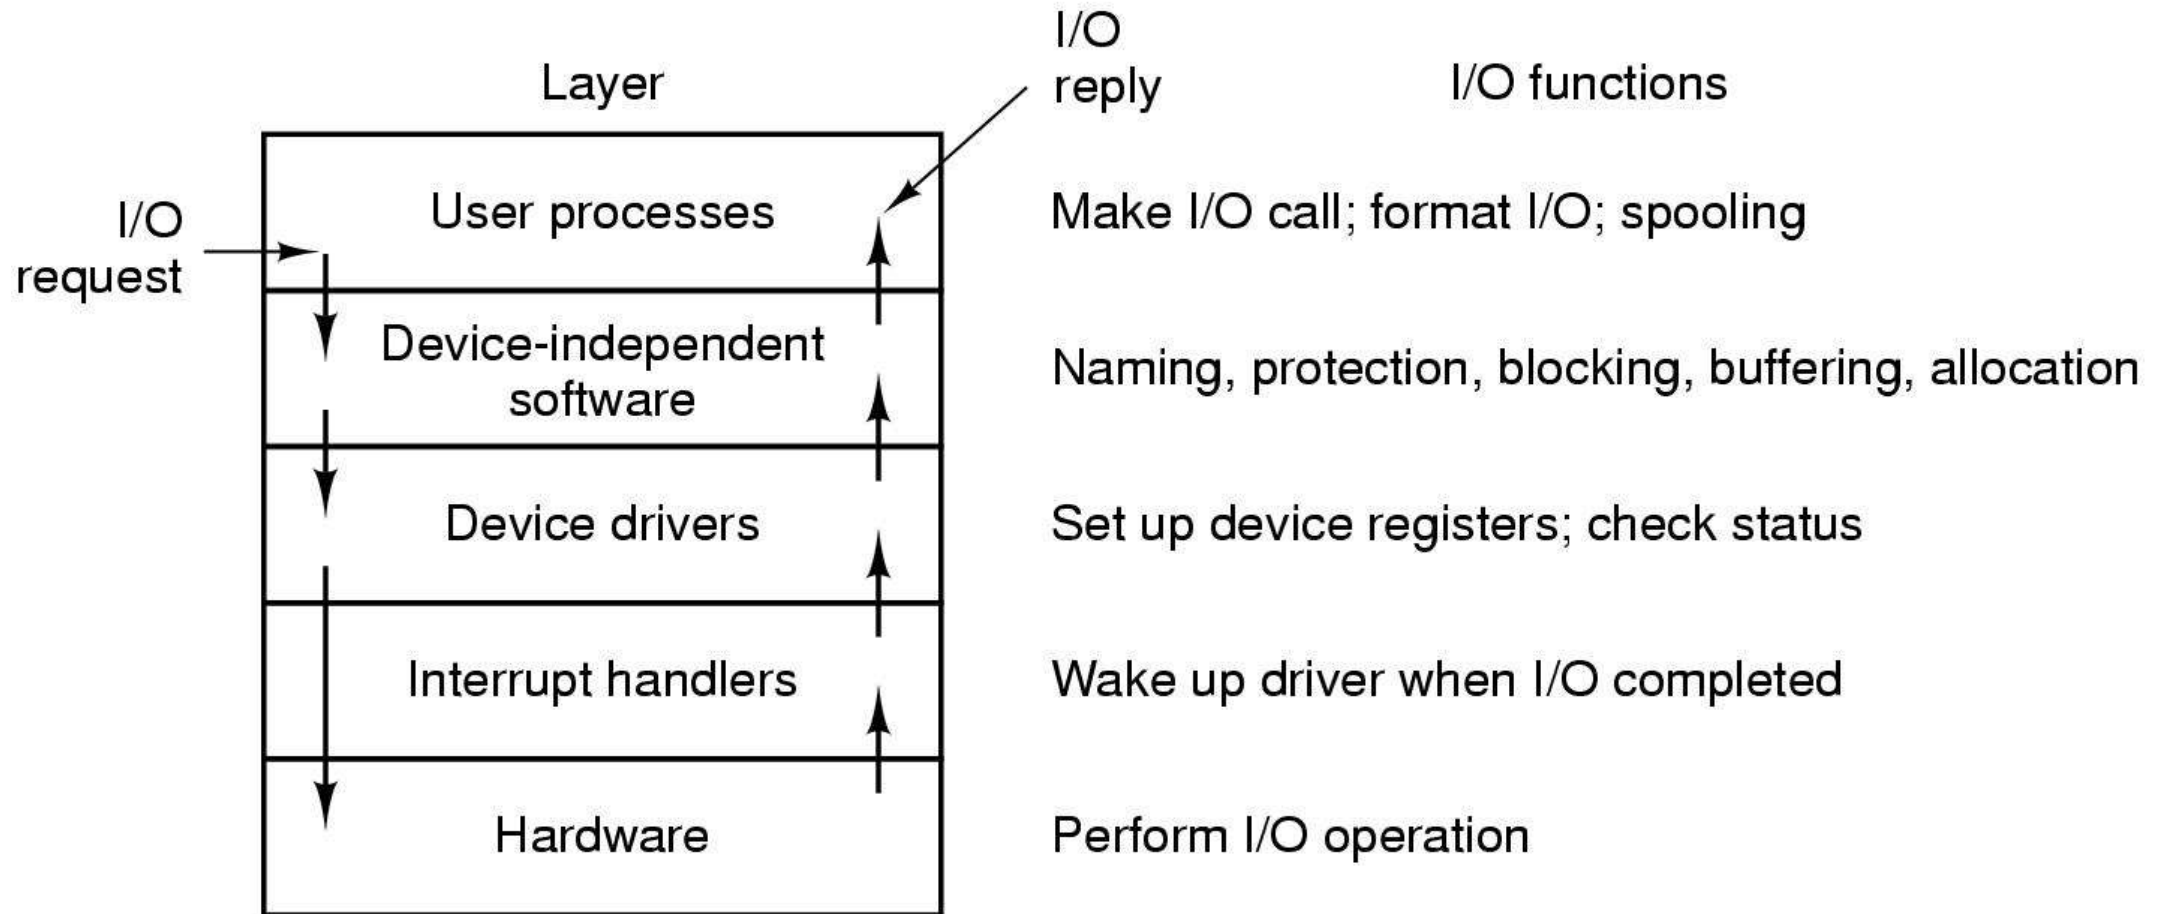
User-Space I/O Software

- In user's (C) program


```
count = write (fd, buffer, nbytes);  
printf ("The value of %s is %d\n", str, i);
```

- Linked with library routines.
- The library routines contain:
 - *Lots of code*
 - *Buffering*
 - *The syscall to trap into the kernel*

Communicating Across I/O Layers



Programmable Timer



- One-shot mode:

- Counter initialized then decremented until zero
- At zero a single interrupt occurs

- Square wave mode:

- At zero the counter is reinitialized with the same value
- Periodic interrupts (called “clock ticks”) occur

Time

- 500 MHz Crystal (oscillates every 2 nanoseconds)
- 32 bit register overflows in 8.6 seconds
 - *So how can we remember what the time is?*
- Backup clock
 - *Similar to digital watch*
 - *Low-power circuitry, battery-powered*
 - *Periodically reset from the internet*
 - *UTC: Universal Coordinated Time*
 - *Unix: Seconds since Jan. 1, 1970*
 - *Windows: Seconds since Jan. 1, 1980*


Goals of Timer Software

- Maintain time of day
 - - *Must update the time-of-day every tick*
- Prevent processes from running too long
- Account for CPU usage
 - - *Separate timer for every process*
 - - *Charge each tick to the current process*
- Handling the “Alarm” syscall
 - - *User programs ask to be sent a signal at a given time*
- Providing watchdog timers for the OS itself
 - - *When to stop the disk, switch to low power mode, etc*
- Doing profiling, monitoring, and statistics gathering

Software Timers

- A process can ask for notification (alarm) at time T
 - *At time T , the OS will signal the process*
- Processes can “go to sleep until time T ”
- Several processes can have active timers
- The CPU has only one clock
 - *Must service the “alarms” in the right order*
- Keep a sorted list of all timers
 - *Each entry tells when the alarm goes off and what to do then*

Software Timers



- Alarms set for 4203, 4207, 4213, 4215 and 4216.
- Each entry tells how many ticks past the previous entry.
- On each tick, decrement the “NextSignal”.
- When it gets to 0, then signal the process.