

R

بسم الله الرحمن الرحيم

سیستم عامل

جلسه بیست و یکم – حافظه‌ی جانبی (۱)

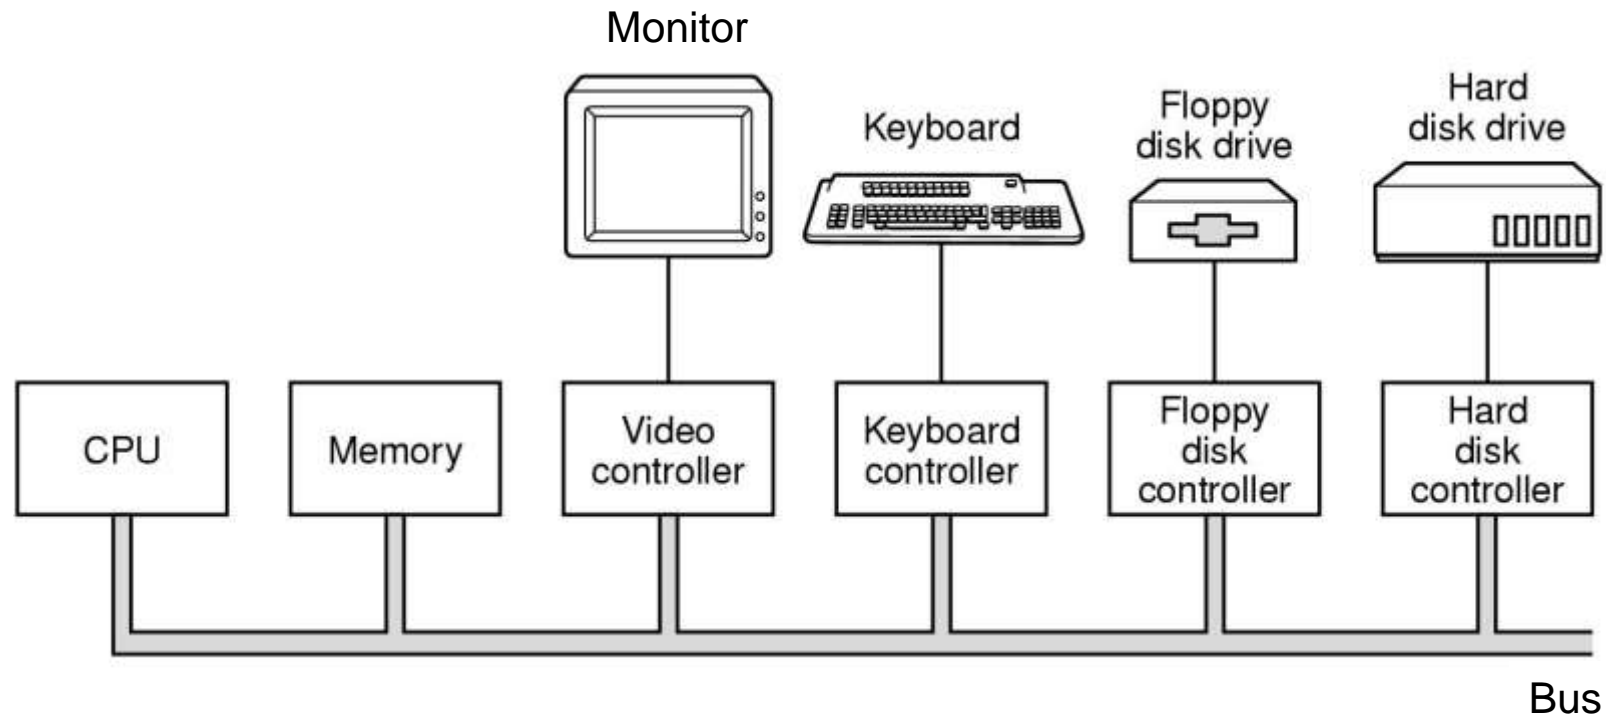
جلسه‌ی گذشته

Device Terminology

- Device (mechanical hardware)
- Device controller (electrical hardware)
- Device driver (software)

Devices & Controllers

■ Components of a simple personal computer



پرسش؟ این دیوایس کنترلر دقیقا کجاست؟

■ مثلا الان به مانیتور وصل کنیم، دیوایس کنترلر توی مانیتور هست؟ یا قبلا توی مادربردمون قرار گرفته؟ موس و پرینتر و ... چطور؟

■ جواب: هر کدام به جوری هستند!

Location Of Device Controller

Device	Location of Device Controller
Mouse/Keyboard	On the motherboard (USB or PS/2 controller)
Speakers	Sound card (or integrated motherboard audio controller)
Hard Disk (HDD/SSD)	Inside the drive (HDD controller or SSD controller)
Printer	Inside the printer (Printer controller)
Network Card (NIC)	On the network card (Ethernet or Wi-Fi adapter controller)
USB Devices	On the motherboard (USB host controller) and inside each USB device

Device Controllers

■ The Device vs. its Controller

■ Some duties of a device controller:

- *Interface between CPU and the Device*
- *Start/Stop device activity*
- *Convert serial bit stream to a block of bytes*
- *Deal with error detection/correction*
- *Move data to/from main memory*

■ Some controllers may handle several (similar) devices

I/O Ports

- Each port has a separate number.
- CPU has special I/O instructions

- *in* *r4, 3*


- *out* *3, r4*

The I/O Port Number

- Port numbers form an “address space” ... separate from main memory
- Contrast with
 - *load* *r4, 3*
 - *store* *3, r4*

Memory-Mapped I/O

- One address space for
 - *main memory*
 - *I/O devices*
- CPU has no special instructions
 - *load r4, addr*
 - *store addr, r4*
- I/O devices are “mapped” into
 - *very high addresses*



I/O Device Speed





Figure 12.11 Common PC and data-center I/O device and interface speeds.

Programmed I/O



Steps in printing a string

Interrupt-Driven I/O

■ Getting the I/O started:

```
CopyFromUser(virtAddr, kernelBuffer,  
byteCount)
```

```
EnableInterrupts()
```

```
while *serialStatusReg != READY
```

```
endWhile
```

```
*serialDataReg = kernelBuffer[0]
```

```
Sleep ()
```

Interrupt-Driven I/O

■The Interrupt Handler:

```
if i == byteCount
```

```
    Wake up the user process
```

```
else
```

```
    *serialDataReg = kernelBuffer[i]
```

```
    i = i + 1
```

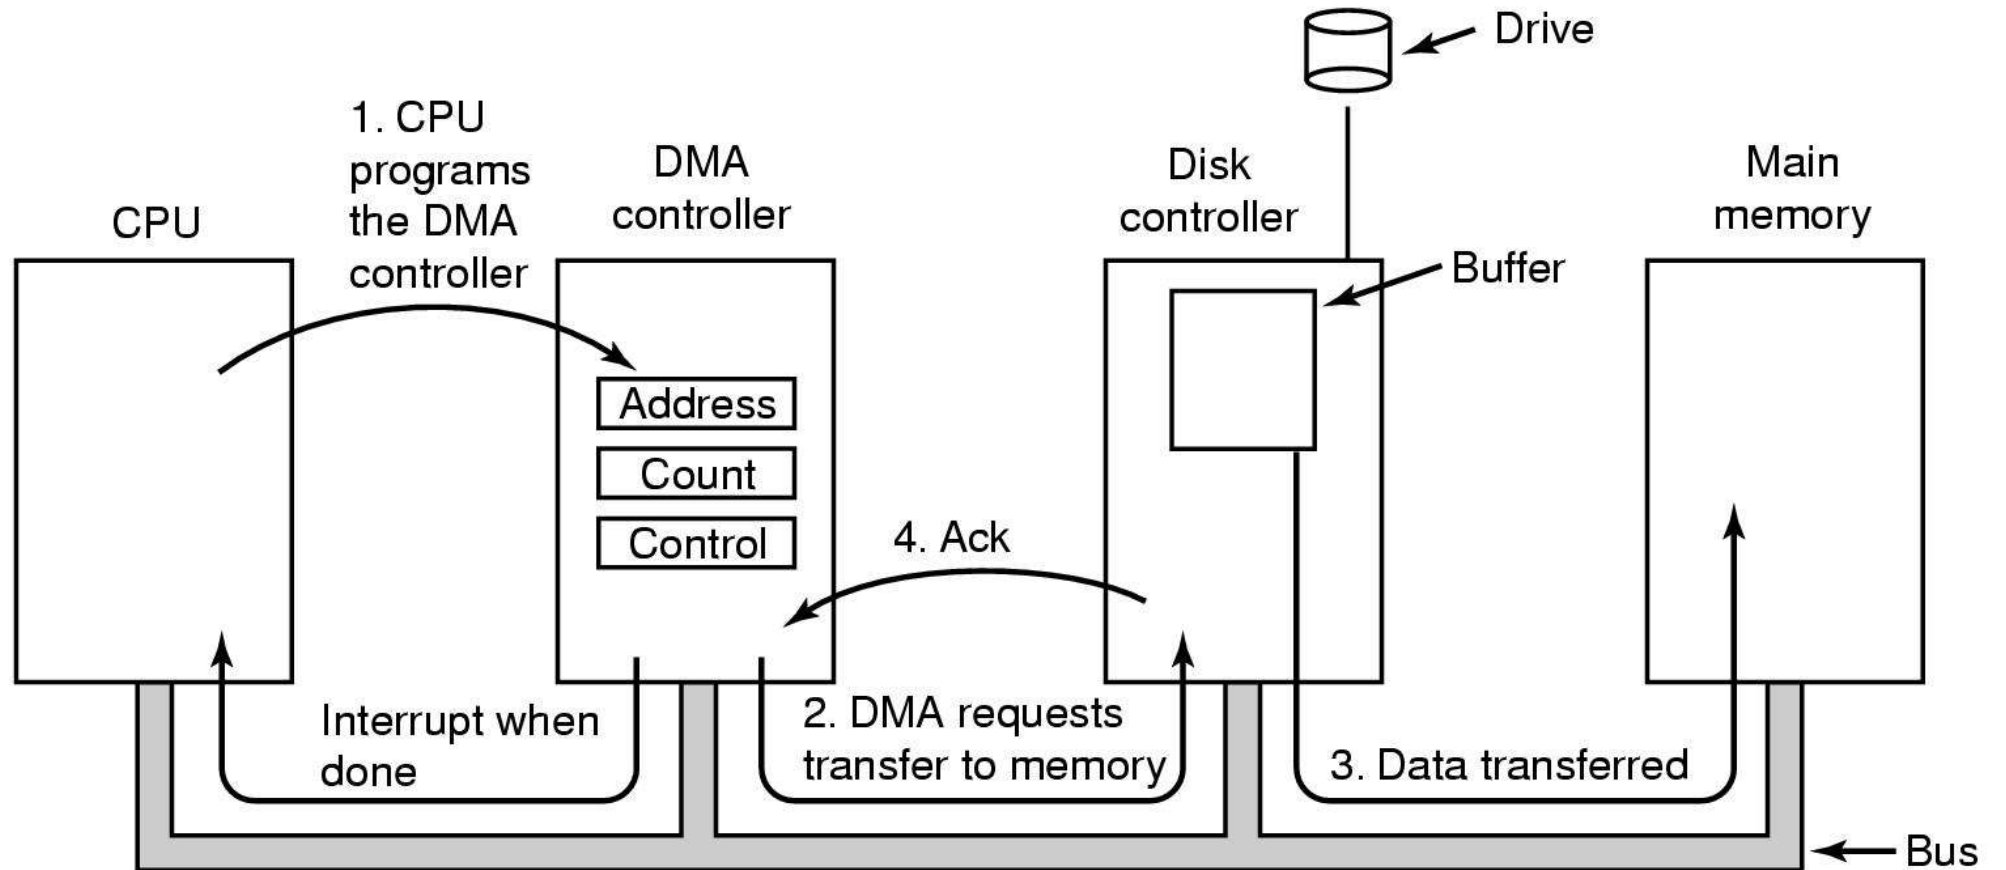
```
endIf
```

```
Return from interrupt
```

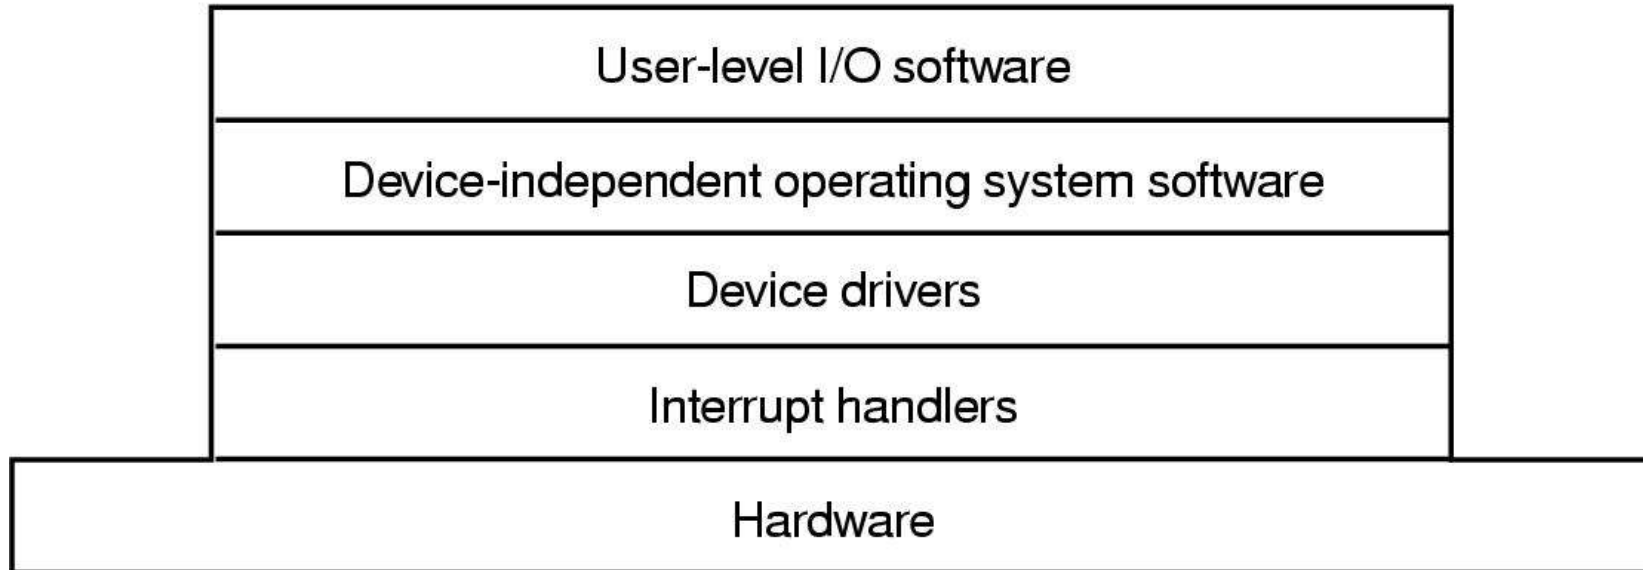
Direct Memory Access (DMA)

- Data transferred from device straight to/from memory
- CPU not involved
- *The DMA controller:*
 - *Does the work of moving the data*
 - *CPU sets up the DMA controller (“programs it ”)*
 - *CPU continues*
 - *The DMA controller moves the bytes*

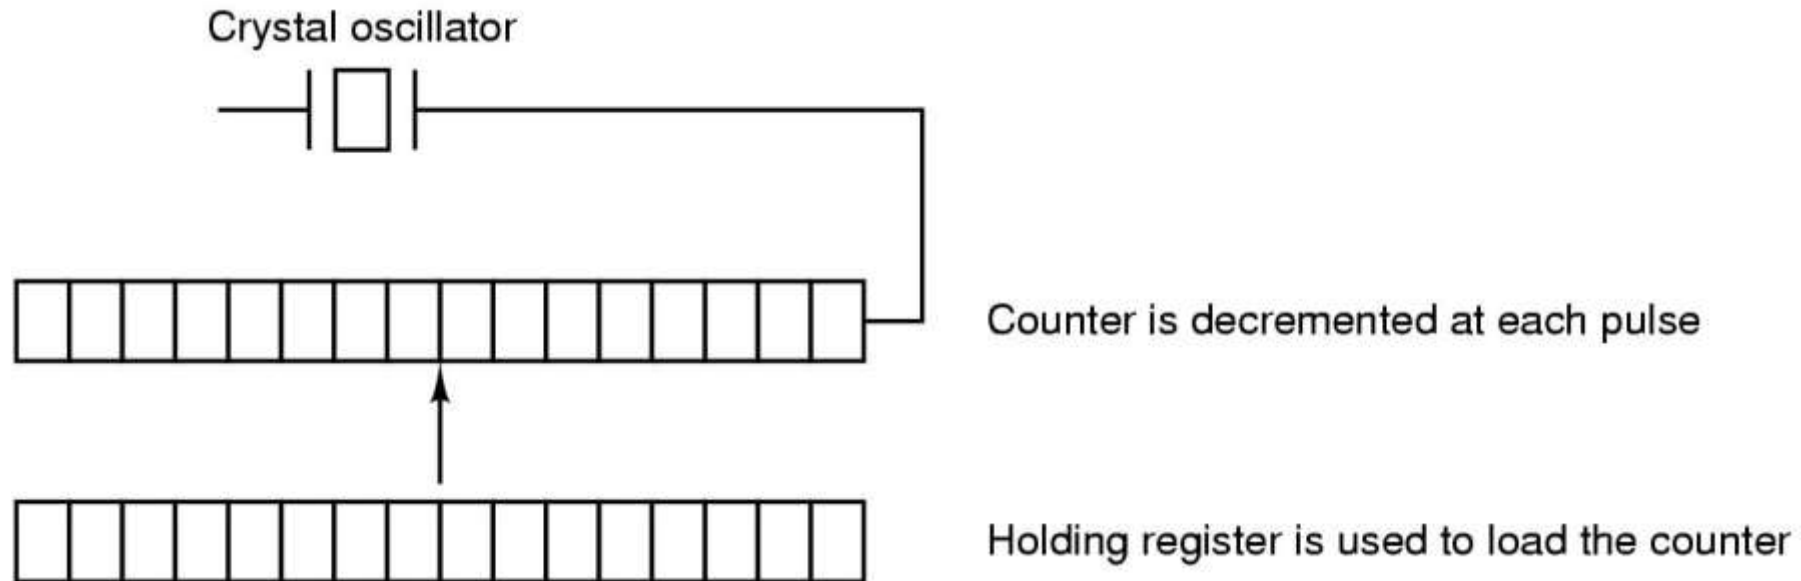
Direct Memory Access (DMA)



I/O Software Layers



Programmable Timer



- One-shot mode:

- Counter initialized then decremented until zero
- At zero a single interrupt occurs

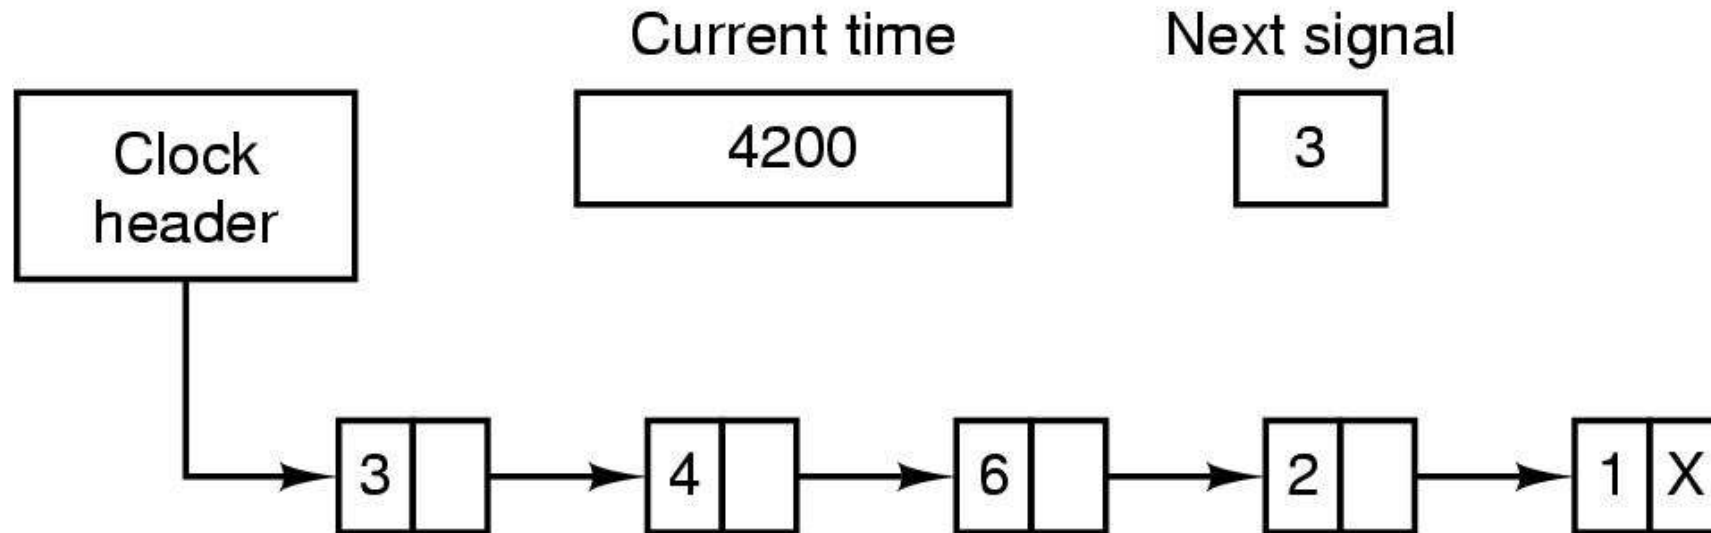
- Square wave mode:

- At zero the counter is reinitialized with the same value
- Periodic interrupts (called “clock ticks”) occur

Goals of Timer Software

- Maintain time of day
 - - *Must update the time-of-day every tick*
- Prevent processes from running too long
- Account for CPU usage
 - - *Separate timer for every process*
 - - *Charge each tick to the current process*
- Handling the “Alarm” syscall
 - - *User programs ask to be sent a signal at a given time*
- Providing watchdog timers for the OS itself
 - - *When to stop the disk, switch to low power mode, etc*
- Doing profiling, monitoring, and statistics gathering

Software Timers



- Alarms set for 4203, 4207, 4213, 4215 and 4216.
- Each entry tells how many ticks past the previous entry.
- On each tick, decrement the “NextSignal”.
- When it gets to 0, then signal the process.



جلسه جدید

Overview of Mass Storage Structure


- Bulk of secondary storage for modern computers is **hard disk drives (HDDs)** and **nonvolatile memory (NVM)** devices
- **HDDs** spin platters of magnetically-coated material under moving read-write heads
 - *Drives rotate at 60 to 250 times per second*
 - **Transfer rate** is rate at which data flow between drive and computer
 - **Positioning time (random-access time)** is time to move disk arm to desired cylinder (**seek time**) and time for desired sector to rotate under the disk head (**rotational latency**)
 - **Head crash** results from disk head making contact with the disk surface -- That's bad
- Disks can be removable

Disk Geometry

- Disk head, surfaces, tracks, sectors ...




Disk Geometry



Example Disk Characteristics

Parameter	IBM 360-KB floppy disk	WD 18300 hard disk
Number of cylinders	40	10601
Tracks per cylinder	2	12
Sectors per track	9	281 (avg)
Sectors per disk	720	35742000
Bytes per sector	512	512
Disk capacity	360 KB	18.3 GB
Seek time (adjacent cylinders)	6 msec	0.8 msec
Seek time (average case)	77 msec	6.9 msec
Rotation time	200 msec	8.33 msec
Motor stop/start time	250 msec	20 sec
Time to transfer 1 sector	22 msec	17 μ sec

Disk Surface Geometry



Constant rotation speed

- Want constant bit density

Inner tracks:

- Fewer sectors per track

Outer tracks:


- More sectors per track

Virtual Geometry


- Physical Geometry

- *The actual layout of sectors on the disk may be complicated*
- *The disk controller does the translation*
- *The CPU sees a “virtual geometry”.*

Virtual Geometry



physical geometry



virtual geometry


Sector Formatting

- A disk sector




- Typically
 - *512 bytes / sector by 2010, (now about 4KB)*
 - *ECC = 16 bytes*


Cylinder Skew




Sector Interleaving



**No
Interleaving**




**Single
Interleaving**



**Double
Interleaving**

Disk Scheduling Algorithms


- Time required to read or write a disk block determined by 3 factors
 - *Seek time*
 - *Rotational delay*
 - *Actual transfer time*
- Seek time dominates
 - *Schedule disk heads to minimize it!*



Disk Scheduling Algorithms

- First-come first serve
- Shortest seek time first
- Scan → back and forth to ends of disk
- C-Scan → only one direction
- Look → back and forth to last request
- C-Look → only one direction

Shortest Seek First (SSF)




Shortest Seek First (SSF)

- Cuts arm motion in half
- Fatal problem:
 - *Starvation is possible!*

The Elevator Algorithm

- Use one bit to track which direction the arm is moving
 - *Up*
 - *Down*
- Keep moving in that direction
- Service the next pending request in that direction
- When there are no more requests in the current direction, reverse direction

The Elevator Algorithm (SCAN)



Other Algorithms

- First-come first serve
- Shortest seek time first
- Scan → back and forth to ends of disk
- C-Scan → only one direction
- Look → back and forth to last request
- C-Look → only one direction